

hw4

November 8, 2023

0.0.1 Lab2 (a) Model preperation

```
[21]: from resnet20 import ResNetCIFAR
      from train_util import train, finetune, test
      import torch
      import numpy as np
      import matplotlib.pyplot as plt

      import time

      import torchvision.transforms as transforms
      import torchvision
      import torch.nn as nn
      import torch.optim as optim

      from FP_layers import *

      device = 'cuda' if torch.cuda.is_available() else 'cpu'
      device
```

```
[21]: 'cuda'
```

```
[22]: net = ResNetCIFAR(num_layers=20, Nbits=None)
      net = net.to(device)
```

```
[3]: # Load the best weight paramters
      net.load_state_dict(torch.load("pretrained_model.pt"))
      test(net)
```

Files already downloaded and verified
Test Loss=0.3231, Test accuracy=0.9150

```
[3]: 0.915
```

0.0.2 Lab2 (b) Prune by percentage

```
[4]: def prune_by_percentage(layer, q):  
    """  
    Pruning the weight paramters by threshold.  
    :param q: pruning percentile. 'q' percent of the least  
    significant weight parameters will be pruned.  
    """  
    with torch.no_grad():  
        # Convert the weight of "layer" to numpy array  
        layer_weight = layer.weight.detach().cpu().numpy()  
        # Compute the q-th percentile of the abs of the converted array  
        percentile = np.percentile(np.abs(layer_weight.flatten()), q)  
        # Generate a binary mask same shape as weight to decide which element  
        ↪to prune  
        masked_obj = np.ma.masked_greater_equal(x=np.abs(layer_weight),  
        ↪value=percentile, copy=True)  
        mask_int = np.ma.getmask(masked_obj).astype(int)  
        # Convert mask to torch tensor and put on GPU  
        mask_tensor = torch.tensor(mask_int).to(device)  
        # Multiply the weight by mask to perform pruning  
        assert mask_int.shape == layer_weight.shape  
        # layer.weight.data = mask_tensor * layer.weight.data  
        layer.weight.data = layer.weight.data.clone().detach().  
        ↪requires_grad_(True) * mask_tensor  
    return
```

```
[5]: q_list = np.array([0.3, 0.5, 0.7]) * 100  
  
test_acc_prune_by_perc = []  
  
for q_val in q_list:  
    net.load_state_dict(torch.load("pretrained_model.pt"))  
    for name, layer in net.named_modules():  
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and  
        ↪'id_mapping' not in name:  
            # change q value  
            prune_by_percentage(layer, q=q_val)  
            # break  
            ## Optional: Check the sparsity you achieve in each layer  
            ## Convert the weight of "layer" to numpy array  
            np_weight = layer.weight.detach().cpu().numpy()  
            ## Count number of zeros  
            zeros = sum((np_weight == 0).flatten())  
            ## Count number of parameters  
            total = len(np_weight.flatten())
```

```

    ## Print sparsity
    # print('Sparsity of ' + name + ': ' + str(zeros/total))
    print('Sparsity of %s: %g' % (name, zeros/total))
test_acc_prune_by_perc.append(test(net))

```

```

Sparsity of head_conv.0.conv: 0.300926
Sparsity of body_op.0.conv1.0.conv: 0.299913
Sparsity of body_op.0.conv2.0.conv: 0.299913
Sparsity of body_op.1.conv1.0.conv: 0.299913
Sparsity of body_op.1.conv2.0.conv: 0.299913
Sparsity of body_op.2.conv1.0.conv: 0.299913
Sparsity of body_op.2.conv2.0.conv: 0.299913
Sparsity of body_op.3.conv1.0.conv: 0.30013
Sparsity of body_op.3.conv2.0.conv: 0.300022
Sparsity of body_op.4.conv1.0.conv: 0.300022
Sparsity of body_op.4.conv2.0.conv: 0.300022
Sparsity of body_op.5.conv1.0.conv: 0.300022
Sparsity of body_op.5.conv2.0.conv: 0.300022
Sparsity of body_op.6.conv1.0.conv: 0.300022
Sparsity of body_op.6.conv2.0.conv: 0.299995
Sparsity of body_op.7.conv1.0.conv: 0.299995
Sparsity of body_op.7.conv2.0.conv: 0.299995
Sparsity of body_op.8.conv1.0.conv: 0.299995
Sparsity of body_op.8.conv2.0.conv: 0.299995
Sparsity of final_fc.linear: 0.3
Files already downloaded and verified
Test Loss=0.3699, Test accuracy=0.9027
Sparsity of head_conv.0.conv: 0.5
Sparsity of body_op.0.conv1.0.conv: 0.5
Sparsity of body_op.0.conv2.0.conv: 0.5
Sparsity of body_op.1.conv1.0.conv: 0.5
Sparsity of body_op.1.conv2.0.conv: 0.5
Sparsity of body_op.2.conv1.0.conv: 0.5
Sparsity of body_op.2.conv2.0.conv: 0.5
Sparsity of body_op.3.conv1.0.conv: 0.5
Sparsity of body_op.3.conv2.0.conv: 0.5
Sparsity of body_op.4.conv1.0.conv: 0.5
Sparsity of body_op.4.conv2.0.conv: 0.5
Sparsity of body_op.5.conv1.0.conv: 0.5
Sparsity of body_op.5.conv2.0.conv: 0.5
Sparsity of body_op.6.conv1.0.conv: 0.5
Sparsity of body_op.6.conv2.0.conv: 0.5
Sparsity of body_op.7.conv1.0.conv: 0.5
Sparsity of body_op.7.conv2.0.conv: 0.5
Sparsity of body_op.8.conv1.0.conv: 0.5
Sparsity of body_op.8.conv2.0.conv: 0.5
Sparsity of final_fc.linear: 0.5

```

```

Files already downloaded and verified
Test Loss=0.6776, Test accuracy=0.8209
Sparsity of head_conv.0.conv: 0.699074
Sparsity of body_op.0.conv1.0.conv: 0.700087
Sparsity of body_op.0.conv2.0.conv: 0.700087
Sparsity of body_op.1.conv1.0.conv: 0.700087
Sparsity of body_op.1.conv2.0.conv: 0.700087
Sparsity of body_op.2.conv1.0.conv: 0.700087
Sparsity of body_op.2.conv2.0.conv: 0.700087
Sparsity of body_op.3.conv1.0.conv: 0.69987
Sparsity of body_op.3.conv2.0.conv: 0.699978
Sparsity of body_op.4.conv1.0.conv: 0.699978
Sparsity of body_op.4.conv2.0.conv: 0.699978
Sparsity of body_op.5.conv1.0.conv: 0.699978
Sparsity of body_op.5.conv2.0.conv: 0.699978
Sparsity of body_op.6.conv1.0.conv: 0.699978
Sparsity of body_op.6.conv2.0.conv: 0.700005
Sparsity of body_op.7.conv1.0.conv: 0.700005
Sparsity of body_op.7.conv2.0.conv: 0.700005
Sparsity of body_op.8.conv1.0.conv: 0.700005
Sparsity of body_op.8.conv2.0.conv: 0.700005
Sparsity of final_fc.linear: 0.7
Files already downloaded and verified
Test Loss=2.4418, Test accuracy=0.4207

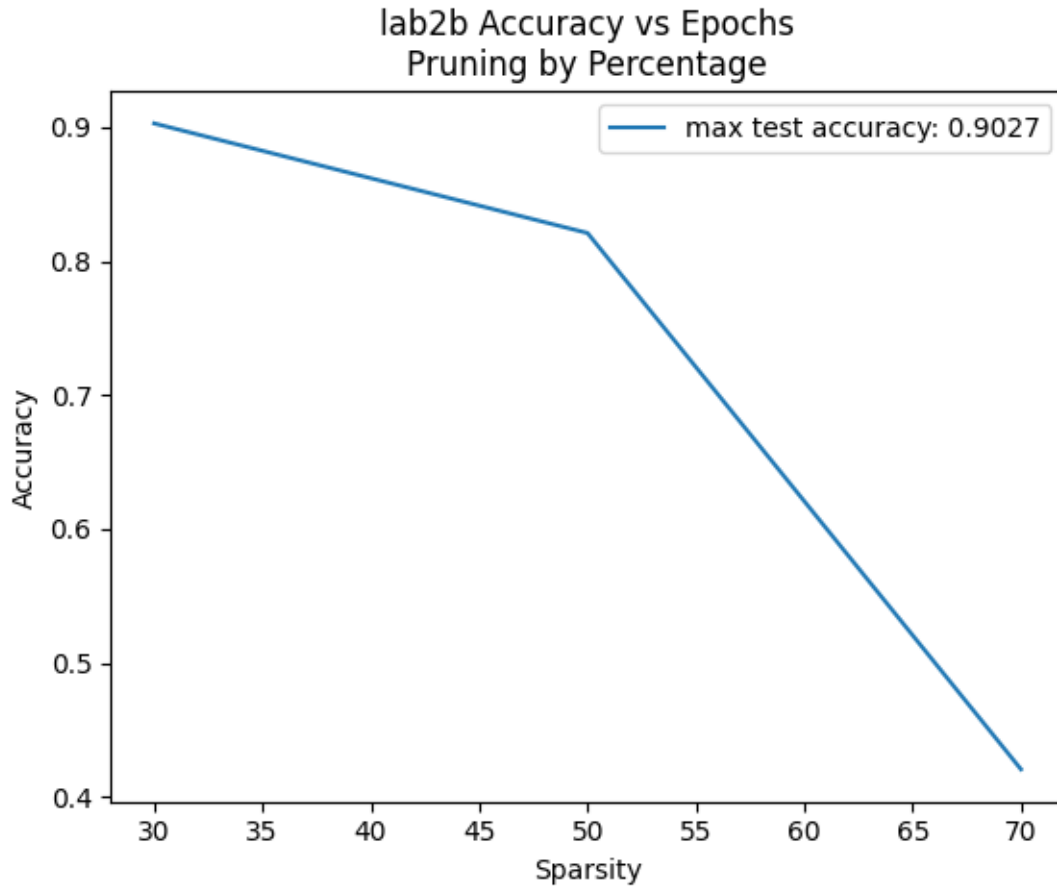
```

```

[6]: fig, ax = plt.subplots(1, 1)
ax.plot(q_list, test_acc_prune_by_perc, label='max test accuracy: %g' % np.
    ↳max(test_acc_prune_by_perc))
ax.set_xlabel('Sparsity')
ax.set_ylabel('Accuracy')
ax.set_title('lab2b Accuracy vs Epochs\nPruning by Percentage')
ax.legend()

plt.savefig('lab2b.pdf', dpi=500, bbox_inches='tight')

```



0.0.3 Lab2 (c) Finetune pruned model

```
[7]: def finetune_after_prune(net, trainloader, criterion, optimizer, prune=True):
    """
    Finetune the pruned model for a single epoch
    Make sure pruned weights are kept as zero
    """
    # Build a dictionary for the nonzero weights
    weight_mask = {}
    for name, layer in net.named_modules():
        if isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear) and
        ↪ 'id_mapping' not in name:
            # Your code here: generate a mask in GPU torch tensor to have 1 for
            ↪ nonzero element and 0 for zero element
            layer_weight = layer.weight.detach().cpu().numpy()
            mask_obj = np.ma.masked_equal(x=layer_weight, value=0) # where
            ↪ it's zero, mask value is True
```

```

        weight_mask[name] = torch.tensor(1 - np.ma.getmask(mask_obj).
↪astype(int)).to(device)
        # zero = 0, nonzero = 1

    global_steps = 0
    train_loss = 0
    correct = 0
    total = 0
    start = time.time()
    for batch_idx, (inputs, targets) in enumerate(trainloader):
        inputs, targets = inputs.to(device), targets.to(device)
        optimizer.zero_grad()
        outputs = net(inputs)
        loss = criterion(outputs, targets)
        loss.backward()
        optimizer.step()

        if prune:
            for name, layer in net.named_modules():
                if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.
↪Linear)) and 'id_mapping' not in name:
                    # Your code here: Use weight_mask to make sure zero
↪elements remains zero
                    layer.weight.data = layer.weight.data.clone().detach().
↪requires_grad_(True) * weight_mask[name]

        train_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()
        global_steps += 1

        if global_steps % 50 == 0:
            end = time.time()
            batch_size = 256
            num_examples_per_second = 50 * batch_size / (end - start)
            print("[Step=%d]\tLoss=%.4f\tacc=%.4f\t%.1f examples/second"
                % (global_steps, train_loss / (batch_idx + 1), (correct /
↪total), num_examples_per_second))
            start = time.time()

```

```

[8]: # Get pruned model
net.load_state_dict(torch.load("pretrained_model.pt"))
for name, layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
↪'id_mapping' not in name:

```

```

        prune_by_percentage(layer, q=70.0)
lab2c_test_b4_finetune = test(net)
# Training setup, do not change
batch_size = 256
lr = 0.002
reg = 1e-4

print('==> Preparing data..')
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

best_acc = 0 # best test accuracy
start_epoch = 0 # start from epoch 0 or last checkpoint epoch
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
    ↪download=True, transform=transform_train)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
    ↪shuffle=True, num_workers=16)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
    ↪download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=100,
    ↪shuffle=False, num_workers=2)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=lr, momentum=0.875,
    ↪weight_decay=reg, nesterov=False)

```

Files already downloaded and verified
Test Loss=2.4418, Test accuracy=0.4207
==> Preparing data..
Files already downloaded and verified
Files already downloaded and verified

```

[9]: # Model finetuning
test_acc_lst = []
n_epochs = 20
for epoch in range(n_epochs):
    print('\nEpoch: %d' % epoch)
    net.train()

```

```

finetune_after_prune(net, trainloader, criterion, optimizer)
#Start the testing code.
net.eval()
test_loss = 0
correct = 0
total = 0
with torch.no_grad():
    for batch_idx, (inputs, targets) in enumerate(testloader):
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = net(inputs)
        loss = criterion(outputs, targets)

        test_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()
num_val_steps = len(testloader)
val_acc = correct / total
print("Test Loss=%.4f, Test acc=%.4f" % (test_loss / (num_val_steps),
↪val_acc))
test_acc_lst.append(val_acc)
if val_acc > best_acc:
    best_acc = val_acc
    print("Saving...")
    torch.save(net.state_dict(), "net_after_finetune.pt")

```

Epoch: 0

[Step=50]	Loss=0.4015	acc=0.8612	6063.1 examples/second
[Step=100]	Loss=0.3621	acc=0.8752	12533.4 examples/second
[Step=150]	Loss=0.3334	acc=0.8853	12582.2 examples/second

Test Loss=0.4290, Test acc=0.8676

Saving..

Epoch: 1

[Step=50]	Loss=0.2458	acc=0.9155	6556.5 examples/second
[Step=100]	Loss=0.2449	acc=0.9152	12105.4 examples/second
[Step=150]	Loss=0.2388	acc=0.9171	12214.2 examples/second

Test Loss=0.3941, Test acc=0.8752

Saving..

Epoch: 2

[Step=50]	Loss=0.2220	acc=0.9273	6774.9 examples/second
[Step=100]	Loss=0.2218	acc=0.9254	12616.7 examples/second
[Step=150]	Loss=0.2150	acc=0.9276	12099.2 examples/second

Test Loss=0.3787, Test acc=0.8789

Saving..

Epoch: 3
[Step=50] Loss=0.1934 acc=0.9318 6663.5 examples/second
[Step=100] Loss=0.2036 acc=0.9284 12776.3 examples/second
[Step=150] Loss=0.2007 acc=0.9292 12756.8 examples/second
Test Loss=0.3710, Test acc=0.8813
Saving...

Epoch: 4
[Step=50] Loss=0.1908 acc=0.9345 6193.6 examples/second
[Step=100] Loss=0.1856 acc=0.9364 13214.0 examples/second
[Step=150] Loss=0.1860 acc=0.9357 13242.3 examples/second
Test Loss=0.3622, Test acc=0.8831
Saving...

Epoch: 5
[Step=50] Loss=0.1849 acc=0.9354 6596.6 examples/second
[Step=100] Loss=0.1848 acc=0.9354 12605.2 examples/second
[Step=150] Loss=0.1826 acc=0.9359 12794.9 examples/second
Test Loss=0.3581, Test acc=0.8853
Saving...

Epoch: 6
[Step=50] Loss=0.1807 acc=0.9342 6581.0 examples/second
[Step=100] Loss=0.1772 acc=0.9370 12638.2 examples/second
[Step=150] Loss=0.1771 acc=0.9373 12780.2 examples/second
Test Loss=0.3544, Test acc=0.8849

Epoch: 7
[Step=50] Loss=0.1696 acc=0.9415 6120.6 examples/second
[Step=100] Loss=0.1741 acc=0.9398 11072.2 examples/second
[Step=150] Loss=0.1734 acc=0.9403 12073.7 examples/second
Test Loss=0.3505, Test acc=0.8871
Saving...

Epoch: 8
[Step=50] Loss=0.1747 acc=0.9388 5946.8 examples/second
[Step=100] Loss=0.1668 acc=0.9422 11545.4 examples/second
[Step=150] Loss=0.1655 acc=0.9422 12119.1 examples/second
Test Loss=0.3488, Test acc=0.8880
Saving...

Epoch: 9
[Step=50] Loss=0.1605 acc=0.9440 6828.1 examples/second
[Step=100] Loss=0.1617 acc=0.9445 13478.8 examples/second
[Step=150] Loss=0.1620 acc=0.9444 13561.5 examples/second
Test Loss=0.3445, Test acc=0.8899
Saving...

Epoch: 10
[Step=50] Loss=0.1613 acc=0.9432 6316.5 examples/second
[Step=100] Loss=0.1633 acc=0.9413 12868.6 examples/second
[Step=150] Loss=0.1629 acc=0.9414 13039.7 examples/second
Test Loss=0.3437, Test acc=0.8906
Saving...

Epoch: 11
[Step=50] Loss=0.1552 acc=0.9457 6349.5 examples/second
[Step=100] Loss=0.1576 acc=0.9442 12022.5 examples/second
[Step=150] Loss=0.1589 acc=0.9437 13227.5 examples/second
Test Loss=0.3418, Test acc=0.8914
Saving...

Epoch: 12
[Step=50] Loss=0.1499 acc=0.9504 6148.7 examples/second
[Step=100] Loss=0.1527 acc=0.9482 12449.1 examples/second
[Step=150] Loss=0.1522 acc=0.9489 12410.9 examples/second
Test Loss=0.3398, Test acc=0.8914

Epoch: 13
[Step=50] Loss=0.1548 acc=0.9458 6635.9 examples/second
[Step=100] Loss=0.1583 acc=0.9450 12242.1 examples/second
[Step=150] Loss=0.1555 acc=0.9462 12662.9 examples/second
Test Loss=0.3404, Test acc=0.8919
Saving...

Epoch: 14
[Step=50] Loss=0.1570 acc=0.9452 6721.6 examples/second
[Step=100] Loss=0.1480 acc=0.9491 12571.6 examples/second
[Step=150] Loss=0.1514 acc=0.9472 12635.6 examples/second
Test Loss=0.3387, Test acc=0.8911

Epoch: 15
[Step=50] Loss=0.1461 acc=0.9513 6599.7 examples/second
[Step=100] Loss=0.1475 acc=0.9502 12253.1 examples/second
[Step=150] Loss=0.1499 acc=0.9492 12467.3 examples/second
Test Loss=0.3371, Test acc=0.8930
Saving...

Epoch: 16
[Step=50] Loss=0.1482 acc=0.9480 5790.8 examples/second
[Step=100] Loss=0.1453 acc=0.9493 12149.8 examples/second
[Step=150] Loss=0.1480 acc=0.9482 12349.3 examples/second
Test Loss=0.3375, Test acc=0.8919

Epoch: 17

```

[Step=50]      Loss=0.1414      acc=0.9503      5905.5 examples/second
[Step=100]     Loss=0.1441      acc=0.9495      12511.9 examples/second
[Step=150]     Loss=0.1428      acc=0.9504      12075.0 examples/second
Test Loss=0.3341, Test acc=0.8940
Saving...

```

```

Epoch: 18
[Step=50]      Loss=0.1442      acc=0.9498      6657.1 examples/second
[Step=100]     Loss=0.1476      acc=0.9484      12269.5 examples/second
[Step=150]     Loss=0.1461      acc=0.9492      12533.2 examples/second
Test Loss=0.3343, Test acc=0.8947
Saving...

```

```

Epoch: 19
[Step=50]      Loss=0.1399      acc=0.9519      6243.7 examples/second
[Step=100]     Loss=0.1443      acc=0.9498      12571.7 examples/second
[Step=150]     Loss=0.1426      acc=0.9509      12707.4 examples/second
Test Loss=0.3344, Test acc=0.8947


```

```

[10]: fig, ax = plt.subplots(1, 1)
      xx = range(n_epochs)
      ax.plot(xx, test_acc_lst, label='last test accuracy: %g' % test_acc_lst[-1])
      ax.set_xlabel('Epochs')
      ax.set_ylabel('Accuracy')
      ax.set_title('lab2c Accuracy vs Epochs\nFine-tune Pruning, before finetuning,
↳test_acc=%.4f'
                  % lab2c_test_b4_finetune)
      ax.legend()

      plt.savefig('Figures/lab2c.pdf', dpi=500, bbox_inches='tight')

```



```
[11]: # Check sparsity of the finetuned model, make sure it's not changed
net.load_state_dict(torch.load("net_after_finetune.pt"))

for name, layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
        'id_mapping' not in name:
        # Your code here:
        ## Convert the weight of "layer" to numpy array
        np_weight = layer.weight.detach().cpu().numpy()
        ## Count number of zeros
        zeros = sum((np_weight == 0).flatten())
        ## Count number of parameters
        total = len(np_weight.flatten())
        # Print sparsity
        print('Sparsity of %s: %g' % (name, zeros/total))

test(net)
```

Sparsity of head_conv.0.conv: 0.699074

```

Sparsity of body_op.0.conv1.0.conv: 0.700087
Sparsity of body_op.0.conv2.0.conv: 0.700087
Sparsity of body_op.1.conv1.0.conv: 0.700087
Sparsity of body_op.1.conv2.0.conv: 0.700087
Sparsity of body_op.2.conv1.0.conv: 0.700087
Sparsity of body_op.2.conv2.0.conv: 0.700087
Sparsity of body_op.3.conv1.0.conv: 0.699978
Sparsity of body_op.3.conv2.0.conv: 0.699978
Sparsity of body_op.4.conv1.0.conv: 0.699978
Sparsity of body_op.4.conv2.0.conv: 0.699978
Sparsity of body_op.5.conv1.0.conv: 0.699978
Sparsity of body_op.5.conv2.0.conv: 0.699978
Sparsity of body_op.6.conv1.0.conv: 0.699978
Sparsity of body_op.6.conv2.0.conv: 0.700005
Sparsity of body_op.7.conv1.0.conv: 0.700005
Sparsity of body_op.7.conv2.0.conv: 0.700005
Sparsity of body_op.8.conv1.0.conv: 0.700005
Sparsity of body_op.8.conv2.0.conv: 0.700005
Sparsity of final_fc.linear: 0.7
Files already downloaded and verified
Test Loss=0.3343, Test accuracy=0.8947

```

[11]: 0.8947

0.0.4 Lab2 (d) Iterative pruning

```

[12]: net.load_state_dict(torch.load("pretrained_model.pt"))
best_acc = 0.
n_epochs = 20

test_acc_iter_prune = []

for epoch in range(n_epochs):
    print('\nEpoch: %d' % epoch)

    net.train()
    if epoch < 10:
        for name, layer in net.named_modules():
            if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and 'id_mapping' not in name:
                # Increase model sparsity
                q = (epoch + 1) * 7
                # Linearly increase the pruning percentage for 10 epochs until reaching 70% in the final epoch
                prune_by_percentage(layer, q=q)
    if epoch < 9:

```

```

        finetune_after_prune(net, trainloader, criterion, optimizer,
↪prune=False)
    else: # starts from epoch==9, where q = 70%
        finetune_after_prune(net, trainloader, criterion, optimizer)

    #Start the testing code.
    net.eval()
    test_loss = 0
    correct = 0
    total = 0
    with torch.no_grad():
        for batch_idx, (inputs, targets) in enumerate(testloader):
            inputs, targets = inputs.to(device), targets.to(device)
            outputs = net(inputs)
            loss = criterion(outputs, targets)

            test_loss += loss.item()
            _, predicted = outputs.max(1)
            total += targets.size(0)
            correct += predicted.eq(targets).sum().item()
    num_val_steps = len(testloader)
    val_acc = correct / total

    test_acc_iter_prune.append(val_acc)

    print("Test Loss=%.4f, Test acc=%.4f" % (test_loss / (num_val_steps),
↪val_acc))

    if epoch>=10:
        if val_acc > best_acc:
            best_acc = val_acc
            print("Saving...")
            torch.save(net.state_dict(), "net_after_iterative_prune.pt")

```

Epoch: 0

[Step=50]	Loss=0.0472	acc=0.9847	6590.0 examples/second
[Step=100]	Loss=0.0482	acc=0.9841	12078.6 examples/second
[Step=150]	Loss=0.0481	acc=0.9843	12069.6 examples/second

Test Loss=0.3262, Test acc=0.9134

Epoch: 1

[Step=50]	Loss=0.0482	acc=0.9840	6651.5 examples/second
[Step=100]	Loss=0.0497	acc=0.9833	12706.2 examples/second
[Step=150]	Loss=0.0481	acc=0.9841	12666.0 examples/second

Test Loss=0.3265, Test acc=0.9139

Epoch: 2

[Step=50]	Loss=0.0475	acc=0.9846	6783.0 examples/second
[Step=100]	Loss=0.0507	acc=0.9833	12782.9 examples/second
[Step=150]	Loss=0.0497	acc=0.9841	12842.8 examples/second

Test Loss=0.3286, Test acc=0.9148

Epoch: 3

[Step=50]	Loss=0.0521	acc=0.9824	5565.0 examples/second
[Step=100]	Loss=0.0516	acc=0.9832	11161.7 examples/second
[Step=150]	Loss=0.0525	acc=0.9824	12087.6 examples/second

Test Loss=0.3304, Test acc=0.9125

Epoch: 4

[Step=50]	Loss=0.0631	acc=0.9784	6435.7 examples/second
[Step=100]	Loss=0.0618	acc=0.9793	12608.6 examples/second
[Step=150]	Loss=0.0608	acc=0.9797	12673.7 examples/second

Test Loss=0.3343, Test acc=0.9117

Epoch: 5

[Step=50]	Loss=0.0740	acc=0.9750	6680.0 examples/second
[Step=100]	Loss=0.0719	acc=0.9757	12194.1 examples/second
[Step=150]	Loss=0.0700	acc=0.9764	12545.0 examples/second

Test Loss=0.3380, Test acc=0.9084

Epoch: 6

[Step=50]	Loss=0.0918	acc=0.9689	6544.7 examples/second
[Step=100]	Loss=0.0897	acc=0.9691	12636.3 examples/second
[Step=150]	Loss=0.0878	acc=0.9701	12305.3 examples/second

Test Loss=0.3330, Test acc=0.9063

Epoch: 7

[Step=50]	Loss=0.1278	acc=0.9573	6697.8 examples/second
[Step=100]	Loss=0.1239	acc=0.9588	12323.9 examples/second
[Step=150]	Loss=0.1204	acc=0.9596	12597.7 examples/second

Test Loss=0.3359, Test acc=0.9022

Epoch: 8

[Step=50]	Loss=0.1709	acc=0.9405	6641.5 examples/second
[Step=100]	Loss=0.1636	acc=0.9422	12175.7 examples/second
[Step=150]	Loss=0.1570	acc=0.9448	12565.0 examples/second

Test Loss=0.3472, Test acc=0.8941

Epoch: 9

[Step=50]	Loss=0.2701	acc=0.9071	6561.1 examples/second
[Step=100]	Loss=0.2503	acc=0.9139	12533.8 examples/second
[Step=150]	Loss=0.2406	acc=0.9166	12643.9 examples/second

Test Loss=0.3901, Test acc=0.8759

Epoch: 10
[Step=50] Loss=0.2112 acc=0.9237 6791.6 examples/second
[Step=100] Loss=0.2032 acc=0.9277 12589.8 examples/second
[Step=150] Loss=0.2042 acc=0.9279 12803.1 examples/second
Test Loss=0.3735, Test acc=0.8804
Saving...

Epoch: 11
[Step=50] Loss=0.1868 acc=0.9356 6858.2 examples/second
[Step=100] Loss=0.1856 acc=0.9359 12521.4 examples/second
[Step=150] Loss=0.1862 acc=0.9343 11670.0 examples/second
Test Loss=0.3674, Test acc=0.8825
Saving...

Epoch: 12
[Step=50] Loss=0.1798 acc=0.9384 6543.2 examples/second
[Step=100] Loss=0.1781 acc=0.9386 12376.5 examples/second
[Step=150] Loss=0.1805 acc=0.9376 13172.6 examples/second
Test Loss=0.3585, Test acc=0.8853
Saving...

Epoch: 13
[Step=50] Loss=0.1798 acc=0.9373 5819.9 examples/second
[Step=100] Loss=0.1740 acc=0.9392 10813.6 examples/second
[Step=150] Loss=0.1709 acc=0.9397 10752.0 examples/second
Test Loss=0.3546, Test acc=0.8866
Saving...

Epoch: 14
[Step=50] Loss=0.1670 acc=0.9427 6128.3 examples/second
[Step=100] Loss=0.1696 acc=0.9418 11742.8 examples/second
[Step=150] Loss=0.1661 acc=0.9427 12339.2 examples/second
Test Loss=0.3511, Test acc=0.8874
Saving...

Epoch: 15
[Step=50] Loss=0.1665 acc=0.9420 6126.5 examples/second
[Step=100] Loss=0.1631 acc=0.9429 12471.9 examples/second
[Step=150] Loss=0.1645 acc=0.9425 12356.1 examples/second
Test Loss=0.3488, Test acc=0.8893
Saving...

Epoch: 16
[Step=50] Loss=0.1545 acc=0.9454 5152.9 examples/second
[Step=100] Loss=0.1574 acc=0.9453 11185.4 examples/second
[Step=150] Loss=0.1581 acc=0.9445 11841.8 examples/second
Test Loss=0.3450, Test acc=0.8899
Saving...


```
Epoch: 17
[Step=50]      Loss=0.1604      acc=0.9439      6179.6 examples/second
[Step=100]     Loss=0.1543      acc=0.9450      11965.1 examples/second
[Step=150]     Loss=0.1543      acc=0.9456      11827.1 examples/second
Test Loss=0.3437, Test acc=0.8911
Saving...
```

```
Epoch: 18
[Step=50]      Loss=0.1557      acc=0.9444      6268.0 examples/second
[Step=100]     Loss=0.1548      acc=0.9464      12063.5 examples/second
[Step=150]     Loss=0.1541      acc=0.9463      12084.5 examples/second
Test Loss=0.3431, Test acc=0.8905
```

```
Epoch: 19
[Step=50]      Loss=0.1500      acc=0.9492      6166.7 examples/second
[Step=100]     Loss=0.1519      acc=0.9481      12370.6 examples/second
[Step=150]     Loss=0.1496      acc=0.9488      12396.4 examples/second
Test Loss=0.3405, Test acc=0.8922
Saving...
```

```
[13]: # Check sparsity of the final model, make sure it's 70%
net.load_state_dict(torch.load("net_after_iterative_prune.pt"))

for name, layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
        'id_mapping' not in name:
        # Your code here: can copy from previous question
        ## Convert the weight of "layer" to numpy array
        np_weight = layer.weight.detach().cpu().numpy()
        ## Count number of zeros
        zeros = sum((np_weight == 0).flatten())
        ## Count number of parameters
        total = len(np_weight.flatten())
        # Print sparsity
        print('Sparsity of %s: %g' % (name, zeros/total))

test(net)
```


```
Sparsity of head_conv.0.conv: 0.699074
Sparsity of body_op.0.conv1.0.conv: 0.700087
Sparsity of body_op.0.conv2.0.conv: 0.700087
Sparsity of body_op.1.conv1.0.conv: 0.700087
Sparsity of body_op.1.conv2.0.conv: 0.700087
Sparsity of body_op.2.conv1.0.conv: 0.700087
Sparsity of body_op.2.conv2.0.conv: 0.700087
Sparsity of body_op.3.conv1.0.conv: 0.69987
```

```
Sparsity of body_op.3.conv2.0.conv: 0.699978
Sparsity of body_op.4.conv1.0.conv: 0.699978
Sparsity of body_op.4.conv2.0.conv: 0.699978
Sparsity of body_op.5.conv1.0.conv: 0.699978
Sparsity of body_op.5.conv2.0.conv: 0.699978
Sparsity of body_op.6.conv1.0.conv: 0.699978
Sparsity of body_op.6.conv2.0.conv: 0.700005
Sparsity of body_op.7.conv1.0.conv: 0.700005
Sparsity of body_op.7.conv2.0.conv: 0.700005
Sparsity of body_op.8.conv1.0.conv: 0.700005
Sparsity of body_op.8.conv2.0.conv: 0.700005
Sparsity of final_fc.linear: 0.7
Files already downloaded and verified
Test Loss=0.3405, Test accuracy=0.8922
```

```
[13]: 0.8922
```

```
[14]: fig, ax = plt.subplots(1, 1)
xx = range(n_epochs)
ax.plot(xx, test_acc_iter_prune, label='last test accuracy: %g' %L
↪test_acc_iter_prune[-1])
ax.set_xlabel('Epochs')
ax.set_ylabel('Accuracy')
ax.set_title('lab2d Accuracy vs Epochs\nIterative Pruning')
ax.legend()

plt.savefig('Figures/lab2d.pdf', dpi=500, bbox_inches='tight')
```



0.0.5 Lab2 (e) Global iterative pruning

```
[15]: import numpy as np

def global_prune_by_percentage(net, q=70.0):
    """
    Pruning the weight paramters by threshold.
    :param q: pruning percentile. 'q' percent of the least
    significant weight parameters will be pruned.
    """
    # A list to gather all the weights
    flattened_weights = []
    # Find global pruning threshold
    for name, layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
        ↪ 'id_mapping' not in name:
            # Convert weight to numpy
            # Flatten the weight and append to flattened_weights
```

```

        flattened_weights.append(layer.weight.detach().cpu().numpy().
↪flatten())

    # Concatenate all weights into a np array
    flattened_weights = np.concatenate(flattened_weights)
    # Find global pruning threshold
    threshold = np.percentile(np.abs(flattened_weights.flatten()), q)

    # Apply pruning threshold to all layers
    for name, layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
↪'id_mapping' not in name:
            # Convert weight to numpy
            layer_weight = layer.weight.detach().cpu().numpy()

            # Generate a binary mask same shape as weight to decide which
↪element to prune
            masked_obj = np.ma.masked_greater_equal(x=np.abs(layer_weight),
↪value=threshold, copy=True)
            mask_int = np.ma.getmask(masked_obj).astype(int)
            # Convert mask to torch tensor and put on GPU
            mask_tensor = torch.tensor(mask_int).to(device)
            # Multiply the weight by mask to perform pruning
            layer.weight.data = layer.weight.data.clone().detach().
↪requires_grad_(True) * mask_tensor

```

```

[16]: net.load_state_dict(torch.load("pretrained_model.pt"))
best_acc = 0.

test_acc_global_prune = []

for epoch in range(20):
    print('\nEpoch: %d' % epoch)
    # q = (epoch + 1) * 8
    q = (epoch + 1) * 7

    net.train()
    # Increase model sparsity
    if epoch < 10:
        global_prune_by_percentage(net, q=q)
    if epoch < 9:
        finetune_after_prune(net, trainloader, criterion, optimizer,
↪prune=False)
    else:
        finetune_after_prune(net, trainloader, criterion, optimizer)

    #Start the testing code.

```

```

net.eval()
test_loss = 0
correct = 0
total = 0
with torch.no_grad():
    for batch_idx, (inputs, targets) in enumerate(testloader):
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = net(inputs)
        loss = criterion(outputs, targets)

        test_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()
num_val_steps = len(testloader)
val_acc = correct / total
print("Test Loss=%.4f, Test acc=%.4f" % (test_loss / (num_val_steps),
↪val_acc))

test_acc_global_prune.append(val_acc)

if epoch >= 10:
    if val_acc > best_acc:
        best_acc = val_acc
        print("Saving...")
        torch.save(net.state_dict(), "net_after_global_iterative_prune.pt")

```

Epoch: 0

[Step=50]	Loss=0.0494	acc=0.9846	6143.9 examples/second
[Step=100]	Loss=0.0484	acc=0.9847	12491.2 examples/second
[Step=150]	Loss=0.0493	acc=0.9843	12542.2 examples/second

Test Loss=0.3235, Test acc=0.9150

Epoch: 1

[Step=50]	Loss=0.0483	acc=0.9848	6045.3 examples/second
[Step=100]	Loss=0.0488	acc=0.9844	12603.3 examples/second
[Step=150]	Loss=0.0488	acc=0.9844	10154.3 examples/second

Test Loss=0.3239, Test acc=0.9143

Epoch: 2

[Step=50]	Loss=0.0480	acc=0.9859	5200.7 examples/second
[Step=100]	Loss=0.0480	acc=0.9857	12315.7 examples/second
[Step=150]	Loss=0.0483	acc=0.9854	12389.8 examples/second

Test Loss=0.3291, Test acc=0.9137

Epoch: 3

[Step=50]	Loss=0.0501	acc=0.9843	6317.4 examples/second
[Step=100]	Loss=0.0496	acc=0.9845	12250.7 examples/second
[Step=150]	Loss=0.0499	acc=0.9841	12418.7 examples/second

Test Loss=0.3274, Test acc=0.9127

Epoch: 4

[Step=50]	Loss=0.0500	acc=0.9841	4792.6 examples/second
[Step=100]	Loss=0.0519	acc=0.9842	12345.4 examples/second
[Step=150]	Loss=0.0524	acc=0.9833	12372.6 examples/second

Test Loss=0.3260, Test acc=0.9141

Epoch: 5

[Step=50]	Loss=0.0617	acc=0.9802	6315.3 examples/second
[Step=100]	Loss=0.0630	acc=0.9797	13431.7 examples/second
[Step=150]	Loss=0.0614	acc=0.9797	12610.3 examples/second

Test Loss=0.3229, Test acc=0.9113

Epoch: 6

[Step=50]	Loss=0.0731	acc=0.9754	6366.3 examples/second
[Step=100]	Loss=0.0719	acc=0.9758	12640.7 examples/second
[Step=150]	Loss=0.0718	acc=0.9759	12408.8 examples/second

Test Loss=0.3234, Test acc=0.9099

Epoch: 7

[Step=50]	Loss=0.0885	acc=0.9702	6106.5 examples/second
[Step=100]	Loss=0.0860	acc=0.9717	12448.0 examples/second
[Step=150]	Loss=0.0872	acc=0.9708	12239.8 examples/second

Test Loss=0.3256, Test acc=0.9068

Epoch: 8

[Step=50]	Loss=0.1173	acc=0.9610	6134.2 examples/second
[Step=100]	Loss=0.1194	acc=0.9599	12445.5 examples/second
[Step=150]	Loss=0.1173	acc=0.9605	12471.5 examples/second

Test Loss=0.3249, Test acc=0.9005

Epoch: 9

[Step=50]	Loss=0.1843	acc=0.9363	6212.4 examples/second
[Step=100]	Loss=0.1780	acc=0.9388	11804.6 examples/second
[Step=150]	Loss=0.1724	acc=0.9407	10877.9 examples/second

Test Loss=0.3410, Test acc=0.8903

Epoch: 10

[Step=50]	Loss=0.1532	acc=0.9495	6142.6 examples/second
[Step=100]	Loss=0.1586	acc=0.9470	13074.3 examples/second
[Step=150]	Loss=0.1582	acc=0.9464	11776.5 examples/second

Test Loss=0.3339, Test acc=0.8934

Saving...

Epoch: 11
[Step=50] Loss=0.1566 acc=0.9456 6226.4 examples/second
[Step=100] Loss=0.1532 acc=0.9479 11704.4 examples/second
[Step=150] Loss=0.1514 acc=0.9485 11097.2 examples/second
Test Loss=0.3300, Test acc=0.8946
Saving...

Epoch: 12
[Step=50] Loss=0.1441 acc=0.9504 5309.6 examples/second
[Step=100] Loss=0.1457 acc=0.9515 12221.5 examples/second
[Step=150] Loss=0.1454 acc=0.9513 12054.7 examples/second
Test Loss=0.3273, Test acc=0.8949
Saving...

Epoch: 13
[Step=50] Loss=0.1325 acc=0.9542 6226.3 examples/second
[Step=100] Loss=0.1376 acc=0.9525 12487.7 examples/second
[Step=150] Loss=0.1391 acc=0.9525 12441.0 examples/second
Test Loss=0.3229, Test acc=0.8970
Saving...

Epoch: 14
[Step=50] Loss=0.1314 acc=0.9566 6036.0 examples/second
[Step=100] Loss=0.1372 acc=0.9543 11868.5 examples/second
[Step=150] Loss=0.1358 acc=0.9542 11846.7 examples/second
Test Loss=0.3231, Test acc=0.8969

Epoch: 15
[Step=50] Loss=0.1355 acc=0.9521 6031.4 examples/second
[Step=100] Loss=0.1347 acc=0.9539 12181.6 examples/second
[Step=150] Loss=0.1353 acc=0.9537 12342.5 examples/second
Test Loss=0.3204, Test acc=0.8985
Saving...

Epoch: 16
[Step=50] Loss=0.1276 acc=0.9568 5985.7 examples/second
[Step=100] Loss=0.1318 acc=0.9554 12347.7 examples/second
[Step=150] Loss=0.1319 acc=0.9547 11193.9 examples/second
Test Loss=0.3189, Test acc=0.8984

Epoch: 17
[Step=50] Loss=0.1329 acc=0.9560 6316.6 examples/second
[Step=100] Loss=0.1312 acc=0.9562 12470.3 examples/second
[Step=150] Loss=0.1286 acc=0.9568 12600.0 examples/second
Test Loss=0.3184, Test acc=0.8981

Epoch: 18
[Step=50] Loss=0.1268 acc=0.9566 6341.9 examples/second

```

[Step=100]      Loss=0.1260      acc=0.9570      12705.8 examples/second
[Step=150]      Loss=0.1268      acc=0.9572      10798.4 examples/second
Test Loss=0.3182, Test acc=0.8986
Saving..

```

```

Epoch: 19
[Step=50]      Loss=0.1266      acc=0.9572      5950.9 examples/second
[Step=100]     Loss=0.1259      acc=0.9574      11822.7 examples/second
[Step=150]     Loss=0.1261      acc=0.9572      12023.3 examples/second
Test Loss=0.3178, Test acc=0.8985

```

```

[23]: net.load_state_dict(torch.load("net_after_global_iterative_prune.pt"))
sparsity_list = []
x_ticks_labels = []
zeros_sum = 0
total_sum = 0
for name, layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
        ↪ 'id_mapping' not in name:
        # Your code here:
        # Convert the weight of "layer" to numpy array
        np_weight = layer.weight.detach().cpu().numpy()
        # Count number of zeros
        zeros = sum((np_weight == 0).flatten())
        # Count number of parameters
        total = len(np_weight.flatten())
        zeros_sum += zeros
        total_sum += total
        sparsity_list.append(zeros / (total / 1.0))
        x_ticks_labels.append(name)
        print('Sparsity of %s: %g' % (name, zeros / (total / 1.0)))
print('Total sparsity: %g' % (zeros_sum / (total_sum / 1.0)))
test(net)

```

```

Sparsity of head_conv.0.conv: 0.240741
Sparsity of body_op.0.conv1.0.conv: 0.550781
Sparsity of body_op.0.conv2.0.conv: 0.528646
Sparsity of body_op.1.conv1.0.conv: 0.518663
Sparsity of body_op.1.conv2.0.conv: 0.552517
Sparsity of body_op.2.conv1.0.conv: 0.516493
Sparsity of body_op.2.conv2.0.conv: 0.565104
Sparsity of body_op.3.conv1.0.conv: 0.52474
Sparsity of body_op.3.conv2.0.conv: 0.583008
Sparsity of body_op.4.conv1.0.conv: 0.615343
Sparsity of body_op.4.conv2.0.conv: 0.677409
Sparsity of body_op.5.conv1.0.conv: 0.611654
Sparsity of body_op.5.conv2.0.conv: 0.702799
Sparsity of body_op.6.conv1.0.conv: 0.615343


```



```
Sparsity of body_op.6.conv2.0.conv: 0.651232
Sparsity of body_op.7.conv1.0.conv: 0.662137
Sparsity of body_op.7.conv2.0.conv: 0.718777
Sparsity of body_op.8.conv1.0.conv: 0.74821
Sparsity of body_op.8.conv2.0.conv: 0.937283
Sparsity of final_fc.linear: 0.11875
Total sparsity: 0.699999
Files already downloaded and verified
Test Loss=0.3182, Test accuracy=0.8986
```


```
[23]: 0.8986
```

```
[25]: fig, ax = plt.subplots(1, 1, figsize=(10, 6))
xx = range(len(sparsity_list))
ax.stem(xx, sparsity_list)
ax.set_xticks(range(len(xx)), x_ticks_labels, rotation='vertical', fontsize=7)
ax.set_xlabel('Layer')
ax.set_ylabel('Sparsity')
ax.set_title('lab2e Sparsity for Each Layer\nGlobal Iterative Pruning')
for (x, y) in zip(xx, sparsity_list):
    plt.annotate('{:.4f}'.format(y), xy=(x, y), xytext=(0, 5),
        ↪textcoords='offset points', ha='center', size=8)
plt.savefig('Figures/lab2e_sparsity.pdf', dpi=500, bbox_inches='tight')
```



```
[18]: fig, ax = plt.subplots(1, 1)
xx = range(n_epochs)
ax.plot(xx, test_acc_global_prune, label='last test accuracy: %g' %_
↳test_acc_global_prune[-1])
ax.set_xlabel('Epochs')
ax.set_ylabel('Accuracy')
ax.set_title('lab2e Accuracy vs Epochs\nGlobal Iterative Pruning')
ax.legend()

plt.savefig('Figures/lab2e.pdf', dpi=500, bbox_inches='tight')
```



0.0.6 Lab 3 (b) and (c): Fixed-point quantization

```
[4]: # load pretrained ResNet-20 model, report the accuracy of the floating-point
    ↪ pretrained model.
net = ResNetCIFAR(num_layers=20)
net = net.to(device)
net.load_state_dict(torch.load("pretrained_model.pt"))
print("before quantisation:")
test(net)
print()

# Define quantized model and load weight
Nbits_arr = [2, 3, 4, 5, 6]
test_acc_lab3b = []
for iii, nb in enumerate(Nbits_arr):
    Nbits = nb #Change this value to finish (b) and (c)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
```

```

net = net.to(device)
net.load_state_dict(torch.load("pretrained_model.pt"))
test_acc_lab3b.append(test(net))

```

before quantisation:

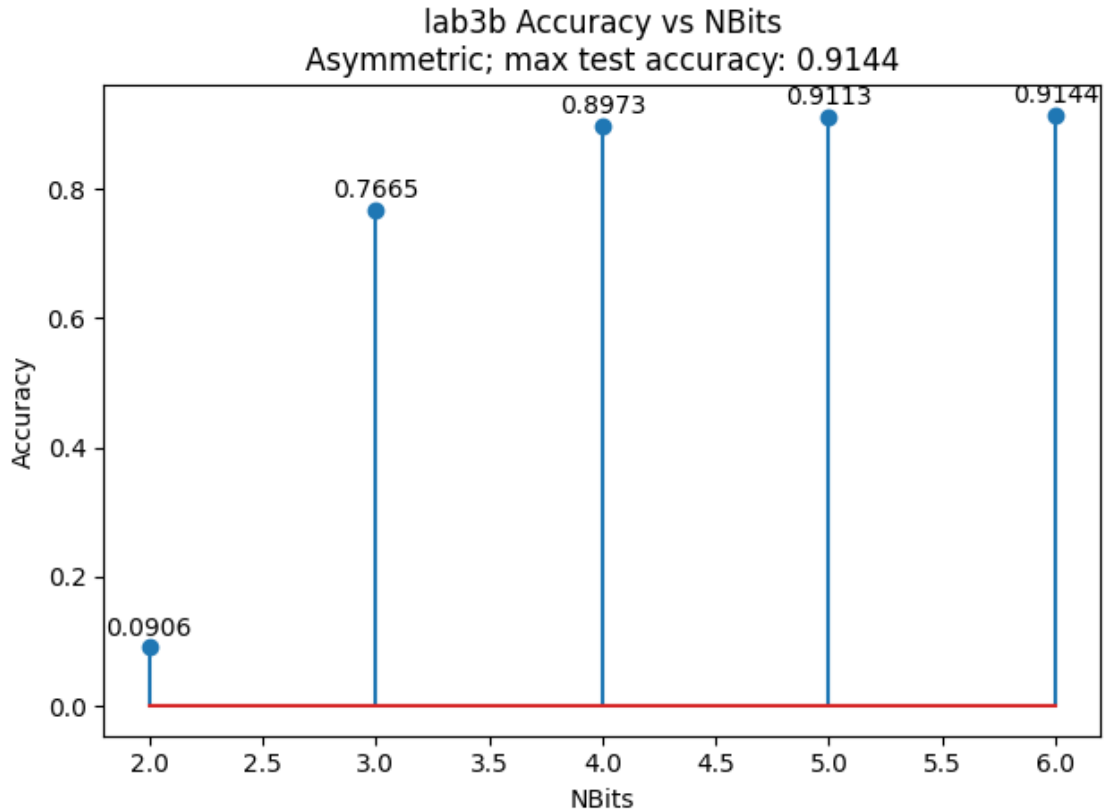
Files already downloaded and verified
Test Loss=0.3231, Test accuracy=0.9150

Files already downloaded and verified
Test Loss=9.6086, Test accuracy=0.0906
Files already downloaded and verified
Test Loss=0.9857, Test accuracy=0.7665
Files already downloaded and verified
Test Loss=0.3860, Test accuracy=0.8973
Files already downloaded and verified
Test Loss=0.3391, Test accuracy=0.9113
Files already downloaded and verified
Test Loss=0.3364, Test accuracy=0.9144

```

[10]: fig, ax = plt.subplots(1, 1)
xx = Nbits_arr
ax.stem(xx, test_acc_lab3b)
# ax.plot(xx, test_acc_lab3b, label='max test accuracy: %g' % np.
↳max(test_acc_lab3b))
ax.set_xlabel('NBits')
ax.set_ylabel('Accuracy')
ax.set_title('lab3b Accuracy vs NBits\nAsymmetric; max test accuracy: %g' % np.
↳max(test_acc_lab3b))
for (x, y) in zip(xx, test_acc_lab3b):
    plt.annotate('{:.4f}'.format(y), xy=(x, y), xytext=(0, 5),
↳textcoords='offset points', ha='center')
fig.tight_layout()
plt.savefig('Figures/lab3b.pdf', dpi=500, bbox_inches='tight')

```



```
[11]: # lab 3 part c is shown below:
```

```
[12]: Nbits_arr_c = [2, 3, 4]
test_acc_lab3c_finetune = {}
# test accuracy for each finetuning as fn of epoch
final_test_acc_lab3c = {} # test acc after finetuning for each Nbit
n_epochs = 20
test_acc_lab3c_b4_finetune = {} # test acc before finetuning

for iii, nb in enumerate(Nbits_arr_c):
    # Define quantized model and load weight
    Nbits = nb #Change this value to finish (b) and (c)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
    net = net.to(device)
    net.load_state_dict(torch.load("pretrained_model.pt"))
    test_acc_lab3c_b4_finetune[Nbits] = test(net)

    # Quantized model finetuning
    test_acc_lab3c_finetune[Nbits] = np.asarray(
        finetune(net, epochs=n_epochs, batch_size=256, lr=0.002, reg=1e-4))
```

```
# finetune() supposedly returns list of val_accuracy values
# Load the model with best accuracy
net.load_state_dict(torch.load("quantized_net_after_finetune.pt"))
final_test_acc_lab3c[Nbits] = test(net)
```

Files already downloaded and verified
 Test Loss=9.6086, Test accuracy=0.0906
 ==> Preparing data..
 Files already downloaded and verified
 Files already downloaded and verified

Epoch: 0

[Step=50]	Loss=1.3199	acc=0.6250	5064.9 examples/second
[Step=100]	Loss=1.0535	acc=0.6870	9559.5 examples/second
[Step=150]	Loss=0.9346	acc=0.7166	9594.5 examples/second

Test Loss=0.7387, Test acc=0.7820
 Saving..

Epoch: 1

[Step=200]	Loss=0.6141	acc=0.8047	3416.9 examples/second
[Step=250]	Loss=0.5912	acc=0.8040	9611.2 examples/second
[Step=300]	Loss=0.5688	acc=0.8098	9720.0 examples/second
[Step=350]	Loss=0.5459	acc=0.8157	9729.8 examples/second

Test Loss=0.6160, Test acc=0.8087
 Saving..

Epoch: 2

[Step=400]	Loss=0.4761	acc=0.8359	2916.8 examples/second
[Step=450]	Loss=0.4810	acc=0.8365	8614.2 examples/second
[Step=500]	Loss=0.4718	acc=0.8390	7992.3 examples/second
[Step=550]	Loss=0.4650	acc=0.8406	9019.1 examples/second

Test Loss=0.5682, Test acc=0.8207
 Saving..

Epoch: 3

[Step=600]	Loss=0.4422	acc=0.8503	2982.4 examples/second
[Step=650]	Loss=0.4296	acc=0.8514	7740.0 examples/second
[Step=700]	Loss=0.4260	acc=0.8522	8216.0 examples/second
[Step=750]	Loss=0.4242	acc=0.8532	9281.9 examples/second

Test Loss=0.5630, Test acc=0.8256
 Saving..

Epoch: 4

[Step=800]	Loss=0.3994	acc=0.8660	3616.9 examples/second
[Step=850]	Loss=0.4016	acc=0.8609	8880.7 examples/second
[Step=900]	Loss=0.3993	acc=0.8597	8941.6 examples/second
[Step=950]	Loss=0.3989	acc=0.8605	8840.8 examples/second

Test Loss=0.5452, Test acc=0.8309

Saving...

Epoch: 5

[Step=1000]	Loss=0.4074	acc=0.8596	3349.6 examples/second
[Step=1050]	Loss=0.3827	acc=0.8658	9304.9 examples/second
[Step=1100]	Loss=0.3851	acc=0.8651	9023.1 examples/second
[Step=1150]	Loss=0.3811	acc=0.8667	9238.3 examples/second

Test Loss=0.5262, Test acc=0.8353

Saving...

Epoch: 6

[Step=1200]	Loss=0.3542	acc=0.8791	3387.8 examples/second
[Step=1250]	Loss=0.3542	acc=0.8769	9156.1 examples/second
[Step=1300]	Loss=0.3613	acc=0.8749	9068.4 examples/second
[Step=1350]	Loss=0.3592	acc=0.8750	9290.3 examples/second

Test Loss=0.5987, Test acc=0.8243

Epoch: 7

[Step=1400]	Loss=0.3612	acc=0.8726	3634.1 examples/second
[Step=1450]	Loss=0.3503	acc=0.8758	9067.2 examples/second
[Step=1500]	Loss=0.3517	acc=0.8771	8964.0 examples/second
[Step=1550]	Loss=0.3503	acc=0.8772	9364.0 examples/second

Test Loss=0.5400, Test acc=0.8394

Saving...

Epoch: 8

[Step=1600]	Loss=0.3455	acc=0.8785	3548.8 examples/second
[Step=1650]	Loss=0.3429	acc=0.8808	9186.2 examples/second
[Step=1700]	Loss=0.3387	acc=0.8811	9122.0 examples/second
[Step=1750]	Loss=0.3402	acc=0.8801	9614.7 examples/second

Test Loss=0.5119, Test acc=0.8431

Saving...

Epoch: 9

[Step=1800]	Loss=0.3359	acc=0.8837	3562.8 examples/second
[Step=1850]	Loss=0.3359	acc=0.8843	9036.8 examples/second
[Step=1900]	Loss=0.3303	acc=0.8847	9079.4 examples/second
[Step=1950]	Loss=0.3288	acc=0.8851	9645.1 examples/second

Test Loss=0.5415, Test acc=0.8349

Epoch: 10

[Step=2000]	Loss=0.3316	acc=0.8798	3638.5 examples/second
[Step=2050]	Loss=0.3348	acc=0.8805	9018.6 examples/second
[Step=2100]	Loss=0.3331	acc=0.8820	9029.8 examples/second
[Step=2150]	Loss=0.3294	acc=0.8833	9577.1 examples/second

Test Loss=0.5093, Test acc=0.8420

Epoch: 11

[Step=2200]	Loss=0.3162	acc=0.8890	3573.9 examples/second
[Step=2250]	Loss=0.3176	acc=0.8871	8949.2 examples/second
[Step=2300]	Loss=0.3180	acc=0.8883	9003.6 examples/second
[Step=2350]	Loss=0.3155	acc=0.8886	9855.7 examples/second

Test Loss=0.4624, Test acc=0.8555
Saving...

Epoch: 12

[Step=2400]	Loss=0.3081	acc=0.8923	3372.2 examples/second
[Step=2450]	Loss=0.3098	acc=0.8915	9674.6 examples/second
[Step=2500]	Loss=0.3116	acc=0.8908	9699.1 examples/second

Test Loss=0.5014, Test acc=0.8462

Epoch: 13

[Step=2550]	Loss=0.2429	acc=0.9199	3746.8 examples/second
[Step=2600]	Loss=0.3055	acc=0.8939	8866.0 examples/second
[Step=2650]	Loss=0.3042	acc=0.8935	9083.3 examples/second
[Step=2700]	Loss=0.3060	acc=0.8927	9076.7 examples/second

Test Loss=0.4841, Test acc=0.8508

Epoch: 14

[Step=2750]	Loss=0.2795	acc=0.9049	3611.9 examples/second
[Step=2800]	Loss=0.2960	acc=0.8965	8951.2 examples/second
[Step=2850]	Loss=0.2997	acc=0.8952	8979.6 examples/second
[Step=2900]	Loss=0.2992	acc=0.8955	9206.7 examples/second

Test Loss=0.4598, Test acc=0.8545

Epoch: 15

[Step=2950]	Loss=0.2990	acc=0.8902	3626.7 examples/second
[Step=3000]	Loss=0.2943	acc=0.8950	9052.5 examples/second
[Step=3050]	Loss=0.2922	acc=0.8971	8802.7 examples/second
[Step=3100]	Loss=0.2892	acc=0.8976	8801.3 examples/second

Test Loss=0.4655, Test acc=0.8568
Saving...

Epoch: 16

[Step=3150]	Loss=0.2909	acc=0.8970	3331.8 examples/second
[Step=3200]	Loss=0.2963	acc=0.8968	8891.1 examples/second
[Step=3250]	Loss=0.2917	acc=0.8972	8881.5 examples/second
[Step=3300]	Loss=0.2872	acc=0.8978	9036.9 examples/second

Test Loss=0.4836, Test acc=0.8509

Epoch: 17

[Step=3350]	Loss=0.2757	acc=0.9028	3560.6 examples/second
[Step=3400]	Loss=0.2887	acc=0.8952	8938.8 examples/second
[Step=3450]	Loss=0.2872	acc=0.8962	8840.4 examples/second
[Step=3500]	Loss=0.2860	acc=0.8984	8871.0 examples/second

Test Loss=0.4483, Test acc=0.8576

Saving...

Epoch: 18

[Step=3550]	Loss=0.2802	acc=0.9034	3519.2 examples/second
[Step=3600]	Loss=0.2809	acc=0.9002	9101.0 examples/second
[Step=3650]	Loss=0.2837	acc=0.8998	8806.1 examples/second
[Step=3700]	Loss=0.2874	acc=0.8980	8418.4 examples/second

Test Loss=0.4596, Test acc=0.8576

Epoch: 19

[Step=3750]	Loss=0.2893	acc=0.8980	3245.6 examples/second
[Step=3800]	Loss=0.2876	acc=0.8983	9225.6 examples/second
[Step=3850]	Loss=0.2844	acc=0.8989	8938.6 examples/second
[Step=3900]	Loss=0.2842	acc=0.8997	9348.4 examples/second

Test Loss=0.5014, Test acc=0.8440

Files already downloaded and verified

Test Loss=0.4483, Test accuracy=0.8576

Files already downloaded and verified

Test Loss=0.9857, Test accuracy=0.7665

==> Preparing data..

Files already downloaded and verified

Files already downloaded and verified

Epoch: 0

[Step=50]	Loss=0.1696	acc=0.9402	5790.4 examples/second
[Step=100]	Loss=0.1587	acc=0.9446	9041.2 examples/second
[Step=150]	Loss=0.1515	acc=0.9473	9184.6 examples/second

Test Loss=0.3917, Test acc=0.8957

Saving...

Epoch: 1

[Step=200]	Loss=0.1164	acc=0.9521	3681.9 examples/second
[Step=250]	Loss=0.1214	acc=0.9543	8930.0 examples/second
[Step=300]	Loss=0.1235	acc=0.9542	8967.8 examples/second
[Step=350]	Loss=0.1225	acc=0.9551	8787.4 examples/second

Test Loss=0.3819, Test acc=0.8987

Saving...

Epoch: 2

[Step=400]	Loss=0.1068	acc=0.9614	3329.4 examples/second
[Step=450]	Loss=0.1119	acc=0.9591	9067.3 examples/second
[Step=500]	Loss=0.1118	acc=0.9593	8954.9 examples/second
[Step=550]	Loss=0.1118	acc=0.9597	9145.0 examples/second

Test Loss=0.3750, Test acc=0.9002

Saving...

Epoch: 3

[Step=600]	Loss=0.1122	acc=0.9606	3373.3 examples/second
[Step=650]	Loss=0.1050	acc=0.9624	8948.5 examples/second
[Step=700]	Loss=0.1057	acc=0.9623	8616.6 examples/second
[Step=750]	Loss=0.1058	acc=0.9619	8644.7 examples/second

Test Loss=0.3728, Test acc=0.9002

Epoch: 4

[Step=800]	Loss=0.0996	acc=0.9629	3277.9 examples/second
[Step=850]	Loss=0.1045	acc=0.9621	9003.6 examples/second
[Step=900]	Loss=0.1031	acc=0.9628	8842.6 examples/second
[Step=950]	Loss=0.1047	acc=0.9622	8735.5 examples/second

Test Loss=0.3686, Test acc=0.9040

Saving...

Epoch: 5

[Step=1000]	Loss=0.0925	acc=0.9656	3289.7 examples/second
[Step=1050]	Loss=0.1019	acc=0.9615	9219.4 examples/second
[Step=1100]	Loss=0.1018	acc=0.9618	8997.6 examples/second
[Step=1150]	Loss=0.1023	acc=0.9624	9011.6 examples/second

Test Loss=0.3672, Test acc=0.9016

Epoch: 6

[Step=1200]	Loss=0.0971	acc=0.9681	3603.5 examples/second
[Step=1250]	Loss=0.1015	acc=0.9638	9218.9 examples/second
[Step=1300]	Loss=0.1008	acc=0.9639	8944.5 examples/second
[Step=1350]	Loss=0.1011	acc=0.9642	9182.3 examples/second

Test Loss=0.3728, Test acc=0.9010

Epoch: 7

[Step=1400]	Loss=0.0914	acc=0.9669	3034.5 examples/second
[Step=1450]	Loss=0.0907	acc=0.9675	8991.2 examples/second
[Step=1500]	Loss=0.0951	acc=0.9664	9560.0 examples/second
[Step=1550]	Loss=0.0975	acc=0.9653	9428.4 examples/second

Test Loss=0.3679, Test acc=0.9015

Epoch: 8

[Step=1600]	Loss=0.0983	acc=0.9641	3363.7 examples/second
[Step=1650]	Loss=0.0958	acc=0.9653	8885.9 examples/second
[Step=1700]	Loss=0.0964	acc=0.9651	8842.4 examples/second
[Step=1750]	Loss=0.0968	acc=0.9654	8737.4 examples/second

Test Loss=0.3743, Test acc=0.9012

Epoch: 9

[Step=1800]	Loss=0.0962	acc=0.9652	3242.6 examples/second
[Step=1850]	Loss=0.0936	acc=0.9663	8963.5 examples/second
[Step=1900]	Loss=0.0942	acc=0.9661	8960.3 examples/second
[Step=1950]	Loss=0.0935	acc=0.9664	8973.6 examples/second

Test Loss=0.3675, Test acc=0.9027

Epoch: 10
 [Step=2000] Loss=0.0971 acc=0.9646 3627.0 examples/second
 [Step=2050] Loss=0.0966 acc=0.9650 9085.6 examples/second
 [Step=2100] Loss=0.0963 acc=0.9651 9032.3 examples/second
 [Step=2150] Loss=0.0952 acc=0.9654 9184.7 examples/second
 Test Loss=0.3578, Test acc=0.9034

Epoch: 11
 [Step=2200] Loss=0.0865 acc=0.9691 3356.4 examples/second
 [Step=2250] Loss=0.0929 acc=0.9664 8594.0 examples/second
 [Step=2300] Loss=0.0932 acc=0.9669 8620.8 examples/second
 [Step=2350] Loss=0.0938 acc=0.9665 8909.3 examples/second
 Test Loss=0.3577, Test acc=0.9046
 Saving...

Epoch: 12
 [Step=2400] Loss=0.0869 acc=0.9689 3574.0 examples/second
 [Step=2450] Loss=0.0901 acc=0.9679 8848.3 examples/second
 [Step=2500] Loss=0.0898 acc=0.9684 8947.4 examples/second
 Test Loss=0.3648, Test acc=0.9022

Epoch: 13
 [Step=2550] Loss=0.0755 acc=0.9766 3390.1 examples/second
 [Step=2600] Loss=0.0876 acc=0.9686 8889.5 examples/second
 [Step=2650] Loss=0.0884 acc=0.9683 8928.4 examples/second
 [Step=2700] Loss=0.0886 acc=0.9682 9048.0 examples/second
 Test Loss=0.3537, Test acc=0.9053
 Saving...

Epoch: 14
 [Step=2750] Loss=0.0935 acc=0.9642 3611.2 examples/second
 [Step=2800] Loss=0.0906 acc=0.9672 9169.0 examples/second
 [Step=2850] Loss=0.0898 acc=0.9681 8868.0 examples/second
 [Step=2900] Loss=0.0891 acc=0.9684 8897.3 examples/second
 Test Loss=0.3601, Test acc=0.9052

Epoch: 15
 [Step=2950] Loss=0.0934 acc=0.9652 3546.7 examples/second
 [Step=3000] Loss=0.0898 acc=0.9673 9089.0 examples/second
 [Step=3050] Loss=0.0890 acc=0.9676 8988.5 examples/second
 [Step=3100] Loss=0.0889 acc=0.9679 9194.3 examples/second
 Test Loss=0.3632, Test acc=0.9032

Epoch: 16
 [Step=3150] Loss=0.0829 acc=0.9696 3325.5 examples/second
 [Step=3200] Loss=0.0890 acc=0.9676 9060.3 examples/second
 [Step=3250] Loss=0.0882 acc=0.9681 9120.1 examples/second

[Step=3300] Loss=0.0897 acc=0.9673 9176.5 examples/second
Test Loss=0.3562, Test acc=0.9032

Epoch: 17

[Step=3350] Loss=0.0849 acc=0.9685 3356.1 examples/second
[Step=3400] Loss=0.0883 acc=0.9688 9136.2 examples/second
[Step=3450] Loss=0.0896 acc=0.9686 9003.3 examples/second
[Step=3500] Loss=0.0896 acc=0.9682 9184.7 examples/second
Test Loss=0.3600, Test acc=0.9016

Epoch: 18

[Step=3550] Loss=0.0824 acc=0.9714 3551.5 examples/second
[Step=3600] Loss=0.0878 acc=0.9692 9263.9 examples/second
[Step=3650] Loss=0.0879 acc=0.9687 8926.6 examples/second
[Step=3700] Loss=0.0887 acc=0.9686 9366.3 examples/second
Test Loss=0.3725, Test acc=0.9014

Epoch: 19

[Step=3750] Loss=0.0858 acc=0.9672 3266.7 examples/second
[Step=3800] Loss=0.0887 acc=0.9679 9136.1 examples/second
[Step=3850] Loss=0.0847 acc=0.9702 8982.6 examples/second
[Step=3900] Loss=0.0840 acc=0.9704 9328.0 examples/second

Test Loss=0.3572, Test acc=0.9050

Files already downloaded and verified

Test Loss=0.3537, Test accuracy=0.9053

Files already downloaded and verified

Test Loss=0.3860, Test accuracy=0.8973

==> Preparing data..

Files already downloaded and verified

Files already downloaded and verified

Epoch: 0

[Step=50] Loss=0.0636 acc=0.9789 5830.0 examples/second
[Step=100] Loss=0.0648 acc=0.9780 8858.1 examples/second
[Step=150] Loss=0.0659 acc=0.9770 9042.0 examples/second

Test Loss=0.3373, Test acc=0.9088

Saving...

Epoch: 1

[Step=200] Loss=0.0409 acc=0.9902 3393.4 examples/second
[Step=250] Loss=0.0663 acc=0.9766 8828.5 examples/second
[Step=300] Loss=0.0650 acc=0.9774 8885.3 examples/second
[Step=350] Loss=0.0647 acc=0.9776 8957.7 examples/second

Test Loss=0.3343, Test acc=0.9094

Saving...

Epoch: 2

[Step=400] Loss=0.0656 acc=0.9766 3312.1 examples/second

[Step=450]	Loss=0.0577	acc=0.9817	8889.2 examples/second
[Step=500]	Loss=0.0608	acc=0.9797	8918.0 examples/second
[Step=550]	Loss=0.0622	acc=0.9791	8865.4 examples/second

Test Loss=0.3353, Test acc=0.9116

Saving...

Epoch: 3

[Step=600]	Loss=0.0627	acc=0.9759	3325.2 examples/second
[Step=650]	Loss=0.0575	acc=0.9797	8984.9 examples/second
[Step=700]	Loss=0.0570	acc=0.9804	8935.8 examples/second
[Step=750]	Loss=0.0581	acc=0.9801	9351.3 examples/second

Test Loss=0.3322, Test acc=0.9107

Epoch: 4

[Step=800]	Loss=0.0512	acc=0.9839	3577.3 examples/second
[Step=850]	Loss=0.0567	acc=0.9816	9041.1 examples/second
[Step=900]	Loss=0.0588	acc=0.9805	8821.9 examples/second
[Step=950]	Loss=0.0588	acc=0.9804	8731.1 examples/second

Test Loss=0.3313, Test acc=0.9108

Epoch: 5

[Step=1000]	Loss=0.0583	acc=0.9824	3322.6 examples/second
[Step=1050]	Loss=0.0599	acc=0.9815	8863.7 examples/second
[Step=1100]	Loss=0.0604	acc=0.9806	8867.6 examples/second
[Step=1150]	Loss=0.0598	acc=0.9806	9044.9 examples/second

Test Loss=0.3308, Test acc=0.9111

Epoch: 6

[Step=1200]	Loss=0.0558	acc=0.9798	3600.8 examples/second
[Step=1250]	Loss=0.0556	acc=0.9813	9778.9 examples/second
[Step=1300]	Loss=0.0555	acc=0.9814	9589.5 examples/second
[Step=1350]	Loss=0.0560	acc=0.9813	9310.6 examples/second

Test Loss=0.3344, Test acc=0.9108

Epoch: 7

[Step=1400]	Loss=0.0567	acc=0.9812	3224.2 examples/second
[Step=1450]	Loss=0.0576	acc=0.9805	9086.4 examples/second
[Step=1500]	Loss=0.0577	acc=0.9804	9005.8 examples/second
[Step=1550]	Loss=0.0599	acc=0.9795	9333.0 examples/second

Test Loss=0.3323, Test acc=0.9119

Saving...

Epoch: 8

[Step=1600]	Loss=0.0539	acc=0.9816	3416.3 examples/second
[Step=1650]	Loss=0.0549	acc=0.9819	9024.3 examples/second
[Step=1700]	Loss=0.0556	acc=0.9812	9039.2 examples/second
[Step=1750]	Loss=0.0562	acc=0.9811	9628.8 examples/second

Test Loss=0.3337, Test acc=0.9100

Epoch: 9

[Step=1800]	Loss=0.0595	acc=0.9809	3369.8 examples/second
[Step=1850]	Loss=0.0582	acc=0.9813	8824.3 examples/second
[Step=1900]	Loss=0.0580	acc=0.9812	8817.4 examples/second
[Step=1950]	Loss=0.0563	acc=0.9817	9178.9 examples/second

Test Loss=0.3334, Test acc=0.9132

Saving...

Epoch: 10

[Step=2000]	Loss=0.0507	acc=0.9844	3437.9 examples/second
[Step=2050]	Loss=0.0541	acc=0.9821	8487.0 examples/second
[Step=2100]	Loss=0.0547	acc=0.9824	8808.1 examples/second
[Step=2150]	Loss=0.0553	acc=0.9820	9507.2 examples/second

Test Loss=0.3339, Test acc=0.9121

Epoch: 11

[Step=2200]	Loss=0.0524	acc=0.9823	3598.0 examples/second
[Step=2250]	Loss=0.0554	acc=0.9810	9539.4 examples/second
[Step=2300]	Loss=0.0549	acc=0.9811	9591.2 examples/second
[Step=2350]	Loss=0.0551	acc=0.9816	9555.4 examples/second

Test Loss=0.3340, Test acc=0.9138

Saving...

Epoch: 12

[Step=2400]	Loss=0.0589	acc=0.9801	3414.0 examples/second
[Step=2450]	Loss=0.0570	acc=0.9804	8979.7 examples/second
[Step=2500]	Loss=0.0568	acc=0.9806	9029.9 examples/second

Test Loss=0.3344, Test acc=0.9129

Epoch: 13

[Step=2550]	Loss=0.0447	acc=0.9844	3638.7 examples/second
[Step=2600]	Loss=0.0516	acc=0.9843	8909.1 examples/second
[Step=2650]	Loss=0.0541	acc=0.9827	8822.3 examples/second
[Step=2700]	Loss=0.0546	acc=0.9820	9170.6 examples/second

Test Loss=0.3373, Test acc=0.9109

Epoch: 14

[Step=2750]	Loss=0.0563	acc=0.9798	3613.7 examples/second
[Step=2800]	Loss=0.0537	acc=0.9818	9061.7 examples/second
[Step=2850]	Loss=0.0528	acc=0.9823	8936.2 examples/second
[Step=2900]	Loss=0.0530	acc=0.9824	8969.6 examples/second

Test Loss=0.3317, Test acc=0.9138

Epoch: 15

[Step=2950]	Loss=0.0498	acc=0.9867	3337.1 examples/second
[Step=3000]	Loss=0.0499	acc=0.9841	9156.2 examples/second
[Step=3050]	Loss=0.0522	acc=0.9831	8975.3 examples/second

```
[Step=3100]      Loss=0.0532      acc=0.9829      9281.6 examples/second
Test Loss=0.3345, Test acc=0.9134
```

Epoch: 16

```
[Step=3150]      Loss=0.0611      acc=0.9780      3484.5 examples/second
[Step=3200]      Loss=0.0600      acc=0.9798      9034.7 examples/second
[Step=3250]      Loss=0.0562      acc=0.9814      8955.5 examples/second
[Step=3300]      Loss=0.0552      acc=0.9818      9386.9 examples/second
Test Loss=0.3330, Test acc=0.9124
```

Epoch: 17

```
[Step=3350]      Loss=0.0602      acc=0.9803      3644.2 examples/second
[Step=3400]      Loss=0.0570      acc=0.9799      9065.9 examples/second
[Step=3450]      Loss=0.0558      acc=0.9808      8923.4 examples/second
[Step=3500]      Loss=0.0531      acc=0.9816      9067.4 examples/second
Test Loss=0.3364, Test acc=0.9141
```

Saving...

Epoch: 18

```
[Step=3550]      Loss=0.0494      acc=0.9837      3300.7 examples/second
[Step=3600]      Loss=0.0499      acc=0.9837      9015.7 examples/second
[Step=3650]      Loss=0.0513      acc=0.9833      8850.9 examples/second
[Step=3700]      Loss=0.0512      acc=0.9832      9148.4 examples/second
Test Loss=0.3413, Test acc=0.9111
```

Epoch: 19

```
[Step=3750]      Loss=0.0495      acc=0.9826      3279.0 examples/second
[Step=3800]      Loss=0.0488      acc=0.9842      9286.5 examples/second
[Step=3850]      Loss=0.0503      acc=0.9835      9694.6 examples/second
[Step=3900]      Loss=0.0513      acc=0.9828      9261.1 examples/second
Test Loss=0.3384, Test acc=0.9111
```


Files already downloaded and verified

Test Loss=0.3364, Test accuracy=0.9141

```
[13]: fig, ax = plt.subplots(1, 3, figsize=(12, 4))

for iii, nbits in enumerate(Nbits_arr_c):
    xx = range(n_epochs)
    ax[iii].plot(xx, test_acc_lab3c_finetune[nbits], label='before finetuning:␣
    ↳%g\max accuracy: %g'
                % (test_acc_lab3c_b4_finetune[nbits], np.
    ↳max(test_acc_lab3c_finetune[nbits])))
    ax[iii].set_xlabel('Epochs')
    ax[iii].set_ylabel('Accuracy')
    ax[iii].set_title('lab3c Accuracy vs Epochs during Finetuning\maxAsymmetric,␣
    ↳NBits=%d' % nbits)
    ax[iii].legend()
```

```
fig.tight_layout()
plt.savefig('Figures/lab3c.pdf', dpi=500, bbox_inches='tight')
```



0.0.7 Lab3 (d) Quantize pruned model

```
[2]: nbits_lab3d = [2, 3, 4]
test_acc_lab3d_b4_finetune = {}
test_acc_arrays_finetune = {}
test_acc_lab3d_after_finetune = {}
n_epochs = 20

for iii, nb in enumerate(nbits_lab3d):
    # Define quantized model and load weight
    Nbits = nb #Change this value to finish (d)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
    net = net.to(device)
    net.load_state_dict(torch.load("net_after_global_iterative_prune.pt"))
    test_acc_lab3d_b4_finetune[Nbits] = test(net)

    # Quantized model finetuning
    test_acc_arrays_finetune[Nbits] = np.array(
        finetune(net, epochs=n_epochs, batch_size=256, lr=0.002, reg=1e-4))

    # Do not load the model with best accuracy since we want the last acc value.
    # net.load_state_dict(torch.load("quantized_net_after_finetune.pt"))
    test_acc_lab3d_after_finetune[Nbits] = test(net)
```

Files already downloaded and verified
Test Loss=8523.9728, Test accuracy=0.1000
==> Preparing data..
Files already downloaded and verified
Files already downloaded and verified

Epoch: 0
 [Step=50] Loss=2.2584 acc=0.1487 2551.6 examples/second
 [Step=100] Loss=2.2225 acc=0.1541 4022.4 examples/second
 [Step=150] Loss=2.2040 acc=0.1603 2889.7 examples/second
 Test Loss=2.1456, Test acc=0.1768
 Saving...

Epoch: 1
 [Step=200] Loss=2.1437 acc=0.1777 1498.7 examples/second
 [Step=250] Loss=2.1324 acc=0.1859 5755.3 examples/second
 [Step=300] Loss=2.1203 acc=0.1960 7315.4 examples/second
 [Step=350] Loss=2.1173 acc=0.2018 6008.7 examples/second
 Test Loss=2.0912, Test acc=0.2219
 Saving...

Epoch: 2
 [Step=400] Loss=2.0949 acc=0.2173 2328.7 examples/second
 [Step=450] Loss=2.0796 acc=0.2308 7142.5 examples/second
 [Step=500] Loss=2.0751 acc=0.2343 6791.7 examples/second
 [Step=550] Loss=2.0707 acc=0.2361 5729.8 examples/second
 Test Loss=2.0630, Test acc=0.2419
 Saving...

Epoch: 3
 [Step=600] Loss=2.0576 acc=0.2409 2538.4 examples/second
 [Step=650] Loss=2.0530 acc=0.2421 7277.8 examples/second
 [Step=700] Loss=2.0493 acc=0.2451 6790.3 examples/second
 [Step=750] Loss=2.0461 acc=0.2457 6486.7 examples/second
 Test Loss=2.0509, Test acc=0.2500
 Saving...

Epoch: 4
 [Step=800] Loss=2.0399 acc=0.2627 2230.6 examples/second
 [Step=850] Loss=2.0310 acc=0.2558 7058.7 examples/second
 [Step=900] Loss=2.0260 acc=0.2560 6324.5 examples/second
 [Step=950] Loss=2.0229 acc=0.2574 6454.7 examples/second
 Test Loss=2.0337, Test acc=0.2449

Epoch: 5
 [Step=1000] Loss=2.0018 acc=0.2617 1983.5 examples/second
 [Step=1050] Loss=2.0080 acc=0.2606 7184.5 examples/second
 [Step=1100] Loss=2.0030 acc=0.2634 6992.6 examples/second
 [Step=1150] Loss=2.0020 acc=0.2630 7130.1 examples/second
 Test Loss=2.0206, Test acc=0.2597
 Saving...

Epoch: 6
 [Step=1200] Loss=1.9806 acc=0.2803 1913.3 examples/second

[Step=1250]	Loss=1.9885	acc=0.2740	5945.2 examples/second
[Step=1300]	Loss=1.9872	acc=0.2716	6223.8 examples/second
[Step=1350]	Loss=1.9816	acc=0.2729	6070.4 examples/second

Test Loss=1.9910, Test acc=0.2674
Saving...

Epoch: 7

[Step=1400]	Loss=1.9636	acc=0.2856	2413.9 examples/second
[Step=1450]	Loss=1.9622	acc=0.2817	5625.7 examples/second
[Step=1500]	Loss=1.9618	acc=0.2820	6169.3 examples/second
[Step=1550]	Loss=1.9566	acc=0.2813	6561.9 examples/second

Test Loss=1.9693, Test acc=0.2750
Saving...

Epoch: 8

[Step=1600]	Loss=1.9517	acc=0.2808	2430.4 examples/second
[Step=1650]	Loss=1.9424	acc=0.2841	5413.7 examples/second
[Step=1700]	Loss=1.9401	acc=0.2849	5684.2 examples/second
[Step=1750]	Loss=1.9356	acc=0.2879	6402.2 examples/second

Test Loss=2.0433, Test acc=0.2379

Epoch: 9

[Step=1800]	Loss=1.9403	acc=0.2780	2304.2 examples/second
[Step=1850]	Loss=1.9254	acc=0.2899	5910.4 examples/second
[Step=1900]	Loss=1.9200	acc=0.2916	6326.7 examples/second
[Step=1950]	Loss=1.9147	acc=0.2928	6645.6 examples/second

Test Loss=1.9648, Test acc=0.2771
Saving...

Epoch: 10

[Step=2000]	Loss=1.8978	acc=0.2988	2007.0 examples/second
[Step=2050]	Loss=1.8933	acc=0.3024	5328.3 examples/second
[Step=2100]	Loss=1.8940	acc=0.3034	6084.0 examples/second
[Step=2150]	Loss=1.8897	acc=0.3032	5355.3 examples/second

Test Loss=1.9089, Test acc=0.2939
Saving...

Epoch: 11

[Step=2200]	Loss=1.8837	acc=0.3026	2341.1 examples/second
[Step=2250]	Loss=1.8721	acc=0.3062	6857.4 examples/second
[Step=2300]	Loss=1.8687	acc=0.3090	5330.3 examples/second
[Step=2350]	Loss=1.8647	acc=0.3116	5505.8 examples/second

Test Loss=1.9105, Test acc=0.2882

Epoch: 12

[Step=2400]	Loss=1.8482	acc=0.3193	2237.5 examples/second
[Step=2450]	Loss=1.8453	acc=0.3168	5464.0 examples/second
[Step=2500]	Loss=1.8443	acc=0.3168	5736.0 examples/second

Test Loss=1.9252, Test acc=0.2729

Epoch: 13

[Step=2550]	Loss=1.8271	acc=0.3008	2240.8 examples/second
[Step=2600]	Loss=1.8203	acc=0.3265	7001.6 examples/second
[Step=2650]	Loss=1.8133	acc=0.3278	5569.1 examples/second
[Step=2700]	Loss=1.8097	acc=0.3282	5655.1 examples/second

Test Loss=1.8428, Test acc=0.3200

Saving...

Epoch: 14

[Step=2750]	Loss=1.8031	acc=0.3340	2244.0 examples/second
[Step=2800]	Loss=1.7921	acc=0.3316	6283.3 examples/second
[Step=2850]	Loss=1.7925	acc=0.3335	6634.1 examples/second
[Step=2900]	Loss=1.7897	acc=0.3358	6687.3 examples/second

Test Loss=1.8365, Test acc=0.2960

Epoch: 15

[Step=2950]	Loss=1.7824	acc=0.3461	2300.1 examples/second
[Step=3000]	Loss=1.7727	acc=0.3393	5836.0 examples/second
[Step=3050]	Loss=1.7672	acc=0.3381	6780.5 examples/second
[Step=3100]	Loss=1.7629	acc=0.3405	6285.5 examples/second

Test Loss=1.8815, Test acc=0.2833

Epoch: 16

[Step=3150]	Loss=1.7534	acc=0.3474	2198.5 examples/second
[Step=3200]	Loss=1.7428	acc=0.3513	6244.0 examples/second
[Step=3250]	Loss=1.7403	acc=0.3530	6646.7 examples/second
[Step=3300]	Loss=1.7354	acc=0.3532	6492.4 examples/second

Test Loss=1.7803, Test acc=0.3290

Saving...

Epoch: 17

[Step=3350]	Loss=1.7069	acc=0.3707	2389.7 examples/second
[Step=3400]	Loss=1.7099	acc=0.3621	6052.0 examples/second
[Step=3450]	Loss=1.7125	acc=0.3596	5945.7 examples/second
[Step=3500]	Loss=1.7076	acc=0.3617	6578.9 examples/second

Test Loss=1.7614, Test acc=0.3492

Saving...

Epoch: 18

[Step=3550]	Loss=1.6993	acc=0.3684	2168.6 examples/second
[Step=3600]	Loss=1.6906	acc=0.3697	5199.1 examples/second
[Step=3650]	Loss=1.6898	acc=0.3688	5951.0 examples/second
[Step=3700]	Loss=1.6859	acc=0.3715	7146.8 examples/second

Test Loss=1.7413, Test acc=0.3392

Epoch: 19

[Step=3750]	Loss=1.6793	acc=0.3733	2451.1 examples/second
[Step=3800]	Loss=1.6771	acc=0.3738	7020.0 examples/second
[Step=3850]	Loss=1.6717	acc=0.3757	6958.3 examples/second
[Step=3900]	Loss=1.6672	acc=0.3772	6609.1 examples/second

Test Loss=1.7031, Test acc=0.3565

Saving...

Files already downloaded and verified

Test Loss=1.7031, Test accuracy=0.3565

Files already downloaded and verified

Test Loss=1.1891, Test accuracy=0.6081

=> Preparing data..

Files already downloaded and verified

Files already downloaded and verified

Epoch: 0

[Step=50]	Loss=1.4277	acc=0.5354	4222.7 examples/second
[Step=100]	Loss=1.1835	acc=0.6058	7022.2 examples/second
[Step=150]	Loss=1.0297	acc=0.6548	6769.1 examples/second

Test Loss=0.6359, Test acc=0.7825

Saving...

Epoch: 1

[Step=200]	Loss=0.5752	acc=0.8027	2173.6 examples/second
[Step=250]	Loss=0.5377	acc=0.8179	6475.0 examples/second
[Step=300]	Loss=0.5151	acc=0.8266	6194.4 examples/second
[Step=350]	Loss=0.4983	acc=0.8321	6088.8 examples/second

Test Loss=0.5133, Test acc=0.8250

Saving...

Epoch: 2

[Step=400]	Loss=0.4255	acc=0.8506	2096.5 examples/second
[Step=450]	Loss=0.4051	acc=0.8598	5950.1 examples/second
[Step=500]	Loss=0.3997	acc=0.8615	7218.1 examples/second
[Step=550]	Loss=0.3895	acc=0.8661	6015.9 examples/second

Test Loss=0.4873, Test acc=0.8383

Saving...

Epoch: 3

[Step=600]	Loss=0.3470	acc=0.8766	2682.0 examples/second
[Step=650]	Loss=0.3460	acc=0.8815	5688.7 examples/second
[Step=700]	Loss=0.3396	acc=0.8844	6365.3 examples/second
[Step=750]	Loss=0.3421	acc=0.8827	6453.5 examples/second

Test Loss=0.4538, Test acc=0.8501

Saving...

Epoch: 4

[Step=800]	Loss=0.3275	acc=0.8860	2154.1 examples/second
[Step=850]	Loss=0.3198	acc=0.8885	5757.1 examples/second

[Step=900]	Loss=0.3167	acc=0.8899	5449.9 examples/second
[Step=950]	Loss=0.3155	acc=0.8904	5455.2 examples/second

Test Loss=0.4395, Test acc=0.8558
Saving..

Epoch: 5

[Step=1000]	Loss=0.2984	acc=0.9020	2189.2 examples/second
[Step=1050]	Loss=0.3011	acc=0.8963	6190.7 examples/second
[Step=1100]	Loss=0.2990	acc=0.8958	7163.5 examples/second
[Step=1150]	Loss=0.2936	acc=0.8980	6157.8 examples/second

Test Loss=0.4294, Test acc=0.8606
Saving..

Epoch: 6

[Step=1200]	Loss=0.2749	acc=0.8991	1772.0 examples/second
[Step=1250]	Loss=0.2748	acc=0.9019	5588.0 examples/second
[Step=1300]	Loss=0.2736	acc=0.9026	5355.3 examples/second
[Step=1350]	Loss=0.2742	acc=0.9025	4679.5 examples/second

Test Loss=0.4084, Test acc=0.8677
Saving..

Epoch: 7

[Step=1400]	Loss=0.2616	acc=0.9081	2051.1 examples/second
[Step=1450]	Loss=0.2589	acc=0.9086	6170.7 examples/second
[Step=1500]	Loss=0.2606	acc=0.9071	5862.7 examples/second
[Step=1550]	Loss=0.2606	acc=0.9075	5641.9 examples/second

Test Loss=0.4032, Test acc=0.8680
Saving..

Epoch: 8

[Step=1600]	Loss=0.2536	acc=0.9092	2380.5 examples/second
[Step=1650]	Loss=0.2542	acc=0.9104	7170.1 examples/second
[Step=1700]	Loss=0.2512	acc=0.9118	6127.2 examples/second
[Step=1750]	Loss=0.2497	acc=0.9122	5993.7 examples/second

Test Loss=0.4039, Test acc=0.8685
Saving..

Epoch: 9

[Step=1800]	Loss=0.2454	acc=0.9147	2019.8 examples/second
[Step=1850]	Loss=0.2425	acc=0.9158	6349.9 examples/second
[Step=1900]	Loss=0.2408	acc=0.9151	6465.4 examples/second
[Step=1950]	Loss=0.2417	acc=0.9149	6120.4 examples/second

Test Loss=0.3963, Test acc=0.8713
Saving..

Epoch: 10

[Step=2000]	Loss=0.2398	acc=0.9127	2246.5 examples/second
[Step=2050]	Loss=0.2355	acc=0.9161	5403.4 examples/second

[Step=2100]	Loss=0.2376	acc=0.9152	5753.6 examples/second
[Step=2150]	Loss=0.2367	acc=0.9162	5892.1 examples/second

Test Loss=0.3939, Test acc=0.8755
Saving..

Epoch: 11

[Step=2200]	Loss=0.2332	acc=0.9165	2256.8 examples/second
[Step=2250]	Loss=0.2312	acc=0.9179	5677.7 examples/second
[Step=2300]	Loss=0.2278	acc=0.9191	6285.6 examples/second
[Step=2350]	Loss=0.2310	acc=0.9180	5419.5 examples/second

Test Loss=0.3791, Test acc=0.8772
Saving..

Epoch: 12

[Step=2400]	Loss=0.2261	acc=0.9202	2164.6 examples/second
[Step=2450]	Loss=0.2236	acc=0.9224	5809.2 examples/second
[Step=2500]	Loss=0.2241	acc=0.9219	6129.8 examples/second

Test Loss=0.3780, Test acc=0.8789
Saving..

Epoch: 13

[Step=2550]	Loss=0.1799	acc=0.9238	1804.3 examples/second
[Step=2600]	Loss=0.2177	acc=0.9222	5828.9 examples/second
[Step=2650]	Loss=0.2184	acc=0.9233	6692.2 examples/second
[Step=2700]	Loss=0.2185	acc=0.9237	6149.5 examples/second

Test Loss=0.3976, Test acc=0.8757

Epoch: 14

[Step=2750]	Loss=0.1851	acc=0.9342	1931.4 examples/second
[Step=2800]	Loss=0.2050	acc=0.9252	4914.1 examples/second
[Step=2850]	Loss=0.2103	acc=0.9252	4737.3 examples/second
[Step=2900]	Loss=0.2133	acc=0.9250	6176.9 examples/second

Test Loss=0.3798, Test acc=0.8790
Saving..

Epoch: 15

[Step=2950]	Loss=0.2011	acc=0.9301	2068.7 examples/second
[Step=3000]	Loss=0.2064	acc=0.9265	6152.8 examples/second
[Step=3050]	Loss=0.2099	acc=0.9259	5452.2 examples/second
[Step=3100]	Loss=0.2087	acc=0.9259	6601.5 examples/second

Test Loss=0.3737, Test acc=0.8813
Saving..

Epoch: 16

[Step=3150]	Loss=0.2170	acc=0.9196	2050.2 examples/second
[Step=3200]	Loss=0.2105	acc=0.9246	6915.4 examples/second
[Step=3250]	Loss=0.2071	acc=0.9266	5686.8 examples/second
[Step=3300]	Loss=0.2051	acc=0.9268	6234.8 examples/second

Test Loss=0.3775, Test acc=0.8803

Epoch: 17

[Step=3350]	Loss=0.1894	acc=0.9368	2062.4 examples/second
[Step=3400]	Loss=0.1977	acc=0.9299	7147.3 examples/second
[Step=3450]	Loss=0.2038	acc=0.9288	5848.8 examples/second
[Step=3500]	Loss=0.2009	acc=0.9296	6100.3 examples/second

Test Loss=0.3755, Test acc=0.8786

Epoch: 18

[Step=3550]	Loss=0.1927	acc=0.9318	2099.4 examples/second
[Step=3600]	Loss=0.1999	acc=0.9315	6471.2 examples/second
[Step=3650]	Loss=0.1972	acc=0.9315	5835.3 examples/second
[Step=3700]	Loss=0.1970	acc=0.9315	6041.3 examples/second

Test Loss=0.3723, Test acc=0.8828

Saving...

Epoch: 19

[Step=3750]	Loss=0.1721	acc=0.9381	2535.8 examples/second
[Step=3800]	Loss=0.1910	acc=0.9331	5669.5 examples/second
[Step=3850]	Loss=0.1923	acc=0.9328	6447.8 examples/second
[Step=3900]	Loss=0.1938	acc=0.9324	5829.4 examples/second

Test Loss=0.3688, Test acc=0.8830

Saving...

Files already downloaded and verified

Test Loss=0.3688, Test accuracy=0.8830

Files already downloaded and verified

Test Loss=0.3924, Test accuracy=0.8793

==> Preparing data..

Files already downloaded and verified

Files already downloaded and verified

Epoch: 0

[Step=50]	Loss=0.6967	acc=0.7641	4052.4 examples/second
[Step=100]	Loss=0.4982	acc=0.8292	5687.7 examples/second
[Step=150]	Loss=0.4115	acc=0.8585	5593.0 examples/second

Test Loss=0.3715, Test acc=0.8811

Saving...

Epoch: 1

[Step=200]	Loss=0.2036	acc=0.9375	2385.5 examples/second
[Step=250]	Loss=0.2031	acc=0.9301	6396.2 examples/second
[Step=300]	Loss=0.2006	acc=0.9309	6802.8 examples/second
[Step=350]	Loss=0.1965	acc=0.9321	6917.3 examples/second

Test Loss=0.3271, Test acc=0.8893

Saving...

Epoch: 2

[Step=400]	Loss=0.1735	acc=0.9409	2093.1 examples/second
[Step=450]	Loss=0.1761	acc=0.9380	6020.2 examples/second
[Step=500]	Loss=0.1766	acc=0.9383	5815.3 examples/second
[Step=550]	Loss=0.1762	acc=0.9397	5984.7 examples/second

Test Loss=0.3157, Test acc=0.8940

Saving...

Epoch: 3

[Step=600]	Loss=0.1613	acc=0.9486	2228.8 examples/second
[Step=650]	Loss=0.1710	acc=0.9437	6017.6 examples/second
[Step=700]	Loss=0.1668	acc=0.9451	6502.6 examples/second
[Step=750]	Loss=0.1691	acc=0.9440	5935.8 examples/second

Test Loss=0.3165, Test acc=0.8964

Saving...

Epoch: 4

[Step=800]	Loss=0.1518	acc=0.9492	2282.8 examples/second
[Step=850]	Loss=0.1561	acc=0.9473	6279.8 examples/second
[Step=900]	Loss=0.1583	acc=0.9467	5865.7 examples/second
[Step=950]	Loss=0.1580	acc=0.9463	5459.3 examples/second

Test Loss=0.3109, Test acc=0.8971

Saving...

Epoch: 5

[Step=1000]	Loss=0.1585	acc=0.9467	1470.3 examples/second
[Step=1050]	Loss=0.1578	acc=0.9458	5756.0 examples/second
[Step=1100]	Loss=0.1571	acc=0.9461	6313.0 examples/second
[Step=1150]	Loss=0.1543	acc=0.9468	5814.4 examples/second

Test Loss=0.3100, Test acc=0.9001

Saving...

Epoch: 6

[Step=1200]	Loss=0.1564	acc=0.9437	2333.5 examples/second
[Step=1250]	Loss=0.1507	acc=0.9463	5865.2 examples/second
[Step=1300]	Loss=0.1472	acc=0.9477	5825.0 examples/second
[Step=1350]	Loss=0.1452	acc=0.9487	5583.1 examples/second

Test Loss=0.3526, Test acc=0.8932

Epoch: 7

[Step=1400]	Loss=0.1462	acc=0.9513	2358.5 examples/second
[Step=1450]	Loss=0.1482	acc=0.9503	6254.3 examples/second
[Step=1500]	Loss=0.1451	acc=0.9511	6041.6 examples/second
[Step=1550]	Loss=0.1449	acc=0.9509	6297.7 examples/second

Test Loss=0.3447, Test acc=0.8991

Epoch: 8

[Step=1600]	Loss=0.1276	acc=0.9556	2116.9 examples/second
[Step=1650]	Loss=0.1265	acc=0.9560	5582.5 examples/second

[Step=1700]	Loss=0.1320	acc=0.9530	5041.7 examples/second
[Step=1750]	Loss=0.1342	acc=0.9529	4866.9 examples/second

Test Loss=0.3123, Test acc=0.9019
Saving..

Epoch: 9

[Step=1800]	Loss=0.1250	acc=0.9539	2431.3 examples/second
[Step=1850]	Loss=0.1270	acc=0.9551	4829.3 examples/second
[Step=1900]	Loss=0.1276	acc=0.9543	5838.1 examples/second
[Step=1950]	Loss=0.1294	acc=0.9538	6848.1 examples/second

Test Loss=0.3368, Test acc=0.9031
Saving..

Epoch: 10

[Step=2000]	Loss=0.1238	acc=0.9568	2374.8 examples/second
[Step=2050]	Loss=0.1251	acc=0.9556	5640.0 examples/second
[Step=2100]	Loss=0.1247	acc=0.9552	6107.3 examples/second
[Step=2150]	Loss=0.1245	acc=0.9551	6087.2 examples/second

Test Loss=0.3415, Test acc=0.9003

Epoch: 11

[Step=2200]	Loss=0.1217	acc=0.9578	2327.4 examples/second
[Step=2250]	Loss=0.1211	acc=0.9566	6317.9 examples/second
[Step=2300]	Loss=0.1195	acc=0.9576	6078.1 examples/second
[Step=2350]	Loss=0.1205	acc=0.9574	6444.7 examples/second

Test Loss=0.3428, Test acc=0.9015

Epoch: 12

[Step=2400]	Loss=0.1196	acc=0.9557	2055.1 examples/second
[Step=2450]	Loss=0.1187	acc=0.9564	6098.4 examples/second
[Step=2500]	Loss=0.1199	acc=0.9563	6342.8 examples/second

Test Loss=0.3553, Test acc=0.9015

Epoch: 13

[Step=2550]	Loss=0.1105	acc=0.9629	1285.4 examples/second
[Step=2600]	Loss=0.1130	acc=0.9620	4766.6 examples/second
[Step=2650]	Loss=0.1148	acc=0.9603	5926.2 examples/second
[Step=2700]	Loss=0.1122	acc=0.9612	5831.9 examples/second

Test Loss=0.3735, Test acc=0.9012

Epoch: 14

[Step=2750]	Loss=0.1233	acc=0.9564	2233.3 examples/second
[Step=2800]	Loss=0.1131	acc=0.9602	7137.4 examples/second
[Step=2850]	Loss=0.1098	acc=0.9621	7470.3 examples/second
[Step=2900]	Loss=0.1101	acc=0.9623	6062.7 examples/second

Test Loss=0.3353, Test acc=0.9035
Saving..

Epoch: 15
 [Step=2950] Loss=0.1143 acc=0.9586 2158.8 examples/second
 [Step=3000] Loss=0.1147 acc=0.9590 6172.7 examples/second
 [Step=3050] Loss=0.1135 acc=0.9602 7081.9 examples/second
 [Step=3100] Loss=0.1115 acc=0.9610 6637.8 examples/second
 Test Loss=0.3364, Test acc=0.9027

Epoch: 16
 [Step=3150] Loss=0.1184 acc=0.9570 2410.7 examples/second
 [Step=3200] Loss=0.1092 acc=0.9616 7257.2 examples/second
 [Step=3250] Loss=0.1094 acc=0.9609 7521.6 examples/second
 [Step=3300] Loss=0.1083 acc=0.9617 5921.7 examples/second
 Test Loss=0.3262, Test acc=0.9011

Epoch: 17
 [Step=3350] Loss=0.1003 acc=0.9642 2214.1 examples/second
 [Step=3400] Loss=0.1039 acc=0.9633 6910.0 examples/second
 [Step=3450] Loss=0.1051 acc=0.9623 6623.9 examples/second
 [Step=3500] Loss=0.1048 acc=0.9627 6385.0 examples/second
 Test Loss=0.3378, Test acc=0.9035

Epoch: 18
 [Step=3550] Loss=0.0929 acc=0.9677 2317.9 examples/second
 [Step=3600] Loss=0.0987 acc=0.9667 6866.8 examples/second
 [Step=3650] Loss=0.1022 acc=0.9648 6104.4 examples/second
 [Step=3700] Loss=0.1022 acc=0.9642 4787.4 examples/second
 Test Loss=0.3430, Test acc=0.9055

Saving...

Epoch: 19
 [Step=3750] Loss=0.0954 acc=0.9674 2287.8 examples/second
 [Step=3800] Loss=0.1015 acc=0.9649 6324.0 examples/second
 [Step=3850] Loss=0.0972 acc=0.9665 5853.9 examples/second
 [Step=3900] Loss=0.0986 acc=0.9656 5434.9 examples/second
 Test Loss=0.3385, Test acc=0.9028
 Files already downloaded and verified
 Test Loss=0.3385, Test accuracy=0.9028

```
[25]: fig, ax = plt.subplots(1, 3, figsize=(12, 4))


for iii, nbits in enumerate(nbits_lab3d):
    xx = range(n_epochs)
    ax[iii].plot(xx, test_acc_arrays_finetune[nbits], label='before finetuning:␣
    ↪%g\after finetuning: %g'
                % (test_acc_lab3d_b4_finetune[nbits],␣
    ↪test_acc_lab3d_after_finetune[nbits]))
    ax[iii].set_xlabel('Epochs')
```

```

ax[iii].set_ylabel('Accuracy')
ax[iii].set_title('lab3d Accuracy vs Epochs during Finetuning\n'
                  'Asymmetric,global iterative prune,NBits=%d' % nbits)

ax[iii].legend()
fig.tight_layout()
plt.savefig('Figures/lab3d.pdf', dpi=500, bbox_inches='tight')

```



0.0.8 Lab3 (e) Symmetric quantization

Implement symmetric quantization in FP_layers.py, and repeat the process in (b)

```

[2]: # check the performance of symmetric quantization with 6, 5, 4, 3, 2 bits
# Define quantized model and load weight
Nbits_arr_lab3e = [2, 3, 4, 5, 6]
test_acc_lab3e = []
for iii, nb in enumerate(Nbits_arr_lab3e):
    Nbits = nb #Change this value

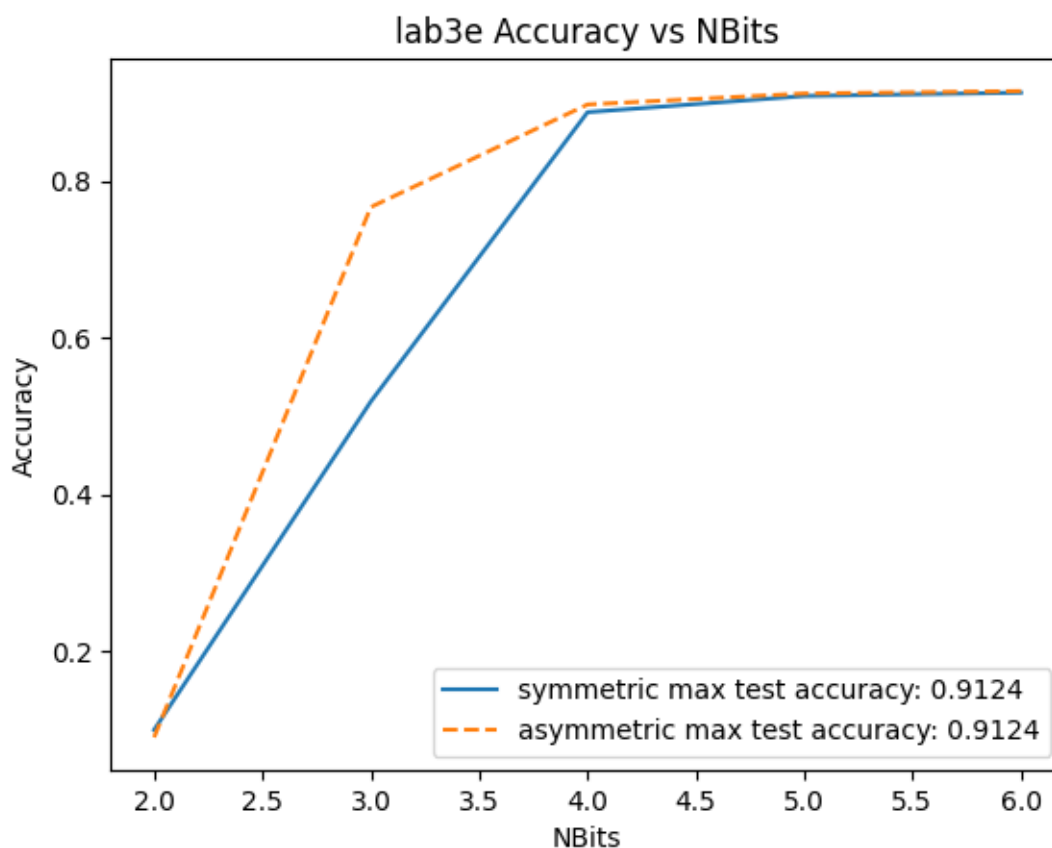
    net = ResNetCIFAR(num_layers=20, Nbits=Nbits, symmetric=True) # symmetric
    net = net.to(device)
    net.load_state_dict(torch.load("pretrained_model.pt"))
    test_acc_lab3e.append(test(net))

```

Files already downloaded and verified
Test Loss=42.7983, Test accuracy=0.1000
Files already downloaded and verified
Test Loss=2.3751, Test accuracy=0.5186
Files already downloaded and verified
Test Loss=0.4227, Test accuracy=0.8875
Files already downloaded and verified
Test Loss=0.3518, Test accuracy=0.9081
Files already downloaded and verified
Test Loss=0.3276, Test accuracy=0.9124

```
[5]: fig, ax = plt.subplots(1, 1)
xx = Nbits_arr_lab3e
ax.plot(xx, test_acc_lab3e, label='symmetric max test accuracy: %g' % np.
    ↳max(test_acc_lab3e))
ax.plot(xx, test_acc_lab3b, '--', label='asymmetric max test accuracy: %g' % np.
    ↳max(test_acc_lab3e))
ax.set_xlabel('NBits')
ax.set_ylabel('Accuracy')
ax.set_title('lab3e Accuracy vs NBits')
ax.legend()

plt.savefig('Figures/lab3e.pdf', dpi=500, bbox_inches='tight')
```



```
[11]: def plot_quant_symmetry(symmetric, NBits=3, bins=100, save=False):

    net = ResNetCIFAR(num_layers=20) # this is for original weights
    net = net.to(device)
    net.load_state_dict(torch.load("pretrained_model.pt"))
    weights_dict_quant = {}
    w_dict_orig = {}
```

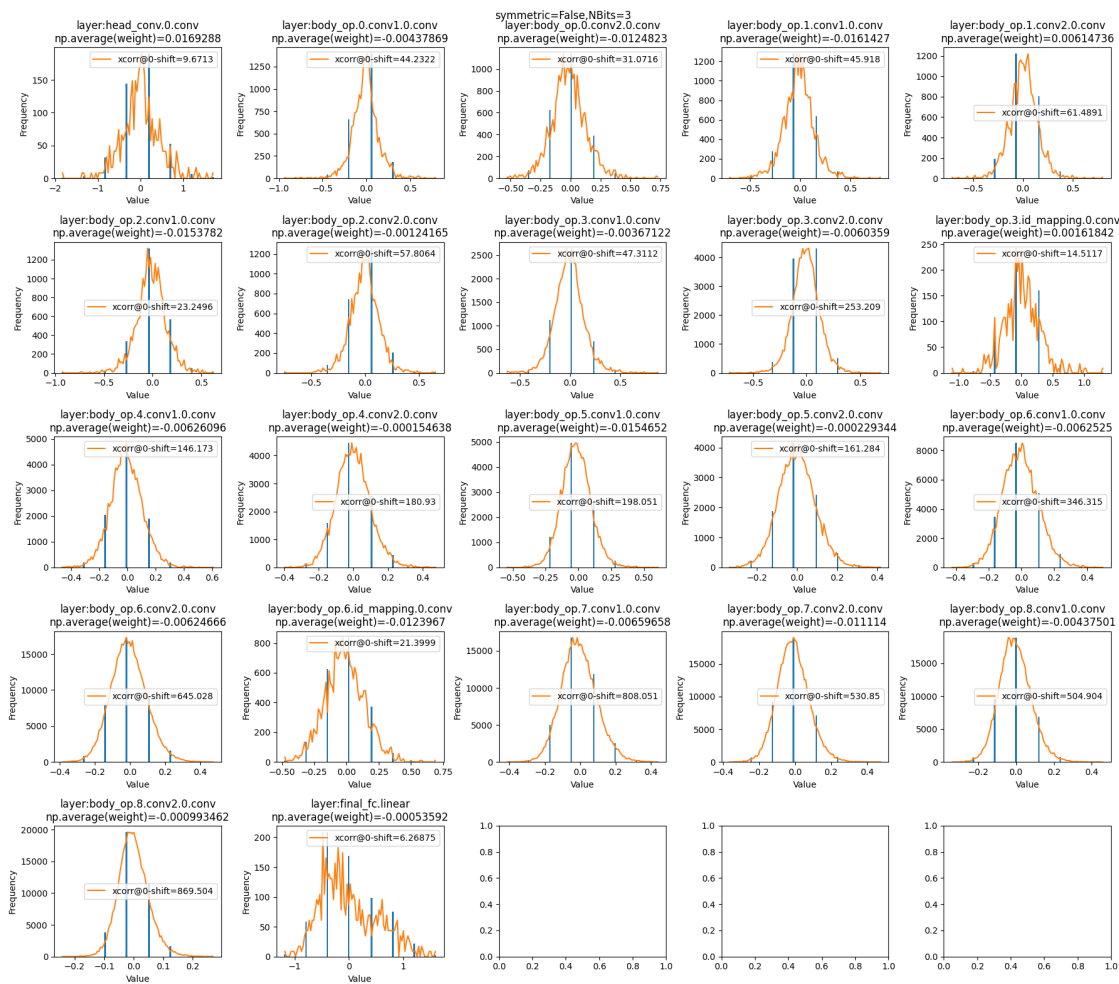
```

conv_list = []
for name, module in net.named_modules():
    if isinstance(module, nn.Conv2d) or isinstance(module, nn.Linear):
        w_dict_orig[str(name)] = module.weight.detach().cpu().numpy().
        flatten()
        www = STE().forward(ctx=None, w=module.weight, bit=NBits,
        symmetric=symmetric)
        np_weights = www.cpu().detach().numpy()
        weights_dict_quant[str(name)] = np_weights.flatten()
        # print(len(weights_dict_quant.keys()))
        # print(weights_dict_quant.keys())
        fig, ax = plt.subplots(5, 5, figsize=(18, 16))
        ax_flat = ax.flatten()
        for iii, layer_name in enumerate(list(weights_dict_quant.keys())):
            y_val, edges = np.histogram(w_dict_orig[layer_name], bins=bins)
            x_val = np.convolve(edges, [0.5, 0.5])[1:-1] # compute average btw
            adjacent elements of edges

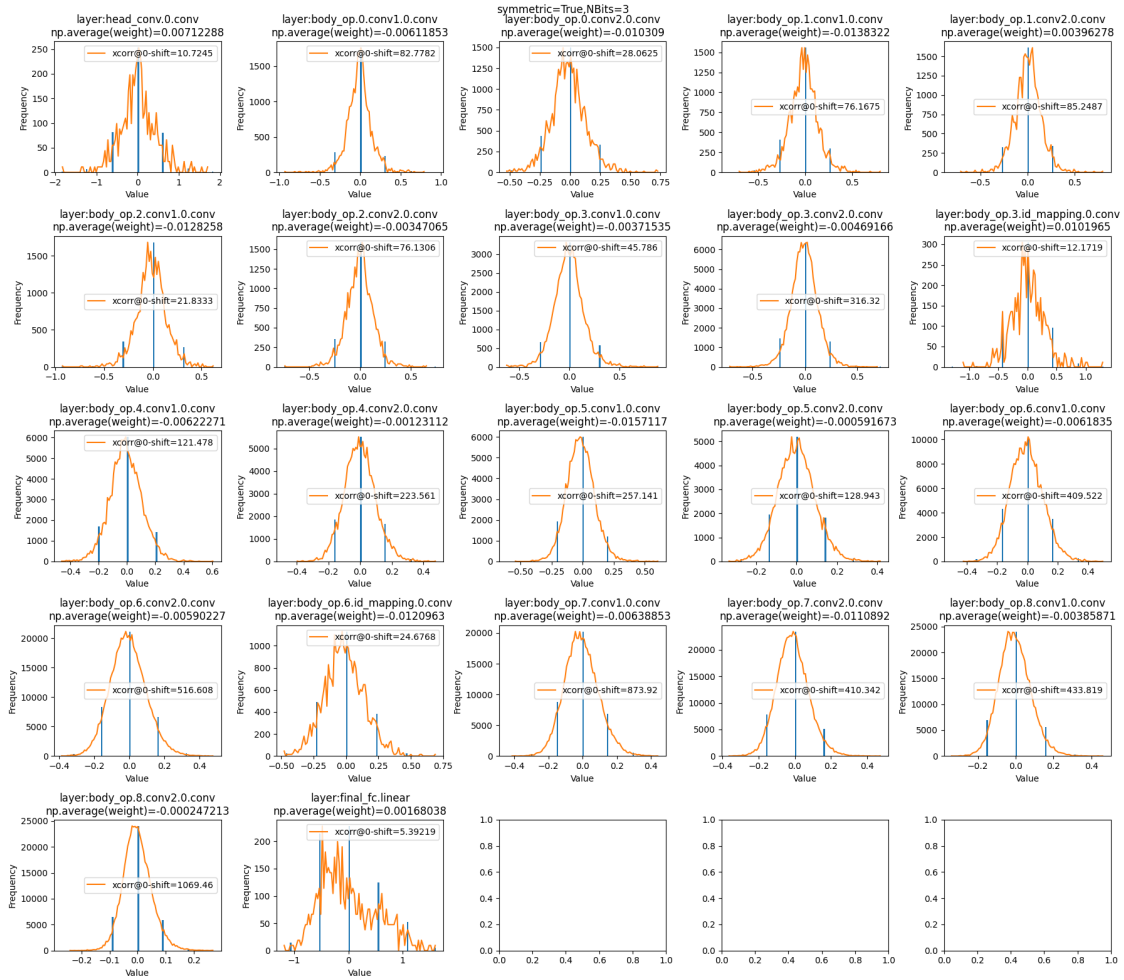
            y_val_quant, _ = np.histogram(weights_dict_quant[layer_name], bins=bins)
            ax_flat[iii].hist(weights_dict_quant[layer_name], bins=bins)
            ax_flat[iii].set_title(str("layer:%s\nnp.average(weight)=%g" %
                                     (layer_name, np.
            average(weights_dict_quant[layer_name]))))
            ax_flat[iii].set_xlabel('Value')
            ax_flat[iii].set_ylabel('Frequency')
            conv = np.convolve(y_val_quant, y_val, 'valid')
            # if len(conv) % 2 == 0:
            #     zeroshift_convolve = 0.5*(conv[int((len(conv)-1)/2)] +
            conv[int((len(conv)+1)/2)]) \
            # / len(weights_dict_quant[layer_name])
            # else:
            #     zeroshift_convolve = (conv[int(len(conv)/2)]) /
            len(weights_dict_quant[layer_name])
            zeroshift_convolve = conv / len(weights_dict_quant[layer_name])
            ax_flat[iii].plot(x_val, y_val * (np.max(y_val_quant) / np.max(y_val)),
                             label='xcorr@0-shift=%g' % zeroshift_convolve)
            ax_flat[iii].legend()
            conv_list.append(zeroshift_convolve)
        fig.suptitle('symmetric=' + str(symmetric) + ',NBits=' + str(NBits))
        fig.tight_layout()
        if save:
            plt.savefig('Figures/lab3e_hist_sym=%s_NBits=%d.pdf' % (symmetric,
            NBits), dpi=500, bbox_inches='tight')
        return conv_list

```

```
c_asym = plot_quant_symmetry(symmetric=False, save=True)
```



```
[12]: c_sym = plot_quant_symmetry(symmetric=True, save=True)
```



```
[20]: difference = np.array(c_asym) - np.array(c_sym)
plt.plot(0.5*difference / (np.array(c_asym) + np.array(c_sym)))
plt.plot([0, 20], [0, 0])
```

```
[20]: [<matplotlib.lines.Line2D at 0x7f6401a07fa0>]
```

