

ECE 661 - Homework 4

Michael Li (z1310)

November 8, 2023

I have adhered to the Duke Community Standard in completing this assignment.

Michael Li

Contents

1 True/False Questions (20 pts)	1
2 Lab 1: Sparse optimization of linear models (30 pts)	3
3 Lab 2: Pruning ResNet-20 model (25 pts)	9
4 Lab 3: Fixed-point quantization and finetuning (25 pts)	13

List of Figures

1 Lab1 part b, plain gradient descent.	4
2 Lab1 part c, iterative pruning.	5
3 Lab1 part d, l_1 -regularisation for $\lambda = \{0.2, 0.5, 1.0, 2.0\}$.	6
4 Lab1 part e, proximal gradient update for threshold = {0.004, 0.01, 0.02, 0.04}.	7
5 Lab1 part f, trimmed l_1 for $\lambda = \{1.0, 2.0, 5.0, 10.0\}$.	8
6 Test accuracy vs pruning percentage q .	9
7 Test accuracy vs epoch for $q = 0.7$.	10
8 Test accuracy vs epoch, iterative pruning.	11
9 Lab2 part e, global iterative pruning.	12
10 Partial code, FP_layers.py.	13
11 Test accuracy vs Nbits.	14
12 Test accuracy vs Epoch for Nbits $\in \{4, 3, 2\}$.	15
13 Test accuracy vs Epoch for Nbits $\in \{4, 3, 2\}$ with model trained using global iterative pruning.	16
14 Final test accuracy Nbits $\in \{4, 3, 2\}$, comparing asymmetric and symmetric pruning.	17

1 True/False Questions (20 pts)

Problem 1.1 (2 pts) Generally speaking, the weight pruning does not intervene the weight quantisation, as they are orthogonal techniques in compressing the size of DNN models.

True. We did this in lab 3 of this assignment. Both can be applied without interfering with each other since the effect on the weight values are different: weight pruning is a form of regularisation (and occasionally can save computation resources with applicable hardware) while weight quantisation attempts to save memory usage.

Problem 1.2 (2 pts) In weight pruning techniques, the distribution of the remaining weights does not affect the inference latency.

False. Non-structured sparsity doesn't bring much speedup on traditional GPUs even with specialised sparse-matrix multiplication algorithms according to Lecture 12 slide 25. Structured sparsity constrains the distribution of the nonzero weights, which may result in shorter inference latency.

Problem 1.3 (2 pts) In deep compression pipeline, even if we skip the quantisation step, the pruned model can still be effectively encoded by the following Huffman coding process as pruning greatly reduces the number of weight variables.

False. Without quantisation, it is unlikely that floats will repeat, so Huffman encoding will not help. Huffman encoding only increases storage efficiency when there are repetitions in the values of the weights. This means that the most commonly appearing weights can be represented in fewer bits to save the most space and a look-up table can be used. With quantisation, there exists repetition of values in the weight matrix, so Huffman encoding will reduce number of bits.

Problem 1.4 (2 pts) Directly using SGD to optimize a sparsity-inducing regulariser (i.e. L-1, Deep-Hoyer etc.) with the training loss will lead to exact zero values in the weight elements, there's no need to apply additional pruning step after the optimization process.

False. The weight values will be *guided* towards 0 and won't be exactly zero. According to lecture 13, slide 19, a final pruning step is required to achieve a sparse model; the pruning step's necessity means that the sparsity-inducing regulariser only makes the weights go very close to zero.

Problem 1.5 (2 pts) Using soft thresholding operator will lead to better results comparing to using L-1 regularisation directly as it solves the "bias" problem of L-1.

False. Soft-thresholding is still "biased" since the larger values always shrink by λ . This causes a loss of variance since the larger weights become uniformly smaller according to lecture 14, slide 9.

Problem 1.6 (2 pts) Group Lasso can lead to structured sparsity on DNNs, which is more hardware-friendly. The idea of Group Lasso comes from applying L-2 regularisation to the L-1 norm of all of the groups.

False. It is the other way round. One should apply L-1 regularisation to the L-2 norms of all the groups according to lecture 13, slide 21. This makes sure that the L-2 norms all shrink to zero due to L-1 regularisation. In any case, applying L-2 regularisation to an expression is worse at making the expression go to zero compared to applying L-1.

Problem 1.7 (2 pts) Proximal gradient descent introduces an additional proximity term to minimize regularisation loss in the proximity of weight parameters. The proximity term allows smoother convergence of the overall objective.

True. According to lecture 14, slide 6, convergence will become smoother. This behaviour is observed in Figure 3 (L-1) and 4 (proximal gradient update). Compare the smoother convergence in Figure 4 with the more squiggly convergence in Figure 3 (vanilla L-1).

Problem 1.8 (2 pts) Models equipped with early exits allows some inputs to be computed with only part of the model, thus overcoming the issue of over-fitting and overthinking.

True. According to lecture 14, slides 27-28, early exits reduces over-fitting and overthinking and improves testing accuracy. Using less layers to identify and predict inputs that are less complicated reduces the chance that the model will mis-predict in the later stages.

Problem 1.9 (2 pts) When implementing quantisation-aware training with STE, gradients are quantised during back-propagation, ensuring that updates are consistent with the quantised weights.

False because the gradients are NOT quantised during back-propagation (lecture 15, slide 10). The full-precision gradients and weight values are always kept when training the model. The quantised weights are used in the forward-propagating step with prediction and loss calculation. The model would become unstable if the gradient values are also quantised.

Problem 1.10 (2 pts) Comparing to quantising all the layers in a DNN to the same precision, mixed-precision quantisation scheme can reach higher accuracy with a similar-sized model.

True, according to lecture 15 slide 28. For example, more "important" layers (those that affect the loss a lot when their weights are changed, quantitatively determined using the Hessian) require more bits to quantise (lecture 15 slide 35) and fewer quantisation bits are required for layers with more weight values.

2 Lab 1: Sparse optimization of linear models (30 pts)

(a) (4 pts) Theoretical analysis: with learning rate μ , and supposing the weight you have after step k is W^k , derive the symbolic formulation of weight W^{k+1} after step $k + 1$ of full-batch gradient descent with $X_i, y_i, i \in \{1, 2, 3\}$.

According to slide 9 lecture 2, the relationship between W^{k+1} and W^k is $W^{k+1} = W^k - \mu \frac{\partial L}{\partial W} \Big|_{W=W^k}$.

$$\begin{aligned} W^{k+1} &= W^k - \mu \frac{\partial L}{\partial W} \Big|_{W=W^k} = W^k - \mu \frac{\partial}{\partial W} \left[\sum_i (X_i W - y_i)^2 \right] \Big|_{W=W^k} \\ &= W^k - \mu \left[\sum_i \frac{\partial}{\partial W} (X_i W - y_i)^2 \Big|_{W=W^k} \right] \end{aligned} \quad (1)$$

Let $u = X_i W - y_i$ and $g(u) = u^2 = (X_i W - y_i)^2$ where $g : \mathbb{R}^1 \rightarrow \mathbb{R}^1$. Using rules of matrix calculus and the numerator layout as explained here (https://en.wikipedia.org/wiki/Matrix_calculus#Numerator-layout_notation),

$$\frac{\partial L}{\partial W} = \sum_i \frac{\partial}{\partial W} (X_i W - y_i)^2 = \sum_i \frac{\partial g(u)}{\partial u} \frac{\partial u}{\partial W} = \sum_i [(2 \cdot (X_i W - y_i) X_i^T)] \quad (2)$$

Combining Equations 1 and 2,

$$W^{k+1} = W^k - \mu \left[\sum_i \frac{\partial}{\partial W} (X_i W - y_i)^2 \Big|_{W=W^k} \right] = \boxed{W^k - \mu \cdot \sum_i [2 \cdot (X_i W^k - y_i) X_i^T]} \quad (3)$$

(b) (3 pts) In Python, directly minimize the objective L without any sparsity-inducing regularisation/constraint. Plot the value of $\log(L)$ vs. number of steps throughout the training, and use another figure to plot how the value of each element in W is changing throughout the training.

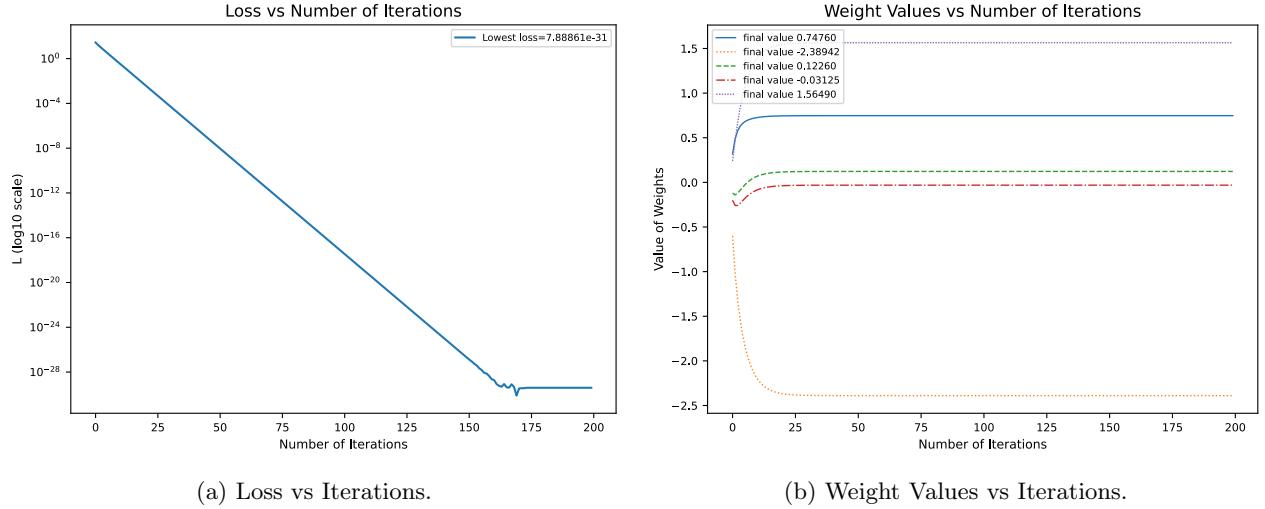


Figure 1: Lab1 part b, plain gradient descent.

```
def no_regularisation_update(in_x, in_y, w_prev, l_1_reg_param=None, lr=None):
    n_r, n_c = in_x.shape
    grad = 0
    for inner in range(n_r):
        grad += 2 * (np.dot(in_x[inner, :], w_prev) - in_y[inner]) * in_x[inner].reshape(-1, 1)
    if lr is not None:
        return w_prev - lr * grad
    else:
        return grad
```

The code snippet above shows my function for calculating gradient descent with no pruning, no regularisation, and no proximal gradient update. The sum in Equation 3 is computed (i.e., the gradient is computed) and the weight is updated using lr . The gradient update function are passed in as the argument `grad_fn` in `run_grad_descent(...)`. Please refer to the PDF of the Jupyter Notebook for more details.

From your result, is W converging to an optimal solution? Is W converging to a sparse solution?

W is converging to an optimal solution since the loss is extremely close to zero at iteration 199 (Figure 1a). However, W isn't converging to a sparse solution since all of the weight values are nonzero (1b).

(c) (6 pts) Since we have the knowledge that the ground-truth weight should have $\|W\|_0 \leq 2$, we can apply projected gradient descent to enforce this sparse constraint. Redo the optimization process in (b), this time prune the elements in W after every gradient descent step to ensure $\|W^l\|_0 \leq 2$. Plot the value of $\log(L)$ throughout the training, and use another figure to plot the value of each element in W in each step.

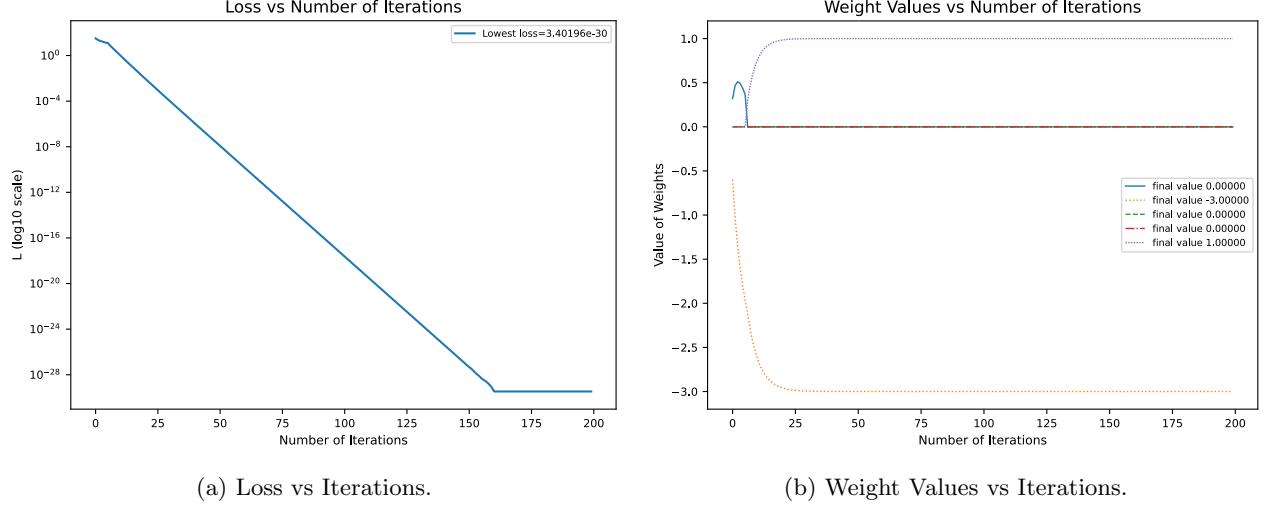


Figure 2: Lab1 part c, iterative pruning.

```
def run_grad_descent(w_init, lr, in_x, in_y, epochs, w0_prune=False, thresh=None, l_1_reg_param=None,
                     grad_fn=no_regularisation_update):
    # not shown ...
    for i in range(epochs):
        if not w0_prune and thresh is not None:  # only for lab 1 f, l_1-trimmed
            w_curr = grad_fn(in_x, in_y, w_prev, l_1_reg_param, lr=lr, thresh=thresh)
        else:  # lab 1b, c, d, e
            w_curr = grad_fn(in_x, in_y, w_prev, l_1_reg_param, lr=lr)

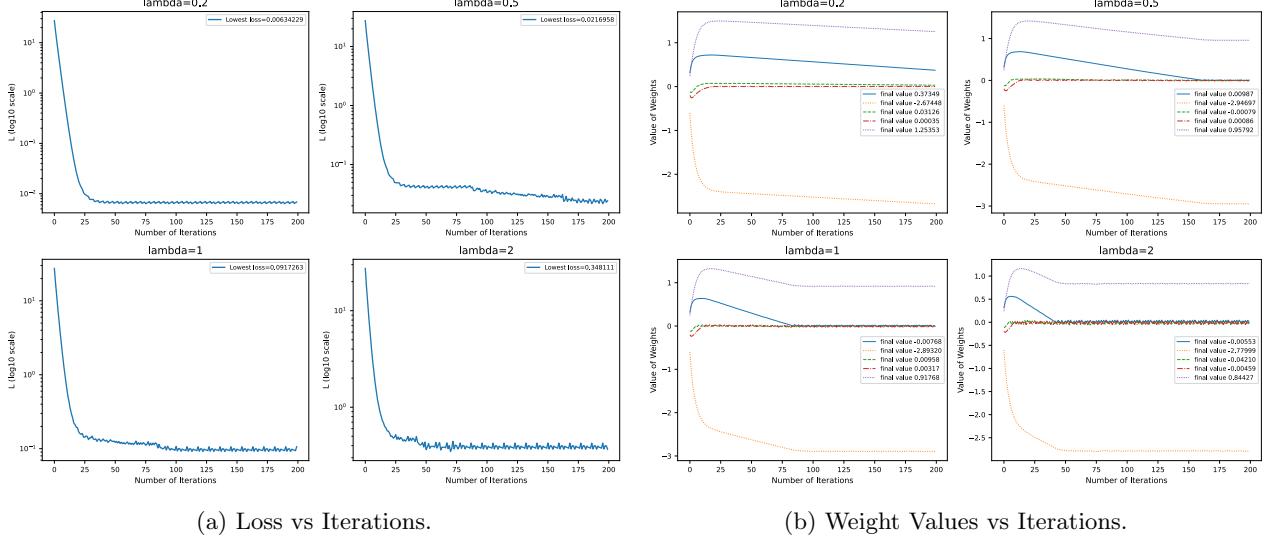
        if w0_prune and thresh is not None:  # only for lab 1 c, pruning
            w_curr_sorted_idx = np.argsort(np.abs(w_curr.flatten()))  # sort ascending
            for inner in range(len(w_init) - thresh):
                # len(w_init) - thresh = 5 - 2 = 3, so three smallest weights (absolute value) are
                # forced to zero
                w_curr[w_curr_sorted_idx[inner], 0] = 0  # zero-ing out smallest elements
    # not shown ...
    return # not shown ...
```

The above code snippet highlights the weight pruning process. The for loop is for gradient update. In lab1 part c `grad_fn == no_regularisation_update` (i.e., using vanilla gradient update) and the condition `if w0_prune and thresh is not None` is satisfied. The coordinates in the weight vector with the smallest absolute values are set to zero. Please refer to the PDF of the Jupyter Notebook for more details.

From your result, is W converging to an optimal solution? Is W converging to a sparse solution?

W is converging to an optimal solution (Figure 2a) and W is also converging to a sparse solution since Figure 2b shows that there are only two coordinates in the weight vector are nonzero.

(d) (5 pts) In this problem we apply l_1 -regularisation to induce the sparse solution. The minimisation objective therefore changes to $L + \lambda ||W||_1$. Please use full-batch gradient descent to minimize this objective, with $\lambda = \{0.2, 0.5, 1.0, 2.0\}$ respectively. For each case, plot the value of $\log(L)$ throughout the training, and use another figure to plot the value of each element in W in each step.



(a) Loss vs Iterations.

(b) Weight Values vs Iterations.

Figure 3: Lab1 part d, l_1 -regularisation for $\lambda = \{0.2, 0.5, 1.0, 2.0\}$.

```
def l_1_norm_update(in_x, in_y, w_prev, l_1_reg_param, lr):
    """
    Args: # not shown ...
    Returns: updated weight
    """
    grad = no_regularisation_update(in_x, in_y, w_prev, l_1_reg_param, lr=None) + \
           l_1_reg_param * np.sign(w_prev)
    return w_prev - lr * grad
```

The above code snippet shows the weight-update function used for lab 1 part d. The l_1 -norm gradient update adds the gradient of the l_1 -norm ($w_{i+1} = \lambda \cdot \text{sgn}(w_i)$) to the gradient of the loss function. Please refer to the PDF of the Jupyter Notebook for more details.

From your result, comment on the convergence performance under different λ .

See Figure 3 for the required plots. The performance of the optimisation process using l_1 -norm becomes worse as λ becomes larger. There is significant oscillation in the loss value as the regularisation parameter increases in magnitude. The loss does not converge to a low value, meaning that the optimisation process is stuck on a local minimum. There are small, sudden decreases in the loss value when $\lambda = \{0.5, 1, 2\}$ in Figure 3a (perhaps because the optimisation process jumped from one local minimum into another local minimum).

The values of the coordinates in the weight vector (part (d)) for $\lambda = 0.5$ are very close to those in the weight vector in part (c) (iterative pruning). However, due to oscillations in the loss, the weight values using l_1 -regularisation never end up being equal to the weight values when using iterative pruning, causing the loss to never decrease beyond 10^{-2} . The coordinates in the weight values go to zero more quickly but oscillate more for larger values of λ .

(e) (6 pts) Here we optimise the same objective as in (d), this time using **proximal gradient update**. Recall that the proximal operator of the l_1 regulariser is the soft thresholding function. Set the threshold in the soft thresholding function to $\{0.004, 0.01, 0.02, 0.04\}$ respectively. Plot the value of $\log(L)$ throughout the training, and use another figure to plot the value of each element in W in each step. (Hint: Optimizing $L + \lambda||W||_1$ using gradient descent with learning rate μ should correspond to proximal gradient update with threshold $\mu\lambda$)

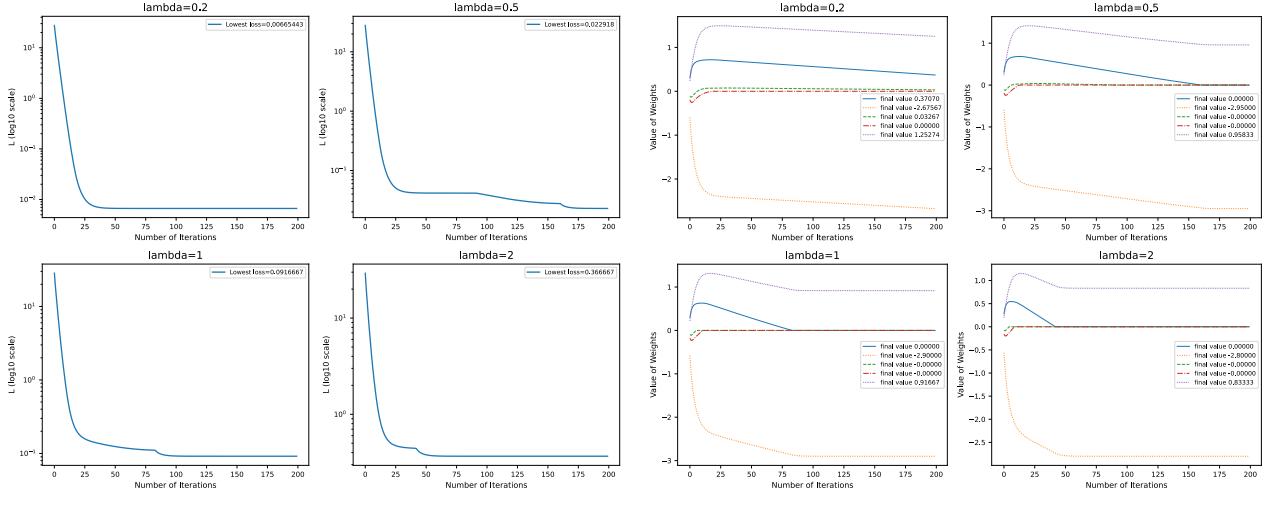


Figure 4: Lab1 part e, proximal gradient update for threshold = $\{0.004, 0.01, 0.02, 0.04\}$.

```
def l_1_norm_proximal_update(in_x, in_y, w_prev, l_1_reg_param, lr, which_weights=None):
    """
    Computes weight using proximal gradient descent (l_1-norm). threshold is l_1_reg_param*lr
    Args: ...
        which_weights: indices of which weights for which proximal update is applied
    Note: threshold = lr * lambda should be in set{0.004, 0.01, 0.02, 0.04}
    Returns: updated weight
    """
    if which_weights is None:
        which_weights = np.array(range(len(w_prev.flatten())))
    weight_tilde = no_regularisation_update(in_x, in_y, w_prev, lr=lr) # main objective
    weight_ret = np.copy(weight_tilde)

    threshold = lr * l_1_reg_param
    for iii in which_weights: # only updates the weights specified in which_weights
        this_weight = weight_tilde[iii, 0]
        weight_ret[iii, 0] = my_relu(np.abs(this_weight) - threshold) * np.sign(this_weight)
    return weight_ret
```

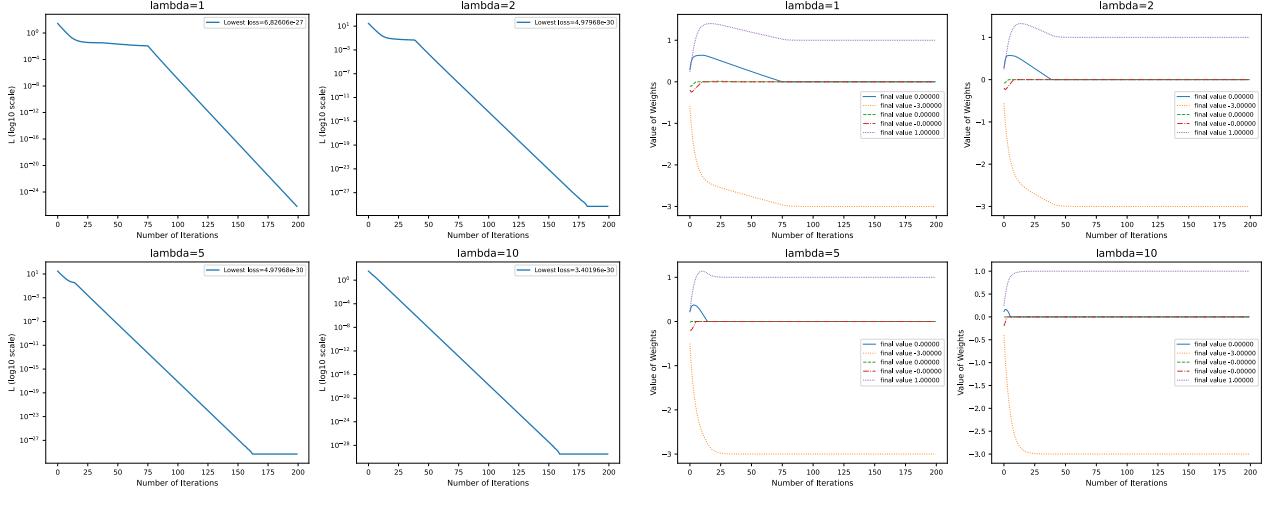
The code snippet above shows the l_1 -norm proximal gradient update. The intermediate weight vector is calculated using vanilla gradient update. Then, each value in the weight vector is edited according to the formula for soft-thresholding. Please refer to the PDF of the Jupyter Notebook for more details.

Compare the convergence performance with the results in (d).

Compared to Figure 3a (l₁), there are less oscillations in Figure 4a (soft-thresholding). However, the shapes and values (ignore the oscillatory behaviour) of the loss curves in Figures 3a and 4a are very similar. Both algorithms fail to find the best and sparsest weight vector in 200 iterations and both manage to exit and then enter local minima.

The values in the weight vector do not oscillate when using proximal gradient update (Figure 4b) compared to part (d) (Figure 3b). Both plots (Figures 4b, 3b) show similar trends in how the weight vector changes over time and neither l_1 nor soft-thresholding manage to produce an optimal sparse solution.

(f) (6 pts) Trimmed l_1 (Tl_1) regularizer is proposed to solve the “bias” problem of l_1 . For simplicity you may implement the T l_1 regularizer as applying a l_1 regularisation with strength λ on the 3 elements of W with the smallest absolute value, with no penalty on other elements. Minimize $L + \lambda Tl_1(W)$ using proximal gradient update with $\lambda = \{1.0, 2.0, 5.0, 10.0\}$ (correspond to the soft thresholding threshold values $\{0.02, 0.04, 0.1, 0.2\}$). Plot the value of $\log(L)$ throughout the training, and use another figure to plot the value of each element in W in each step. Compare the convergence behaviour of the Trimmed l_1 and the l_1 . Also compare the behavior of the early iterations (e.g. first 20) between Tl_1 and iterative pruning.



(a) Loss vs Iterations.

(b) Weight Values vs Iterations.

Figure 5: Lab1 part f, trimmed l_1 for $\lambda = \{1.0, 2.0, 5.0, 10.0\}$.

```
def l1_norm_trimmed_update(in_x, in_y, w_prev, l1_reg_param, lr, thresh):
    """
    thresh: number of elements to keep unchanged when performing proximal l1, returns updated weight
    """
    ascend_idx = np.argsort(np.abs(w_prev.flatten()))
    weight_l1_proximal = l1_norm_proximal_update(in_x, in_y, w_prev, l1_reg_param, lr,
                                                which_weights=ascend_idx[0:w_prev.shape[0]-thresh])
    return weight_l1_proximal
```

The Tl_1 weight update function is implemented in the code snippet above. Proximal gradient update is applied to the largest values in the weight vector as specified by the parameter `thresh`. Please refer to the PDF of the Jupyter Notebook for more details.

Compare the convergence behaviour of the Trimmed l_1 and the l_1 algorithms (part (d)):

In Figure 5a, we see that the Tl_1 algorithm converges to the optimal weight value, while the l_1 algorithm fails to converge (Figure 3a). There is no oscillation using Tl_1 and the algorithm doesn't get stuck on local minima. The weight values for Tl_1 (Figure 5b) do not oscillate; however, the weight value plots with Tl_1 bear strong resemblance to the corresponding plots with the l_1 algorithm (Figure 3b, look at $\lambda = 1, 2$). This means that Tl_1 does not get stuck and goes to the global minimum without oscillating.

Compare the behavior of the early steps (e.g. first 20) between the Trimmed l_1 and the iterative pruning algorithms (part (c)).

In Figure 2a (pruning), the weight vector spends very little time in a local minimum since the loss value starts decreasing quickly in the first few iterations and heads towards the global minimum. In contrast, the weight vector spends more epochs navigating local minima for lower values of λ in Figure 5a (Tl_1) as evinced by the slow initial decrease of the loss value when $\lambda = 1, 2$ compared to pruning. When $\lambda = 10$ and using Tl_1 , (bottom-right of Figure 5a), the weight vector spends less time in a local minimum than with weight pruning since there's no relatively flat part of the loss vs number of iterations plot for $\lambda = 10$ using Tl_1 .

For small λ 's, the weight values using Tl_1 (top 2 plots in Figure 5b) change more slowly compared to weight pruning. Also, weight values overshoot more with Tl_1 and small λ (purple line of the plots with $\lambda = \{1, 2, 5\}$ in Figure 5b). When λ becomes larger, the weight values with Tl_1 (bottom-right of Figure 5b) approach the sparsest optimal solution more rapidly than with weight pruning (Figure 2b). The purple line in Figure 2b (pruning) starts off at zero, while the purple lines in Figure 5b (Tl_1) do not start zero.

3 Lab 2: Pruning ResNet-20 model (25 pts)

(a) (2 pts) In `hw4.ipynb`, run through the first three code block, report the accuracy of the floating-point pre-trained model.

Test Loss is 0.3231, Test accuracy is 0.9150.

(b) (6 pts) Complete the implementation of *pruning by percentage function* in the notebook. Here we determines the pruning threshold in each DNN layer by the ' q -th percentile' value in the absolute value of layer's weight element. Use the next block to call your implemented pruning by percentage. Try pruning percentage $q = 0.3, 0.5, 0.7$. Report the test accuracy q .

```
with torch.no_grad():
    # Convert the weight of "layer" to numpy array
    layer_weight = layer.weight.detach().cpu().numpy()
    # Compute the q-th percentile of the abs of the converted array
    percentile = np.percentile(np.abs(layer_weight.flatten()), q)
    # Generate a binary mask same shape as weight to decide which element to prune
    masked_obj = np.ma.masked_greater_equal(x=np.abs(layer_weight), value=percentile, copy=True)
    mask_int = np.ma.getmask(masked_obj).astype(int)
    # Convert mask to torch tensor and put on GPU
    mask_tensor = torch.tensor(mask_int).to(device)
    # Multiply the weight by mask to perform pruning
    assert mask_int.shape == layer_weight.shape
    # layer.weight.data = mask_tensor * layer.weight.data
    layer.weight.data = layer.weight.data.clone().detach().requires_grad_(True) * mask_tensor
```

The code snippet shown above is the body of the `prune_by_percentage` function. First, the percentile is found based on q . Then, the mask is generated. Finally, the mask is applied to the weight. Please refer to the PDF of the Jupyter Notebook for more details.

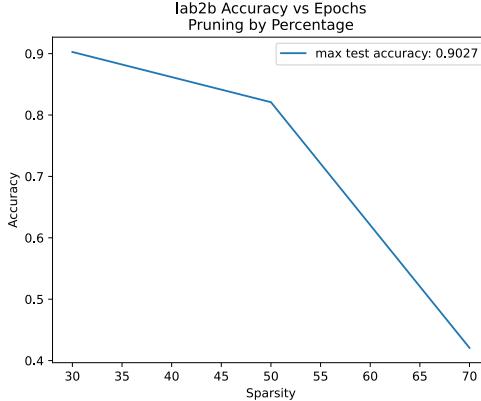


Figure 6: Test accuracy vs pruning percentage q .

The test accuracy values are reported in Figure 6; I assumed that we prune q percent of all weights (instead of keep q percent of all weights un-pruned). For $q = 0.3$, `test_acc = 0.9027`; for $q = 0.5$, `test_acc = 0.8209`; for $q = 0.7$, `test_acc = 0.4207`.

(c) (6 pts) Fill in the `finetune_after_prune` function for pruned model finetuning. Make sure the pruned away elements in previous step are kept as 0 throughout the finetuning process. Finetune the pruned model with $q = 0.7$ for 20 epochs with the provided training pipeline. Report the best accuracy achieved during finetuning. Finish the code for sparsity evaluation to check if the finetuned model preserves the sparsity.

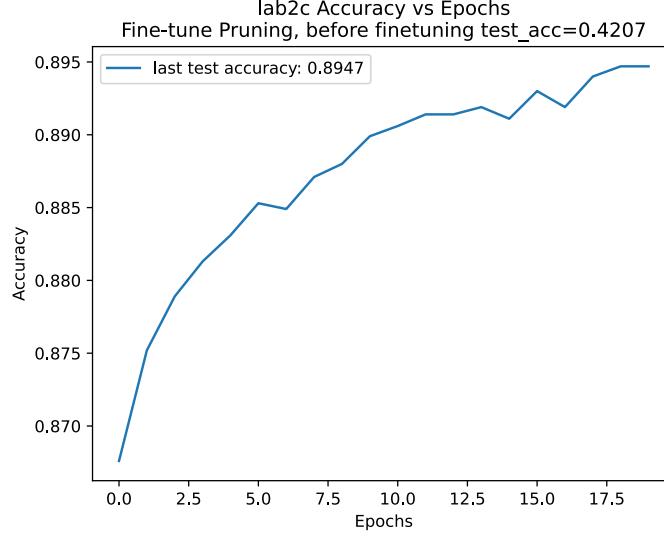


Figure 7: Test accuracy vs epoch for $q = 0.7$.

The best accuracy achieved during fine-tuning is 0.8947 in the final two epochs (Figure 7).

```
def finetune_after_prune(net, trainloader, criterion, optimizer, prune=True):
...
# Your code here: generate a mask in GPU torch tensor to have 1 for nonzero element and 0 for
# zero element
    layer_weight = layer.weight.detach().cpu().numpy()
    mask_obj = np.ma.masked_equal(x=layer_weight, value=0) # where it's zero, mask value is True
    weight_mask[name] = torch.tensor(1 - np.ma.getmask(mask_obj).astype(int)).to(device)
...
# Your code here: Use weight_mask to make sure zero elements remains zero
    layer.weight.data = layer.weight.data.clone().detach().requires_grad_(True) * weight_mask[name]
...
```

The mask generation process and the actual masking process in part (c) are very similar to the code that I explained in part (b). To generate a mask, I used the `np.ma.masked_equal` and moved the mask onto the GPU (note: the making function generates `True` values where the input is zero, meaning that I need to invert the condition by computing `1 - mask`). A different mask is generated for each layer. I ensured that the weight elements remained at zero after back-propagation in the second part of the code snippet above.

```
net.load_state_dict(torch.load("net_after_finetune.pt"))
for name, layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and 'id_mapping' not in name:
        # Your code here:
        ## Convert the weight of "layer" to numpy array
        np_weight = layer.weight.detach().cpu().numpy()
        ## Count number of zeros
        zeros = sum((np_weight == 0).flatten())
        ## Count number of parameters
        total = len(np_weight.flatten())
        # Print sparsity
        print('Sparsity of %s: %g' % (name, zeros/total))
test(net)
```

The code snippet above shows how I made sure that all the layers still had the same number of nonzero elements. I changed the tensor into a numpy array and computed the ratio between the number of zeros and the total length in that numpy array. Please refer to the PDF of the Jupyter Notebook for more details. Please see the next page to confirm that sparsity is preserved.

```
Sparsity of head.conv.0.conv: 0.699074
Sparsity of body.op.0.conv1.0.conv: 0.700087
Sparsity of body.op.0.conv2.0.conv: 0.700087
Sparsity of body.op.1.conv1.0.conv: 0.700087
Sparsity of body.op.1.conv2.0.conv: 0.700087
Sparsity of body.op.2.conv1.0.conv: 0.700087
Sparsity of body.op.2.conv2.0.conv: 0.700087
```

```

Sparsity of body_op.3.conv1.0.conv: 0.699978
Sparsity of body_op.3.conv2.0.conv: 0.699978
Sparsity of body_op.4.conv1.0.conv: 0.699978
Sparsity of body_op.4.conv2.0.conv: 0.699978
Sparsity of body_op.5.conv1.0.conv: 0.699978
Sparsity of body_op.5.conv2.0.conv: 0.699978
Sparsity of body_op.6.conv1.0.conv: 0.699978
Sparsity of body_op.6.conv2.0.conv: 0.700005
Sparsity of body_op.7.conv1.0.conv: 0.700005
Sparsity of body_op.7.conv2.0.conv: 0.700005
Sparsity of body_op.8.conv1.0.conv: 0.700005
Sparsity of body_op.8.conv2.0.conv: 0.700005
Sparsity of final_fc.linear: 0.7

```

(d) (5 pts) Implement iterative pruning. Instead of applying single step pruning before finetuning, try to iteratively increase the sparsity of the model before each epoch of finetuning. Linearly increase the pruning percentage for 10 epochs until reaching 70% in the final epoch (prune by $(7 \times e)\%$ before epoch e) then continue finetune for 10 epochs. Pruned weight can be recovered during the iterative pruning process before the final pruning step. Compare performance with (c).

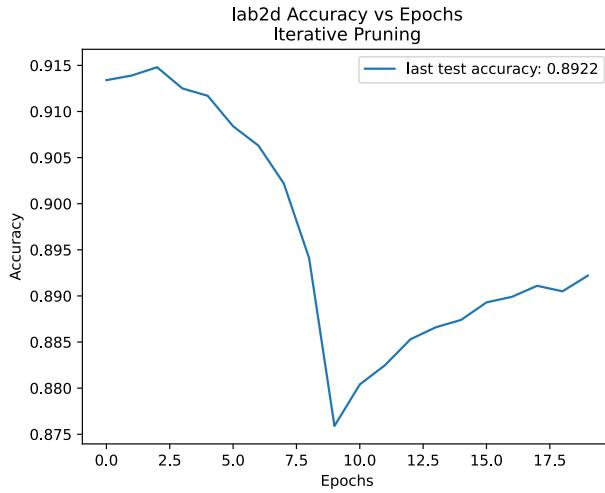


Figure 8: Test accuracy vs epoch, iterative pruning.

```

if epoch < 10:
    for name,layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and 'id_mapping' not in name:
            # Increase model sparsity
            q = (epoch + 1) * 7
            # Linearly increase the pruning percentage for 10 epochs until reaching 70% in the final epoch
            prune_by_percentage(layer, q=q)
if epoch < 9:
    finetune_after_prune(net, trainloader, criterion, optimizer, prune=False)
else: # starts from epoch==9, where q == 70%
    finetune_after_prune(net, trainloader, criterion, optimizer)

```

The code snippet shown above shows the changes to q for the first 10 epochs.

Compare performance with (c):

The performance using iterative pruning (Figure 8) is similar to the performance in (c) (Figure 7). The highest accuracy in part (c) is around the same as the accuracy using iterative pruning.

The weights that were zero before every epoch in part (c) stay zero after the training step, while the weights that were zero before every epoch in iterative finetuning aren't necessarily zero after each training step from epoch 0 to epoch 8 (since `prune=False` in `finetune_after_prune(...)` before epoch 9). The model in (c) isn't able to change the mask while iterative tuning allows the model to learn the optimal location of the nonzero weights. However, the ability to learn the location of the nonzero values doesn't seem to help the model's performance after 20 epochs.

The main effect of iterative pruning is that the accuracy initially decreases as q increases and then increases when q levels off at 70%. Since the model has a different mask for each epoch from epoch 0 to 8, the location of the zero weights are unstable and the model isn't able to learn well in the first few epochs. However, we see that the model's testing accuracy in epoch 9 (70% pruning, iterative pruning) is higher than the model's testing accuracy in epoch 0 (same percentage pruned, part (c)). This suggests that the model in (d) learned a bit when undergoing gradual iterative pruning. The model in (d) not having enough epochs to converge could explain the slightly lower final accuracy compared to (c).

(e) (6 pts) Perform magnitude-based global iterative pruning. Previously we set the pruning threshold of each layer following the weight distribution of the layer and prune all layers to the same sparsity. This will constrain the flexibility in the final sparsity pattern across layers. In this question, fill in the `global_prune_by_percentage` function to perform a global ranking of the weight magnitude from all the layers, and determine a single pruning threshold by percentage for all the layers. Repeat iterative pruning to 70% sparsity, and report final accuracy and the percentage of zeros in each layer.

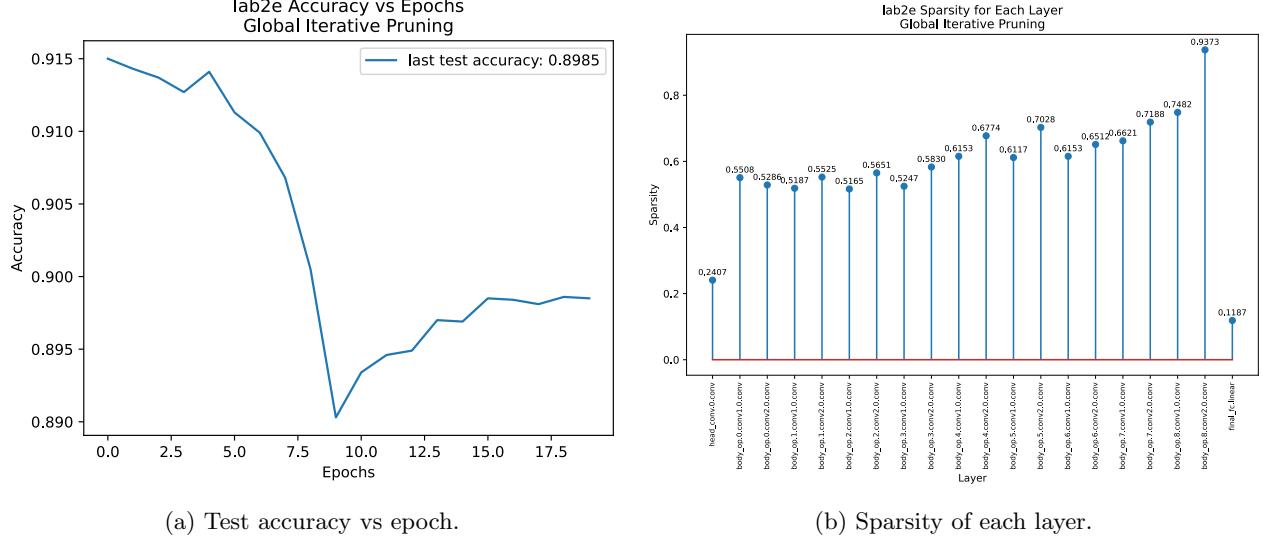


Figure 9: Lab2 part e, global iterative pruning.

```
def global_prune_by_percentage(net, q=70.0):
    # A list to gather all the weights
    flattened_weights = []
    # Find global pruning threshold
    for name,layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and 'id_mapping' not in name:
            # Convert weight to numpy
            # Flatten the weight and append to flattened_weights
            flattened_weights.append(layer.weight.detach().cpu().numpy().flatten())

    # Concat all weights into a np array
    flattened_weights = np.concatenate(flattened_weights)
    # Find global pruning threshold
    threshold = np.percentile(np.abs(flattened_weights.flatten()), q)

    # Apply pruning threshold to all layers
    for name,layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and 'id_mapping' not in name:
            # Convert weight to numpy
            layer_weight = layer.weight.detach().cpu().numpy()

            # Generate a binary mask same shape as weight to decide which element to prune
            masked_obj = np.ma.masked_greater_equal(x=np.abs(layer_weight), value=threshold, copy=True)
            mask_int = np.ma.getmask(masked_obj).astype(int)
            # Convert mask to torch tensor and put on GPU
            mask_tensor = torch.tensor(mask_int).to(device)
            # Multiply the weight by mask to perform pruning
            layer.weight.data = layer.weight.data.clone().detach_.requires_grad_(True) * mask_tensor
```

The above code snippet shows the global pruning function. First, I get a list of all the weights and find the threshold based on the percentile q . Then, I create a mask for each layer and apply the mask to the weights of that layer.

Global iterative pruning at $q = 70\%$ achieves test accuracy of 0.8985 in the last epoch according to Figure 9a. The sparsity for each layer is presented in Figure 9b. Please refer to the PDF of the Jupyter Notebook for more details.

4 Lab 3: Fixed-point quantization and finetuning (25 pts)

(a) (10 pts) As is mentioned in lecture 15, to train a quantized model we need to use floating-point weight as trainable variable while use a straight-through estimator (STE) in forward and backward pass to convert the weight into quantized value. Intuitively, the forward pass of STE converts a float weight into fixed-point, while the backward pass passes the gradient straightly through the quantizer to the float weight.

To start with, implement the STE forward function in `FP_layers.py`, so that it serves as a linear quantizer with dynamic scaling, as introduced on page 9 of lecture 15. Please follow the comments in the code to figure out the expected functionality of each line. Take a screenshot of the finished STE class and paste it into the report. Submission of the `FP_layers.py` file is not required.

```

hw4.ipynb M   FP_layers.py X
.. ece661-hw-fa23 > HW4 > FP_layers.py

20     # Build a mask to record position of zero weights
21     np_weights = w.detach().cpu().numpy()
22     weight_mask = np.ma.masked_equal(x=np_weights, value=0) # zero will be marked true
23     weight_mask = torch.tensor(1 - np.ma.getmask(weight_mask).astype(int)).to(device)
24     # zero will be marked zero, weight_mask onto GPU
25
26     flattened_np_weights = np_weights.flatten()
27
28     # Lab3-(a), Your code here:
29     if symmetric == False:
30         max_, min_ = np.max(flattened_np_weights), np.min(flattened_np_weights)
31         # Compute alpha (scale) for dynamic scaling
32         alpha = max_ - min_
33         # Compute beta (bias) for dynamic scaling
34         beta = min_
35         # Scale w with alpha and beta so that all elements in ws are between 0 and 1
36         ws = (w - beta) / alpha
37
38         step = (2 ** (bit)) - 1
39         # Quantize ws with a linear quantizer to "bit" bits
40         R = (1/step) * torch.round(step * ws) # r_o = 1/(2^k-1) * round((2^k-1)*r_i), assume r_i=r_in=ws
41         # Scale the quantized weight R back with alpha and beta
42         wq = alpha * R + beta
43
44     # Lab3 (e), Your code here:
45     else: # if symmetric == True
46         # In symmetric quantization, the quantization levels are symmetric to zero.
47         abs_max_ = np.max(np.abs(flattened_np_weights))
48         alpha = abs_max_
49         beta = 0 # do not shift
50
51         # Scale w with alpha and beta=0 so that all elements in ws are between -0.5 and 0.5
52         ws = ((w - beta) / alpha)
53
54         step = (2 ** (bit - 1)) - 1
55
56         # Quantize ws with a linear quantizer to "bit" bits
57         R = (1/step) * torch.round(step * ws)
58         # Scale the quantized weight R back with alpha and beta
59         wq = alpha * R + beta
60         pass
61
62     # Restore zero elements in wq
63     wq = wq * weight_mask
64
65
66     return wq

```

Figure 10: Partial code, `FP_layers.py`.

The `symmetric == False` case implements the algorithm shown on Slide 9, Lecture 15. First, I compute `alpha` and `beta`. Then, I compute the shifted and scaled weight `ws` whose values should all be between $\{0, 1\}$. The scaled weight is then quantised on line 40 and the original values are recovered on line 42. Finally, the mask is applied on line 63.

The code block under the `else` condition completes the requirements for lab 3 (e).

(b) (2 pts) In `hw4.ipynb`, load pretrained ResNet-20 model, **report the accuracy of the floating-point pretrained model**. Then set `Nbits` in the first line of block 4 to 6, 5, 4, 3, and 2 respectively, run it and **report the test accuracy** you got.

The floating point pretrained model has the following accuracy:

```
before quantisation: Test accuracy=0.9150
```

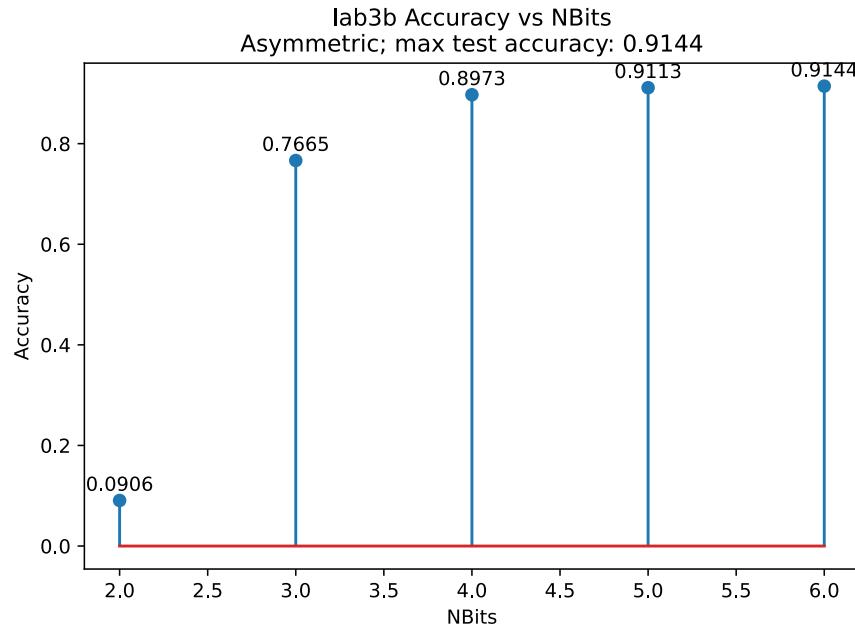


Figure 11: Test accuracy vs Nbits.

Please refer to Figure 11 for the test accuracy quantisation. The test accuracy becomes lower as the number of bits used in the quantisation decreases. This is expected because when `Nbits` becomes smaller, the number of distinct values in the weights become smaller. Smaller `Nbits` values mean that the model is less likely to record the original distribution of the weights accurately, especially for weights that are of high magnitudes.

(c) (5 pts) With `Nbits` set to 4, 3, and 2, run code blocks 4 and 5 to finetune the quantised model for 20 epochs. You do not need to change other parameter in the finetune function. For each precision, report the highest testing accuracy you get during finetuning.

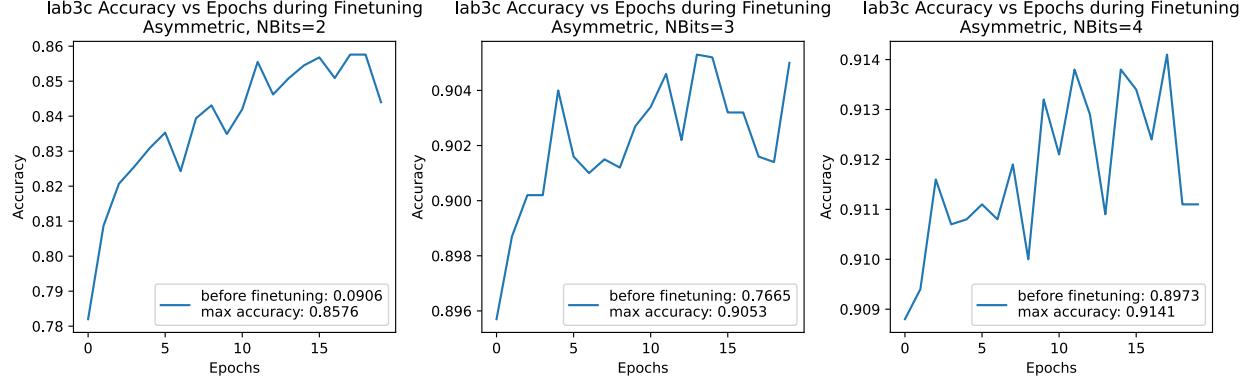


Figure 12: Test accuracy vs Epoch for $\text{Nbits} \in \{4, 3, 2\}$.

The highest test accuracy for each quantisation (shown in Figure 12) is higher than the test accuracy before fine-tuning. Even after a one epoch the testing accuracy becomes higher than the testing accuracy before finetuning for `Nbits == 2` and `Nbits == 3`, which means that the quantised model benefits a great deal from finetuning when `Nbits` is small.

Comment on the relationship between precision and accuracy, and on the effectiveness of finetuning.

As the precision increases (higher value for `Nbits`), the highest test accuracy increases. With more bits, the model is able to change more flexibly during the training process and thus reach a state where the loss is lower. Too much information is lost about the model when the number of bits used to represent weight values is too low, causing the model to be less accurate. There's also too much constraints on what the values of each weight take on, meaning that convolution and fully-connected layers aren't fine-grained enough.

The absolute difference between the final test accuracy and the test accuracy before finetuning decreases as `Nbits` increases. When `Nbits` is large, the model already performs quite well before finetuning occurs, meaning that there's limited room for further improvement. When `Nbits` is small (`Nbits == 2` for example), there's a lot of room for the accuracy to increase given that the initial accuracy before finetuning with `Nbits == 2` is so low. **So, we can say that finetuning is less effective when Nbits is large.**

```

Nbits_arr_c = [2, 3, 4]
test_acc_lab3c_finetune = {}
# test accuracy for each finetuning as fn of epoch
final_test_acc_lab3c = {} # test acc after finetuning for each Nbit
n_epochs = 20
test_acc_lab3c_b4_finetune = {} # test acc before finetuning

for iii, nb in enumerate(Nbits_arr_c):
    # Define quantized model and load weight
    Nbits = nb #Change this value to finish (b) and (c)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
    net = net.to(device)
    net.load_state_dict(torch.load("pretrained_model.pt"))
    test_acc_lab3c_b4_finetune[Nbits] = test(net)

    # Quantized model finetuning
    test_acc_lab3c_finetune[Nbits] = np.asarray(
        finetune(net, epochs=n_epochs, batch_size=256, lr=0.002, reg=1e-4))
    # finetune() returns list of val_accuracy values because I edited it
    # Load the model with best accuracy
    net.load_state_dict(torch.load("quantized_net_after_finetune.pt"))
    final_test_acc_lab3c[Nbits] = test(net)

```

The code snippet above shows the process for fine-tuning the quantised model. In each iteration, `Nbits` is changed, the model is re-created & re-loaded with the pretrained model, and its accuracy is recorded before finetuning. The best model that resulted from finetuning is loaded and its test accuracy is recorded.

(d) (4 pts) In practice, we want to apply both pruning and quantization on the DNN model. Here we explore how pruning will affect quantization performance. Please load the checkpoint of the 70% sparsity model with the best accuracy from Lab 2, repeat the process in (c), and report the accuracy before and after finetuning.

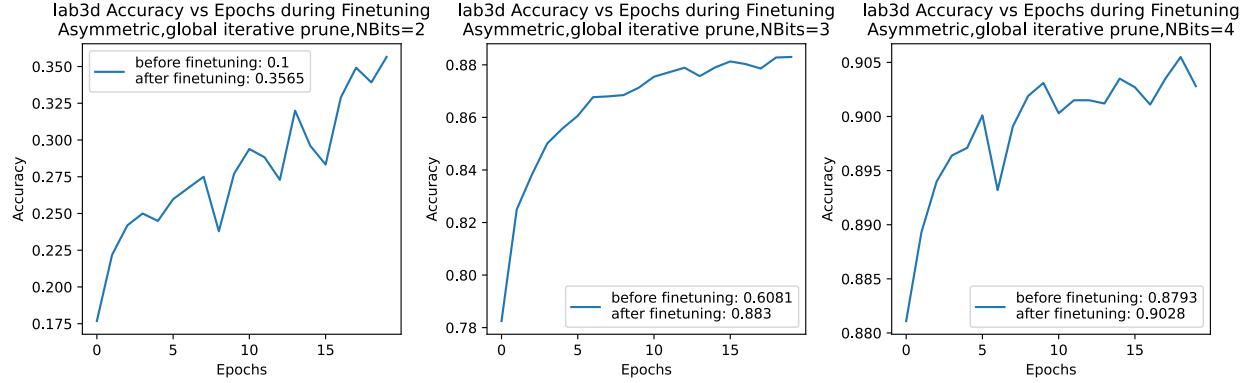


Figure 13: Test accuracy vs Epoch for $\text{Nbits} \in \{4, 3, 2\}$ with model trained using global iterative pruning.

The test accuracy values before and after finetuning are reported in Figure 13. Note that the accuracy before finetuning is reported for the quantised version of the best global iteratively pruned model.

Discuss your observations comparing to (c)'s results:

Before finetuning, the model in (d) performs worse than the models in (c). This makes sense because the model in (d) loses more information than the model in (c) does before any finetuning occurs. The test accuracy values after finetuning for the model in (d) is also lower than the corresponding values for the model in (c). This also makes sense because the model in part (d) started off in a less accurate state compared to the model in part (c) since the model is quantised *and* the model was pruned at the beginning. Although the `finetune(...)` function does not appear to force zero values to remain at zero after each training step (this means that the model will become less sparse over time), the model that we started with in part (d) is imprecise enough that it takes more time to train the model.

These effects are especially strong with `Nbits == 2` in part (d). Although the corresponding model in part (c) was able to achieve an accuracy as high as 0.86 during finetuning, part (d)'s model with `Nbits == 2` only achieves an accuracy of 0.36 after finetuning.

```

nbits_lab3d = [2, 3, 4]
test_acc_lab3d_b4_finetune = {}
test_acc_arrays_finetune = {}
test_acc_lab3d_after_finetune = {}
n_epochs = 20

for iii, nb in enumerate(nbits_lab3d):
    # Define quantized model and load weight
    Nbits = nb #Change this value to finish (d)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
    net = net.to(device)
    net.load_state_dict(torch.load("net_after_global_iterative_prune.pt"))
    test_acc_lab3d_b4_finetune[Nbits] = test(net)

    # Quantized model finetuning
    test_acc_arrays_finetune[Nbits] = np.array(
        finetune(net, epochs=n_epochs, batch_size=256, lr=0.002, reg=1e-4))

    # Do not load the model with best accuracy since we want the last acc value.
    # net.load_state_dict(torch.load("quantized_net_after_finetune.pt"))
    test_acc_lab3d_after_finetune[Nbits] = test(net)

```

The above code snippet shows the finetuning process for the model that came about using global iterative prune. In each iteration of the loop, the net is re-loaded, tested, and then fine-tuned. Different from the previous page (lab 3 part (c)), the net is tested again without loading the net with the best result. This records the accuracy both before and after finetuning.

(e) (4 pts) Symmetric quantization is a commonly used and hardware-friendly quantization approach. In symmetric quantization, the quantization levels are symmetric to zero. Implement symmetric quantization in the STE class and repeat the process in (b). Compare and analyze the performance of symmetric quantization and asymmetric quantization.

Please refer to the bottom half of Figure 10 for details on how I implemented quantisation where `symmetric == True`. I computed the quantised weights using this formula:

$$w_q = w_{\text{mask}} \odot \left[\frac{1}{2^{\text{bit}-1} - 1} \cdot \text{round} \left[(2^{\text{bit}-1} - 1) \cdot w \div \max(\text{abs}(w)) \right] \right] \cdot \max(\text{abs}(w))$$

Figure 14 shows that the accuracy of asymmetric quantisation is higher than that of symmetric quantisation.

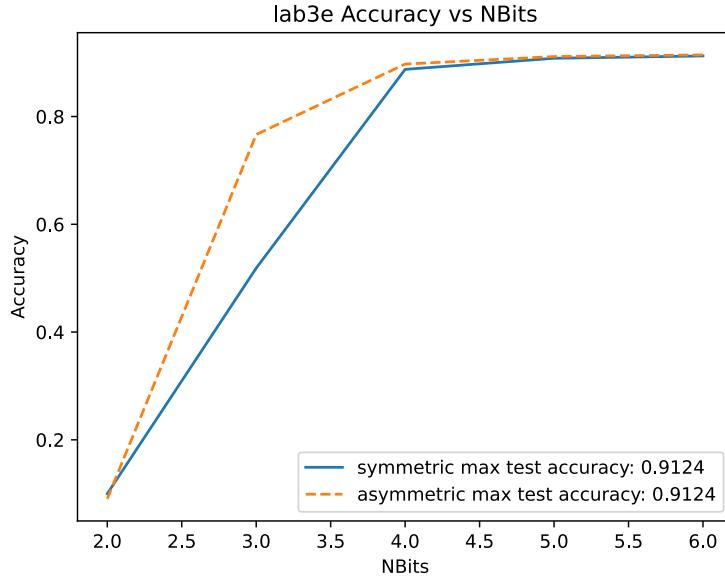


Figure 14: Final test accuracy $\text{Nbits} \in \{4, 3, 2\}$, comparing asymmetric and symmetric pruning.

The range of values in the quantisation for the symmetric case is larger than or equal to that of the asymmetric case since $\alpha = \max |w|$ in the symmetric case ($\alpha = \max w - \min w$ in the asymmetric case). This means that the numerical differences between adjacent quantisation levels are larger. Since symmetric quantisation results in quantisation levels that are spaced further apart, the weight values are mapped less finely when using symmetric quantisation. When a larger-than-necessary range is used in symmetric quantisation, it results in higher loss of information compared to the asymmetric case.

When $\text{Nbits} = 3$, asymmetric quantisation has just enough bits to capture the distribution of high-magnitude weights and predict with passable accuracy, while the symmetric quantisation loses just enough precision to predict less accurately. With $\text{Nbits} = 2$, the number of bits in both the asymmetric and symmetric cases are too low to capture the higher-magnitude weights, while with $\text{Nbits} \in \{4, 5, 6\}$, the number of bits is high enough that the small-ish loss in accuracy (loss in accuracy compares symmetric and asymmetric quantisation) caused by using symmetric quantisation doesn't decrease the accuracy significantly.

tl;dr: The higher information loss in the symmetric case means that the distributions of the high-magnitude weights are captured less well, causing the performance to be lower.

The piece of text below prints the test accuracy for symmetric quantisation.

```
Nbits=2, Test accuracy=0.1000
Nbits=3, Test accuracy=0.5186
Nbits=4, Test accuracy=0.8875
Nbits=5, Test accuracy=0.9081
Nbits=6, Test accuracy=0.9124
```

lab1

November 8, 2023

```
[294]: import sys  
print(sys.executable)
```

/Users/zeyuli/anaconda3/envs/ECE661_HW/bin/python

```
[295]: import numpy as np  
import matplotlib.pyplot as plt  
from typing import Tuple  
  
x1, y1 = np.array([ 1, -2, -1, -1,  1]), 7  
x2, y2 = np.array([ 2, -1,  2,  0, -2]), 1  
x3, y3 = np.array([-1,  0,  2,  2,  1]), 1  
  
x_input = np.array([x1, x2, x3]).reshape(3, 5)  
y_input = np.array([y1, y2, y3])  
  
w0 = np.zeros(5).reshape(-1, 1)  
mu = 0.02  
n_epochs = 200  
print(x_input, y_input)  
  
LEGEND_size = 7  
TITLE_size = 12  
AXLABEL_size = 9  
TICK_size = 8
```

[[1 -2 -1 -1 1]
 [2 -1 2 0 -2]
 [-1 0 2 2 1]] [7 1 1]

```
[296]: def my_relu(a):  
    return np.abs(a) * (a > 0)  
  
def l1_norm_trimmed_update(in_x, in_y, w_prev, l1_reg_param, lr, thresh):  
    """  
    thresh: number of elements to keep unchanged when performing proximal l1  
    """
```

```

    ascend_idx = np.argsort(np.abs(w_prev.flatten()))
    weight_l1_proximal = l1_norm_proximal_update(in_x, in_y, w_prev,
    ↪l1_reg_param, lr,
                                         which_weights=ascend_idx[0:
    ↪w_prev.shape[0] - thresh])
    return weight_l1_proximal

def l1_norm_proximal_update(in_x, in_y, w_prev, l1_reg_param, lr, ↪
    ↪which_weights=None):
    """
        Computes weight using proximal gradient descent (l1-norm). threshold is ↪
    ↪l1_reg_param*lr
    Args:
        in_x: x values (data)
        in_y: y values (target)
        w_prev: previous weight
        l1_reg_param: l1 regularisation (lambda)
        lr: learning rate, must be not None.
        which_weights: indices of which weights for which proximal update is ↪
    ↪applied
    Note: threshold = lr * lambda should be in set{0.004, 0.01, 0.02, 0.04}
    Returns: updated weight
    """
    if which_weights is None:
        which_weights = np.array(range(len(w_prev.flatten())))
        # which_weights = [0,1,2,3,4] if w_prev.shape == (5, 1)
    weight_tilde = no_regularisation_update(in_x, in_y, w_prev, lr=lr) # main ↪
    ↪objective
    weight_ret = np.copy(weight_tilde)

    threshold = lr * l1_reg_param
    assert threshold >= 0

    for iii in which_weights: # only updates the weights specified in ↪
    ↪which_weights
        this_weight = weight_tilde[iii, 0]
        weight_ret[iii, 0] = my_relu(np.abs(this_weight) - threshold) * np.
    ↪sign(this_weight)
    return weight_ret

def l1_norm_update(in_x, in_y, w_prev, l1_reg_param, lr):
    """
    Args:
        in_x: x values (data)

```

```

    in_y: y values (target)
    w_prev: previous weight
    l1_reg_param: l1 regularisation (lambda)
    lr: learning rate, must be not None
    Returns: updated weight
    """
grad = no_regularisation_update(in_x, in_y, w_prev, l1_reg_param, lr=None) ↵
    ↵+ \
        l1_reg_param * np.sign(w_prev)
    assert grad.shape == w_prev.shape
    return w_prev - lr * grad

def no_regularisation_update(in_x, in_y, w_prev, l1_reg_param=None, lr=None):
    """
Args:
    in_x: x values (data)
    in_y: y values (target)
    w_prev: previous weight
    l1_reg_param: dummy
    lr: if None, it means some fn other than run_grad_descent called it; if ↵
    ↵run_grad_descent called it,
        must be (not None)
    Returns: returns updated weight if lr is not None; returns gradient if lr ↵
    ↵is None
    """
n_r, n_c = in_x.shape
grad = 0
for inner in range(n_r):
    grad += 2 * (np.dot(in_x[inner, :], w_prev) - in_y[inner]) * ↵
in_x[inner].reshape(-1, 1)
if lr is not None:
    return w_prev - lr * grad
else:
    return grad

def run_grad_descent(w_init, lr, in_x, in_y, epochs, w0_prune=False, ↵
    ↵thresh=None, l1_reg_param=None,
    grad_fn=no_regularisation_update) -> Tuple[np.array, np. ↵
array, np.array]:
    """
Args:
    w_init: w0 = np.zeros(5)
    lr: mu=0.02
    in_x: x1, x2, x3

```

```

in_y: y1, y2, y3
epochs: 200
w0_prune: prune or not (lab 1 c). should be False in all except for lab_1c
    thresh: threshold for pruning or trimmed l1. should be True only in lab_1c, lab 1 f
    l1_reg_param: lambda, used in update weight
    grad_fn: function for updating weight
    Returns: optimal weight (np.array), loss for each epoch (np.array), and w_arr matrix to
    record the progression of each weight through time
    """
w_prev = w_init
l_arr = []
w_arr = np.zeros((len(w_init), epochs))
n_r, n_c = in_x.shape
for i in range(epochs):
    if not w0_prune and thresh is not None: # only for lab 1 f, l1-trimmed
        w_curr = grad_fn(in_x, in_y, w_prev, l1_reg_param, lr=lr, thresh=thresh)
    else:
        w_curr = grad_fn(in_x, in_y, w_prev, l1_reg_param, lr=lr)

    if w0_prune and thresh is not None: # only for lab 1 c, pruning
        w_curr_sorted_idx = np.argsort(np.abs(w_curr.flatten())) # sort ascending
        for inner in range(len(w_init) - thresh):
            # len(w_init) - thresh = 5 - 2 = 3, so three smallest weights (absolute value) are
            # forced to zero
            w_curr[w_curr_sorted_idx[inner], 0] = 0 # zero-ing out smallest elements
    l_curr = 0
    for inner in range(n_r):
        l_curr += (np.dot(in_x[inner, :], w_curr) - in_y[inner]) ** 2
    l_arr.append(l_curr)
    w_arr[:, i] = w_curr.flatten()
    w_prev = w_curr
return w_prev, np.array(l_arr), w_arr

def plot_log_loss_vs_n_iters(num_epochs, loss_array, save_name_, save=False,
                             plt_axis=None):
    if plt_axis is None:
        fig, ax = plt.subplots(1, 1)
    else:

```

```

        fig, ax = None, plt_axis
        ax.semilogy(range(num_epochs), loss_array, label='Lowest loss=%g' % np.
        ↪min(loss_array))
        ax.set_xlabel('Number of Iterations', fontsize=AXLABEL_size)
        ax.set_ylabel('L (log10 scale)', fontsize=AXLABEL_size)
        ax.set_title('Loss vs Number of Iterations', fontsize=TITLE_size)
        ax.tick_params(axis='both', labelsize=TICK_size)
        ax.legend(loc='best', fontsize=LEGEND_size)
        if fig is not None:
            fig.tight_layout()
        if save:
            plt.savefig('%s.pdf' % save_name_, dpi=700, bbox_inches='tight')

def plot_weight_val_vs_n_iters(num_epochs, w_init, final_weights, weight_array, ↪
    ↪save_name_, save=False,
                           plt_axis=None):
    if plt_axis is None:
        fig, ax = plt.subplots(1, 1)
    else:
        fig, ax = None, plt_axis
    line_styles = [ '-', ':', '--', '-.', (0, (1, 1))]
    for iii in range(len(w_init)):
        ax.plot(range(num_epochs), weight_array[iii, :], ↪
    ↪linestyle=line_styles[iii],
                linewidth=1., label='final value %.5f' % final_weights[iii])
    ax.legend(loc='best', fontsize=LEGEND_size)
    ax.tick_params(axis='both', labelsize=TICK_size)
    ax.set_xlabel('Number of Iterations', fontsize=AXLABEL_size)
    ax.set_ylabel('Value of Weights', fontsize=AXLABEL_size)
    ax.set_title('Weight Values vs Number of Iterations', fontsize=TITLE_size)
    if fig is not None:
        fig.tight_layout()
    if save:
        plt.savefig('%s.pdf' % save_name_, dpi=700, bbox_inches='tight')

def plot_many_w_vals_or_loss_vs_n_iters(num_epochs, w_init, l1_reg_params, ↪
    ↪save_name_, save=False,
                           loss_arr_lst=None, final_weights_list=None, ↪
    ↪weight_array_list=None):
    """
    Args:
        num_epochs: number of iterations (200)
        w_init: initial weights
        final_weights_list: list of weights achieved by gradient descent by
    ↪default None

```

```

    weight_array_list: list of np array that describe weight progression by_
    ↵default None
    loss_arr_lst: list of loss arrays (len(loss_arr_lst) corresponds to_
    ↵len(l1_reg_params))
        by default, loss_arr_lst=None. If not None,
        l1_reg_params: reg params (lambda)
        save_name_: name of file
        save: save or not
    Either (final_weights_list and weight_array_list are not None) or_
    ↵(loss_arr_lst is not None)
    Returns: None
    """
    assert (final_weights_list is not None and weight_array_list is not None) or \
    ↵or \
        loss_arr_lst is not None, \
        "Either (final_weights_list and weight_array_list are not None) or_
    ↵(loss_arr_lst is not None)"
    assert not (final_weights_list is not None and weight_array_list is not_
    ↵None and
        loss_arr_lst is not None), "all three are not None."
    fig, ax = plt.subplots(2, 2, figsize=(10, 8))
    ax_flat = ax.flatten()
    for iii, axis in enumerate(ax_flat):
        if loss_arr_lst is None:
            plot_weight_val_vs_n_iters(num_epochs, w_init,_
            ↵final_weights=final_weights_list[iii],_
                weight_array=weight_array_list[iii],_
            ↵save_name_=None, save=False,
                plt_axis=axis)
        else:
            plot_log_loss_vs_n_iters(num_epochs, loss_array=loss_arr_lst[iii],_
            ↵save_name_=None,
                save=False, plt_axis=axis)
            ax_flat[iii].set_title("lambda=%g" % l1_reg_params[iii])
        if fig is not None:
            fig.tight_layout()
        if save_name_ is not None and save:
            plt.savefig('%s.pdf' % save_name_, dpi=700, bbox_inches='tight')
    return

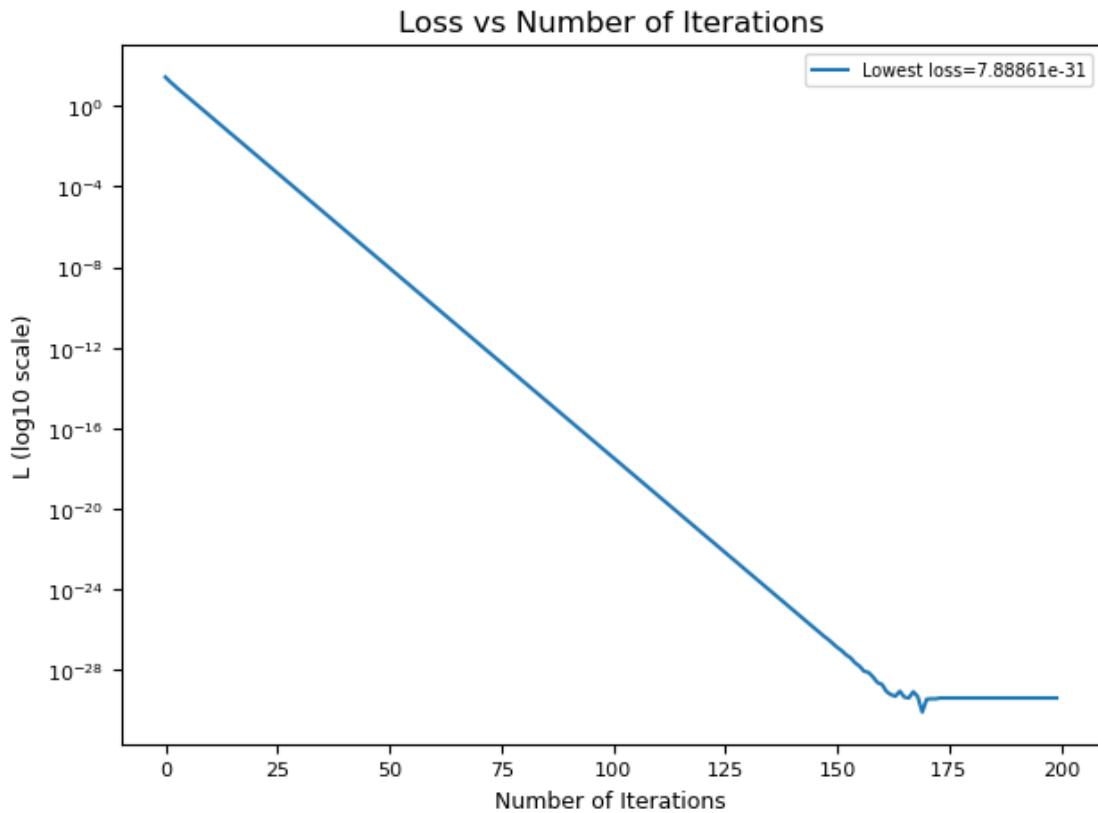
```

[297]: """ part b (3 pts) In Python, directly minimize the objective L without any sparsity-inducing regularisation/constraint. Plot the value of $\log(L)$ vs. #steps throughout the training, and use another figure to plot how the value of each element in W is changing throughout the training. From your result, is W converging to an optimal solution? Is W converging to a sparse solution? """

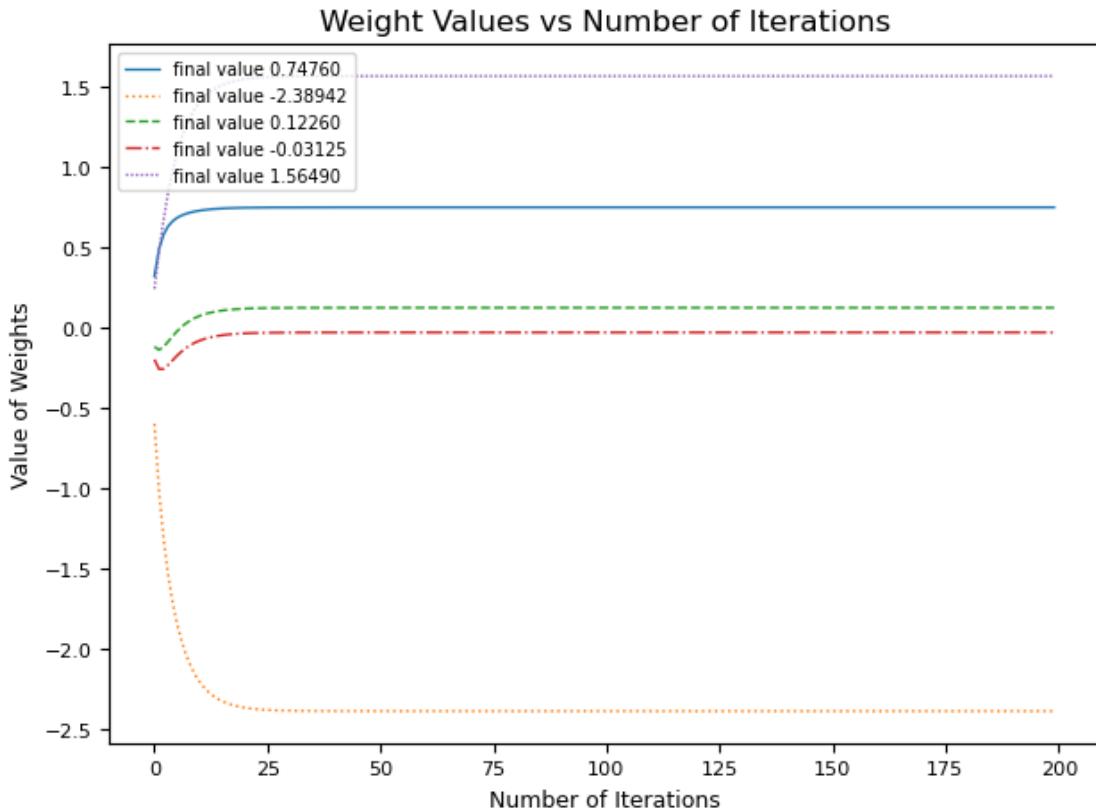
```
weight, loss_arr, weight_arr = run_grad_descent(w0, mu, x_input, y_input, epochs=n_epochs)
print(weight.shape, loss_arr.shape, weight_arr.shape)
```

(5, 1) (200, 1) (5, 200)

```
[298]: # part b
plot_log_loss_vs_n_iters(num_epochs=n_epochs, loss_array=loss_arr, save_name_='lab1b_loss', save=True)
```

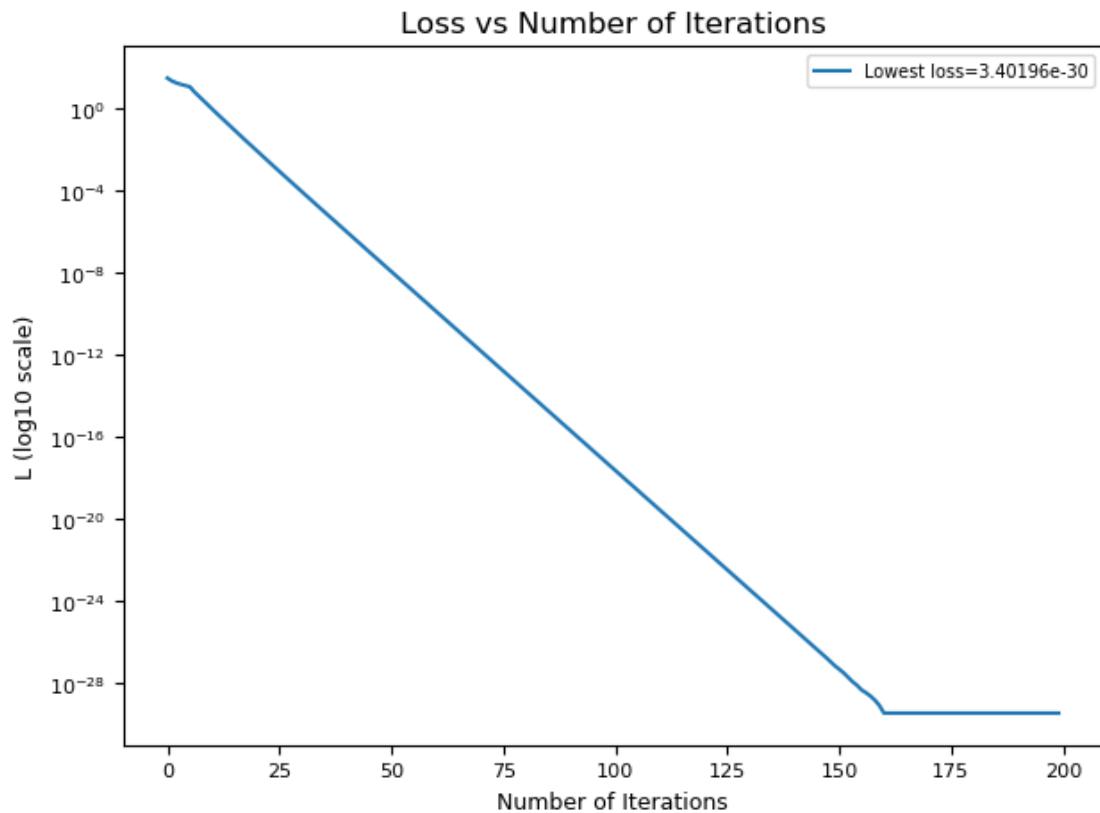


```
[299]: # part b
plot_weight_val_vs_n_iters(num_epochs=n_epochs, w_init=w0, final_weights=weight,
                           weight_array=weight_arr, save_name_='lab1b_weights', save=True)
```

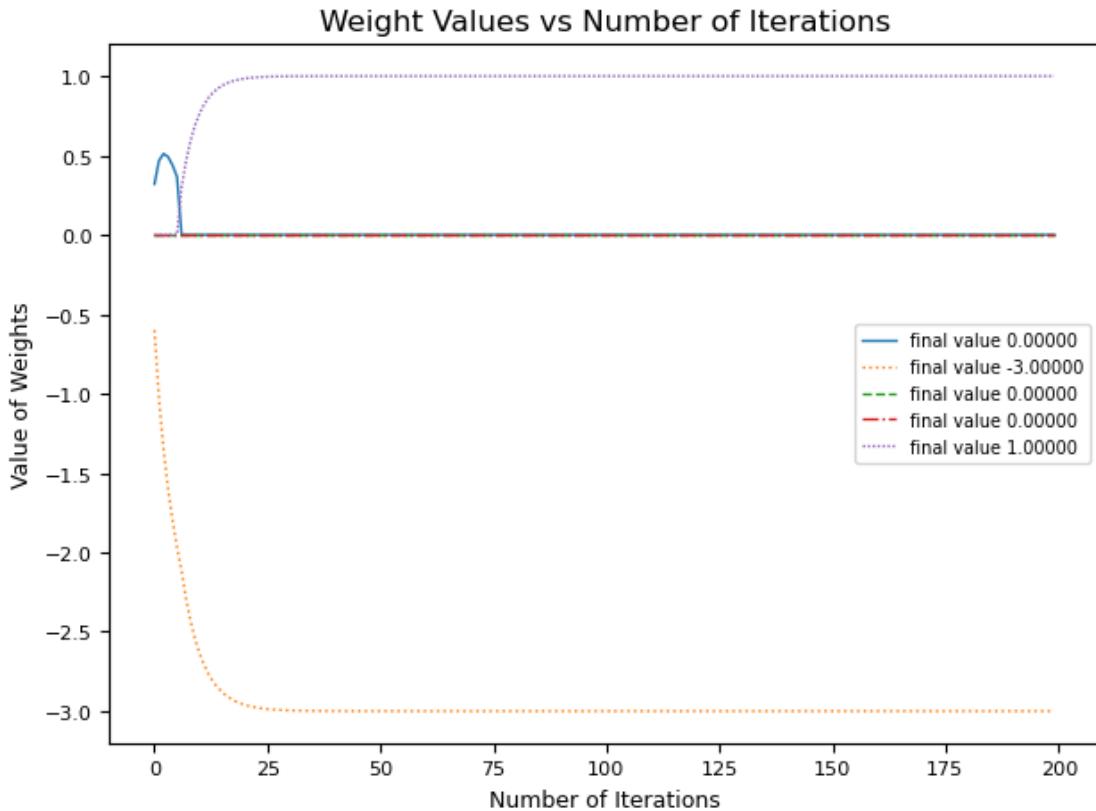


[300]:
 """part (c) (6 pts) Since we have the knowledge that the ground-truth weight should have $\|W\|_0 \leq 2$, we can apply projected gradient descent to enforce this sparse constraint. Redo the optimization process in (b), this time prune the elements in W after every gradient descent step to ensure $\|W\|_0 \leq 2$. Plot the value of $\log(L)$ throughout the training, and use another figure to plot the value of each element in W in each step. From your result, is W converging to an optimal solution? Is W converging to a sparse solution?"""
 weight_pruned, loss_arr_pruned, weight_arr_pruned = \
 run_grad_descent(w0, mu, x_input, y_input, n_epochs, w0_prune=True, \
 thresh=2)

[301]:
 plot_log_loss_vs_n_iters(num_epochs=n_epochs, loss_array=loss_arr_pruned, \
 save_name_='lab1c_loss',
 save=True)



```
[302]: # part c
plot_weight_val_vs_n_iters(num_epochs=n_epochs, w_init=w0,
                            final_weights=weight_pruned,
                            weight_array=weight_arr_pruned,
                            save_name_='lab1c_weights', save=True)
```



[303]: *"""part (d) (5 pts) In this problem we apply l1 regularization to induce the sparse solution. The minimization objective therefore changes to $L + \lambda ||W||_1$. Please use full-batch gradient descent to minimize this objective, with $\lambda = \{0.2, 0.5, 1.0, 2.0\}$ respectively. For each case, plot the value of $\log(L)$ throughout the training, and use another figure to plot the value of each element in W in each step. From your result, comment on the convergence performance under different λ ."""*

```
lambda_list = [0.2, 0.5, 1.0, 2.0]
weight_l1_list, loss_l1_list, weight_arr_l1_list = [], [], []
for idx, l11 in enumerate(lambda_list):
    weight_l11, loss_arr_l11, weight_arr_l11 = run_grad_descent(w0, mu, x_input,
                                                               y_input, n_epochs,
                                                               w0_prune=False,
                                                               thresh=None,
                                                               l1_reg_param=lambda_list[idx],
                                                               grad_fn=l1_norm_update)
    weight_l1_list.append(weight_l11)
```

```

loss_l1_list.append(loss_arr_l1)
weight_arr_l1_list.append(weight_arr_l1)

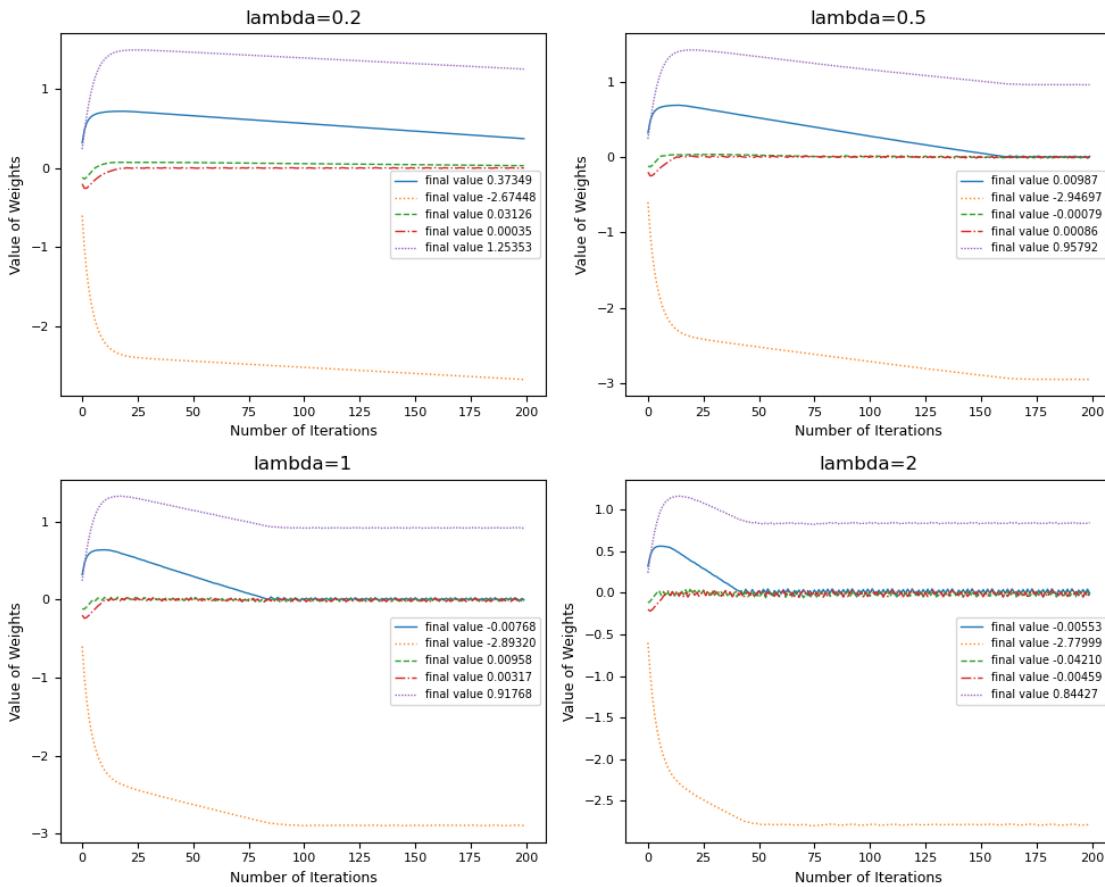
```

[304]:

```

plot_many_w_vals_or_loss_vs_n_iters(num_epochs=n_epochs, w_init=w0,
                                    final_weights_list=weight_l1_list,
                                    weight_array_list=weight_arr_l1_list,
                                    loss_arr_lst=None,
                                    l1_reg_params=lambda_list,
                                    save_name_='lab1d_weight', save=True)

```

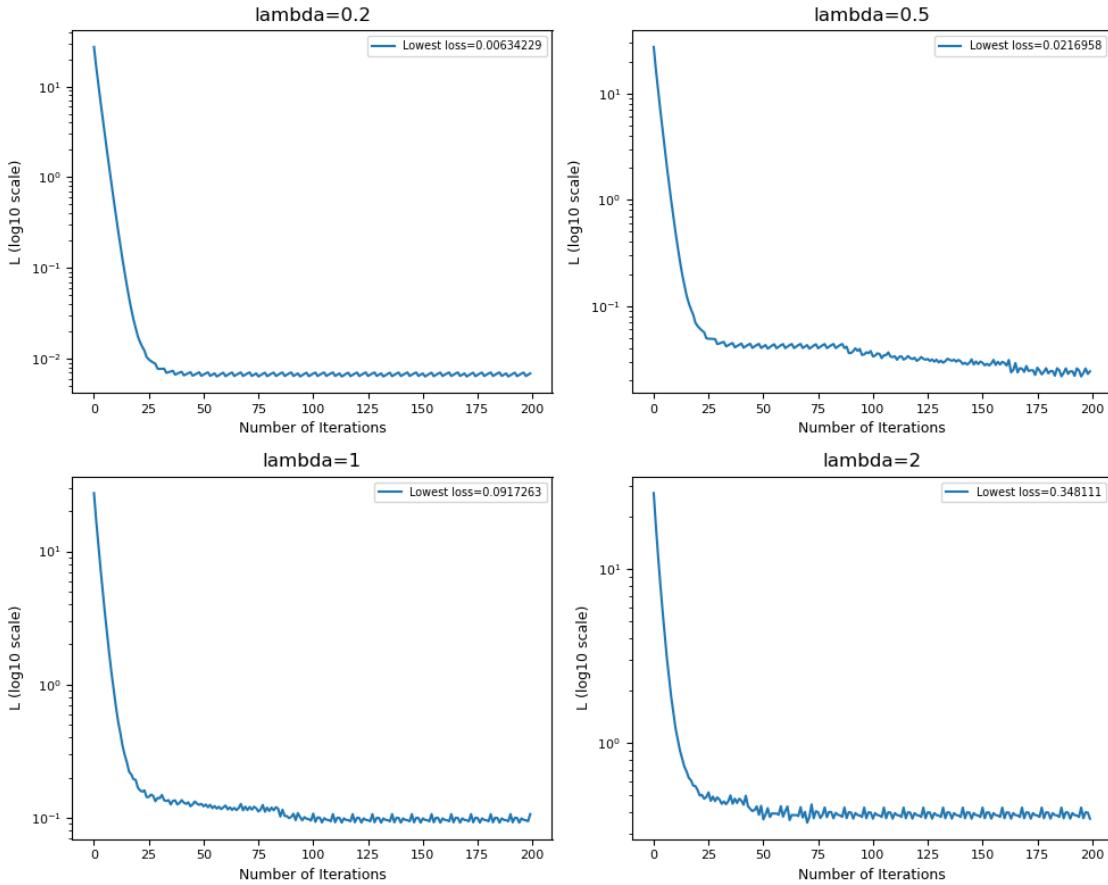


[305]:

```

plot_many_w_vals_or_loss_vs_n_iters(num_epochs=n_epochs, w_init=w0,
                                    final_weights_list=None,
                                    weight_array_list=None,
                                    loss_arr_lst=loss_l1_list,
                                    l1_reg_params=lambda_list,
                                    save_name_='lab1d_loss', save=True)

```



[306] :

"""
 part (e) (6pts) Here we optimize the same objective as in (d), this time using
 ↪ proximal gradient update. Recall that the proximal operator of the l_1 ↪
 ↪ regulariser is the soft thresholding function. Set the threshold in the soft ↪
 ↪ thresholding function to {0.004, 0.01, 0.02, 0.04} respectively. Plot the ↪
 ↪ value of $\log(L)$ throughout the training, and use another figure to plot the ↪
 ↪ value of each element in W in each step. Compare the convergence performance ↪
 ↪ with the results in (d). (Hint: Optimizing $L + \lambda \|\mathbf{w}\|_1$ using gradient ↪
 ↪ descent with learning rate μ should correspond to proximal gradient update ↪
 ↪ with threshold $\mu * \lambda$) """

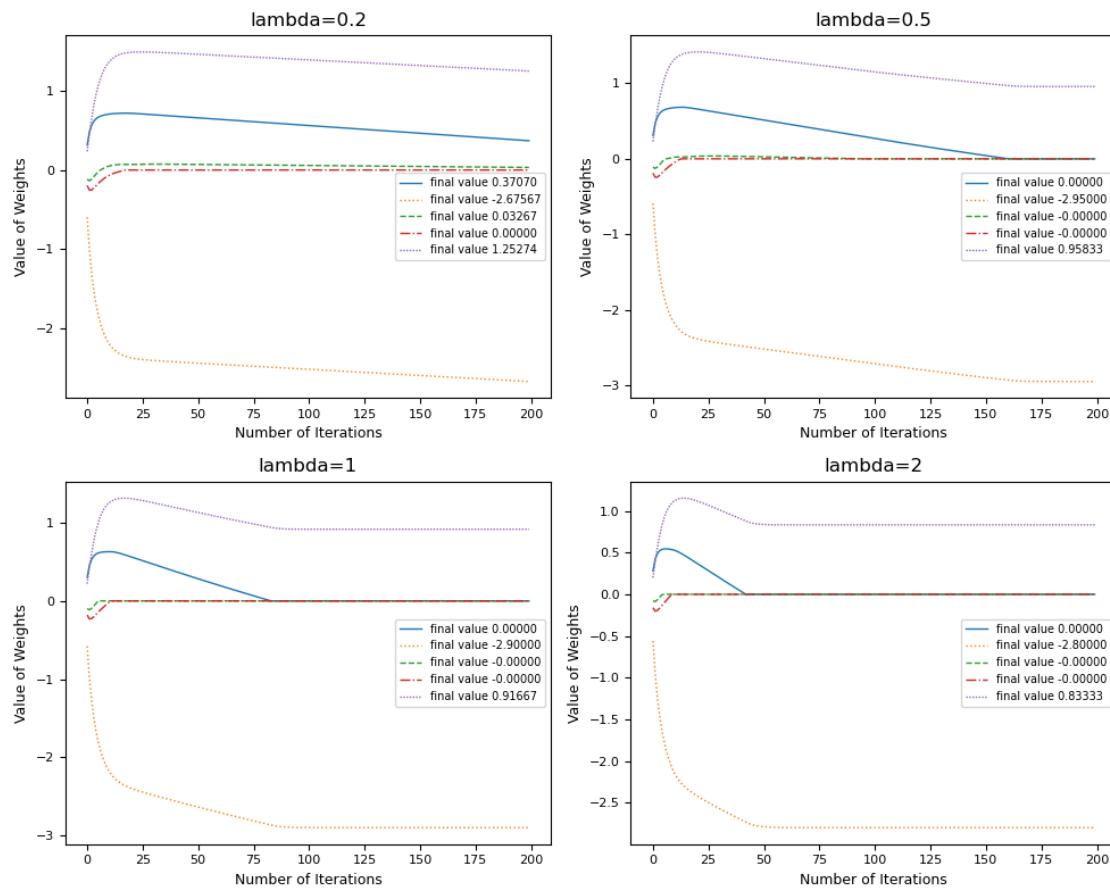
```
lambda_list = [0.2, 0.5, 1.0, 2.0]
weight_prox_list, loss_prox_list, weight_arr_prox_list = [], [], []
for idx, lll in enumerate(lambda_list):
    weight_prox, loss_arr_prox, weight_arr_prox = run_grad_descent(w0, mu,
        ↪ x_input, y_input, n_epochs,
        ↪ w0_prune=False, thresh=None,
```

```

↳ l1_reg_param=lll,
↳ grad_fn=l1_norm_proximal_update)
weight_prox_list.append(weight_prox)
loss_prox_list.append(loss_arr_prox)
weight_arr_prox_list.append(weight_arr_prox)

```

[307]: `plot_many_w_vals_or_loss_vs_n_iters(num_epochs=n_epochs, w_init=w0, final_weights_list=weight_prox_list, weight_array_list=weight_arr_prox_list, loss_arr_lst=None, l1_reg_params=lambda_list, save_name_='label_weight', save=True)`

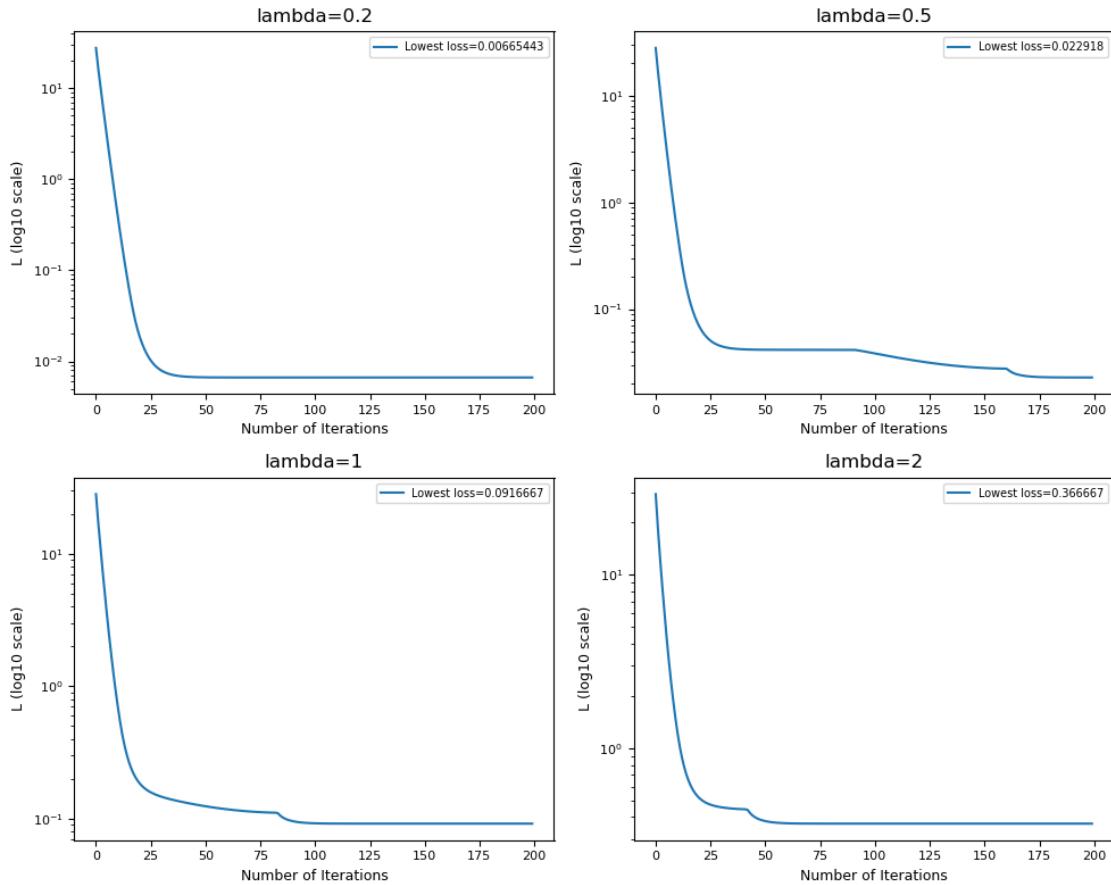


[308]: `plot_many_w_vals_or_loss_vs_n_iters(num_epochs=n_epochs, w_init=w0, final_weights_list=None, weight_array_list=None, loss_arr_lst=loss_prox_list,`

```

l1_reg_params=lambda_list,
↪save_name_='label1e_loss', save=True)

```



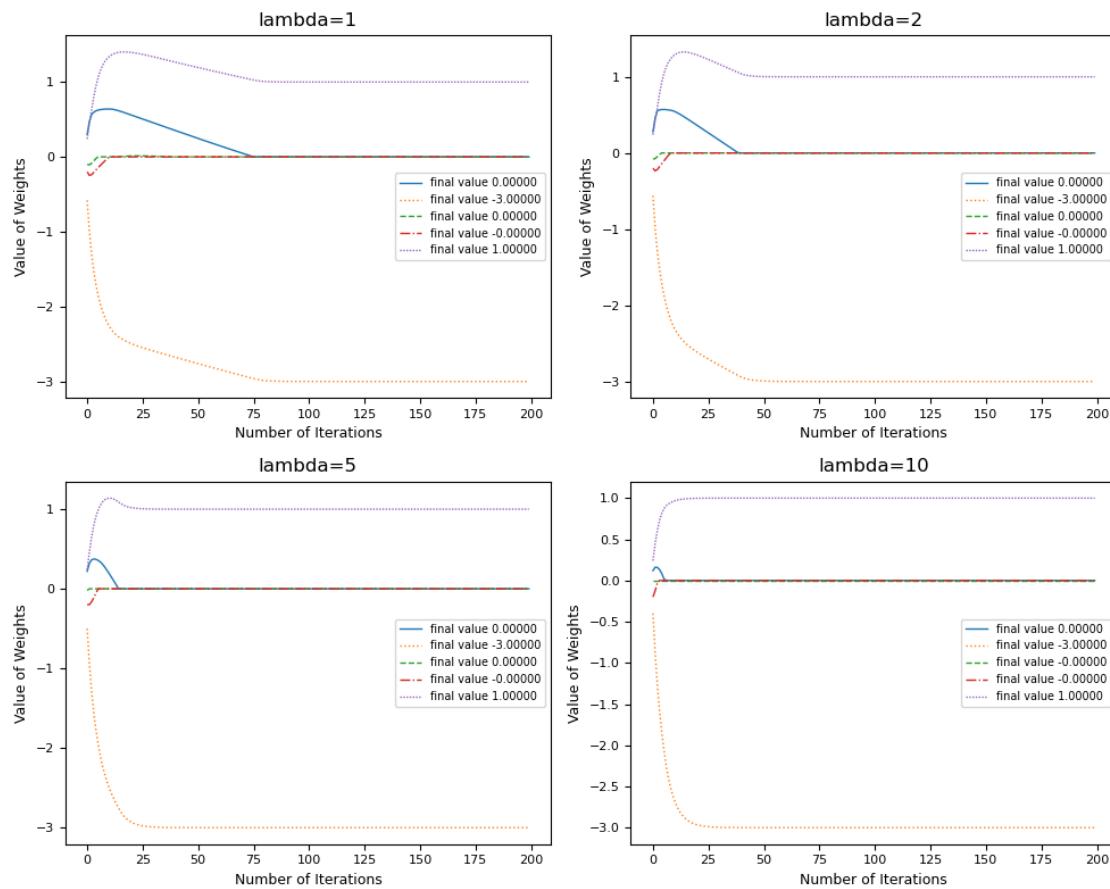
[309]: *"""(f) (6 pts) Trimmed l1 (Tl1) regulariser is proposed to solve the "bias" problem of l1. For simplicity you may implement the T l1 regulariser as applying a l1 regularization with strength on the 3 elements of W with the smallest absolute value, with no penalty on other elements. Minimize $L + T_{l1}(W)$ using proximal gradient update with $\gamma = \{1.0, 2.0, 5.0, 10.0\}$ (correspond to the soft thresholding threshold $\{0.02, 0.04, 0.1, 0.2\}$). Plot the value of $\log(L)$ throughout the training, and use another figure to plot the value of each element in W in each step. Comment on the convergence comparison of the Trimmed l1 and the l1. Also compare the behavior of the early steps (e.g. first 20) between the Trimmed l1 and the iterative pruning.*

```
lambda_list_f = np.array([1.0, 2.0, 5.0, 10.0]) * 1
```

```
weight_trim_list, loss_trim_list, weight_arr_trim_list = [], [], []
for idx, lll in enumerate(lambda_list_f):
```

```
w, l, w_array = run_grad_descent(w0, mu, x_input, y_input, n_epochs,
    ↪w0_prune=False, thresh=2,
                                l1_reg_param=111, □
    ↪grad_fn=l1_norm_trimmed_update)
    weight_trim_list.append(w)
    loss_trim_list.append(l)
    weight_arr_trim_list.append(w_array)
```

[310]: `plot_many_w_vals_or_loss_vs_n_iters(num_epochs=n_epochs, w_init=w0, □
 ↪final_weights_list=weight_trim_list,
 weight_array_list=weight_arr_trim_list, □
 ↪loss_arr_lst=None,
 l1_reg_params=lambda_list_f, □
 ↪save_name_='lab1f_weight', save=True)`

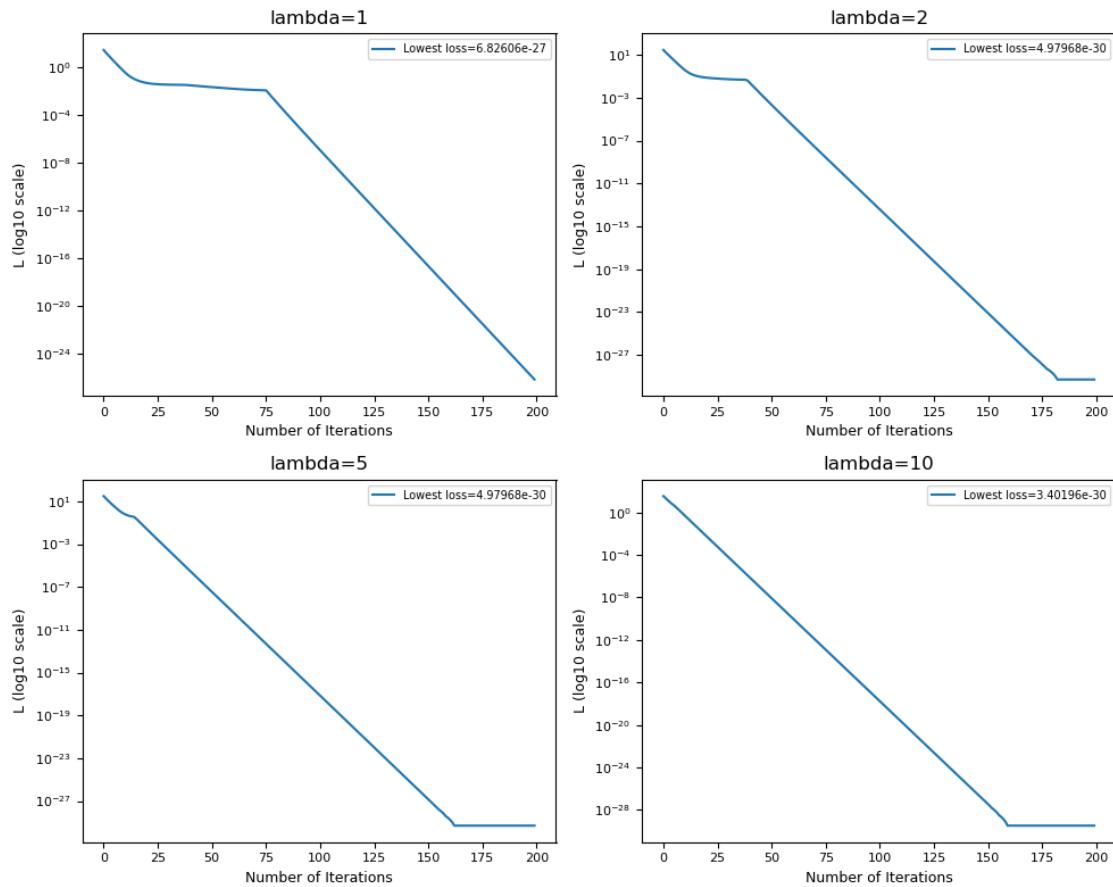


[311]: `plot_many_w_vals_or_loss_vs_n_iters(num_epochs=n_epochs, w_init=w0, □
 ↪final_weights_list=None,
 weight_array_list=None, □
 ↪loss_arr_lst=loss_trim_list,`

```

    l1_reg_params=lambda_list_f, u
    ↪save_name_='lab1f_loss', save=True)

```



hw4

November 8, 2023

0.0.1 Lab2 (a) Model preparation

```
[21]: from resnet20 import ResNetCIFAR
from train_util import train, finetune, test
import torch
import numpy as np
import matplotlib.pyplot as plt

import time

import torchvision.transforms as transforms
import torchvision
import torch.nn as nn
import torch.optim as optim

from FP_layers import *

device = 'cuda' if torch.cuda.is_available() else 'cpu'
device
```

[21]: 'cuda'

```
[22]: net = ResNetCIFAR(num_layers=20, Nbits=None)
net = net.to(device)
```

```
[3]: # Load the best weight paramters
net.load_state_dict(torch.load("pretrained_model.pt"))
test(net)
```

Files already downloaded and verified
Test Loss=0.3231, Test accuracy=0.9150

[3]: 0.915

0.0.2 Lab2 (b) Prune by percentage

```
[4]: def prune_by_percentage(layer, q):
    """
    Pruning the weight parameters by threshold.
    :param q: pruning percentile. 'q' percent of the least
    significant weight parameters will be pruned.
    """
    with torch.no_grad():
        # Convert the weight of "layer" to numpy array
        layer_weight = layer.weight.detach().cpu().numpy()
        # Compute the q-th percentile of the abs of the converted array
        percentile = np.percentile(np.abs(layer_weight.flatten()), q)
        # Generate a binary mask same shape as weight to decide which element
        # to prune
        masked_obj = np.ma.masked_greater_equal(x=np.abs(layer_weight), m
        value=percentile, copy=True)
        mask_int = np.ma.getmask(masked_obj).astype(int)
        # Convert mask to torch tensor and put on GPU
        mask_tensor = torch.tensor(mask_int).to(device)
        # Multiply the weight by mask to perform pruning
        assert mask_int.shape == layer_weight.shape
        # layer.weight.data = mask_tensor * layer.weight.data
        layer.weight.data = layer.weight.data.clone().detach() *
        requires_grad_(True) * mask_tensor
    return
```

```
[5]: q_list = np.array([0.3, 0.5, 0.7]) * 100
test_acc_prune_by_perc = []

for q_val in q_list:
    net.load_state_dict(torch.load("pretrained_model.pt"))
    for name, layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
        'id_mapping' not in name:
            # change q value
            prune_by_percentage(layer, q=q_val)
            # break
            ## Optional: Check the sparsity you achieve in each layer
            ## Convert the weight of "layer" to numpy array
            np_weight = layer.weight.detach().cpu().numpy()
            ## Count number of zeros
            zeros = sum((np_weight == 0).flatten())
            ## Count number of parameters
            total = len(np_weight.flatten())
    test_acc_prune_by_perc.append(test_accuracy(net, test_loader))
```

```

## Print sparsity
# print('Sparsity of ' + name + ': ' + str(zeros/total))
print('Sparsity of %s: %g' % (name, zeros/total))
test_acc_prune_by_perc.append(test(net))

```

```

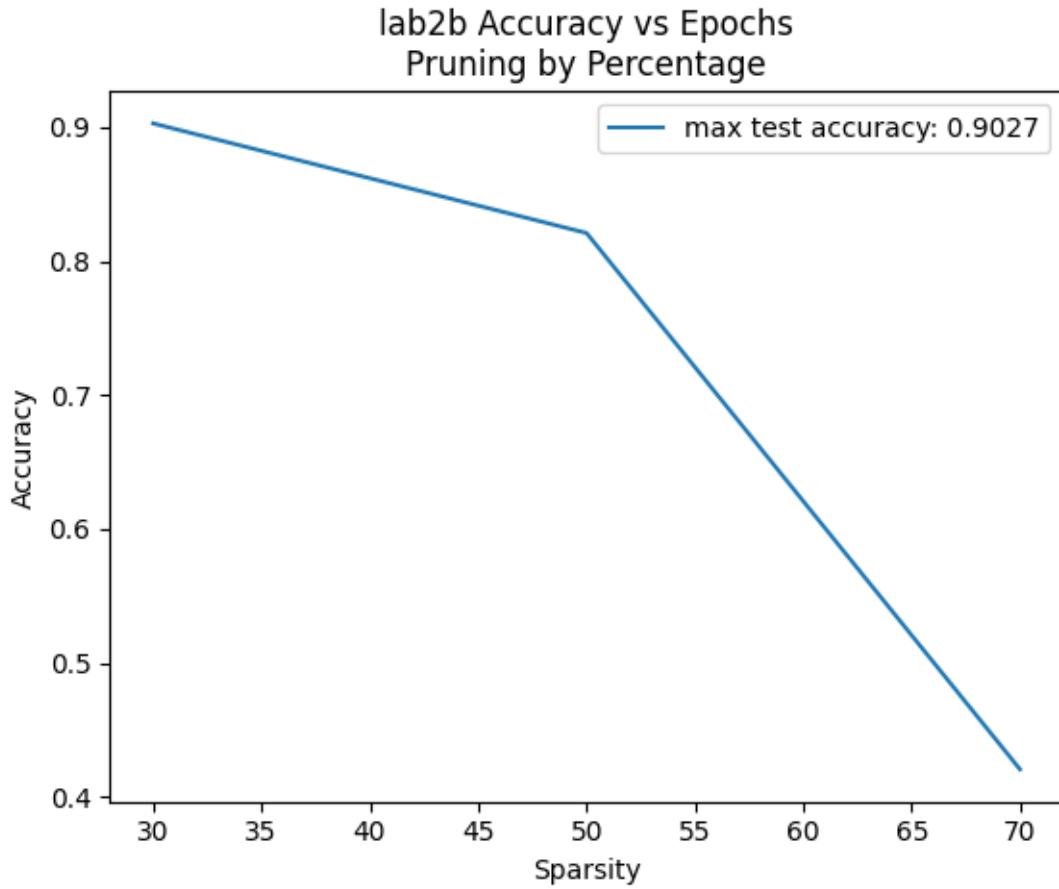
Sparsity of head_conv.0.conv: 0.300926
Sparsity of body_op.0.conv1.0.conv: 0.299913
Sparsity of body_op.0.conv2.0.conv: 0.299913
Sparsity of body_op.1.conv1.0.conv: 0.299913
Sparsity of body_op.1.conv2.0.conv: 0.299913
Sparsity of body_op.2.conv1.0.conv: 0.299913
Sparsity of body_op.2.conv2.0.conv: 0.299913
Sparsity of body_op.3.conv1.0.conv: 0.30013
Sparsity of body_op.3.conv2.0.conv: 0.300022
Sparsity of body_op.4.conv1.0.conv: 0.300022
Sparsity of body_op.4.conv2.0.conv: 0.300022
Sparsity of body_op.5.conv1.0.conv: 0.300022
Sparsity of body_op.5.conv2.0.conv: 0.300022
Sparsity of body_op.6.conv1.0.conv: 0.300022
Sparsity of body_op.6.conv2.0.conv: 0.299995
Sparsity of body_op.7.conv1.0.conv: 0.299995
Sparsity of body_op.7.conv2.0.conv: 0.299995
Sparsity of body_op.8.conv1.0.conv: 0.299995
Sparsity of body_op.8.conv2.0.conv: 0.299995
Sparsity of final_fc.linear: 0.3
Files already downloaded and verified
Test Loss=0.3699, Test accuracy=0.9027
Sparsity of head_conv.0.conv: 0.5
Sparsity of body_op.0.conv1.0.conv: 0.5
Sparsity of body_op.0.conv2.0.conv: 0.5
Sparsity of body_op.1.conv1.0.conv: 0.5
Sparsity of body_op.1.conv2.0.conv: 0.5
Sparsity of body_op.2.conv1.0.conv: 0.5
Sparsity of body_op.2.conv2.0.conv: 0.5
Sparsity of body_op.3.conv1.0.conv: 0.5
Sparsity of body_op.3.conv2.0.conv: 0.5
Sparsity of body_op.4.conv1.0.conv: 0.5
Sparsity of body_op.4.conv2.0.conv: 0.5
Sparsity of body_op.5.conv1.0.conv: 0.5
Sparsity of body_op.5.conv2.0.conv: 0.5
Sparsity of body_op.6.conv1.0.conv: 0.5
Sparsity of body_op.6.conv2.0.conv: 0.5
Sparsity of body_op.7.conv1.0.conv: 0.5
Sparsity of body_op.7.conv2.0.conv: 0.5
Sparsity of body_op.8.conv1.0.conv: 0.5
Sparsity of body_op.8.conv2.0.conv: 0.5
Sparsity of final_fc.linear: 0.5

```

```
Files already downloaded and verified
Test Loss=0.6776, Test accuracy=0.8209
Sparsity of head_conv.0.conv: 0.699074
Sparsity of body_op.0.conv1.0.conv: 0.700087
Sparsity of body_op.0.conv2.0.conv: 0.700087
Sparsity of body_op.1.conv1.0.conv: 0.700087
Sparsity of body_op.1.conv2.0.conv: 0.700087
Sparsity of body_op.2.conv1.0.conv: 0.700087
Sparsity of body_op.2.conv2.0.conv: 0.700087
Sparsity of body_op.3.conv1.0.conv: 0.69987
Sparsity of body_op.3.conv2.0.conv: 0.699978
Sparsity of body_op.4.conv1.0.conv: 0.699978
Sparsity of body_op.4.conv2.0.conv: 0.699978
Sparsity of body_op.5.conv1.0.conv: 0.699978
Sparsity of body_op.5.conv2.0.conv: 0.699978
Sparsity of body_op.6.conv1.0.conv: 0.699978
Sparsity of body_op.6.conv2.0.conv: 0.700005
Sparsity of body_op.7.conv1.0.conv: 0.700005
Sparsity of body_op.7.conv2.0.conv: 0.700005
Sparsity of body_op.8.conv1.0.conv: 0.700005
Sparsity of body_op.8.conv2.0.conv: 0.700005
Sparsity of final_fc.linear: 0.7
Files already downloaded and verified
Test Loss=2.4418, Test accuracy=0.4207
```

```
[6]: fig, ax = plt.subplots(1, 1)
ax.plot(q_list, test_acc_prune_by_perc, label='max test accuracy: %g' % np.
        ↪max(test_acc_prune_by_perc))
ax.set_xlabel('Sparsity')
ax.set_ylabel('Accuracy')
ax.set_title('lab2b Accuracy vs Epochs\nPruning by Percentage')
ax.legend()

plt.savefig('lab2b.pdf', dpi=500, bbox_inches='tight')
```



0.0.3 Lab2 (c) Finetune pruned model

```
[7]: def finetune_after_prune(net, trainloader, criterion, optimizer, prune=True):
    """
    Finetune the pruned model for a single epoch
    Make sure pruned weights are kept as zero
    """
    # Build a dictionary for the nonzero weights
    weight_mask = {}
    for name, layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
           'id_mapping' not in name:
            # Your code here: generate a mask in GPU torch tensor to have 1 for
            # nonzero element and 0 for zero element
            layer_weight = layer.weight.detach().cpu().numpy()
            mask_obj = np.ma.masked_equal(x=layer_weight, value=0) # where
            # it's zero, mask value is True
```

```

        weight_mask[name] = torch.tensor(1 - np.ma.getmask(mask_obj)).
˓→astype(int)).to(device)
        # zero = 0, nonzero = 1

global_steps = 0
train_loss = 0
correct = 0
total = 0
start = time.time()
for batch_idx, (inputs, targets) in enumerate(trainloader):
    inputs, targets = inputs.to(device), targets.to(device)
    optimizer.zero_grad()
    outputs = net(inputs)
    loss = criterion(outputs, targets)
    loss.backward()
    optimizer.step()

    if prune:
        for name, layer in net.named_modules():
            if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.
˓→Linear)) and 'id_mapping' not in name:
                # Your code here: Use weight_mask to make sure zero
˓→elements remains zero
                layer.weight.data = layer.weight.data.clone().detach().
˓→requires_grad_(True) * weight_mask[name]

    train_loss += loss.item()
    _, predicted = outputs.max(1)
    total += targets.size(0)
    correct += predicted.eq(targets).sum().item()
    global_steps += 1

    if global_steps % 50 == 0:
        end = time.time()
        batch_size = 256
        num_examples_per_second = 50 * batch_size / (end - start)
        print("[Step=%d]\tLoss=%.4f\tacc=%.4f\t%.1f examples/second"
              % (global_steps, train_loss / (batch_idx + 1), (correct /
˓→total), num_examples_per_second))
        start = time.time()

```

[8]: # Get pruned model

```

net.load_state_dict(torch.load("pretrained_model.pt"))
for name, layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
˓→'id_mapping' not in name:

```

```

        prune_by_percentage(layer, q=70.0)
lab2c_test_b4_finetune = test(net)
# Training setup, do not change
batch_size = 256
lr = 0.002
reg = 1e-4

print('==> Preparing data..')
transform_train = transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])

transform_test = transforms.Compose([
    transforms.ToTensor(),
    transforms.Normalize((0.4914, 0.4822, 0.4465), (0.2023, 0.1994, 0.2010)),
])
best_acc = 0 # best test accuracy
start_epoch = 0 # start from epoch 0 or last checkpoint epoch
trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform_train)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
                                         shuffle=True, num_workers=16)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform_test)
testloader = torch.utils.data.DataLoader(testset, batch_size=100,
                                         shuffle=False, num_workers=2)

criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=lr, momentum=0.875,
                      weight_decay=reg, nesterov=False)

```

Files already downloaded and verified
Test Loss=2.4418, Test accuracy=0.4207
==> Preparing data..
Files already downloaded and verified
Files already downloaded and verified

[9]: # Model finetuning

```

test_acc_lst = []
n_epochs = 20
for epoch in range(n_epochs):
    print('\nEpoch: %d' % epoch)
    net.train()

```

```

finetune_after_prune(net, trainloader, criterion, optimizer)
#Start the testing code.
net.eval()
test_loss = 0
correct = 0
total = 0
with torch.no_grad():
    for batch_idx, (inputs, targets) in enumerate(testloader):
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = net(inputs)
        loss = criterion(outputs, targets)

        test_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()
num_val_steps = len(testloader)
val_acc = correct / total
print("Test Loss=% .4f, Test acc=% .4f" % (test_loss / (num_val_steps), val_acc))
    test_acc_lst.append(val_acc)
    if val_acc > best_acc:
        best_acc = val_acc
        print("Saving...")
        torch.save(net.state_dict(), "net_after_finetune.pt")

```

Epoch: 0

[Step=50]	Loss=0.4015	acc=0.8612	6063.1 examples/second
[Step=100]	Loss=0.3621	acc=0.8752	12533.4 examples/second
[Step=150]	Loss=0.3334	acc=0.8853	12582.2 examples/second

Test Loss=0.4290, Test acc=0.8676

Saving...

Epoch: 1

[Step=50]	Loss=0.2458	acc=0.9155	6556.5 examples/second
[Step=100]	Loss=0.2449	acc=0.9152	12105.4 examples/second
[Step=150]	Loss=0.2388	acc=0.9171	12214.2 examples/second

Test Loss=0.3941, Test acc=0.8752

Saving...

Epoch: 2

[Step=50]	Loss=0.2220	acc=0.9273	6774.9 examples/second
[Step=100]	Loss=0.2218	acc=0.9254	12616.7 examples/second
[Step=150]	Loss=0.2150	acc=0.9276	12099.2 examples/second

Test Loss=0.3787, Test acc=0.8789

Saving...

Epoch: 3
[Step=50] Loss=0.1934 acc=0.9318 6663.5 examples/second
[Step=100] Loss=0.2036 acc=0.9284 12776.3 examples/second
[Step=150] Loss=0.2007 acc=0.9292 12756.8 examples/second

Test Loss=0.3710, Test acc=0.8813

Saving...

Epoch: 4
[Step=50] Loss=0.1908 acc=0.9345 6193.6 examples/second
[Step=100] Loss=0.1856 acc=0.9364 13214.0 examples/second
[Step=150] Loss=0.1860 acc=0.9357 13242.3 examples/second

Test Loss=0.3622, Test acc=0.8831

Saving...

Epoch: 5
[Step=50] Loss=0.1849 acc=0.9354 6596.6 examples/second
[Step=100] Loss=0.1848 acc=0.9354 12605.2 examples/second
[Step=150] Loss=0.1826 acc=0.9359 12794.9 examples/second

Test Loss=0.3581, Test acc=0.8853

Saving...

Epoch: 6
[Step=50] Loss=0.1807 acc=0.9342 6581.0 examples/second
[Step=100] Loss=0.1772 acc=0.9370 12638.2 examples/second
[Step=150] Loss=0.1771 acc=0.9373 12780.2 examples/second

Test Loss=0.3544, Test acc=0.8849

Epoch: 7
[Step=50] Loss=0.1696 acc=0.9415 6120.6 examples/second
[Step=100] Loss=0.1741 acc=0.9398 11072.2 examples/second
[Step=150] Loss=0.1734 acc=0.9403 12073.7 examples/second

Test Loss=0.3505, Test acc=0.8871

Saving...

Epoch: 8
[Step=50] Loss=0.1747 acc=0.9388 5946.8 examples/second
[Step=100] Loss=0.1668 acc=0.9422 11545.4 examples/second
[Step=150] Loss=0.1655 acc=0.9422 12119.1 examples/second

Test Loss=0.3488, Test acc=0.8880

Saving...

Epoch: 9
[Step=50] Loss=0.1605 acc=0.9440 6828.1 examples/second
[Step=100] Loss=0.1617 acc=0.9445 13478.8 examples/second
[Step=150] Loss=0.1620 acc=0.9444 13561.5 examples/second

Test Loss=0.3445, Test acc=0.8899

Saving...

Epoch: 10
[Step=50] Loss=0.1613 acc=0.9432 6316.5 examples/second
[Step=100] Loss=0.1633 acc=0.9413 12868.6 examples/second
[Step=150] Loss=0.1629 acc=0.9414 13039.7 examples/second

Test Loss=0.3437, Test acc=0.8906

Saving...

Epoch: 11
[Step=50] Loss=0.1552 acc=0.9457 6349.5 examples/second
[Step=100] Loss=0.1576 acc=0.9442 12022.5 examples/second
[Step=150] Loss=0.1589 acc=0.9437 13227.5 examples/second

Test Loss=0.3418, Test acc=0.8914

Saving...

Epoch: 12
[Step=50] Loss=0.1499 acc=0.9504 6148.7 examples/second
[Step=100] Loss=0.1527 acc=0.9482 12449.1 examples/second
[Step=150] Loss=0.1522 acc=0.9489 12410.9 examples/second

Test Loss=0.3398, Test acc=0.8914

Epoch: 13
[Step=50] Loss=0.1548 acc=0.9458 6635.9 examples/second
[Step=100] Loss=0.1583 acc=0.9450 12242.1 examples/second
[Step=150] Loss=0.1555 acc=0.9462 12662.9 examples/second

Test Loss=0.3404, Test acc=0.8919

Saving...

Epoch: 14
[Step=50] Loss=0.1570 acc=0.9452 6721.6 examples/second
[Step=100] Loss=0.1480 acc=0.9491 12571.6 examples/second
[Step=150] Loss=0.1514 acc=0.9472 12635.6 examples/second

Test Loss=0.3387, Test acc=0.8911

Epoch: 15
[Step=50] Loss=0.1461 acc=0.9513 6599.7 examples/second
[Step=100] Loss=0.1475 acc=0.9502 12253.1 examples/second
[Step=150] Loss=0.1499 acc=0.9492 12467.3 examples/second

Test Loss=0.3371, Test acc=0.8930

Saving...

Epoch: 16
[Step=50] Loss=0.1482 acc=0.9480 5790.8 examples/second
[Step=100] Loss=0.1453 acc=0.9493 12149.8 examples/second
[Step=150] Loss=0.1480 acc=0.9482 12349.3 examples/second

Test Loss=0.3375, Test acc=0.8919

Epoch: 17

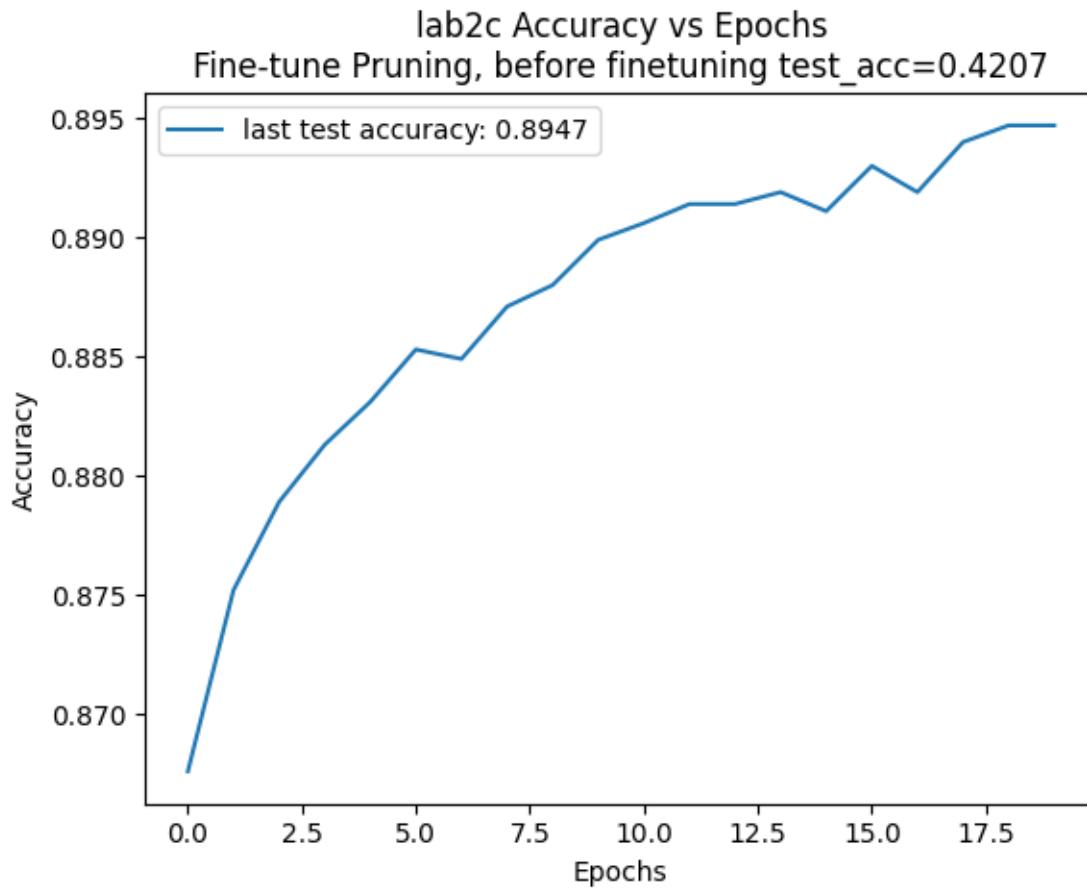
```
[Step=50]      Loss=0.1414      acc=0.9503      5905.5 examples/second
[Step=100]     Loss=0.1441      acc=0.9495      12511.9 examples/second
[Step=150]     Loss=0.1428      acc=0.9504      12075.0 examples/second
Test Loss=0.3341, Test acc=0.8940
Saving...
```

```
Epoch: 18
[Step=50]      Loss=0.1442      acc=0.9498      6657.1 examples/second
[Step=100]     Loss=0.1476      acc=0.9484      12269.5 examples/second
[Step=150]     Loss=0.1461      acc=0.9492      12533.2 examples/second
Test Loss=0.3343, Test acc=0.8947
Saving...
```

```
Epoch: 19
[Step=50]      Loss=0.1399      acc=0.9519      6243.7 examples/second
[Step=100]     Loss=0.1443      acc=0.9498      12571.7 examples/second
[Step=150]     Loss=0.1426      acc=0.9509      12707.4 examples/second
Test Loss=0.3344, Test acc=0.8947
```

```
[10]: fig, ax = plt.subplots(1, 1)
xx = range(n_epochs)
ax.plot(xx, test_acc_lst, label='last test accuracy: %g' % test_acc_lst[-1])
ax.set_xlabel('Epochs')
ax.set_ylabel('Accuracy')
ax.set_title('lab2c Accuracy vs Epochs\nFine-tune Pruning, before finetuning')
    ↵test_acc=%.4f'
    % lab2c_test_b4_finetune)
ax.legend()

plt.savefig('Figures/lab2c.pdf', dpi=500, bbox_inches='tight')
```



```
[11]: # Check sparsity of the finetuned model, make sure it's not changed
net.load_state_dict(torch.load("net_after_finetune.pt"))

for name,layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
       'id_mapping' not in name:
        # Your code here:
        ## Convert the weight of "layer" to numpy array
        np_weight = layer.weight.detach().cpu().numpy()
        ## Count number of zeros
        zeros = sum((np_weight == 0).flatten())
        ## Count number of parameters
        total = len(np_weight.flatten())
        # Print sparsity
        print('Sparsity of %s: %g' % (name, zeros/total))

test(net)
```

Sparsity of head_conv.0.conv: 0.699074

```

Sparsity of body_op.0.conv1.0.conv: 0.700087
Sparsity of body_op.0.conv2.0.conv: 0.700087
Sparsity of body_op.1.conv1.0.conv: 0.700087
Sparsity of body_op.1.conv2.0.conv: 0.700087
Sparsity of body_op.2.conv1.0.conv: 0.700087
Sparsity of body_op.2.conv2.0.conv: 0.700087
Sparsity of body_op.3.conv1.0.conv: 0.69987
Sparsity of body_op.3.conv2.0.conv: 0.699978
Sparsity of body_op.4.conv1.0.conv: 0.699978
Sparsity of body_op.4.conv2.0.conv: 0.699978
Sparsity of body_op.5.conv1.0.conv: 0.699978
Sparsity of body_op.5.conv2.0.conv: 0.699978
Sparsity of body_op.6.conv1.0.conv: 0.699978
Sparsity of body_op.6.conv2.0.conv: 0.700005
Sparsity of body_op.7.conv1.0.conv: 0.700005
Sparsity of body_op.7.conv2.0.conv: 0.700005
Sparsity of body_op.8.conv1.0.conv: 0.700005
Sparsity of body_op.8.conv2.0.conv: 0.700005
Sparsity of final_fc.linear: 0.7
Files already downloaded and verified
Test Loss=0.3343, Test accuracy=0.8947

```

[11]: 0.8947

0.0.4 Lab2 (d) Iterative pruning

```

[12]: net.load_state_dict(torch.load("pretrained_model.pt"))
best_acc = 0.
n_epochs = 20

test_acc_iter_prune = []

for epoch in range(n_epochs):
    print('\nEpoch: %d' % epoch)

    net.train()
    if epoch < 10:
        for name,layer in net.named_modules():
            if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and 'id_mapping' not in name:
                # Increase model sparsity
                q = (epoch + 1) * 7
                # Linearly increase the pruning percentage for 10 epochs until reaching 70% in the final epoch
                prune_by_percentage(layer, q=q)
    if epoch < 9:

```

```

        finetune_after_prune(net, trainloader, criterion, optimizer, ↴
↳prune=False)
    else: # starts from epoch==9, where q = 70%
        finetune_after_prune(net, trainloader, criterion, optimizer)

#Start the testing code.
net.eval()
test_loss = 0
correct = 0
total = 0
with torch.no_grad():
    for batch_idx, (inputs, targets) in enumerate(testloader):
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = net(inputs)
        loss = criterion(outputs, targets)

        test_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()
num_val_steps = len(testloader)
val_acc = correct / total

test_acc_iter_prune.append(val_acc)

print("Test Loss=%.4f, Test acc=%.4f" % (test_loss / (num_val_steps), ↴
↳val_acc))

if epoch>=10:
    if val_acc > best_acc:
        best_acc = val_acc
        print("Saving...")
        torch.save(net.state_dict(), "net_after_iterative_prune.pt")

```

Epoch: 0

[Step=50]	Loss=0.0472	acc=0.9847	6590.0 examples/second
[Step=100]	Loss=0.0482	acc=0.9841	12078.6 examples/second
[Step=150]	Loss=0.0481	acc=0.9843	12069.6 examples/second

Test Loss=0.3262, Test acc=0.9134

Epoch: 1

[Step=50]	Loss=0.0482	acc=0.9840	6651.5 examples/second
[Step=100]	Loss=0.0497	acc=0.9833	12706.2 examples/second
[Step=150]	Loss=0.0481	acc=0.9841	12666.0 examples/second

Test Loss=0.3265, Test acc=0.9139

Epoch: 2

[Step=50]	Loss=0.0475	acc=0.9846	6783.0 examples/second
[Step=100]	Loss=0.0507	acc=0.9833	12782.9 examples/second
[Step=150]	Loss=0.0497	acc=0.9841	12842.8 examples/second

Test Loss=0.3286, Test acc=0.9148

Epoch: 3

[Step=50]	Loss=0.0521	acc=0.9824	5565.0 examples/second
[Step=100]	Loss=0.0516	acc=0.9832	11161.7 examples/second
[Step=150]	Loss=0.0525	acc=0.9824	12087.6 examples/second

Test Loss=0.3304, Test acc=0.9125

Epoch: 4

[Step=50]	Loss=0.0631	acc=0.9784	6435.7 examples/second
[Step=100]	Loss=0.0618	acc=0.9793	12608.6 examples/second
[Step=150]	Loss=0.0608	acc=0.9797	12673.7 examples/second

Test Loss=0.3343, Test acc=0.9117

Epoch: 5

[Step=50]	Loss=0.0740	acc=0.9750	6680.0 examples/second
[Step=100]	Loss=0.0719	acc=0.9757	12194.1 examples/second
[Step=150]	Loss=0.0700	acc=0.9764	12545.0 examples/second

Test Loss=0.3380, Test acc=0.9084

Epoch: 6

[Step=50]	Loss=0.0918	acc=0.9689	6544.7 examples/second
[Step=100]	Loss=0.0897	acc=0.9691	12636.3 examples/second
[Step=150]	Loss=0.0878	acc=0.9701	12305.3 examples/second

Test Loss=0.3330, Test acc=0.9063

Epoch: 7

[Step=50]	Loss=0.1278	acc=0.9573	6697.8 examples/second
[Step=100]	Loss=0.1239	acc=0.9588	12323.9 examples/second
[Step=150]	Loss=0.1204	acc=0.9596	12597.7 examples/second

Test Loss=0.3359, Test acc=0.9022

Epoch: 8

[Step=50]	Loss=0.1709	acc=0.9405	6641.5 examples/second
[Step=100]	Loss=0.1636	acc=0.9422	12175.7 examples/second
[Step=150]	Loss=0.1570	acc=0.9448	12565.0 examples/second

Test Loss=0.3472, Test acc=0.8941

Epoch: 9

[Step=50]	Loss=0.2701	acc=0.9071	6561.1 examples/second
[Step=100]	Loss=0.2503	acc=0.9139	12533.8 examples/second
[Step=150]	Loss=0.2406	acc=0.9166	12643.9 examples/second

Test Loss=0.3901, Test acc=0.8759

Epoch: 10
 [Step=50] Loss=0.2112 acc=0.9237 6791.6 examples/second
 [Step=100] Loss=0.2032 acc=0.9277 12589.8 examples/second
 [Step=150] Loss=0.2042 acc=0.9279 12803.1 examples/second
 Test Loss=0.3735, Test acc=0.8804
 Saving...

Epoch: 11
 [Step=50] Loss=0.1868 acc=0.9356 6858.2 examples/second
 [Step=100] Loss=0.1856 acc=0.9359 12521.4 examples/second
 [Step=150] Loss=0.1862 acc=0.9343 11670.0 examples/second
 Test Loss=0.3674, Test acc=0.8825
 Saving...

Epoch: 12
 [Step=50] Loss=0.1798 acc=0.9384 6543.2 examples/second
 [Step=100] Loss=0.1781 acc=0.9386 12376.5 examples/second
 [Step=150] Loss=0.1805 acc=0.9376 13172.6 examples/second
 Test Loss=0.3585, Test acc=0.8853
 Saving...

Epoch: 13
 [Step=50] Loss=0.1798 acc=0.9373 5819.9 examples/second
 [Step=100] Loss=0.1740 acc=0.9392 10813.6 examples/second
 [Step=150] Loss=0.1709 acc=0.9397 10752.0 examples/second
 Test Loss=0.3546, Test acc=0.8866
 Saving...

Epoch: 14
 [Step=50] Loss=0.1670 acc=0.9427 6128.3 examples/second
 [Step=100] Loss=0.1696 acc=0.9418 11742.8 examples/second
 [Step=150] Loss=0.1661 acc=0.9427 12339.2 examples/second
 Test Loss=0.3511, Test acc=0.8874
 Saving...

Epoch: 15
 [Step=50] Loss=0.1665 acc=0.9420 6126.5 examples/second
 [Step=100] Loss=0.1631 acc=0.9429 12471.9 examples/second
 [Step=150] Loss=0.1645 acc=0.9425 12356.1 examples/second
 Test Loss=0.3488, Test acc=0.8893
 Saving...

Epoch: 16
 [Step=50] Loss=0.1545 acc=0.9454 5152.9 examples/second
 [Step=100] Loss=0.1574 acc=0.9453 11185.4 examples/second
 [Step=150] Loss=0.1581 acc=0.9445 11841.8 examples/second
 Test Loss=0.3450, Test acc=0.8899
 Saving...

```

Epoch: 17
[Step=50]      Loss=0.1604      acc=0.9439      6179.6 examples/second
[Step=100]     Loss=0.1543      acc=0.9450      11965.1 examples/second
[Step=150]     Loss=0.1543      acc=0.9456      11827.1 examples/second
Test Loss=0.3437, Test acc=0.8911
Saving...

Epoch: 18
[Step=50]      Loss=0.1557      acc=0.9444      6268.0 examples/second
[Step=100]     Loss=0.1548      acc=0.9464      12063.5 examples/second
[Step=150]     Loss=0.1541      acc=0.9463      12084.5 examples/second
Test Loss=0.3431, Test acc=0.8905

Epoch: 19
[Step=50]      Loss=0.1500      acc=0.9492      6166.7 examples/second
[Step=100]     Loss=0.1519      acc=0.9481      12370.6 examples/second
[Step=150]     Loss=0.1496      acc=0.9488      12396.4 examples/second
Test Loss=0.3405, Test acc=0.8922
Saving...

```

```
[13]: # Check sparsity of the final model, make sure it's 70%
net.load_state_dict(torch.load("net_after_iterative_prune.pt"))

for name,layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
       'id_mapping' not in name:
        # Your code here: can copy from previous question
        ## Convert the weight of "layer" to numpy array
        np_weight = layer.weight.detach().cpu().numpy()
        ## Count number of zeros
        zeros = sum((np_weight == 0).flatten())
        ## Count number of parameters
        total = len(np_weight.flatten())
        # Print sparsity
        print('Sparsity of %s: %g' % (name, zeros/total))

test(net)
```

```

Sparsity of head_conv.0.conv: 0.699074
Sparsity of body_op.0.conv1.0.conv: 0.700087
Sparsity of body_op.0.conv2.0.conv: 0.700087
Sparsity of body_op.1.conv1.0.conv: 0.700087
Sparsity of body_op.1.conv2.0.conv: 0.700087
Sparsity of body_op.2.conv1.0.conv: 0.700087
Sparsity of body_op.2.conv2.0.conv: 0.700087
Sparsity of body_op.3.conv1.0.conv: 0.69987

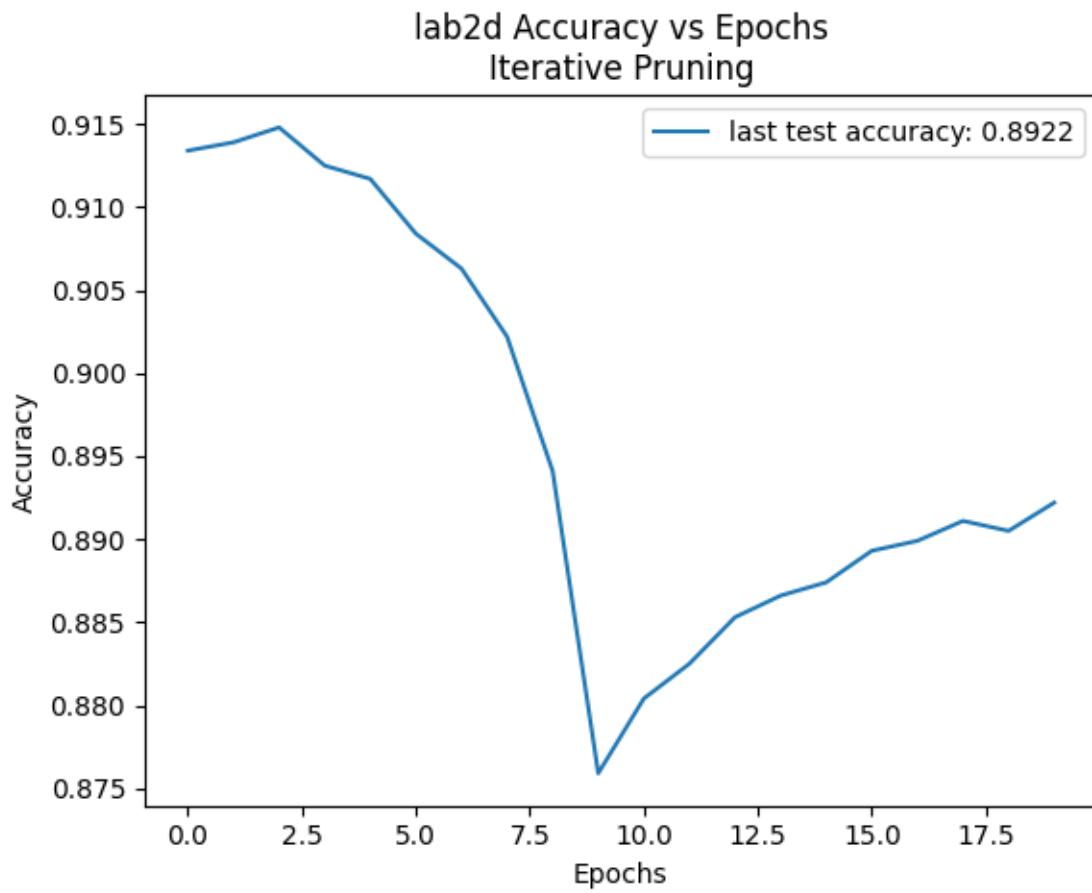
```

```
Sparsity of body_op.3.conv2.0.conv: 0.699978
Sparsity of body_op.4.conv1.0.conv: 0.699978
Sparsity of body_op.4.conv2.0.conv: 0.699978
Sparsity of body_op.5.conv1.0.conv: 0.699978
Sparsity of body_op.5.conv2.0.conv: 0.699978
Sparsity of body_op.6.conv1.0.conv: 0.699978
Sparsity of body_op.6.conv2.0.conv: 0.700005
Sparsity of body_op.7.conv1.0.conv: 0.700005
Sparsity of body_op.7.conv2.0.conv: 0.700005
Sparsity of body_op.8.conv1.0.conv: 0.700005
Sparsity of body_op.8.conv2.0.conv: 0.700005
Sparsity of final_fc.linear: 0.7
Files already downloaded and verified
Test Loss=0.3405, Test accuracy=0.8922
```

[13]: 0.8922

```
[14]: fig, ax = plt.subplots(1, 1)
xx = range(n_epochs)
ax.plot(xx, test_acc_iter_prune, label='last test accuracy: %g' %_
        test_acc_iter_prune[-1])
ax.set_xlabel('Epochs')
ax.set_ylabel('Accuracy')
ax.set_title('lab2d Accuracy vs Epochs\nIterative Pruning')
ax.legend()

plt.savefig('Figures/lab2d.pdf', dpi=500, bbox_inches='tight')
```



0.0.5 Lab2 (e) Global iterative pruning

```
[15]: import numpy as np

def global_prune_by_percentage(net, q=70.0):
    """
    Pruning the weight parameters by threshold.
    :param q: pruning percentile. 'q' percent of the least
    significant weight parameters will be pruned.
    """
    # A list to gather all the weights
    flattened_weights = []
    # Find global pruning threshold
    for name,layer in net.named_modules():
        if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
        ↵'id_mapping' not in name:
            # Convert weight to numpy
            # Flatten the weight and append to flattened_weights
```

```

        flattened_weights.append(layer.weight.detach().cpu().numpy() .
flatten())

# Concat all weights into a np array
flattened_weights = np.concatenate(flattened_weights)
# Find global pruning threshold
threshold = np.percentile(np.abs(flattened_weights.flatten()), q)

# Apply pruning threshold to all layers
for name,layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
'__id_mapping' not in name:
        # Convert weight to numpy
        layer_weight = layer.weight.detach().cpu().numpy()

        # Generate a binary mask same shape as weight to decide which
element to prune
        masked_obj = np.ma.masked_greater_equal(x=np.abs(layer_weight),u
value=threshold, copy=True)
        mask_int = np.ma.getmask(masked_obj).astype(int)
        # Convert mask to torch tensor and put on GPU
        mask_tensor = torch.tensor(mask_int).to(device)
        # Multiply the weight by mask to perform pruning
        layer.weight.data = layer.weight.data.clone().detach() .
requires_grad_(True) * mask_tensor

```

```

[16]: net.load_state_dict(torch.load("pretrained_model.pt"))
best_acc = 0.

test_acc_global_prune = []

for epoch in range(20):
    print('\nEpoch: %d' % epoch)
    # q = (epoch + 1) * 8
    q = (epoch + 1) * 7

    net.train()
    # Increase model sparsity
    if epoch < 10:
        global_prune_by_percentage(net, q=q)
    if epoch < 9:
        finetune_after_prune(net, trainloader, criterion, optimizer,
prune=False)
    else:
        finetune_after_prune(net, trainloader, criterion, optimizer)

#Start the testing code.

```

```

net.eval()
test_loss = 0
correct = 0
total = 0
with torch.no_grad():
    for batch_idx, (inputs, targets) in enumerate(testloader):
        inputs, targets = inputs.to(device), targets.to(device)
        outputs = net(inputs)
        loss = criterion(outputs, targets)

        test_loss += loss.item()
        _, predicted = outputs.max(1)
        total += targets.size(0)
        correct += predicted.eq(targets).sum().item()
num_val_steps = len(testloader)
val_acc = correct / total
print("Test Loss=%.4f, Test acc=%.4f" % (test_loss / (num_val_steps), val_acc))

test_acc_global_prune.append(val_acc)

if epoch >= 10:
    if val_acc > best_acc:
        best_acc = val_acc
    print("Saving...")
    torch.save(net.state_dict(), "net_after_global_iterative_prune.pt")

```

Epoch: 0

[Step=50]	Loss=0.0494	acc=0.9846	6143.9 examples/second
[Step=100]	Loss=0.0484	acc=0.9847	12491.2 examples/second
[Step=150]	Loss=0.0493	acc=0.9843	12542.2 examples/second

Test Loss=0.3235, Test acc=0.9150

Epoch: 1

[Step=50]	Loss=0.0483	acc=0.9848	6045.3 examples/second
[Step=100]	Loss=0.0488	acc=0.9844	12603.3 examples/second
[Step=150]	Loss=0.0488	acc=0.9844	10154.3 examples/second

Test Loss=0.3239, Test acc=0.9143

Epoch: 2

[Step=50]	Loss=0.0480	acc=0.9859	5200.7 examples/second
[Step=100]	Loss=0.0480	acc=0.9857	12315.7 examples/second
[Step=150]	Loss=0.0483	acc=0.9854	12389.8 examples/second

Test Loss=0.3291, Test acc=0.9137

Epoch: 3

[Step=50]	Loss=0.0501	acc=0.9843	6317.4 examples/second
[Step=100]	Loss=0.0496	acc=0.9845	12250.7 examples/second
[Step=150]	Loss=0.0499	acc=0.9841	12418.7 examples/second
Test Loss=0.3274, Test acc=0.9127			
Epoch: 4			
[Step=50]	Loss=0.0500	acc=0.9841	4792.6 examples/second
[Step=100]	Loss=0.0519	acc=0.9842	12345.4 examples/second
[Step=150]	Loss=0.0524	acc=0.9833	12372.6 examples/second
Test Loss=0.3260, Test acc=0.9141			
Epoch: 5			
[Step=50]	Loss=0.0617	acc=0.9802	6315.3 examples/second
[Step=100]	Loss=0.0630	acc=0.9797	13431.7 examples/second
[Step=150]	Loss=0.0614	acc=0.9797	12610.3 examples/second
Test Loss=0.3229, Test acc=0.9113			
Epoch: 6			
[Step=50]	Loss=0.0731	acc=0.9754	6366.3 examples/second
[Step=100]	Loss=0.0719	acc=0.9758	12640.7 examples/second
[Step=150]	Loss=0.0718	acc=0.9759	12408.8 examples/second
Test Loss=0.3234, Test acc=0.9099			
Epoch: 7			
[Step=50]	Loss=0.0885	acc=0.9702	6106.5 examples/second
[Step=100]	Loss=0.0860	acc=0.9717	12448.0 examples/second
[Step=150]	Loss=0.0872	acc=0.9708	12239.8 examples/second
Test Loss=0.3256, Test acc=0.9068			
Epoch: 8			
[Step=50]	Loss=0.1173	acc=0.9610	6134.2 examples/second
[Step=100]	Loss=0.1194	acc=0.9599	12445.5 examples/second
[Step=150]	Loss=0.1173	acc=0.9605	12471.5 examples/second
Test Loss=0.3249, Test acc=0.9005			
Epoch: 9			
[Step=50]	Loss=0.1843	acc=0.9363	6212.4 examples/second
[Step=100]	Loss=0.1780	acc=0.9388	11804.6 examples/second
[Step=150]	Loss=0.1724	acc=0.9407	10877.9 examples/second
Test Loss=0.3410, Test acc=0.8903			
Epoch: 10			
[Step=50]	Loss=0.1532	acc=0.9495	6142.6 examples/second
[Step=100]	Loss=0.1586	acc=0.9470	13074.3 examples/second
[Step=150]	Loss=0.1582	acc=0.9464	11776.5 examples/second
Test Loss=0.3339, Test acc=0.8934			
Saving...			

Epoch: 11
 [Step=50] Loss=0.1566 acc=0.9456 6226.4 examples/second
 [Step=100] Loss=0.1532 acc=0.9479 11704.4 examples/second
 [Step=150] Loss=0.1514 acc=0.9485 11097.2 examples/second
 Test Loss=0.3300, Test acc=0.8946
 Saving...

Epoch: 12
 [Step=50] Loss=0.1441 acc=0.9504 5309.6 examples/second
 [Step=100] Loss=0.1457 acc=0.9515 12221.5 examples/second
 [Step=150] Loss=0.1454 acc=0.9513 12054.7 examples/second
 Test Loss=0.3273, Test acc=0.8949
 Saving...

Epoch: 13
 [Step=50] Loss=0.1325 acc=0.9542 6226.3 examples/second
 [Step=100] Loss=0.1376 acc=0.9525 12487.7 examples/second
 [Step=150] Loss=0.1391 acc=0.9525 12441.0 examples/second
 Test Loss=0.3229, Test acc=0.8970
 Saving...

Epoch: 14
 [Step=50] Loss=0.1314 acc=0.9566 6036.0 examples/second
 [Step=100] Loss=0.1372 acc=0.9543 11868.5 examples/second
 [Step=150] Loss=0.1358 acc=0.9542 11846.7 examples/second
 Test Loss=0.3231, Test acc=0.8969

Epoch: 15
 [Step=50] Loss=0.1355 acc=0.9521 6031.4 examples/second
 [Step=100] Loss=0.1347 acc=0.9539 12181.6 examples/second
 [Step=150] Loss=0.1353 acc=0.9537 12342.5 examples/second
 Test Loss=0.3204, Test acc=0.8985
 Saving...

Epoch: 16
 [Step=50] Loss=0.1276 acc=0.9568 5985.7 examples/second
 [Step=100] Loss=0.1318 acc=0.9554 12347.7 examples/second
 [Step=150] Loss=0.1319 acc=0.9547 11193.9 examples/second
 Test Loss=0.3189, Test acc=0.8984

Epoch: 17
 [Step=50] Loss=0.1329 acc=0.9560 6316.6 examples/second
 [Step=100] Loss=0.1312 acc=0.9562 12470.3 examples/second
 [Step=150] Loss=0.1286 acc=0.9568 12600.0 examples/second
 Test Loss=0.3184, Test acc=0.8981

Epoch: 18
 [Step=50] Loss=0.1268 acc=0.9566 6341.9 examples/second

```

[Step=100]      Loss=0.1260      acc=0.9570      12705.8 examples/second
[Step=150]      Loss=0.1268      acc=0.9572      10798.4 examples/second
Test Loss=0.3182, Test acc=0.8986
Saving...

```

```

Epoch: 19
[Step=50]      Loss=0.1266      acc=0.9572      5950.9 examples/second
[Step=100]     Loss=0.1259      acc=0.9574      11822.7 examples/second
[Step=150]     Loss=0.1261      acc=0.9572      12023.3 examples/second
Test Loss=0.3178, Test acc=0.8985

```

```

[23]: net.load_state_dict(torch.load("net_after_global_iterative_prune.pt"))
sparsity_list = []
x_ticks_labels = []
zeros_sum = 0
total_sum = 0
for name, layer in net.named_modules():
    if (isinstance(layer, nn.Conv2d) or isinstance(layer, nn.Linear)) and
       'id_mapping' not in name:
        # Your code here:
        # Convert the weight of "layer" to numpy array
        np_weight = layer.weight.detach().cpu().numpy()
        # Count number of zeros
        zeros = sum((np_weight == 0).flatten())
        # Count number of parameters
        total = len(np_weight.flatten())
        zeros_sum += zeros
        total_sum += total
        sparsity_list.append(zeros / (total / 1.0))
        x_ticks_labels.append(name)
        print('Sparsity of %s: %g' % (name, zeros / (total / 1.0)))
print('Total sparsity: %g' % (zeros_sum / (total_sum / 1.0)))
test(net)

```

```

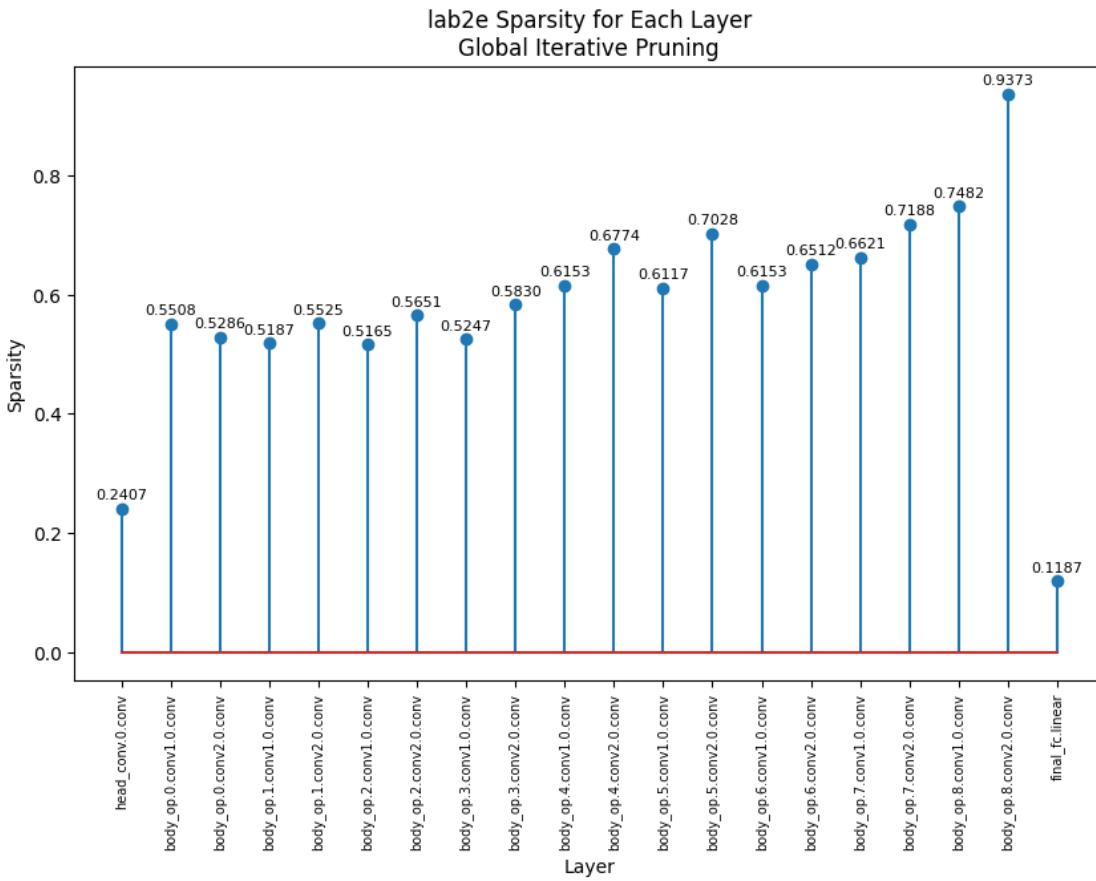
Sparsity of head_conv.0.conv: 0.240741
Sparsity of body_op.0.conv1.0.conv: 0.550781
Sparsity of body_op.0.conv2.0.conv: 0.528646
Sparsity of body_op.1.conv1.0.conv: 0.518663
Sparsity of body_op.1.conv2.0.conv: 0.552517
Sparsity of body_op.2.conv1.0.conv: 0.516493
Sparsity of body_op.2.conv2.0.conv: 0.565104
Sparsity of body_op.3.conv1.0.conv: 0.52474
Sparsity of body_op.3.conv2.0.conv: 0.583008
Sparsity of body_op.4.conv1.0.conv: 0.615343
Sparsity of body_op.4.conv2.0.conv: 0.677409
Sparsity of body_op.5.conv1.0.conv: 0.611654
Sparsity of body_op.5.conv2.0.conv: 0.702799
Sparsity of body_op.6.conv1.0.conv: 0.615343

```

```
Sparsity of body_op.6.conv2.0.conv: 0.651232
Sparsity of body_op.7.conv1.0.conv: 0.662137
Sparsity of body_op.7.conv2.0.conv: 0.718777
Sparsity of body_op.8.conv1.0.conv: 0.74821
Sparsity of body_op.8.conv2.0.conv: 0.937283
Sparsity of final_fc.linear: 0.11875
Total sparsity: 0.699999
Files already downloaded and verified
Test Loss=0.3182, Test accuracy=0.8986
```

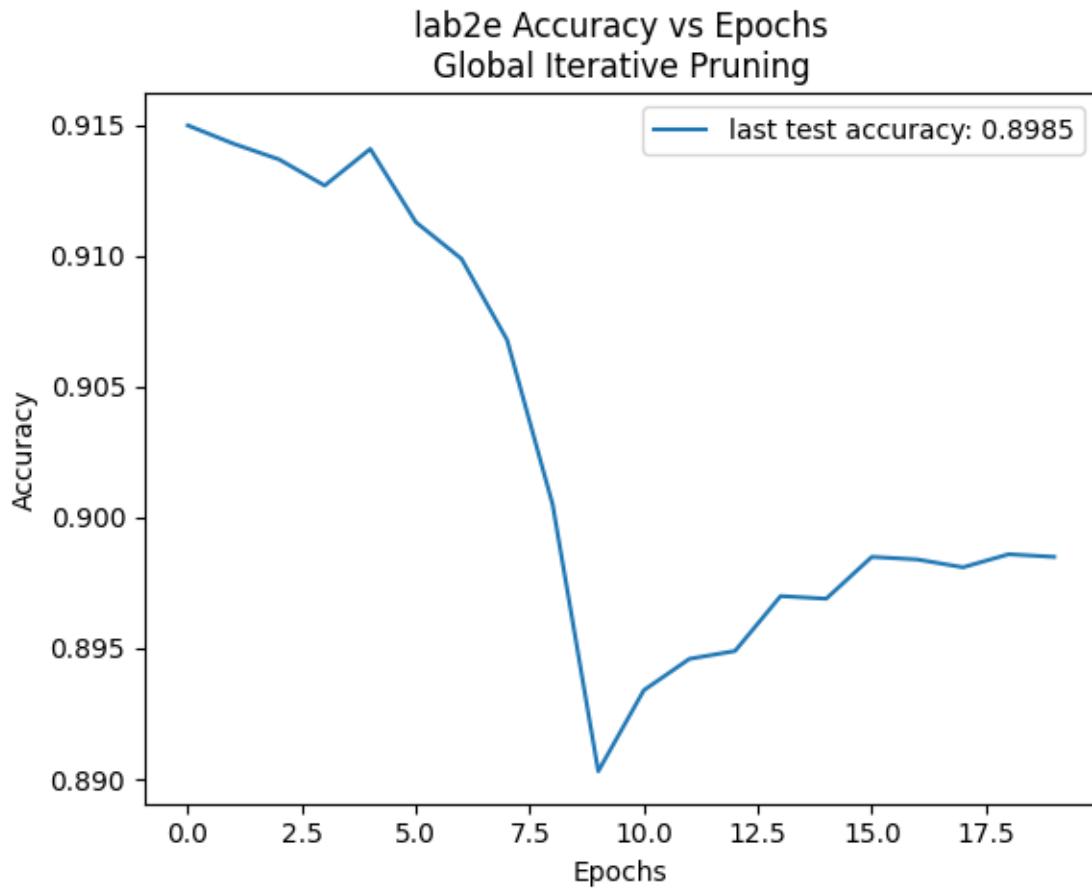
[23]: 0.8986

```
[25]: fig, ax = plt.subplots(1, 1, figsize=(10, 6))
xx = range(len(sparsity_list))
ax.stem(xx, sparsity_list)
ax.set_xticks(range(len(xx)), x_ticks_labels, rotation='vertical', fontsize=7)
ax.set_xlabel('Layer')
ax.set_ylabel('Sparsity')
ax.set_title('lab2e Sparsity for Each Layer\nGlobal Iterative Pruning')
for (x, y) in zip(xx, sparsity_list):
    plt.annotate('{:.4f}'.format(y), xy=(x, y), xytext=(0, 5),
    textcoords='offset points', ha='center', size=8)
plt.savefig('Figures/lab2e_sparsity.pdf', dpi=500, bbox_inches='tight')
```



```
[18]: fig, ax = plt.subplots(1, 1)
xx = range(n_epochs)
ax.plot(xx, test_acc_global_prune, label='last test accuracy: %g' % test_acc_global_prune[-1])
ax.set_xlabel('Epochs')
ax.set_ylabel('Accuracy')
ax.set_title('lab2e Accuracy vs Epochs\nGlobal Iterative Pruning')
ax.legend()

plt.savefig('Figures/lab2e.pdf', dpi=500, bbox_inches='tight')
```



0.0.6 Lab 3 (b) and (c): Fixed-point quantization

```
[4]: # load pretrained ResNet-20 model, report the accuracy of the floating-point
      ↪pretrained model.

net = ResNetCIFAR(num_layers=20)
net = net.to(device)
net.load_state_dict(torch.load("pretrained_model.pt"))
print("before quantisation:")
test(net)
print()

# Define quantized model and load weight
Nbits_arr = [2, 3, 4, 5, 6]
test_acc_lab3b = []
for iii, nb in enumerate(Nbits_arr):
    Nbits = nb #Change this value to finish (b) and (c)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
```

```

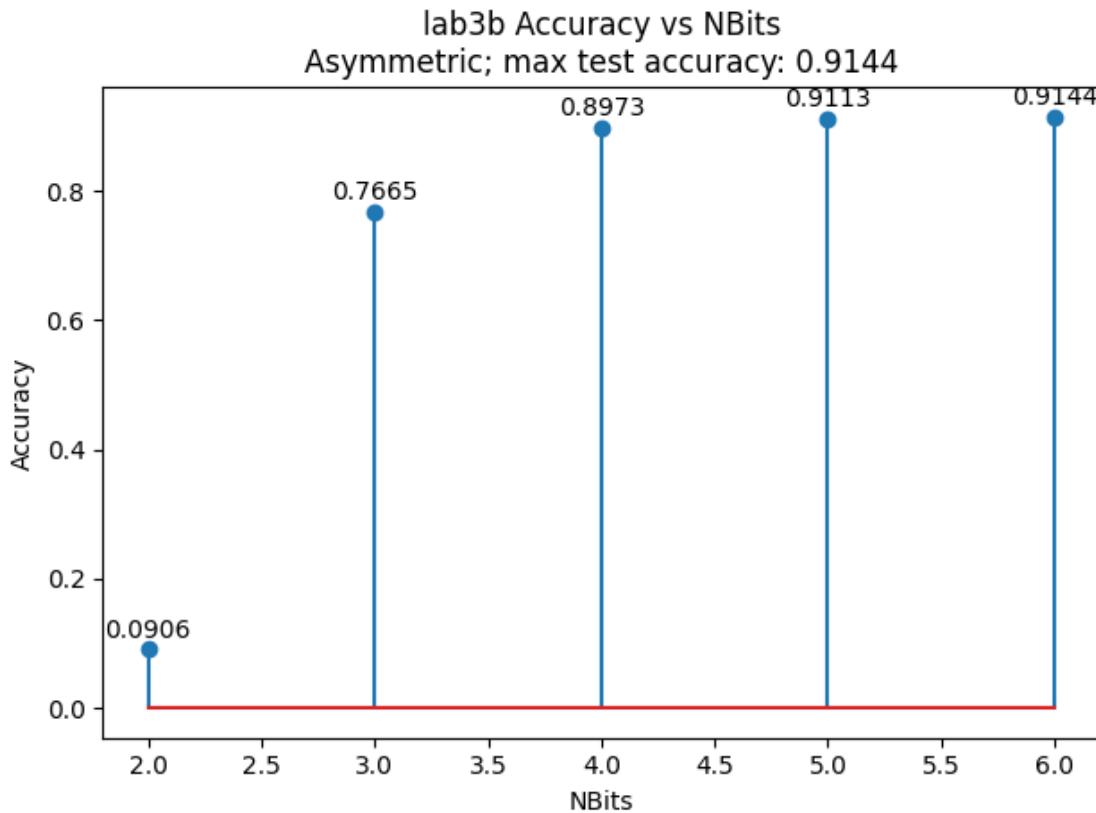
net = net.to(device)
net.load_state_dict(torch.load("pretrained_model.pt"))
test_acc_lab3b.append(test(net))

```

before quantisation:
 Files already downloaded and verified
 Test Loss=0.3231, Test accuracy=0.9150

Files already downloaded and verified
 Test Loss=9.6086, Test accuracy=0.0906
 Files already downloaded and verified
 Test Loss=0.9857, Test accuracy=0.7665
 Files already downloaded and verified
 Test Loss=0.3860, Test accuracy=0.8973
 Files already downloaded and verified
 Test Loss=0.3391, Test accuracy=0.9113
 Files already downloaded and verified
 Test Loss=0.3364, Test accuracy=0.9144

```
[10]: fig, ax = plt.subplots(1, 1)
xx = Nbits_arr
ax.stem(xx, test_acc_lab3b)
# ax.plot(xx, test_acc_lab3b, label='max test accuracy: %g' % np.
#         ↪max(test_acc_lab3b))
ax.set_xlabel('NBits')
ax.set_ylabel('Accuracy')
ax.set_title('lab3b Accuracy vs NBits\nAsymmetric; max test accuracy: %g' % np.
             ↪max(test_acc_lab3b))
for (x, y) in zip(xx, test_acc_lab3b):
    plt.annotate('{:.4f}'.format(y), xy=(x, y), xytext=(0, 5), ↪
                 ↪textcoords='offset points', ha='center')
fig.tight_layout()
plt.savefig('Figures/lab3b.pdf', dpi=500, bbox_inches='tight')
```



```
[11]: # lab 3 part c is shown below:
```

```
[12]: Nbits_arr_c = [2, 3, 4]
test_acc_lab3c_finetune = []
# test accuracy for each finetuning as fn of epoch
final_test_acc_lab3c = {} # test acc after finetuning for each Nbit
n_epochs = 20
test_acc_lab3c_b4_finetune = {} # test acc before finetuning

for iii, nb in enumerate(Nbits_arr_c):
    # Define quantized model and load weight
    Nbits = nb #Change this value to finish (b) and (c)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
    net = net.to(device)
    net.load_state_dict(torch.load("pretrained_model.pt"))
    test_acc_lab3c_b4_finetune[Nbits] = test(net)

    # Quantized model finetuning
    test_acc_lab3c_finetune[Nbits] = np.asarray(
        finetune(net, epochs=n_epochs, batch_size=256, lr=0.002, reg=1e-4))
```

```

# finetune() supposedly returns list of val_accuracy values
# Load the model with best accuracy
net.load_state_dict(torch.load("quantized_net_after_finetune.pt"))
final_test_acc_lab3c[Nbits] = test(net)

```

Files already downloaded and verified
Test Loss=9.6086, Test accuracy=0.0906
==> Preparing data..
Files already downloaded and verified
Files already downloaded and verified

Epoch: 0
[Step=50] Loss=1.3199 acc=0.6250 5064.9 examples/second
[Step=100] Loss=1.0535 acc=0.6870 9559.5 examples/second
[Step=150] Loss=0.9346 acc=0.7166 9594.5 examples/second
Test Loss=0.7387, Test acc=0.7820
Saving...

Epoch: 1
[Step=200] Loss=0.6141 acc=0.8047 3416.9 examples/second
[Step=250] Loss=0.5912 acc=0.8040 9611.2 examples/second
[Step=300] Loss=0.5688 acc=0.8098 9720.0 examples/second
[Step=350] Loss=0.5459 acc=0.8157 9729.8 examples/second
Test Loss=0.6160, Test acc=0.8087
Saving...

Epoch: 2
[Step=400] Loss=0.4761 acc=0.8359 2916.8 examples/second
[Step=450] Loss=0.4810 acc=0.8365 8614.2 examples/second
[Step=500] Loss=0.4718 acc=0.8390 7992.3 examples/second
[Step=550] Loss=0.4650 acc=0.8406 9019.1 examples/second
Test Loss=0.5682, Test acc=0.8207
Saving...

Epoch: 3
[Step=600] Loss=0.4422 acc=0.8503 2982.4 examples/second
[Step=650] Loss=0.4296 acc=0.8514 7740.0 examples/second
[Step=700] Loss=0.4260 acc=0.8522 8216.0 examples/second
[Step=750] Loss=0.4242 acc=0.8532 9281.9 examples/second
Test Loss=0.5630, Test acc=0.8256
Saving...

Epoch: 4
[Step=800] Loss=0.3994 acc=0.8660 3616.9 examples/second
[Step=850] Loss=0.4016 acc=0.8609 8880.7 examples/second
[Step=900] Loss=0.3993 acc=0.8597 8941.6 examples/second
[Step=950] Loss=0.3989 acc=0.8605 8840.8 examples/second

Test Loss=0.5452, Test acc=0.8309

Saving...

Epoch: 5

[Step=1000]	Loss=0.4074	acc=0.8596	3349.6 examples/second
[Step=1050]	Loss=0.3827	acc=0.8658	9304.9 examples/second
[Step=1100]	Loss=0.3851	acc=0.8651	9023.1 examples/second
[Step=1150]	Loss=0.3811	acc=0.8667	9238.3 examples/second

Test Loss=0.5262, Test acc=0.8353

Saving...

Epoch: 6

[Step=1200]	Loss=0.3542	acc=0.8791	3387.8 examples/second
[Step=1250]	Loss=0.3542	acc=0.8769	9156.1 examples/second
[Step=1300]	Loss=0.3613	acc=0.8749	9068.4 examples/second
[Step=1350]	Loss=0.3592	acc=0.8750	9290.3 examples/second

Test Loss=0.5987, Test acc=0.8243

Epoch: 7

[Step=1400]	Loss=0.3612	acc=0.8726	3634.1 examples/second
[Step=1450]	Loss=0.3503	acc=0.8758	9067.2 examples/second
[Step=1500]	Loss=0.3517	acc=0.8771	8964.0 examples/second
[Step=1550]	Loss=0.3503	acc=0.8772	9364.0 examples/second

Test Loss=0.5400, Test acc=0.8394

Saving...

Epoch: 8

[Step=1600]	Loss=0.3455	acc=0.8785	3548.8 examples/second
[Step=1650]	Loss=0.3429	acc=0.8808	9186.2 examples/second
[Step=1700]	Loss=0.3387	acc=0.8811	9122.0 examples/second
[Step=1750]	Loss=0.3402	acc=0.8801	9614.7 examples/second

Test Loss=0.5119, Test acc=0.8431

Saving...

Epoch: 9

[Step=1800]	Loss=0.3359	acc=0.8837	3562.8 examples/second
[Step=1850]	Loss=0.3359	acc=0.8843	9036.8 examples/second
[Step=1900]	Loss=0.3303	acc=0.8847	9079.4 examples/second
[Step=1950]	Loss=0.3288	acc=0.8851	9645.1 examples/second

Test Loss=0.5415, Test acc=0.8349

Epoch: 10

[Step=2000]	Loss=0.3316	acc=0.8798	3638.5 examples/second
[Step=2050]	Loss=0.3348	acc=0.8805	9018.6 examples/second
[Step=2100]	Loss=0.3331	acc=0.8820	9029.8 examples/second
[Step=2150]	Loss=0.3294	acc=0.8833	9577.1 examples/second

Test Loss=0.5093, Test acc=0.8420

Epoch: 11
 [Step=2200] Loss=0.3162 acc=0.8890 3573.9 examples/second
 [Step=2250] Loss=0.3176 acc=0.8871 8949.2 examples/second
 [Step=2300] Loss=0.3180 acc=0.8883 9003.6 examples/second
 [Step=2350] Loss=0.3155 acc=0.8886 9855.7 examples/second
 Test Loss=0.4624, Test acc=0.8555
 Saving...

Epoch: 12
 [Step=2400] Loss=0.3081 acc=0.8923 3372.2 examples/second
 [Step=2450] Loss=0.3098 acc=0.8915 9674.6 examples/second
 [Step=2500] Loss=0.3116 acc=0.8908 9699.1 examples/second
 Test Loss=0.5014, Test acc=0.8462

Epoch: 13
 [Step=2550] Loss=0.2429 acc=0.9199 3746.8 examples/second
 [Step=2600] Loss=0.3055 acc=0.8939 8866.0 examples/second
 [Step=2650] Loss=0.3042 acc=0.8935 9083.3 examples/second
 [Step=2700] Loss=0.3060 acc=0.8927 9076.7 examples/second
 Test Loss=0.4841, Test acc=0.8508

Epoch: 14
 [Step=2750] Loss=0.2795 acc=0.9049 3611.9 examples/second
 [Step=2800] Loss=0.2960 acc=0.8965 8951.2 examples/second
 [Step=2850] Loss=0.2997 acc=0.8952 8979.6 examples/second
 [Step=2900] Loss=0.2992 acc=0.8955 9206.7 examples/second
 Test Loss=0.4598, Test acc=0.8545

Epoch: 15
 [Step=2950] Loss=0.2990 acc=0.8902 3626.7 examples/second
 [Step=3000] Loss=0.2943 acc=0.8950 9052.5 examples/second
 [Step=3050] Loss=0.2922 acc=0.8971 8802.7 examples/second
 [Step=3100] Loss=0.2892 acc=0.8976 8801.3 examples/second
 Test Loss=0.4655, Test acc=0.8568
 Saving...

Epoch: 16
 [Step=3150] Loss=0.2909 acc=0.8970 3331.8 examples/second
 [Step=3200] Loss=0.2963 acc=0.8968 8891.1 examples/second
 [Step=3250] Loss=0.2917 acc=0.8972 8881.5 examples/second
 [Step=3300] Loss=0.2872 acc=0.8978 9036.9 examples/second
 Test Loss=0.4836, Test acc=0.8509

Epoch: 17
 [Step=3350] Loss=0.2757 acc=0.9028 3560.6 examples/second
 [Step=3400] Loss=0.2887 acc=0.8952 8938.8 examples/second
 [Step=3450] Loss=0.2872 acc=0.8962 8840.4 examples/second
 [Step=3500] Loss=0.2860 acc=0.8984 8871.0 examples/second

Test Loss=0.4483, Test acc=0.8576

Saving...

Epoch: 18

[Step=3550]	Loss=0.2802	acc=0.9034	3519.2 examples/second
[Step=3600]	Loss=0.2809	acc=0.9002	9101.0 examples/second
[Step=3650]	Loss=0.2837	acc=0.8998	8806.1 examples/second
[Step=3700]	Loss=0.2874	acc=0.8980	8418.4 examples/second

Test Loss=0.4596, Test acc=0.8576

Epoch: 19

[Step=3750]	Loss=0.2893	acc=0.8980	3245.6 examples/second
[Step=3800]	Loss=0.2876	acc=0.8983	9225.6 examples/second
[Step=3850]	Loss=0.2844	acc=0.8989	8938.6 examples/second
[Step=3900]	Loss=0.2842	acc=0.8997	9348.4 examples/second

Test Loss=0.5014, Test acc=0.8440

Files already downloaded and verified

Test Loss=0.4483, Test accuracy=0.8576

Files already downloaded and verified

Test Loss=0.9857, Test accuracy=0.7665

==> Preparing data..

Files already downloaded and verified

Files already downloaded and verified

Epoch: 0

[Step=50]	Loss=0.1696	acc=0.9402	5790.4 examples/second
[Step=100]	Loss=0.1587	acc=0.9446	9041.2 examples/second
[Step=150]	Loss=0.1515	acc=0.9473	9184.6 examples/second

Test Loss=0.3917, Test acc=0.8957

Saving...

Epoch: 1

[Step=200]	Loss=0.1164	acc=0.9521	3681.9 examples/second
[Step=250]	Loss=0.1214	acc=0.9543	8930.0 examples/second
[Step=300]	Loss=0.1235	acc=0.9542	8967.8 examples/second
[Step=350]	Loss=0.1225	acc=0.9551	8787.4 examples/second

Test Loss=0.3819, Test acc=0.8987

Saving...

Epoch: 2

[Step=400]	Loss=0.1068	acc=0.9614	3329.4 examples/second
[Step=450]	Loss=0.1119	acc=0.9591	9067.3 examples/second
[Step=500]	Loss=0.1118	acc=0.9593	8954.9 examples/second
[Step=550]	Loss=0.1118	acc=0.9597	9145.0 examples/second

Test Loss=0.3750, Test acc=0.9002

Saving...

Epoch: 3

[Step=600]	Loss=0.1122	acc=0.9606	3373.3 examples/second
[Step=650]	Loss=0.1050	acc=0.9624	8948.5 examples/second
[Step=700]	Loss=0.1057	acc=0.9623	8616.6 examples/second
[Step=750]	Loss=0.1058	acc=0.9619	8644.7 examples/second
Test Loss=0.3728, Test acc=0.9002			

Epoch: 4			
[Step=800]	Loss=0.0996	acc=0.9629	3277.9 examples/second
[Step=850]	Loss=0.1045	acc=0.9621	9003.6 examples/second
[Step=900]	Loss=0.1031	acc=0.9628	8842.6 examples/second
[Step=950]	Loss=0.1047	acc=0.9622	8735.5 examples/second
Test Loss=0.3686, Test acc=0.9040			
Saving...			

Epoch: 5			
[Step=1000]	Loss=0.0925	acc=0.9656	3289.7 examples/second
[Step=1050]	Loss=0.1019	acc=0.9615	9219.4 examples/second
[Step=1100]	Loss=0.1018	acc=0.9618	8997.6 examples/second
[Step=1150]	Loss=0.1023	acc=0.9624	9011.6 examples/second
Test Loss=0.3672, Test acc=0.9016			

Epoch: 6			
[Step=1200]	Loss=0.0971	acc=0.9681	3603.5 examples/second
[Step=1250]	Loss=0.1015	acc=0.9638	9218.9 examples/second
[Step=1300]	Loss=0.1008	acc=0.9639	8944.5 examples/second
[Step=1350]	Loss=0.1011	acc=0.9642	9182.3 examples/second
Test Loss=0.3728, Test acc=0.9010			

Epoch: 7			
[Step=1400]	Loss=0.0914	acc=0.9669	3034.5 examples/second
[Step=1450]	Loss=0.0907	acc=0.9675	8991.2 examples/second
[Step=1500]	Loss=0.0951	acc=0.9664	9560.0 examples/second
[Step=1550]	Loss=0.0975	acc=0.9653	9428.4 examples/second
Test Loss=0.3679, Test acc=0.9015			

Epoch: 8			
[Step=1600]	Loss=0.0983	acc=0.9641	3363.7 examples/second
[Step=1650]	Loss=0.0958	acc=0.9653	8885.9 examples/second
[Step=1700]	Loss=0.0964	acc=0.9651	8842.4 examples/second
[Step=1750]	Loss=0.0968	acc=0.9654	8737.4 examples/second
Test Loss=0.3743, Test acc=0.9012			

Epoch: 9			
[Step=1800]	Loss=0.0962	acc=0.9652	3242.6 examples/second
[Step=1850]	Loss=0.0936	acc=0.9663	8963.5 examples/second
[Step=1900]	Loss=0.0942	acc=0.9661	8960.3 examples/second
[Step=1950]	Loss=0.0935	acc=0.9664	8973.6 examples/second
Test Loss=0.3675, Test acc=0.9027			

Epoch: 10

[Step=2000]	Loss=0.0971	acc=0.9646	3627.0 examples/second
[Step=2050]	Loss=0.0966	acc=0.9650	9085.6 examples/second
[Step=2100]	Loss=0.0963	acc=0.9651	9032.3 examples/second
[Step=2150]	Loss=0.0952	acc=0.9654	9184.7 examples/second

Test Loss=0.3578, Test acc=0.9034

Epoch: 11

[Step=2200]	Loss=0.0865	acc=0.9691	3356.4 examples/second
[Step=2250]	Loss=0.0929	acc=0.9664	8594.0 examples/second
[Step=2300]	Loss=0.0932	acc=0.9669	8620.8 examples/second
[Step=2350]	Loss=0.0938	acc=0.9665	8909.3 examples/second

Test Loss=0.3577, Test acc=0.9046

Saving...

Epoch: 12

[Step=2400]	Loss=0.0869	acc=0.9689	3574.0 examples/second
[Step=2450]	Loss=0.0901	acc=0.9679	8848.3 examples/second
[Step=2500]	Loss=0.0898	acc=0.9684	8947.4 examples/second

Test Loss=0.3648, Test acc=0.9022

Epoch: 13

[Step=2550]	Loss=0.0755	acc=0.9766	3390.1 examples/second
[Step=2600]	Loss=0.0876	acc=0.9686	8889.5 examples/second
[Step=2650]	Loss=0.0884	acc=0.9683	8928.4 examples/second
[Step=2700]	Loss=0.0886	acc=0.9682	9048.0 examples/second

Test Loss=0.3537, Test acc=0.9053

Saving...

Epoch: 14

[Step=2750]	Loss=0.0935	acc=0.9642	3611.2 examples/second
[Step=2800]	Loss=0.0906	acc=0.9672	9169.0 examples/second
[Step=2850]	Loss=0.0898	acc=0.9681	8868.0 examples/second
[Step=2900]	Loss=0.0891	acc=0.9684	8897.3 examples/second

Test Loss=0.3601, Test acc=0.9052

Epoch: 15

[Step=2950]	Loss=0.0934	acc=0.9652	3546.7 examples/second
[Step=3000]	Loss=0.0898	acc=0.9673	9089.0 examples/second
[Step=3050]	Loss=0.0890	acc=0.9676	8988.5 examples/second
[Step=3100]	Loss=0.0889	acc=0.9679	9194.3 examples/second

Test Loss=0.3632, Test acc=0.9032

Epoch: 16

[Step=3150]	Loss=0.0829	acc=0.9696	3325.5 examples/second
[Step=3200]	Loss=0.0890	acc=0.9676	9060.3 examples/second
[Step=3250]	Loss=0.0882	acc=0.9681	9120.1 examples/second

```

[Step=3300]      Loss=0.0897      acc=0.9673      9176.5 examples/second
Test Loss=0.3562, Test acc=0.9032

Epoch: 17
[Step=3350]      Loss=0.0849      acc=0.9685      3356.1 examples/second
[Step=3400]      Loss=0.0883      acc=0.9688      9136.2 examples/second
[Step=3450]      Loss=0.0896      acc=0.9686      9003.3 examples/second
[Step=3500]      Loss=0.0896      acc=0.9682      9184.7 examples/second
Test Loss=0.3600, Test acc=0.9016

Epoch: 18
[Step=3550]      Loss=0.0824      acc=0.9714      3551.5 examples/second
[Step=3600]      Loss=0.0878      acc=0.9692      9263.9 examples/second
[Step=3650]      Loss=0.0879      acc=0.9687      8926.6 examples/second
[Step=3700]      Loss=0.0887      acc=0.9686      9366.3 examples/second
Test Loss=0.3725, Test acc=0.9014

Epoch: 19
[Step=3750]      Loss=0.0858      acc=0.9672      3266.7 examples/second
[Step=3800]      Loss=0.0887      acc=0.9679      9136.1 examples/second
[Step=3850]      Loss=0.0847      acc=0.9702      8982.6 examples/second
[Step=3900]      Loss=0.0840      acc=0.9704      9328.0 examples/second
Test Loss=0.3572, Test acc=0.9050
Files already downloaded and verified
Test Loss=0.3537, Test accuracy=0.9053
Files already downloaded and verified
Test Loss=0.3860, Test accuracy=0.8973
==> Preparing data..
Files already downloaded and verified
Files already downloaded and verified

Epoch: 0
[Step=50]        Loss=0.0636      acc=0.9789      5830.0 examples/second
[Step=100]       Loss=0.0648      acc=0.9780      8858.1 examples/second
[Step=150]       Loss=0.0659      acc=0.9770      9042.0 examples/second
Test Loss=0.3373, Test acc=0.9088
Saving...

Epoch: 1
[Step=200]       Loss=0.0409      acc=0.9902      3393.4 examples/second
[Step=250]       Loss=0.0663      acc=0.9766      8828.5 examples/second
[Step=300]       Loss=0.0650      acc=0.9774      8885.3 examples/second
[Step=350]       Loss=0.0647      acc=0.9776      8957.7 examples/second
Test Loss=0.3343, Test acc=0.9094
Saving...

Epoch: 2
[Step=400]       Loss=0.0656      acc=0.9766      3312.1 examples/second

```

[Step=450]	Loss=0.0577	acc=0.9817	8889.2 examples/second
[Step=500]	Loss=0.0608	acc=0.9797	8918.0 examples/second
[Step=550]	Loss=0.0622	acc=0.9791	8865.4 examples/second
Test Loss=0.3353, Test acc=0.9116			
Saving...			
Epoch: 3			
[Step=600]	Loss=0.0627	acc=0.9759	3325.2 examples/second
[Step=650]	Loss=0.0575	acc=0.9797	8984.9 examples/second
[Step=700]	Loss=0.0570	acc=0.9804	8935.8 examples/second
[Step=750]	Loss=0.0581	acc=0.9801	9351.3 examples/second
Test Loss=0.3322, Test acc=0.9107			
Epoch: 4			
[Step=800]	Loss=0.0512	acc=0.9839	3577.3 examples/second
[Step=850]	Loss=0.0567	acc=0.9816	9041.1 examples/second
[Step=900]	Loss=0.0588	acc=0.9805	8821.9 examples/second
[Step=950]	Loss=0.0588	acc=0.9804	8731.1 examples/second
Test Loss=0.3313, Test acc=0.9108			
Epoch: 5			
[Step=1000]	Loss=0.0583	acc=0.9824	3322.6 examples/second
[Step=1050]	Loss=0.0599	acc=0.9815	8863.7 examples/second
[Step=1100]	Loss=0.0604	acc=0.9806	8867.6 examples/second
[Step=1150]	Loss=0.0598	acc=0.9806	9044.9 examples/second
Test Loss=0.3308, Test acc=0.9111			
Epoch: 6			
[Step=1200]	Loss=0.0558	acc=0.9798	3600.8 examples/second
[Step=1250]	Loss=0.0556	acc=0.9813	9778.9 examples/second
[Step=1300]	Loss=0.0555	acc=0.9814	9589.5 examples/second
[Step=1350]	Loss=0.0560	acc=0.9813	9310.6 examples/second
Test Loss=0.3344, Test acc=0.9108			
Epoch: 7			
[Step=1400]	Loss=0.0567	acc=0.9812	3224.2 examples/second
[Step=1450]	Loss=0.0576	acc=0.9805	9086.4 examples/second
[Step=1500]	Loss=0.0577	acc=0.9804	9005.8 examples/second
[Step=1550]	Loss=0.0599	acc=0.9795	9333.0 examples/second
Test Loss=0.3323, Test acc=0.9119			
Saving...			
Epoch: 8			
[Step=1600]	Loss=0.0539	acc=0.9816	3416.3 examples/second
[Step=1650]	Loss=0.0549	acc=0.9819	9024.3 examples/second
[Step=1700]	Loss=0.0556	acc=0.9812	9039.2 examples/second
[Step=1750]	Loss=0.0562	acc=0.9811	9628.8 examples/second
Test Loss=0.3337, Test acc=0.9100			

Epoch: 9

[Step=1800]	Loss=0.0595	acc=0.9809	3369.8 examples/second
[Step=1850]	Loss=0.0582	acc=0.9813	8824.3 examples/second
[Step=1900]	Loss=0.0580	acc=0.9812	8817.4 examples/second
[Step=1950]	Loss=0.0563	acc=0.9817	9178.9 examples/second

Test Loss=0.3334, Test acc=0.9132

Saving...

Epoch: 10

[Step=2000]	Loss=0.0507	acc=0.9844	3437.9 examples/second
[Step=2050]	Loss=0.0541	acc=0.9821	8487.0 examples/second
[Step=2100]	Loss=0.0547	acc=0.9824	8808.1 examples/second
[Step=2150]	Loss=0.0553	acc=0.9820	9507.2 examples/second

Test Loss=0.3339, Test acc=0.9121

Epoch: 11

[Step=2200]	Loss=0.0524	acc=0.9823	3598.0 examples/second
[Step=2250]	Loss=0.0554	acc=0.9810	9539.4 examples/second
[Step=2300]	Loss=0.0549	acc=0.9811	9591.2 examples/second
[Step=2350]	Loss=0.0551	acc=0.9816	9555.4 examples/second

Test Loss=0.3340, Test acc=0.9138

Saving...

Epoch: 12

[Step=2400]	Loss=0.0589	acc=0.9801	3414.0 examples/second
[Step=2450]	Loss=0.0570	acc=0.9804	8979.7 examples/second
[Step=2500]	Loss=0.0568	acc=0.9806	9029.9 examples/second

Test Loss=0.3344, Test acc=0.9129

Epoch: 13

[Step=2550]	Loss=0.0447	acc=0.9844	3638.7 examples/second
[Step=2600]	Loss=0.0516	acc=0.9843	8909.1 examples/second
[Step=2650]	Loss=0.0541	acc=0.9827	8822.3 examples/second
[Step=2700]	Loss=0.0546	acc=0.9820	9170.6 examples/second

Test Loss=0.3373, Test acc=0.9109

Epoch: 14

[Step=2750]	Loss=0.0563	acc=0.9798	3613.7 examples/second
[Step=2800]	Loss=0.0537	acc=0.9818	9061.7 examples/second
[Step=2850]	Loss=0.0528	acc=0.9823	8936.2 examples/second
[Step=2900]	Loss=0.0530	acc=0.9824	8969.6 examples/second

Test Loss=0.3317, Test acc=0.9138

Epoch: 15

[Step=2950]	Loss=0.0498	acc=0.9867	3337.1 examples/second
[Step=3000]	Loss=0.0499	acc=0.9841	9156.2 examples/second
[Step=3050]	Loss=0.0522	acc=0.9831	8975.3 examples/second

[Step=3100] Loss=0.0532 acc=0.9829 9281.6 examples/second
Test Loss=0.3345, Test acc=0.9134

Epoch: 16
[Step=3150] Loss=0.0611 acc=0.9780 3484.5 examples/second
[Step=3200] Loss=0.0600 acc=0.9798 9034.7 examples/second
[Step=3250] Loss=0.0562 acc=0.9814 8955.5 examples/second
[Step=3300] Loss=0.0552 acc=0.9818 9386.9 examples/second
Test Loss=0.3330, Test acc=0.9124

Epoch: 17
[Step=3350] Loss=0.0602 acc=0.9803 3644.2 examples/second
[Step=3400] Loss=0.0570 acc=0.9799 9065.9 examples/second
[Step=3450] Loss=0.0558 acc=0.9808 8923.4 examples/second
[Step=3500] Loss=0.0531 acc=0.9816 9067.4 examples/second
Test Loss=0.3364, Test acc=0.9141
Saving...

Epoch: 18
[Step=3550] Loss=0.0494 acc=0.9837 3300.7 examples/second
[Step=3600] Loss=0.0499 acc=0.9837 9015.7 examples/second
[Step=3650] Loss=0.0513 acc=0.9833 8850.9 examples/second
[Step=3700] Loss=0.0512 acc=0.9832 9148.4 examples/second
Test Loss=0.3413, Test acc=0.9111

Epoch: 19
[Step=3750] Loss=0.0495 acc=0.9826 3279.0 examples/second
[Step=3800] Loss=0.0488 acc=0.9842 9286.5 examples/second
[Step=3850] Loss=0.0503 acc=0.9835 9694.6 examples/second
[Step=3900] Loss=0.0513 acc=0.9828 9261.1 examples/second
Test Loss=0.3384, Test acc=0.9111
Files already downloaded and verified
Test Loss=0.3364, Test accuracy=0.9141

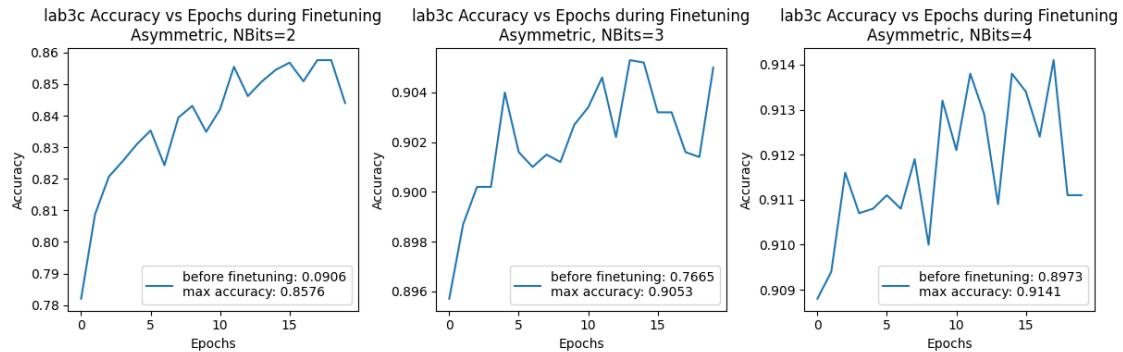
```
[13]: fig, ax = plt.subplots(1, 3, figsize=(12, 4))

for iii, nbits in enumerate(Nbits_arr_c):
    xx = range(n_epochs)
    ax[iii].plot(xx, test_acc_lab3c_finetune[nbits], label='before finetuning: %g\nmax accuracy: %g'
                 % (test_acc_lab3c_b4_finetune[nbits], np.
                    max(test_acc_lab3c_finetune[nbits])))
    ax[iii].set_xlabel('Epochs')
    ax[iii].set_ylabel('Accuracy')
    ax[iii].set_title('lab3c Accuracy vs Epochs during Finetuning\nAsymmetric, NBits=%d' % nbits)
    ax[iii].legend()
```

```

fig.tight_layout()
plt.savefig('Figures/lab3c.pdf', dpi=500, bbox_inches='tight')

```



0.0.7 Lab3 (d) Quantize pruned model

```

[2]: nbits_lab3d = [2, 3, 4]
test_acc_lab3d_b4_finetune = []
test_acc_arrays_finetune = []
test_acc_lab3d_after_finetune = []
n_epochs = 20

for iii, nb in enumerate(nbits_lab3d):
    # Define quantized model and load weight
    Nbits = nb #Change this value to finish (d)

    net = ResNetCIFAR(num_layers=20, Nbits=Nbits)
    net = net.to(device)
    net.load_state_dict(torch.load("net_after_global_iterative_prune.pt"))
    test_acc_lab3d_b4_finetune[Nbits] = test(net)

    # Quantized model finetuning
    test_acc_arrays_finetune[Nbits] = np.array(
        finetune(net, epochs=n_epochs, batch_size=256, lr=0.002, reg=1e-4))

    # Do not load the model with best accuracy since we want the last acc value.
    # net.load_state_dict(torch.load("quantized_net_after_finetune.pt"))
    test_acc_lab3d_after_finetune[Nbits] = test(net)

```

Files already downloaded and verified
Test Loss=8523.9728, Test accuracy=0.1000
==> Preparing data..
Files already downloaded and verified
Files already downloaded and verified

Epoch: 0
 [Step=50] Loss=2.2584 acc=0.1487 2551.6 examples/second
 [Step=100] Loss=2.2225 acc=0.1541 4022.4 examples/second
 [Step=150] Loss=2.2040 acc=0.1603 2889.7 examples/second
 Test Loss=2.1456, Test acc=0.1768
 Saving...

Epoch: 1
 [Step=200] Loss=2.1437 acc=0.1777 1498.7 examples/second
 [Step=250] Loss=2.1324 acc=0.1859 5755.3 examples/second
 [Step=300] Loss=2.1203 acc=0.1960 7315.4 examples/second
 [Step=350] Loss=2.1173 acc=0.2018 6008.7 examples/second
 Test Loss=2.0912, Test acc=0.2219
 Saving...

Epoch: 2
 [Step=400] Loss=2.0949 acc=0.2173 2328.7 examples/second
 [Step=450] Loss=2.0796 acc=0.2308 7142.5 examples/second
 [Step=500] Loss=2.0751 acc=0.2343 6791.7 examples/second
 [Step=550] Loss=2.0707 acc=0.2361 5729.8 examples/second
 Test Loss=2.0630, Test acc=0.2419
 Saving...

Epoch: 3
 [Step=600] Loss=2.0576 acc=0.2409 2538.4 examples/second
 [Step=650] Loss=2.0530 acc=0.2421 7277.8 examples/second
 [Step=700] Loss=2.0493 acc=0.2451 6790.3 examples/second
 [Step=750] Loss=2.0461 acc=0.2457 6486.7 examples/second
 Test Loss=2.0509, Test acc=0.2500
 Saving...

Epoch: 4
 [Step=800] Loss=2.0399 acc=0.2627 2230.6 examples/second
 [Step=850] Loss=2.0310 acc=0.2558 7058.7 examples/second
 [Step=900] Loss=2.0260 acc=0.2560 6324.5 examples/second
 [Step=950] Loss=2.0229 acc=0.2574 6454.7 examples/second
 Test Loss=2.0337, Test acc=0.2449

Epoch: 5
 [Step=1000] Loss=2.0018 acc=0.2617 1983.5 examples/second
 [Step=1050] Loss=2.0080 acc=0.2606 7184.5 examples/second
 [Step=1100] Loss=2.0030 acc=0.2634 6992.6 examples/second
 [Step=1150] Loss=2.0020 acc=0.2630 7130.1 examples/second
 Test Loss=2.0206, Test acc=0.2597
 Saving...

Epoch: 6
 [Step=1200] Loss=1.9806 acc=0.2803 1913.3 examples/second

[Step=1250] Loss=1.9885 acc=0.2740 5945.2 examples/second
[Step=1300] Loss=1.9872 acc=0.2716 6223.8 examples/second
[Step=1350] Loss=1.9816 acc=0.2729 6070.4 examples/second

Test Loss=1.9910, Test acc=0.2674

Saving...

Epoch: 7

[Step=1400] Loss=1.9636 acc=0.2856 2413.9 examples/second
[Step=1450] Loss=1.9622 acc=0.2817 5625.7 examples/second
[Step=1500] Loss=1.9618 acc=0.2820 6169.3 examples/second
[Step=1550] Loss=1.9566 acc=0.2813 6561.9 examples/second

Test Loss=1.9693, Test acc=0.2750

Saving...

Epoch: 8

[Step=1600] Loss=1.9517 acc=0.2808 2430.4 examples/second
[Step=1650] Loss=1.9424 acc=0.2841 5413.7 examples/second
[Step=1700] Loss=1.9401 acc=0.2849 5684.2 examples/second
[Step=1750] Loss=1.9356 acc=0.2879 6402.2 examples/second

Test Loss=2.0433, Test acc=0.2379

Epoch: 9

[Step=1800] Loss=1.9403 acc=0.2780 2304.2 examples/second
[Step=1850] Loss=1.9254 acc=0.2899 5910.4 examples/second
[Step=1900] Loss=1.9200 acc=0.2916 6326.7 examples/second
[Step=1950] Loss=1.9147 acc=0.2928 6645.6 examples/second

Test Loss=1.9648, Test acc=0.2771

Saving...

Epoch: 10

[Step=2000] Loss=1.8978 acc=0.2988 2007.0 examples/second
[Step=2050] Loss=1.8933 acc=0.3024 5328.3 examples/second
[Step=2100] Loss=1.8940 acc=0.3034 6084.0 examples/second
[Step=2150] Loss=1.8897 acc=0.3032 5355.3 examples/second

Test Loss=1.9089, Test acc=0.2939

Saving...

Epoch: 11

[Step=2200] Loss=1.8837 acc=0.3026 2341.1 examples/second
[Step=2250] Loss=1.8721 acc=0.3062 6857.4 examples/second
[Step=2300] Loss=1.8687 acc=0.3090 5330.3 examples/second
[Step=2350] Loss=1.8647 acc=0.3116 5505.8 examples/second

Test Loss=1.9105, Test acc=0.2882

Epoch: 12

[Step=2400] Loss=1.8482 acc=0.3193 2237.5 examples/second
[Step=2450] Loss=1.8453 acc=0.3168 5464.0 examples/second
[Step=2500] Loss=1.8443 acc=0.3168 5736.0 examples/second

Test Loss=1.9252, Test acc=0.2729

Epoch: 13

[Step=2550]	Loss=1.8271	acc=0.3008	2240.8 examples/second
[Step=2600]	Loss=1.8203	acc=0.3265	7001.6 examples/second
[Step=2650]	Loss=1.8133	acc=0.3278	5569.1 examples/second
[Step=2700]	Loss=1.8097	acc=0.3282	5655.1 examples/second

Test Loss=1.8428, Test acc=0.3200

Saving...

Epoch: 14

[Step=2750]	Loss=1.8031	acc=0.3340	2244.0 examples/second
[Step=2800]	Loss=1.7921	acc=0.3316	6283.3 examples/second
[Step=2850]	Loss=1.7925	acc=0.3335	6634.1 examples/second
[Step=2900]	Loss=1.7897	acc=0.3358	6687.3 examples/second

Test Loss=1.8365, Test acc=0.2960

Epoch: 15

[Step=2950]	Loss=1.7824	acc=0.3461	2300.1 examples/second
[Step=3000]	Loss=1.7727	acc=0.3393	5836.0 examples/second
[Step=3050]	Loss=1.7672	acc=0.3381	6780.5 examples/second
[Step=3100]	Loss=1.7629	acc=0.3405	6285.5 examples/second

Test Loss=1.8815, Test acc=0.2833

Epoch: 16

[Step=3150]	Loss=1.7534	acc=0.3474	2198.5 examples/second
[Step=3200]	Loss=1.7428	acc=0.3513	6244.0 examples/second
[Step=3250]	Loss=1.7403	acc=0.3530	6646.7 examples/second
[Step=3300]	Loss=1.7354	acc=0.3532	6492.4 examples/second

Test Loss=1.7803, Test acc=0.3290

Saving...

Epoch: 17

[Step=3350]	Loss=1.7069	acc=0.3707	2389.7 examples/second
[Step=3400]	Loss=1.7099	acc=0.3621	6052.0 examples/second
[Step=3450]	Loss=1.7125	acc=0.3596	5945.7 examples/second
[Step=3500]	Loss=1.7076	acc=0.3617	6578.9 examples/second

Test Loss=1.7614, Test acc=0.3492

Saving...

Epoch: 18

[Step=3550]	Loss=1.6993	acc=0.3684	2168.6 examples/second
[Step=3600]	Loss=1.6906	acc=0.3697	5199.1 examples/second
[Step=3650]	Loss=1.6898	acc=0.3688	5951.0 examples/second
[Step=3700]	Loss=1.6859	acc=0.3715	7146.8 examples/second

Test Loss=1.7413, Test acc=0.3392

Epoch: 19

```

[Step=3750]    Loss=1.6793      acc=0.3733      2451.1 examples/second
[Step=3800]    Loss=1.6771      acc=0.3738      7020.0 examples/second
[Step=3850]    Loss=1.6717      acc=0.3757      6958.3 examples/second
[Step=3900]    Loss=1.6672      acc=0.3772      6609.1 examples/second
Test Loss=1.7031, Test acc=0.3565
Saving...
Files already downloaded and verified
Test Loss=1.7031, Test accuracy=0.3565
Files already downloaded and verified
Test Loss=1.1891, Test accuracy=0.6081
==> Preparing data..
Files already downloaded and verified
Files already downloaded and verified

Epoch: 0
[Step=50]       Loss=1.4277      acc=0.5354      4222.7 examples/second
[Step=100]      Loss=1.1835      acc=0.6058      7022.2 examples/second
[Step=150]      Loss=1.0297      acc=0.6548      6769.1 examples/second
Test Loss=0.6359, Test acc=0.7825
Saving...

Epoch: 1
[Step=200]      Loss=0.5752      acc=0.8027      2173.6 examples/second
[Step=250]      Loss=0.5377      acc=0.8179      6475.0 examples/second
[Step=300]      Loss=0.5151      acc=0.8266      6194.4 examples/second
[Step=350]      Loss=0.4983      acc=0.8321      6088.8 examples/second
Test Loss=0.5133, Test acc=0.8250
Saving...

Epoch: 2
[Step=400]      Loss=0.4255      acc=0.8506      2096.5 examples/second
[Step=450]      Loss=0.4051      acc=0.8598      5950.1 examples/second
[Step=500]      Loss=0.3997      acc=0.8615      7218.1 examples/second
[Step=550]      Loss=0.3895      acc=0.8661      6015.9 examples/second
Test Loss=0.4873, Test acc=0.8383
Saving...

Epoch: 3
[Step=600]      Loss=0.3470      acc=0.8766      2682.0 examples/second
[Step=650]      Loss=0.3460      acc=0.8815      5688.7 examples/second
[Step=700]      Loss=0.3396      acc=0.8844      6365.3 examples/second
[Step=750]      Loss=0.3421      acc=0.8827      6453.5 examples/second
Test Loss=0.4538, Test acc=0.8501
Saving...

Epoch: 4
[Step=800]      Loss=0.3275      acc=0.8860      2154.1 examples/second
[Step=850]      Loss=0.3198      acc=0.8885      5757.1 examples/second

```

[Step=900]	Loss=0.3167	acc=0.8899	5449.9 examples/second
[Step=950]	Loss=0.3155	acc=0.8904	5455.2 examples/second
Test Loss=0.4395, Test acc=0.8558			
Saving...			
Epoch: 5			
[Step=1000]	Loss=0.2984	acc=0.9020	2189.2 examples/second
[Step=1050]	Loss=0.3011	acc=0.8963	6190.7 examples/second
[Step=1100]	Loss=0.2990	acc=0.8958	7163.5 examples/second
[Step=1150]	Loss=0.2936	acc=0.8980	6157.8 examples/second
Test Loss=0.4294, Test acc=0.8606			
Saving...			
Epoch: 6			
[Step=1200]	Loss=0.2749	acc=0.8991	1772.0 examples/second
[Step=1250]	Loss=0.2748	acc=0.9019	5588.0 examples/second
[Step=1300]	Loss=0.2736	acc=0.9026	5355.3 examples/second
[Step=1350]	Loss=0.2742	acc=0.9025	4679.5 examples/second
Test Loss=0.4084, Test acc=0.8677			
Saving...			
Epoch: 7			
[Step=1400]	Loss=0.2616	acc=0.9081	2051.1 examples/second
[Step=1450]	Loss=0.2589	acc=0.9086	6170.7 examples/second
[Step=1500]	Loss=0.2606	acc=0.9071	5862.7 examples/second
[Step=1550]	Loss=0.2606	acc=0.9075	5641.9 examples/second
Test Loss=0.4032, Test acc=0.8680			
Saving...			
Epoch: 8			
[Step=1600]	Loss=0.2536	acc=0.9092	2380.5 examples/second
[Step=1650]	Loss=0.2542	acc=0.9104	7170.1 examples/second
[Step=1700]	Loss=0.2512	acc=0.9118	6127.2 examples/second
[Step=1750]	Loss=0.2497	acc=0.9122	5993.7 examples/second
Test Loss=0.4039, Test acc=0.8685			
Saving...			
Epoch: 9			
[Step=1800]	Loss=0.2454	acc=0.9147	2019.8 examples/second
[Step=1850]	Loss=0.2425	acc=0.9158	6349.9 examples/second
[Step=1900]	Loss=0.2408	acc=0.9151	6465.4 examples/second
[Step=1950]	Loss=0.2417	acc=0.9149	6120.4 examples/second
Test Loss=0.3963, Test acc=0.8713			
Saving...			
Epoch: 10			
[Step=2000]	Loss=0.2398	acc=0.9127	2246.5 examples/second
[Step=2050]	Loss=0.2355	acc=0.9161	5403.4 examples/second

[Step=2100]	Loss=0.2376	acc=0.9152	5753.6 examples/second
[Step=2150]	Loss=0.2367	acc=0.9162	5892.1 examples/second
Test Loss=0.3939, Test acc=0.8755			
Saving...			
Epoch: 11			
[Step=2200]	Loss=0.2332	acc=0.9165	2256.8 examples/second
[Step=2250]	Loss=0.2312	acc=0.9179	5677.7 examples/second
[Step=2300]	Loss=0.2278	acc=0.9191	6285.6 examples/second
[Step=2350]	Loss=0.2310	acc=0.9180	5419.5 examples/second
Test Loss=0.3791, Test acc=0.8772			
Saving...			
Epoch: 12			
[Step=2400]	Loss=0.2261	acc=0.9202	2164.6 examples/second
[Step=2450]	Loss=0.2236	acc=0.9224	5809.2 examples/second
[Step=2500]	Loss=0.2241	acc=0.9219	6129.8 examples/second
Test Loss=0.3780, Test acc=0.8789			
Saving...			
Epoch: 13			
[Step=2550]	Loss=0.1799	acc=0.9238	1804.3 examples/second
[Step=2600]	Loss=0.2177	acc=0.9222	5828.9 examples/second
[Step=2650]	Loss=0.2184	acc=0.9233	6692.2 examples/second
[Step=2700]	Loss=0.2185	acc=0.9237	6149.5 examples/second
Test Loss=0.3976, Test acc=0.8757			
Epoch: 14			
[Step=2750]	Loss=0.1851	acc=0.9342	1931.4 examples/second
[Step=2800]	Loss=0.2050	acc=0.9252	4914.1 examples/second
[Step=2850]	Loss=0.2103	acc=0.9252	4737.3 examples/second
[Step=2900]	Loss=0.2133	acc=0.9250	6176.9 examples/second
Test Loss=0.3798, Test acc=0.8790			
Saving...			
Epoch: 15			
[Step=2950]	Loss=0.2011	acc=0.9301	2068.7 examples/second
[Step=3000]	Loss=0.2064	acc=0.9265	6152.8 examples/second
[Step=3050]	Loss=0.2099	acc=0.9259	5452.2 examples/second
[Step=3100]	Loss=0.2087	acc=0.9259	6601.5 examples/second
Test Loss=0.3737, Test acc=0.8813			
Saving...			
Epoch: 16			
[Step=3150]	Loss=0.2170	acc=0.9196	2050.2 examples/second
[Step=3200]	Loss=0.2105	acc=0.9246	6915.4 examples/second
[Step=3250]	Loss=0.2071	acc=0.9266	5686.8 examples/second
[Step=3300]	Loss=0.2051	acc=0.9268	6234.8 examples/second

Test Loss=0.3775, Test acc=0.8803

Epoch: 17

[Step=3350]	Loss=0.1894	acc=0.9368	2062.4 examples/second
[Step=3400]	Loss=0.1977	acc=0.9299	7147.3 examples/second
[Step=3450]	Loss=0.2038	acc=0.9288	5848.8 examples/second
[Step=3500]	Loss=0.2009	acc=0.9296	6100.3 examples/second

Test Loss=0.3755, Test acc=0.8786

Epoch: 18

[Step=3550]	Loss=0.1927	acc=0.9318	2099.4 examples/second
[Step=3600]	Loss=0.1999	acc=0.9315	6471.2 examples/second
[Step=3650]	Loss=0.1972	acc=0.9315	5835.3 examples/second
[Step=3700]	Loss=0.1970	acc=0.9315	6041.3 examples/second

Test Loss=0.3723, Test acc=0.8828

Saving...

Epoch: 19

[Step=3750]	Loss=0.1721	acc=0.9381	2535.8 examples/second
[Step=3800]	Loss=0.1910	acc=0.9331	5669.5 examples/second
[Step=3850]	Loss=0.1923	acc=0.9328	6447.8 examples/second
[Step=3900]	Loss=0.1938	acc=0.9324	5829.4 examples/second

Test Loss=0.3688, Test acc=0.8830

Saving...

Files already downloaded and verified

Test Loss=0.3688, Test accuracy=0.8830

Files already downloaded and verified

Test Loss=0.3924, Test accuracy=0.8793

==> Preparing data..

Files already downloaded and verified

Files already downloaded and verified

Epoch: 0

[Step=50]	Loss=0.6967	acc=0.7641	4052.4 examples/second
[Step=100]	Loss=0.4982	acc=0.8292	5687.7 examples/second
[Step=150]	Loss=0.4115	acc=0.8585	5593.0 examples/second

Test Loss=0.3715, Test acc=0.8811

Saving...

Epoch: 1

[Step=200]	Loss=0.2036	acc=0.9375	2385.5 examples/second
[Step=250]	Loss=0.2031	acc=0.9301	6396.2 examples/second
[Step=300]	Loss=0.2006	acc=0.9309	6802.8 examples/second
[Step=350]	Loss=0.1965	acc=0.9321	6917.3 examples/second

Test Loss=0.3271, Test acc=0.8893

Saving...

Epoch: 2

[Step=400]	Loss=0.1735	acc=0.9409	2093.1 examples/second
[Step=450]	Loss=0.1761	acc=0.9380	6020.2 examples/second
[Step=500]	Loss=0.1766	acc=0.9383	5815.3 examples/second
[Step=550]	Loss=0.1762	acc=0.9397	5984.7 examples/second
Test Loss=0.3157, Test acc=0.8940			
Saving...			
Epoch: 3			
[Step=600]	Loss=0.1613	acc=0.9486	2228.8 examples/second
[Step=650]	Loss=0.1710	acc=0.9437	6017.6 examples/second
[Step=700]	Loss=0.1668	acc=0.9451	6502.6 examples/second
[Step=750]	Loss=0.1691	acc=0.9440	5935.8 examples/second
Test Loss=0.3165, Test acc=0.8964			
Saving...			
Epoch: 4			
[Step=800]	Loss=0.1518	acc=0.9492	2282.8 examples/second
[Step=850]	Loss=0.1561	acc=0.9473	6279.8 examples/second
[Step=900]	Loss=0.1583	acc=0.9467	5865.7 examples/second
[Step=950]	Loss=0.1580	acc=0.9463	5459.3 examples/second
Test Loss=0.3109, Test acc=0.8971			
Saving...			
Epoch: 5			
[Step=1000]	Loss=0.1585	acc=0.9467	1470.3 examples/second
[Step=1050]	Loss=0.1578	acc=0.9458	5756.0 examples/second
[Step=1100]	Loss=0.1571	acc=0.9461	6313.0 examples/second
[Step=1150]	Loss=0.1543	acc=0.9468	5814.4 examples/second
Test Loss=0.3100, Test acc=0.9001			
Saving...			
Epoch: 6			
[Step=1200]	Loss=0.1564	acc=0.9437	2333.5 examples/second
[Step=1250]	Loss=0.1507	acc=0.9463	5865.2 examples/second
[Step=1300]	Loss=0.1472	acc=0.9477	5825.0 examples/second
[Step=1350]	Loss=0.1452	acc=0.9487	5583.1 examples/second
Test Loss=0.3526, Test acc=0.8932			
Epoch: 7			
[Step=1400]	Loss=0.1462	acc=0.9513	2358.5 examples/second
[Step=1450]	Loss=0.1482	acc=0.9503	6254.3 examples/second
[Step=1500]	Loss=0.1451	acc=0.9511	6041.6 examples/second
[Step=1550]	Loss=0.1449	acc=0.9509	6297.7 examples/second
Test Loss=0.3447, Test acc=0.8991			
Epoch: 8			
[Step=1600]	Loss=0.1276	acc=0.9556	2116.9 examples/second
[Step=1650]	Loss=0.1265	acc=0.9560	5582.5 examples/second

[Step=1700]	Loss=0.1320	acc=0.9530	5041.7 examples/second
[Step=1750]	Loss=0.1342	acc=0.9529	4866.9 examples/second
Test Loss=0.3123, Test acc=0.9019			
Saving...			
Epoch: 9			
[Step=1800]	Loss=0.1250	acc=0.9539	2431.3 examples/second
[Step=1850]	Loss=0.1270	acc=0.9551	4829.3 examples/second
[Step=1900]	Loss=0.1276	acc=0.9543	5838.1 examples/second
[Step=1950]	Loss=0.1294	acc=0.9538	6848.1 examples/second
Test Loss=0.3368, Test acc=0.9031			
Saving...			
Epoch: 10			
[Step=2000]	Loss=0.1238	acc=0.9568	2374.8 examples/second
[Step=2050]	Loss=0.1251	acc=0.9556	5640.0 examples/second
[Step=2100]	Loss=0.1247	acc=0.9552	6107.3 examples/second
[Step=2150]	Loss=0.1245	acc=0.9551	6087.2 examples/second
Test Loss=0.3415, Test acc=0.9003			
Epoch: 11			
[Step=2200]	Loss=0.1217	acc=0.9578	2327.4 examples/second
[Step=2250]	Loss=0.1211	acc=0.9566	6317.9 examples/second
[Step=2300]	Loss=0.1195	acc=0.9576	6078.1 examples/second
[Step=2350]	Loss=0.1205	acc=0.9574	6444.7 examples/second
Test Loss=0.3428, Test acc=0.9015			
Epoch: 12			
[Step=2400]	Loss=0.1196	acc=0.9557	2055.1 examples/second
[Step=2450]	Loss=0.1187	acc=0.9564	6098.4 examples/second
[Step=2500]	Loss=0.1199	acc=0.9563	6342.8 examples/second
Test Loss=0.3553, Test acc=0.9015			
Epoch: 13			
[Step=2550]	Loss=0.1105	acc=0.9629	1285.4 examples/second
[Step=2600]	Loss=0.1130	acc=0.9620	4766.6 examples/second
[Step=2650]	Loss=0.1148	acc=0.9603	5926.2 examples/second
[Step=2700]	Loss=0.1122	acc=0.9612	5831.9 examples/second
Test Loss=0.3735, Test acc=0.9012			
Epoch: 14			
[Step=2750]	Loss=0.1233	acc=0.9564	2233.3 examples/second
[Step=2800]	Loss=0.1131	acc=0.9602	7137.4 examples/second
[Step=2850]	Loss=0.1098	acc=0.9621	7470.3 examples/second
[Step=2900]	Loss=0.1101	acc=0.9623	6062.7 examples/second
Test Loss=0.3353, Test acc=0.9035			
Saving...			

Epoch: 15

[Step=2950]	Loss=0.1143	acc=0.9586	2158.8 examples/second
[Step=3000]	Loss=0.1147	acc=0.9590	6172.7 examples/second
[Step=3050]	Loss=0.1135	acc=0.9602	7081.9 examples/second
[Step=3100]	Loss=0.1115	acc=0.9610	6637.8 examples/second

Test Loss=0.3364, Test acc=0.9027

Epoch: 16

[Step=3150]	Loss=0.1184	acc=0.9570	2410.7 examples/second
[Step=3200]	Loss=0.1092	acc=0.9616	7257.2 examples/second
[Step=3250]	Loss=0.1094	acc=0.9609	7521.6 examples/second
[Step=3300]	Loss=0.1083	acc=0.9617	5921.7 examples/second

Test Loss=0.3262, Test acc=0.9011

Epoch: 17

[Step=3350]	Loss=0.1003	acc=0.9642	2214.1 examples/second
[Step=3400]	Loss=0.1039	acc=0.9633	6910.0 examples/second
[Step=3450]	Loss=0.1051	acc=0.9623	6623.9 examples/second
[Step=3500]	Loss=0.1048	acc=0.9627	6385.0 examples/second

Test Loss=0.3378, Test acc=0.9035

Epoch: 18

[Step=3550]	Loss=0.0929	acc=0.9677	2317.9 examples/second
[Step=3600]	Loss=0.0987	acc=0.9667	6866.8 examples/second
[Step=3650]	Loss=0.1022	acc=0.9648	6104.4 examples/second
[Step=3700]	Loss=0.1022	acc=0.9642	4787.4 examples/second

Test Loss=0.3430, Test acc=0.9055

Saving...

Epoch: 19

[Step=3750]	Loss=0.0954	acc=0.9674	2287.8 examples/second
[Step=3800]	Loss=0.1015	acc=0.9649	6324.0 examples/second
[Step=3850]	Loss=0.0972	acc=0.9665	5853.9 examples/second
[Step=3900]	Loss=0.0986	acc=0.9656	5434.9 examples/second

Test Loss=0.3385, Test acc=0.9028

Files already downloaded and verified

Test Loss=0.3385, Test accuracy=0.9028

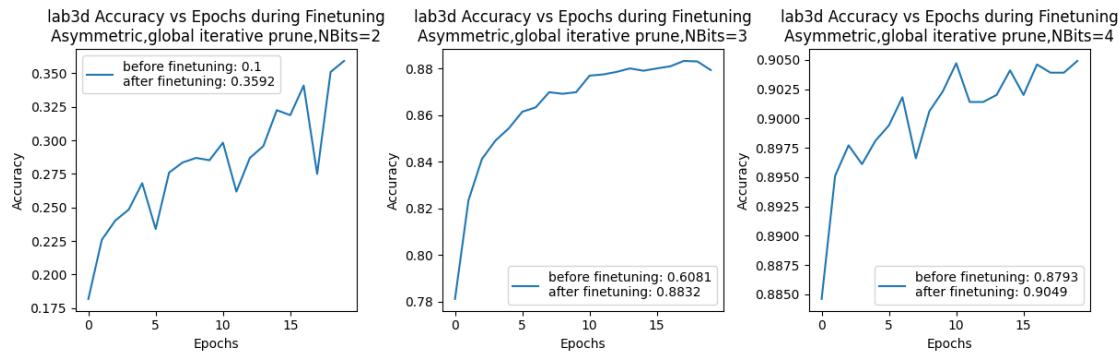
```
[25]: fig, ax = plt.subplots(1, 3, figsize=(12, 4))

for iii, nbits in enumerate(nbites_lab3d):
    xx = range(n_epochs)
    ax[iii].plot(xx, test_acc_arrays_finetune[nbits], label='before finetuning: %g\nafter finetuning: %g'
                 % (test_acc_lab3d_b4_finetune[nbits], test_acc_lab3d_after_finetune[nbits]))
    ax[iii].set_xlabel('Epochs')
```

```

ax[iii].set_ylabel('Accuracy')
ax[iii].set_title('lab3d Accuracy vs Epochs during Finetuning\n'
                  'Asymmetric,global iterative prune,NBits=%d' % nbits)
ax[iii].legend()
fig.tight_layout()
plt.savefig('Figures/lab3d.pdf', dpi=500, bbox_inches='tight')

```



0.0.8 Lab3 (e) Symmetric quantization

Implement symmetric quantization in FP_layers.py, and repeat the process in (b)

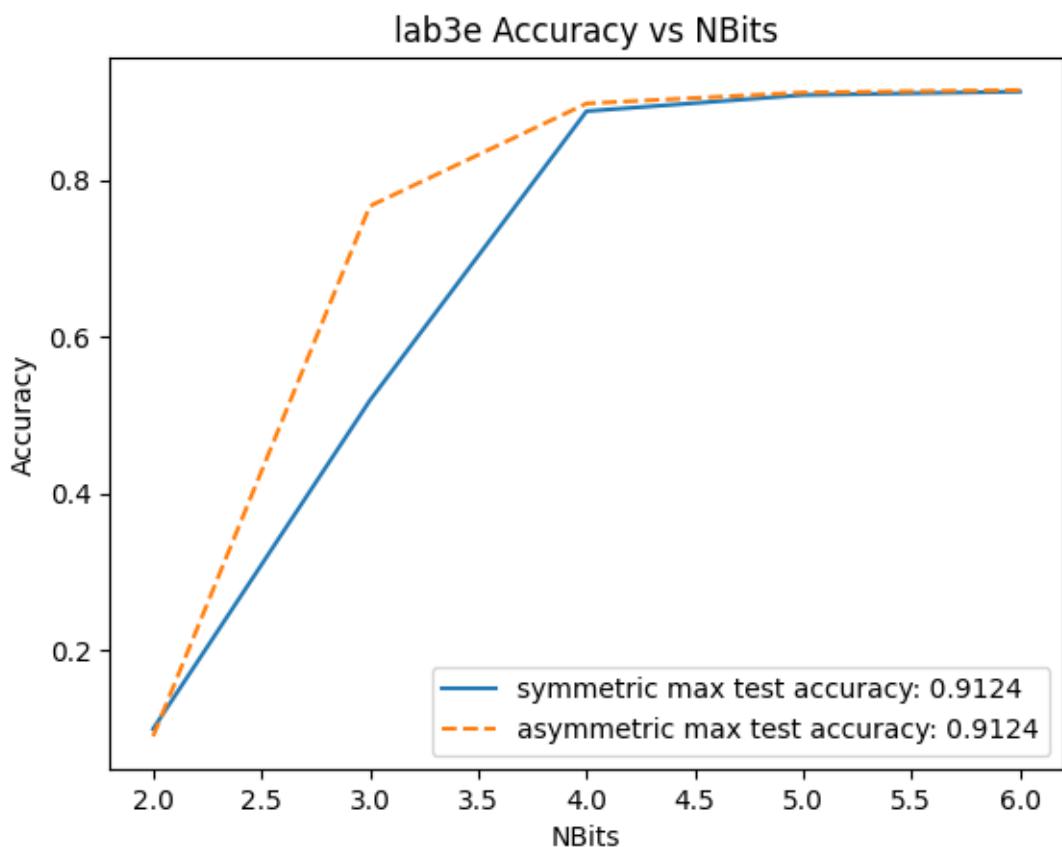
```
[2]: # check the performance of symmetric quantization with 6, 5, 4, 3, 2 bits
# Define quantized model and load weight
Nbts_arr_lab3e = [2, 3, 4, 5, 6]
test_acc_lab3e = []
for iii, nb in enumerate(Nbts_arr_lab3e):
    Nbts = nb #Change this value

    net = ResNetCIFAR(num_layers=20, Nbts=Nbts, symmetric=True) # symmetric
    net = net.to(device)
    net.load_state_dict(torch.load("pretrained_model.pt"))
    test_acc_lab3e.append(test(net))
```

Files already downloaded and verified
Test Loss=42.7983, Test accuracy=0.1000
Files already downloaded and verified
Test Loss=2.3751, Test accuracy=0.5186
Files already downloaded and verified
Test Loss=0.4227, Test accuracy=0.8875
Files already downloaded and verified
Test Loss=0.3518, Test accuracy=0.9081
Files already downloaded and verified
Test Loss=0.3276, Test accuracy=0.9124

```
[5]: fig, ax = plt.subplots(1, 1)
xx = Nbits_arr_lab3e
ax.plot(xx, test_acc_lab3e, label='symmetric max test accuracy: %g' % np.
        ↪max(test_acc_lab3e))
ax.plot(xx, test_acc_lab3b, '--', label='asymmetric max test accuracy: %g' % np.
        ↪max(test_acc_lab3e))
ax.set_xlabel('NBits')
ax.set_ylabel('Accuracy')
ax.set_title('lab3e Accuracy vs NBits')
ax.legend()

plt.savefig('Figures/lab3e.pdf', dpi=500, bbox_inches='tight')
```



```
[11]: def plot_quant_symmetry(symmetric, NBits=3, bins=100, save=False):

    net = ResNetCIFAR(num_layers=20) # this is for original weights
    net = net.to(device)
    net.load_state_dict(torch.load("pretrained_model.pt"))
    weights_dict_quant = {}
    w_dict_orig = {}
```

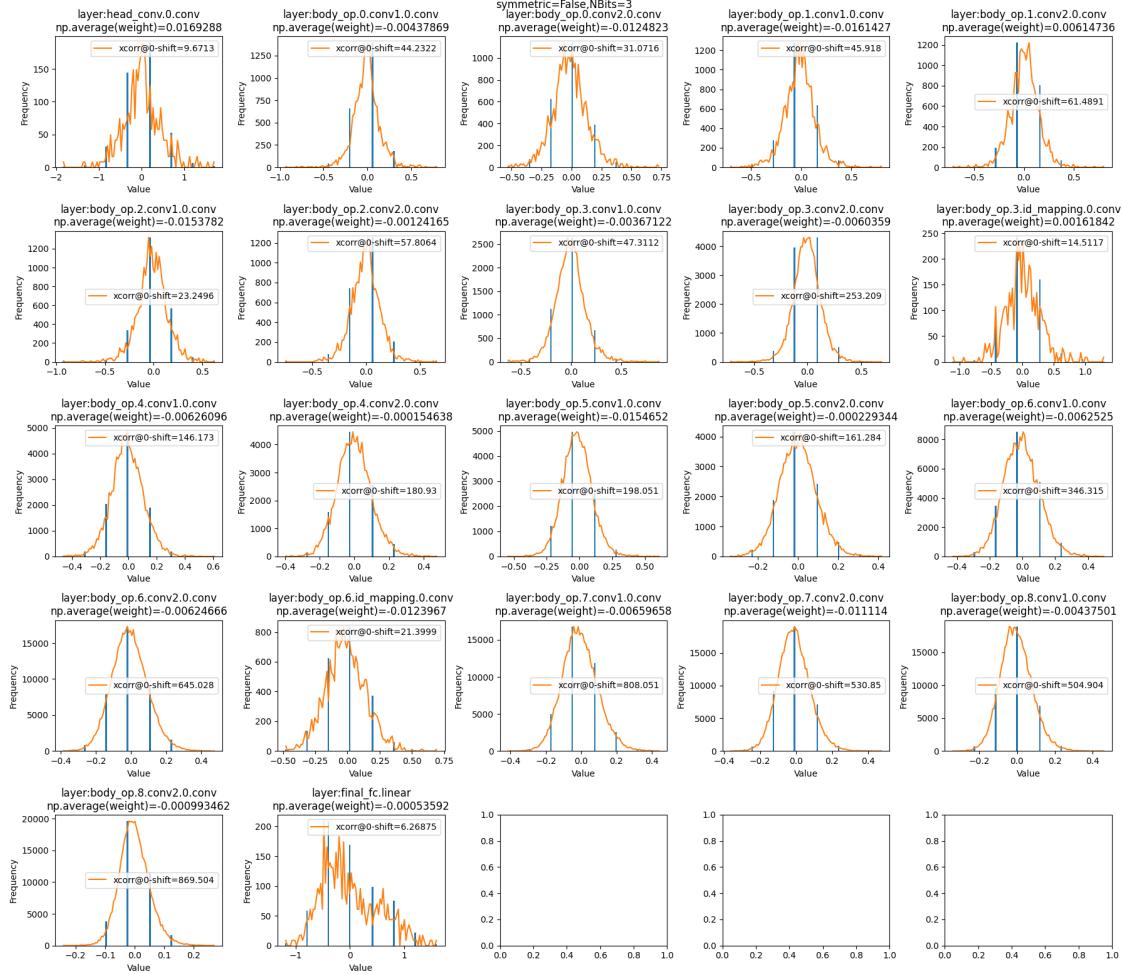
```

conv_list = []
for name, module in net.named_modules():
    if isinstance(module, nn.Conv2d) or isinstance(module, nn.Linear):
        w_dict_orig[str(name)] = module.weight.detach().cpu().numpy().
        flatten()
        www = STE().forward(ctx=None, w=module.weight, bit=NBits,
        ↪symmetric=symmetric)
        np_weights = www.cpu().detach().numpy()
        weights_dict_quant[str(name)] = np_weights.flatten()
# print(len(weights_dict_quant.keys()))
# print(weights_dict_quant.keys())
fig, ax = plt.subplots(5, 5, figsize=(18, 16))
ax_flat = ax.flatten()
for iii, layer_name in enumerate(list(weights_dict_quant.keys())):
    y_val, edges = np.histogram(w_dict_orig[layer_name], bins=bins)
    x_val = np.convolve(edges, [0.5, 0.5])[1:-1] # compute average btw
    ↪adjacent elements of edges

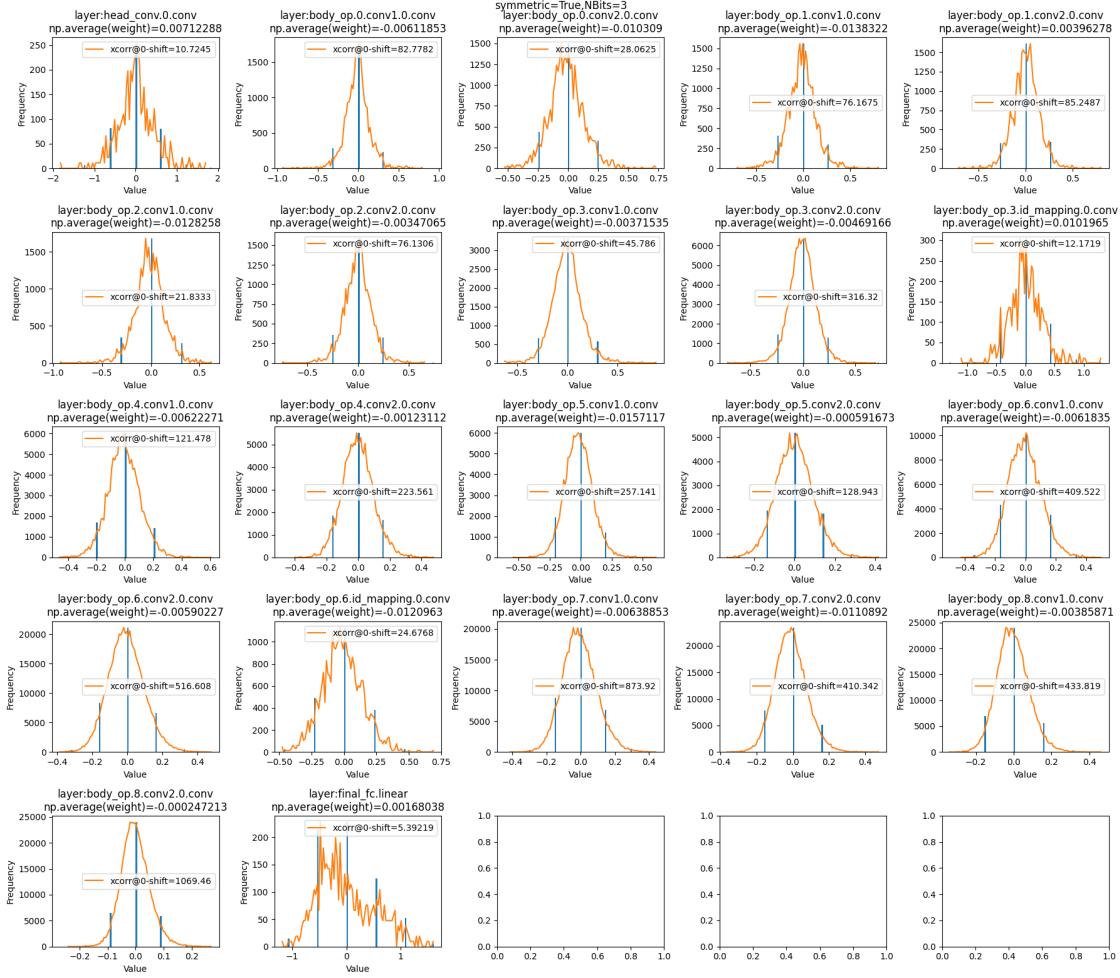
    y_val_quant, _ = np.histogram(weights_dict_quant[layer_name], bins=bins)
    ax_flat[iii].hist(weights_dict_quant[layer_name], bins=bins)
    ax_flat[iii].set_title(str("layer:%s\nnp.average(weight)=%g" %
        (layer_name, np.
        ↪average(weights_dict_quant[layer_name]))))
    ax_flat[iii].set_xlabel('Value')
    ax_flat[iii].set_ylabel('Frequency')
    conv = np.convolve(y_val_quant, y_val, 'valid')
    # if len(conv) % 2 == 0:
    #     zeroshift_convolve = 0.5*(conv[int((len(conv)-1)/2)] +
    ↪conv[int((len(conv)+1)/2)]) \
        # / len(weights_dict_quant[layer_name])
    # else:
    #     zeroshift_convolve = (conv[int(len(conv)/2)]) /
    ↪len(weights_dict_quant[layer_name])
    zeroshift_convolve = conv / len(weights_dict_quant[layer_name])
    ax_flat[iii].plot(x_val, y_val * (np.max(y_val_quant) / np.max(y_val)),
        label='xcorr@0-shift=%g' % zeroshift_convolve)
    ax_flat[iii].legend()
    conv_list.append(zeroshift_convolve)
fig.suptitle('symmetric=' + str(symmetric) + ',NBits=' + str(NBits))
fig.tight_layout()
if save:
    plt.savefig('Figures/lab3e_hist_sym=%s_NBits=%d.pdf' % (symmetric,
    ↪NBits), dpi=500, bbox_inches='tight')
return conv_list

```

```
c_asym = plot_quant_symmetry(symmetric=False, save=True)
```



```
[12]: c_sym = plot_quant_symmetry(symmetric=True, save=True)
```



```
[20]: difference = np.array(c_asym) - np.array(c_sym)
plt.plot(0.5*difference / (np.array(c_asym) + np.array(c_sym)))
plt.plot([0, 20], [0, 0])
```

```
[20]: [<matplotlib.lines.Line2D at 0x7f6401a07fa0>]
```

