

# HWK5\_main

November 22, 2023

## 1 Homework 5: Adversarial Attacks and Defenses

Duke University

ECE661 Fall 2022

### 1.1 Setup

You shouldn't have to change anything in these cells

```
[1]: import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import datasets, transforms
import matplotlib.pyplot as plt
import numpy as np
import random
import os

# Custom
import models
import attacks

device = "cuda" if torch.cuda.is_available() else "cpu"
print("device:", device)
```

device: cuda

```
[2]: train_loader = torch.utils.data.DataLoader(
    datasets.FashionMNIST('./data', train=True, download=True,
    ↪transform=transforms.ToTensor()),
    batch_size = 64, shuffle=True, )
test_loader = torch.utils.data.DataLoader(
    datasets.FashionMNIST('./data', train=False, download=True,
    ↪transform=transforms.ToTensor()),
    batch_size = 64, shuffle=False, )
```

```
[3]: def test_model mdl, loader, device):
    mdl.eval()
    running_correct = 0.
    running_loss = 0.
    running_total = 0.
    with torch.no_grad():
        for batch_idx, (data, labels) in enumerate(loader):
            data = data.to(device); labels = labels.to(device)
            clean_outputs = mdl(data)
            clean_loss = F.cross_entropy(clean_outputs, labels)
            _, clean_preds = clean_outputs.max(1)
            running_correct += clean_preds.eq(labels).sum().item()
            running_loss += clean_loss.item()
            running_total += labels.size(0)
    clean_acc = running_correct/running_total
    clean_loss = running_loss/len(loader)
    mdl.train()
    return clean_acc, clean_loss
```

## 1.2 Model training - Lab 1 a

Train a model and save the checkpoint. This cell is used in Lab-1 (for Lab-3, please see a cell below)

```
[ ]: ## Pick a model architecture
which_net = 'B'
test_acc_arr_lab1a = []

if which_net == 'A':
    net = models.NetA().to(device)
    ## Checkpoint name for this model
    model_checkpoint = "netA_standard.pt"
if which_net == 'B':
    net = models.NetB().to(device)
    model_checkpoint = "netB_standard.pt"

## Basic training params
num_epochs = 20
initial_lr = 0.001
lr_decay_epoch = 15

optimizer = torch.optim.Adam(net.parameters(), lr=initial_lr)

## Training Loop
for epoch in range(num_epochs):
    net.train()
    train_correct = 0.
```

```

train_loss = 0.
train_total = 0.
for batch_idx, (data, labels) in enumerate(train_loader):
    data = data.to(device); labels = labels.to(device)

    # Forward pass
    outputs = net(data)
    net.zero_grad()
    optimizer.zero_grad()
    # Compute loss, gradients, and update params
    loss = F.cross_entropy(outputs, labels)
    loss.backward()
    optimizer.step()
    # Update stats
    _, preds = outputs.max(1)
    train_correct += preds.eq(labels).sum().item()
    train_loss += loss.item()
    train_total += labels.size(0)

# End of training epoch
test_acc, test_loss = test_model(net, test_loader, device)
test_acc_arr_lab1a.append(test_acc)
print("Epoch: [ {} / {} ]; TrainAcc: {:.5f}; TrainLoss: {:.5f}; TestAcc: {:.
↪5f}; TestLoss: {:.5f}".format(
    epoch, num_epochs, train_correct/train_total, train_loss/
↪len(train_loader),
    test_acc, test_loss,
))
# Save model
torch.save(net.state_dict(), model_checkpoint)

# Update LR
if epoch == lr_decay_epoch:
    for param_group in optimizer.param_groups:
        param_group['lr'] = initial_lr*0.1

print("Done!")

```

```

[ ]: fig, ax = plt.subplots(1, 1)
xx = range(num_epochs)
ax.plot(xx, test_acc_arr_lab1a, label='final test accuracy: %g' %
↪(test_acc_arr_lab1a[-1]))
ax.set_xlabel('Epochs')
ax.set_ylabel('Accuracy')
ax.set_title('lab1a net%s Accuracy vs Epochs' % which_net)
ax.legend()

```

```
# plt.savefig('Figures/lab1a_net%s.pdf' % which_net, dpi=500,
↳bbox_inches='tight')
```

Visualize some perturbed samples - Lab-1 b/c/d

```
[12]: def lab1_return_adv_data(model, device, dat, lbl, eps, alpha, iters,
↳rand_start, which_part):
    if which_part == 'b':
        return attacks.PGD_attack(model, device, dat, lbl, eps, alpha, iters,
↳rand_start)[0]
    elif which_part == 'c1':
        return attacks.FGSM_attack(model, device, dat, lbl, eps)[0]
    elif which_part == 'c2':
        return attacks.rFGSM_attack(model, device, dat, lbl, eps)[0]
    elif which_part == 'd':
        return attacks.FGM_L2_attack(model, device, dat, lbl, eps)[0]
    else:
        raise KeyError

def lab1_plot_visualisations(dat, how_many, indices, clean, classes, preds,
↳lab1_part, eps, lab1_save_name, save=False):
    fig, ax = plt.subplots(1, how_many, figsize=(15, 0.58*how_many))
    for jj in range(how_many):
        ax[jj].imshow(dat[inds[jj],0].cpu().numpy(), cmap='gray')
        ax[jj].axis("off")
        if clean:
            ax[jj].set_title("clean. pred={}".format(classes[preds[inds[jj]]]))
        else:
            ax[jj].set_title("adv. pred={}".format(classes[preds[inds[jj]]]))

    fig.suptitle("eps=%g,%s" % (eps, lab1_save_name[lab1_part]))
    plt.tight_layout()
    # plt.show()
    if save:
        if not clean:
            plt.savefig('Figures/lab1%s_netA_eps_%g_%s.pdf' % (lab1_part, eps,
↳lab1_save_name[lab1_part]), dpi=500, bbox_inches='tight')
        else:
            plt.savefig('Figures/lab1_visualisation_data.pdf', dpi=500,
↳bbox_inches='tight')
    # plt.close()
    return
```

```
[ ]: classes = ["t-shirt",
↳"trouser", "pullover", "dress", "coat", "sandal", "shirt", "sneaker", "bag", "boot"]
lab1_parts = ['b', 'c1', 'c2', 'd']
```

```

lab1_save_name = {'b': 'PGD_attack', 'c1': 'FGSM_attack', 'c2': 'rFGSM_attack',
↳ 'd': 'FGM_L2_attack'}
plt_data = False

for data, labels in test_loader:
    data, labels = data.to(device), labels.to(device)
    inds = random.sample(list(range(data.size(0))),6)
    for lab1_part in lab1_parts:
        net = models.NetA().to(device)
        net.load_state_dict(torch.load("netA_standard.pt"))
        # lab1_part = 'd' # possible values: 'b', 'c1', 'c2', 'd'; change this
↳ to plot lab1 b/c/d
        EPS_list_lab1 = np.array([0, 0.005, 0.02, 0.05, 0.075, 0.1, 0.15, 0.2])
        if lab1_part == 'd':
            EPS_list_lab1 = np.array([0, 0.3, 1, 1.5, 2, 3, 3.5, 4])
        print('EPS_list_lab1', EPS_list_lab1)
        for epsilon in EPS_list_lab1:
            ###
            # Compute and apply adversarial perturbation to data
            # EPS in [0.0, 0.2]
            EPS = epsilon
            if lab1_part == 'b' or lab1_part == 'c1' or lab1_part == 'c2':
                assert EPS <= 0.2 and EPS >= 0.0
            elif lab1_part == 'd':
                assert EPS <= 4 and EPS >= 0.0
            else:
                raise KeyError("check lab1_part param")
            ITS = 10
            ALP = 1.85 * (EPS/ITS)
            adv_data = lab1_return_adv_data(model=net, device=device, dat=data,
↳ lbl=labels, eps=EPS, alpha=ALP, iters=ITS,
                                rand_start=True,
↳ which_part=lab1_part)
            ###

            # Compute preds
            with torch.no_grad():
                clean_outputs = net(data)
                _, clean_preds = clean_outputs.max(1)
                clean_preds = clean_preds.cpu().squeeze().numpy()
                adv_outputs = net(adv_data)
                _, adv_preds = adv_outputs.max(1)
                adv_preds = adv_preds.cpu().squeeze().numpy()
            # if not plt_data:
            #     lab1_plot_visualisations(dat=data, how_many=6, indices=inds,
↳ clean=True,

```

```

#                                     classes=classes, preds=clean_preds,␣
↪lab1_part=lab1_part,
#                                     eps=EPS,␣
↪lab1_save_name=lab1_save_name, save=False)
#     plt_data = True
# else:
#     lab1_plot_visualisations(dat=adv_data, how_many=6,␣
↪indices=inds, clean=False,
#                                     classes=classes, preds=adv_preds,␣
↪lab1_part=lab1_part,
#                                     eps=EPS,␣
↪lab1_save_name=lab1_save_name, save=False)
# Plot some samples
plt.figure(figsize=(15,5))
for jj in range(6):
    plt.subplot(2, 6, jj+1)
    plt.imshow(data[inds[jj],0].cpu().numpy(), cmap='gray')
    plt.axis("off")
    plt.title("clean. pred={}".format(classes[clean_preds[inds[jj]]]))
↪format(classes[clean_preds[inds[jj]]]))
    for jj in range(6):
        plt.subplot(2, 6, 6+jj+1)
        plt.imshow(adv_data[inds[jj],0].cpu().numpy(), cmap='gray')
        plt.axis("off")
        plt.title("adv. pred={}".format(classes[adv_preds[inds[jj]]]))
    plt.suptitle("eps=%g,%s" % (EPS, lab1_save_name[lab1_part]))
    plt.tight_layout()
    # plt.show()
    plt.savefig('Figures/lab1%s_netA_eps_%g_%s.pdf' % (lab1_part, EPS,␣
↪lab1_save_name[lab1_part]), dpi=500, bbox_inches='tight')
    plt.close()

break

```

### 1.3 Test Attacks - Whitebox & Blackbox, lab 2 b/c/d

Don't forget to plot accuracy vs. epsilon curves!

```

[ ]: def plot_lab2_white_black(epsilon_arr, white_acc_d, black_acc_d, label_names,␣
↪save=False):
    fig, ax = plt.subplots(1, 2, figsize=(12, 5))
    for label_name in label_names:
        ax[0].plot(epsilon_arr, white_acc_d[label_name], label=label_name)
        ax[0].set_xlabel('eps')
        ax[0].set_ylabel('Accuracy')
        ax[0].set_title('Whitebox Attack')

        ax[1].plot(epsilon_arr, black_acc_d[label_name], label=label_name)

```

```

        ax[1].set_xlabel('eps')
        ax[1].set_ylabel('Accuracy')
        ax[1].set_title('Blackbox Attack')
    ax[0].legend()
    ax[1].legend()
    fig.tight_layout()
    if save:
        plt.savefig('Figures/lab2bcd_attacks.pdf', dpi=500, bbox_inches='tight')
    return fig, ax

def lab2_bcd_return_adv_data(model, device, dat, lbl, eps, alpha, iters,
    ↪rand_start, question_label):
    if question_label == 'Random':
        return attacks.random_noise_attack(model=None, device=device, dat=dat,
    ↪eps=eps)[0]
    elif question_label == 'FGSM':
        return attacks.FGSM_attack(model, device, dat, lbl, eps)[0]
    elif question_label == 'rFGSM':
        return attacks.rFGSM_attack(model, device, dat, lbl, eps)[0]
    elif question_label == 'PGD':
        return attacks.PGD_attack(model, device, dat, lbl, eps, alpha, iters,
    ↪rand_start)[0]
    else:
        raise KeyError

```

```

[ ]: EPS_list_lab2 = np.linspace(0, 0.1, 11)
print('EPS_list_lab2', EPS_list_lab2)
lab2_label = ['Random', 'FGSM', 'rFGSM', 'PGD']
# lab2_label = {'FGSM'}
white_acc_dict, black_acc_dict = {'Random': [], 'FGSM': [], 'rFGSM': [], 'PGD':
    ↪[]}, {'Random': [], 'FGSM': [], 'rFGSM': [], 'PGD': []}

for epsilon in EPS_list_lab2:
    print('epsilon', epsilon)
    for q_label in lab2_label:
        print('q_label', q_label)
        white_acc_lst, black_acc_lst = [], []
        ## Load pretrained models
        whitebox = models.NetA()
        blackbox = models.NetB()

        whitebox.load_state_dict(torch.load("netA_standard.pt")) # TODO
        blackbox.load_state_dict(torch.load("netB_standard.pt")) # TODO

        whitebox, blackbox = whitebox.to(device), blackbox.to(device)
        whitebox.eval()

```

```

blackbox.eval()

test_acc,_ = test_model(whitebox, test_loader, device)
print("Initial Accuracy of Whitebox Model: ",test_acc)
test_acc,_ = test_model(blackbox, test_loader, device)
print("Initial Accuracy of Blackbox Model: ", test_acc)

## Test the models against an adversarial attack

# TODO: Set attack parameters here
ATK_EPS = epsilon
ATK_ITERS = 10
ATK_ALPHA = 1.85 * (ATK_EPS/ATK_ITERS)

whitebox_correct = 0.
blackbox_correct = 0.
running_total = 0.
for batch_idx, (data, labels) in enumerate(test_loader):
    data, labels = data.to(device), labels.to(device)
    # TODO: Perform adversarial attack here
    adv_data = lab2_bcd_return_adv_data(model=whitebox, device=device,
    ↪ dat=data, lbl=labels, eps=ATK_EPS,
                                alpha=ATK_ALPHA,
    ↪ iters=ATK_ITERS, rand_start=True, question_label=q_label)
    # Sanity checking if adversarial example is "legal"
    assert(torch.max(torch.abs(adv_data-data)) <= (ATK_EPS + 1e-5)), \
        "torch.max(torch.abs(adv_data-data)) = %g, %s, ATK_EPS=%g" %
    ↪ (torch.max(torch.abs(adv_data-data)), q_label, ATK_EPS)
    assert(adv_data.max() == 1.), "adv_data.max() = %g, %s, ATK_EPS=%g"
    ↪ % (adv_data.max(), q_label, ATK_EPS)
    assert(adv_data.min() == 0.), "adv_data.min() = %g, %s, ATK_EPS=%g"
    ↪ % (adv_data.min(), q_label, ATK_EPS)

    # Compute accuracy on perturbed data
    with torch.no_grad():
        # Stat keeping - whitebox
        whitebox_outputs = whitebox(adv_data)
        _,whitebox_preds = whitebox_outputs.max(1)
        whitebox_correct += whitebox_preds.eq(labels).sum().item()
        # Stat keeping - blackbox
        blackbox_outputs = blackbox(adv_data)
        _,blackbox_preds = blackbox_outputs.max(1)
        blackbox_correct += blackbox_preds.eq(labels).sum().item()
        running_total += labels.size(0)

## Plot some samples
# if batch_idx == 1:

```



```

# plt.figure(figsize=(15,5))
# for jj in range(12):
#     plt.subplot(2,6,jj+1);plt.imshow(adv_data[jj,0].cpu().
    numpy(), cmap='gray');plt.axis("off")
# plt.tight_layout()
# plt.show()

# Print final
whitebox_acc = whitebox_correct/running_total
blackbox_acc = blackbox_correct/running_total


white_acc_dict[q_label].append(whitebox_acc)
black_acc_dict[q_label].append(blackbox_acc)
print("Attack Epsilon: {}; Whitebox Accuracy: {}; Blackbox Accuracy: {}
    {}".format(ATK_EPS, whitebox_acc, blackbox_acc))
print("Done!")

```

```

[ ]: figure, axis = plot_lab2_white_black(epsilon_arr=EPS_list_lab2,
    white_acc_d=white_acc_dict, black_acc_d=black_acc_dict,
    label_names=lab2_label)
plt.show()

```



## 1.4 Train Robust Models, Lab 3 a/b

Plotting accuracy vs epochs.

```

[6]: def lab3_ab_return_adv_data(model, device, dat, lbl, eps, alpha, iters,
    rand_start, which_method):
    if which_method == 'FGSM':
        return attacks.FGSM_attack(model, device, dat, lbl, eps)[0]
    elif which_method == 'rFGSM':

```

```

        return attacks.rFGSM_attack(model, device, dat, lbl, eps)[0]
    elif which_method == 'PGD':
        return attacks.PGD_attack(model, device, dat, lbl, eps, alpha, iters,
↪rand_start)[0]
    else:
        raise KeyError

def plot_lab3_epoch_vs_test_acc(n_epochs, in_dict, which_part, save=False):
    fig, ax = plt.subplots(1, 1)
    ax.plot(range(n_epochs), in_dict[which_part][2], label='Last Test Acc_
↪(clean data) %.4f' % in_dict[which_part][2][-1])
    ax.plot(range(n_epochs), in_dict[which_part][3], label='Last Train Acc (adv.
↪data) %.4f' % in_dict[which_part][3][-1])
    ax.set_xlabel('Epoch')
    ax.set_ylabel('Accuracy')
    ax.set_title('Adversarial Training (%s attack)' % in_dict[which_part][0])
    ax.legend()

    fig.tight_layout()
    if save:
        plt.savefig('Figures/lab3_%s_%s.pdf' % (which_part,
↪in_dict[which_part][0]), dpi=500, bbox_inches='tight')
        # save name example: Figures/lab3_a1_FSGM.pdf
    return fig, ax

```

```

[ ]: ## lab 3 version of the training code
lab3_parts = ['a1', 'a2', 'b']
# lab3_parts = ['a2']

lab3_adv_training = {'a1': ['FGSM', 'netA_advtrain_fgsm0p1.pt', [], []],
                     'a2': ['rFGSM', 'netA_advtrain_rfgsm0p1.pt', [], []],
                     'b': ['PGD', 'netA_advtrain_pgd0p1.pt', [], []]}

for lab3_part in lab3_parts:
    # lab3_adv_training stores name of attack, name of saved model, and
↪test_acc list
    ATK_EPS = 0.1
    ATK_ITERS = 4 # only for PGD
    ATK_ALPHA = 1.85 * (ATK_EPS/ATK_ITERS) # only for PGD

    ## Pick a model architecture, picked NetA and train from scratch

    net = models.NetA().to(device)
    ## Checkpoint name for this model
    model_checkpoint = lab3_adv_training[lab3_part][1]
    which_method = lab3_adv_training[lab3_part][0]
    ## Basic training params

```

```

num_epochs = 20
initial_lr = 0.001
lr_decay_epoch = 15

optimizer = torch.optim.Adam(net.parameters(), lr=initial_lr)

## Training Loop
for epoch in range(num_epochs):
    net.train()
    train_correct = 0.
    train_loss = 0.
    train_total = 0.
    for batch_idx, (data, labels) in enumerate(train_loader):
        data = data.to(device); labels = labels.to(device)

        adv_data = lab3_ab_return_adv_data(model=net, device=device,
        ↪ dat=data, lbl=labels, eps=ATK_EPS, alpha=ATK_ALPHA,
                                iters=ATK_ITERS,
        ↪ rand_start=True, which_method=which_method)
        # Forward pass
        adv_outputs = net(adv_data)
        net.zero_grad()
        optimizer.zero_grad()
        # Compute loss, gradients, and update params
        loss = F.cross_entropy(adv_outputs, labels)
        loss.backward()
        optimizer.step()
        # Update stats
        _, preds = adv_outputs.max(1)
        train_correct += preds.eq(labels).sum().item()
        train_loss += loss.item()
        train_total += labels.size(0)

    # End of training epoch
    test_acc, test_loss = test_model(net, test_loader, device) # using
    ↪ clean data
    lab3_adv_training[lab3_part][2].append(test_acc)
    lab3_adv_training[lab3_part][3].append(train_correct/train_total)

    print("Epoch: [ {} / {} ]; TrainAcc: {:.5f}; TrainLoss: {:.5f}; TestAcc:
    ↪ {:.5f}; TestLoss: {:.5f}").format(
        epoch, num_epochs, train_correct/train_total, train_loss/
    ↪ len(train_loader),
        test_acc, test_loss,
    ))
    # Save model
    torch.save(net.state_dict(), model_checkpoint)

```

```

    # Update LR
    if epoch == lr_decay_epoch:
        for param_group in optimizer.param_groups:
            param_group['lr'] = initial_lr*0.1

    plot_lab3_epoch_vs_test_acc(n_epochs=num_epochs, in_dict=lab3_adv_training,
↪which_part=lab3_part, save=True)

print("Done!")

```

## 1.5 Test Robust Models, Lab 3 c/d

Don't forget to plot accuracy vs. epsilon curves!

```

[7]: def plot_lab3_eps_vs_acc(epsilon_arr, accuracy_dict, label_names_lst,
↪save=False):
    fig, ax = plt.subplots(1, len(label_names_lst), figsize=(16, 5))

    for iii, model_name in enumerate(label_names_lst):
        for jjj, line_name in enumerate(label_names_lst):
            ax[iii].plot(epsilon_arr, accuracy_dict[model_name][line_name],
↪label=line_name)
            ax[iii].set_xlabel('eps')
            ax[iii].set_ylabel('Accuracy')
            ax[iii].set_title('Model Trained with %s Attack' % model_name)
            ax[iii].legend()

    fig.tight_layout()
    if save:
        plt.savefig('Figures/lab3cd_attacks.pdf', dpi=500, bbox_inches='tight')
    return fig, ax

```

```

[8]: ## lab 3 testing model with adversarial data
EPS_list_lab3cd = np.linspace(0, 0.07, 8) * 2
lab3_labels = ['FGSM', 'rFGSM', 'PGD']
lab3_adv_train_checkpoints = {'FGSM': 'netA_advtrain_fgsm0p1.pt',
                              'rFGSM': 'netA_advtrain_rfgsm0p1.pt',
                              'PGD': 'netA_advtrain_pgd0p1.pt'}

# outer label: name of attack that was used to train the model; inner label:
↪name of attack
lab3_acc_dict = {'FGSM': {'FGSM': [], 'rFGSM': [], 'PGD': []},
                 'rFGSM': {'FGSM': [], 'rFGSM': [], 'PGD': []},
                 'PGD': {'FGSM': [], 'rFGSM': [], 'PGD': []}}

print(EPS_list_lab3cd)

```

```

for epsilon in EPS_list_lab3cd:
    print('epsilon', epsilon)
    for which_model in lab3_labels: # select model using dict in cell above
        model_checkpoint = lab3_adv_train_checkpoints[which_model] # name of
        ↪checkpoint
        # which method was used to train the model, can be "rFGSM, FGSM, PGD"
        print('    which method used to train model', which_model)
        whitebox = models.NetA()
        whitebox.load_state_dict(torch.load(model_checkpoint)) # TODO: Load
        ↪your robust models
        whitebox = whitebox.to(device)
        whitebox.eval()

        test_acc, _ = test_model(whitebox, test_loader, device)
        print("    Initial Accuracy of Whitebox Model: ", test_acc)

        ## Test the model against an adversarial attack

        # TODO: Set attack parameters here
        ATK_EPS = epsilon
        ATK_ITERS = 10 # for testing, use ATK_ITERS = 10
        ATK_ALPHA = 1.85 * (ATK_EPS/ATK_ITERS)
        for which_attack in lab3_labels:
            print('    which attack', which_attack)
            whitebox_correct = 0.
            running_total = 0.
            for batch_idx, (data, labels) in enumerate(test_loader):
                data = data.to(device)
                labels = labels.to(device)

                # TODO: Perform adversarial attack here
                adv_data = lab3_ab_return_adv_data(model=whitebox,
                ↪device=device, dat=data, lbl=labels, eps=ATK_EPS, alpha=ATK_ALPHA,
                iters=ATK_ITERS,
                ↪rand_start=True, which_method=which_attack)
                # Sanity checking if adversarial example is "legal"
                assert(torch.max(torch.abs(adv_data-data)) <= (ATK_EPS + 1e-5) )
                assert(adv_data.max() == 1.)
                assert(adv_data.min() == 0.)

                # Compute accuracy on perturbed data
                with torch.no_grad():
                    whitebox_outputs = whitebox(adv_data)
                    _,whitebox_preds = whitebox_outputs.max(1)
                    whitebox_correct += whitebox_preds.eq(labels).sum().item()
                    running_total += labels.size(0)

```

```

        # Plot some samples
        # if batch_idx == 1:
        #     plt.figure(figsize=(15,5))
        #     for jj in range(12):
        #         plt.subplot(2,6,jj+1);plt.imshow(adv_data[jj,0].cpu().
→numpy(),cmap='gray');plt.axis("off")
        #     plt.tight_layout()
        #     plt.show()
    # Print final
    whitebox_acc = whitebox_correct/running_total
    lab3_acc_dict[which_model][which_attack].append(whitebox_acc)
    print("        Attack Epsilon: {}; Whitebox Accuracy: {}".
→format(ATK_EPS, whitebox_acc))

    print("    Done with this model!")
    print("Done with this epsilon!")

```

```

[0.  0.02 0.04 0.06 0.08 0.1  0.12 0.14]
epsilon 0.0
which method used to train model FGSM
Initial Accuracy of Whitebox Model:  0.6063
    which attack FGSM
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.6063
    which attack rFGSM
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.6063
    which attack PGD
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.6063
Done with this model!
which method used to train model rFGSM
Initial Accuracy of Whitebox Model:  0.8852
    which attack FGSM
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.8852
    which attack rFGSM
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.8852
    which attack PGD
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.8852
Done with this model!
which method used to train model PGD
Initial Accuracy of Whitebox Model:  0.8723
    which attack FGSM
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.8723
    which attack rFGSM
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.8723
    which attack PGD
    Attack Epsilon: 0.0; Whitebox Accuracy: 0.8723
Done with this model!
Done with this epsilon!

```

epsilon 0.02

which method used to train model FGSM

Initial Accuracy of Whitebox Model: 0.6063

which attack FGSM

Attack Epsilon: 0.02; Whitebox Accuracy: 0.5507

which attack rFGSM

Attack Epsilon: 0.02; Whitebox Accuracy: 0.5234

which attack PGD

Attack Epsilon: 0.02; Whitebox Accuracy: 0.433

Done with this model!

which method used to train model rFGSM

Initial Accuracy of Whitebox Model: 0.8852

which attack FGSM

Attack Epsilon: 0.02; Whitebox Accuracy: 0.8633

which attack rFGSM

Attack Epsilon: 0.02; Whitebox Accuracy: 0.8689

which attack PGD

Attack Epsilon: 0.02; Whitebox Accuracy: 0.8605

Done with this model!

which method used to train model PGD

Initial Accuracy of Whitebox Model: 0.8723

which attack FGSM

Attack Epsilon: 0.02; Whitebox Accuracy: 0.8569

which attack rFGSM

Attack Epsilon: 0.02; Whitebox Accuracy: 0.8603

which attack PGD

Attack Epsilon: 0.02; Whitebox Accuracy: 0.8555

Done with this model!

Done with this epsilon!

epsilon 0.04

which method used to train model FGSM

Initial Accuracy of Whitebox Model: 0.6063

which attack FGSM

Attack Epsilon: 0.04; Whitebox Accuracy: 0.5934

which attack rFGSM

Attack Epsilon: 0.04; Whitebox Accuracy: 0.4921

which attack PGD

Attack Epsilon: 0.04; Whitebox Accuracy: 0.2948

Done with this model!

which method used to train model rFGSM

Initial Accuracy of Whitebox Model: 0.8852

which attack FGSM

Attack Epsilon: 0.04; Whitebox Accuracy: 0.8463

which attack rFGSM

Attack Epsilon: 0.04; Whitebox Accuracy: 0.8539

which attack PGD

Attack Epsilon: 0.04; Whitebox Accuracy: 0.8372

Done with this model!

which method used to train model PGD  
 Initial Accuracy of Whitebox Model: 0.8723  
     which attack FGSM  
         Attack Epsilon: 0.04; Whitebox Accuracy: 0.8424  
     which attack rFGSM  
         Attack Epsilon: 0.04; Whitebox Accuracy: 0.85  
     which attack PGD  
         Attack Epsilon: 0.04; Whitebox Accuracy: 0.838  
 Done with this model!  
 Done with this epsilon!  
 epsilon 0.06  
     which method used to train model FGSM  
     Initial Accuracy of Whitebox Model: 0.6063  
         which attack FGSM  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.7214  
         which attack rFGSM  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.4251  
         which attack PGD  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.1361  
 Done with this model!  
     which method used to train model rFGSM  
     Initial Accuracy of Whitebox Model: 0.8852  
         which attack FGSM  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.8313  
         which attack rFGSM  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.8426  
         which attack PGD  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.8139  
 Done with this model!  
     which method used to train model PGD  
     Initial Accuracy of Whitebox Model: 0.8723  
         which attack FGSM  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.83  
         which attack rFGSM  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.8399  
         which attack PGD  
             Attack Epsilon: 0.06; Whitebox Accuracy: 0.8197  
 Done with this model!  
 Done with this epsilon!  
 epsilon 0.08  
     which method used to train model FGSM  
     Initial Accuracy of Whitebox Model: 0.6063  
         which attack FGSM  
             Attack Epsilon: 0.08; Whitebox Accuracy: 0.9366  
         which attack rFGSM  
             Attack Epsilon: 0.08; Whitebox Accuracy: 0.351  
         which attack PGD  
             Attack Epsilon: 0.08; Whitebox Accuracy: 0.0561



Done with this model!  
 which method used to train model rFGSM  
 Initial Accuracy of Whitebox Model: 0.8852  
     which attack FGSM  
         Attack Epsilon: 0.08; Whitebox Accuracy: 0.8194  
     which attack rFGSM  
         Attack Epsilon: 0.08; Whitebox Accuracy: 0.8316  
     which attack PGD  
         Attack Epsilon: 0.08; Whitebox Accuracy: 0.7843  
 Done with this model!  
 which method used to train model PGD  
 Initial Accuracy of Whitebox Model: 0.8723  
     which attack FGSM  
         Attack Epsilon: 0.08; Whitebox Accuracy: 0.8183  
     which attack rFGSM  
         Attack Epsilon: 0.08; Whitebox Accuracy: 0.8302  
     which attack PGD  
         Attack Epsilon: 0.08; Whitebox Accuracy: 0.7992  
 Done with this model!  
 Done with this epsilon!  
 epsilon 0.1  
     which method used to train model FGSM  
     Initial Accuracy of Whitebox Model: 0.6063  
         which attack FGSM  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.9749  
         which attack rFGSM  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.2449  
         which attack PGD  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.0275  
 Done with this model!  
     which method used to train model rFGSM  
     Initial Accuracy of Whitebox Model: 0.8852  
         which attack FGSM  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.8043  
         which attack rFGSM  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.8219  
         which attack PGD  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.7506  
 Done with this model!  
     which method used to train model PGD  
     Initial Accuracy of Whitebox Model: 0.8723  
         which attack FGSM  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.8068  
         which attack rFGSM  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.8224  
         which attack PGD  
             Attack Epsilon: 0.1; Whitebox Accuracy: 0.7801  
 Done with this model!

Done with this epsilon!

epsilon 0.12

which method used to train model FGSM

Initial Accuracy of Whitebox Model: 0.6063

which attack FGSM

Attack Epsilon: 0.12; Whitebox Accuracy: 0.943

which attack rFGSM

Attack Epsilon: 0.12; Whitebox Accuracy: 0.1965

which attack PGD

Attack Epsilon: 0.12; Whitebox Accuracy: 0.0149

Done with this model!

which method used to train model rFGSM

Initial Accuracy of Whitebox Model: 0.8852

which attack FGSM

Attack Epsilon: 0.12; Whitebox Accuracy: 0.7374

which attack rFGSM

Attack Epsilon: 0.12; Whitebox Accuracy: 0.7235

which attack PGD

Attack Epsilon: 0.12; Whitebox Accuracy: 0.4457

Done with this model!

which method used to train model PGD

Initial Accuracy of Whitebox Model: 0.8723

which attack FGSM

Attack Epsilon: 0.12; Whitebox Accuracy: 0.7865

which attack rFGSM

Attack Epsilon: 0.12; Whitebox Accuracy: 0.7777

which attack PGD

Attack Epsilon: 0.12; Whitebox Accuracy: 0.6218

Done with this model!

Done with this epsilon!

epsilon 0.14

which method used to train model FGSM

Initial Accuracy of Whitebox Model: 0.6063

which attack FGSM

Attack Epsilon: 0.14; Whitebox Accuracy: 0.8118

which attack rFGSM

Attack Epsilon: 0.14; Whitebox Accuracy: 0.2299

which attack PGD

Attack Epsilon: 0.14; Whitebox Accuracy: 0.0091

Done with this model!

which method used to train model rFGSM

Initial Accuracy of Whitebox Model: 0.8852

which attack FGSM

Attack Epsilon: 0.14; Whitebox Accuracy: 0.6242

which attack rFGSM


Attack Epsilon: 0.14; Whitebox Accuracy: 0.5502

which attack PGD

Attack Epsilon: 0.14; Whitebox Accuracy: 0.1729

Done with this model!  
 which method used to train model PGD  
 Initial Accuracy of Whitebox Model: 0.8723  
 which attack FGSM  
 Attack Epsilon: 0.14; Whitebox Accuracy: 0.6945  
 which attack rFGSM  
 Attack Epsilon: 0.14; Whitebox Accuracy: 0.573  
 which attack PGD  
 Attack Epsilon: 0.14; Whitebox Accuracy: 0.2582  
 Done with this model!  
 Done with this epsilon!

```
[9]: _, _ = plot_lab3_eps_vs_acc(epsilon_arr=EPS_list_lab3cd,
    ↪ accuracy_dict=lab3_acc_dict, label_names_lst=lab3_labels, save=True)
```



## 1.6 Lab-3e Bonus (train models using PGD AT with different epsilon)

```
[9]: def plot_lab3_bonus_epoch_vs_test_acc(n_epochs, in_dict, eps_list, save=False):
    ## in_dict now maps eps to test_acc_arr
    fig, ax = plt.subplots(1, len(in_dict.keys()), figsize=(16, 5))
    for iii, eps in enumerate(eps_list):
        ax[iii].plot(range(n_epochs), in_dict[eps][0], label='Last Test Acc_
    ↪ (clean data) %.4f' % (in_dict[eps][0][-1]))
        ax[iii].plot(range(n_epochs), in_dict[eps][1], label='Last Train Acc_
    ↪ (adv. data) %.4f' % (in_dict[eps][1][-1]))
        ax[iii].set_xlabel('Epoch')
        ax[iii].set_ylabel('Accuracy')
        ax[iii].set_title('Adversarial Training (PGD attack), AT eps=%g' % eps)
        ax[iii].legend()
    fig.tight_layout()
    if save:
        plt.savefig('Figures/lab3e_testAcc.pdf', dpi=500, bbox_inches='tight')
    return fig, ax
```

```

[7]: ## lab 3 bonus version of the training code, part (e)
eps_lab3_bonus = [0.05, 0.2, 0.4]
lab3_bonus_AT_dict = {eps_val: [[], []] for eps_val in eps_lab3_bonus}
# first list = test_acc, second list = train_acc
## Basic training params

num_epochs = 20
initial_lr = 0.001
lr_decay_epoch = 15

for epsilon in eps_lab3_bonus:
    print('epsilon', epsilon)
    # lab3_adv_training stores name of attack, name of saved model, and
    ↪ test_acc list
    ATK_EPS = epsilon
    ATK_ITERS = 4 # only for PGD
    ATK_ALPHA = 1.85 * (ATK_EPS/ATK_ITERS) # only for PGD

    ## Pick a model architecture, picked NetA and train from scratch

    net = models.NetA().to(device)
    ## Checkpoint name for this model
    model_checkpoint = 'netA_advtrain_pgd_eps_%g.pt' % epsilon

    optimizer = torch.optim.Adam(net.parameters(), lr=initial_lr)

    ## Training Loop
    for epoch in range(num_epochs):
        net.train()
        train_correct = 0.
        train_loss = 0.
        train_total = 0.
        for batch_idx, (data, labels) in enumerate(train_loader):
            data = data.to(device); labels = labels.to(device)

            adv_data = attacks.PGD_attack(model=net, device=device, dat=data,
            ↪ lbl=labels, eps=ATK_EPS, alpha=ATK_ALPHA,
                                   iters=ATK_ITERS, rand_start=True)[0]

            # Forward pass
            adv_outputs = net(adv_data)
            net.zero_grad()
            optimizer.zero_grad()
            # Compute loss, gradients, and update params
            loss = F.cross_entropy(adv_outputs, labels)
            loss.backward()
            optimizer.step()
            # Update stats

```

```

        _, preds = adv_outputs.max(1)
        train_correct += preds.eq(labels).sum().item()
        train_loss += loss.item()
        train_total += labels.size(0)

    # End of training epoch
    test_acc, test_loss = test_model(net, test_loader, device) # using
    ↪ clean data
    lab3_bonus_AT_dict[epsilon][0].append(test_acc)
    lab3_bonus_AT_dict[epsilon][1].append(train_correct/train_total)

    print("Epoch: [ {} / {} ]; TrainAcc: {:.5f}; TrainLoss: {:.5f}; TestAcc:
    ↪ {:.5f}; TestLoss: {:.5f}").format(
        epoch, num_epochs, train_correct/train_total, train_loss/
    ↪ len(train_loader),
        test_acc, test_loss,
    ))
    # Save model
    torch.save(net.state_dict(), model_checkpoint)

    # Update LR
    if epoch == lr_decay_epoch:
        for param_group in optimizer.param_groups:
            param_group['lr'] = initial_lr*0.1
    print("Done, epsilon %g" % epsilon)
print("Done!")

# write eps vals in to a file
f_ptr = open('lab3_bonus_eps_vals.txt', 'w')
for eps in eps_lab3_bonus:
    f_ptr.write('%g\n' % eps)
f_ptr.close()

```

```

epsilon 0.05
Epoch: [ 0 / 20 ]; TrainAcc: 0.73450; TrainLoss: 0.67171; TestAcc: 0.84420;
TestLoss: 0.40490
Epoch: [ 1 / 20 ]; TrainAcc: 0.78738; TrainLoss: 0.53200; TestAcc: 0.86760;
TestLoss: 0.35963
Epoch: [ 2 / 20 ]; TrainAcc: 0.80158; TrainLoss: 0.49526; TestAcc: 0.86780;
TestLoss: 0.35068
Epoch: [ 3 / 20 ]; TrainAcc: 0.80940; TrainLoss: 0.47151; TestAcc: 0.87190;
TestLoss: 0.34070
Epoch: [ 4 / 20 ]; TrainAcc: 0.81770; TrainLoss: 0.45434; TestAcc: 0.87610;
TestLoss: 0.33318
Epoch: [ 5 / 20 ]; TrainAcc: 0.82280; TrainLoss: 0.44043; TestAcc: 0.87400;
TestLoss: 0.33037
Epoch: [ 6 / 20 ]; TrainAcc: 0.82738; TrainLoss: 0.42892; TestAcc: 0.87930;

```

TestLoss: 0.31664  
 Epoch: [ 7 / 20 ]; TrainAcc: 0.82920; TrainLoss: 0.42050; TestAcc: 0.88150;  
 TestLoss: 0.30749  
 Epoch: [ 8 / 20 ]; TrainAcc: 0.83325; TrainLoss: 0.41069; TestAcc: 0.88020;  
 TestLoss: 0.30810  
 Epoch: [ 9 / 20 ]; TrainAcc: 0.83573; TrainLoss: 0.40491; TestAcc: 0.88390;  
 TestLoss: 0.30630  
 Epoch: [ 10 / 20 ]; TrainAcc: 0.83795; TrainLoss: 0.39746; TestAcc: 0.88460;  
 TestLoss: 0.30458  
 Epoch: [ 11 / 20 ]; TrainAcc: 0.84070; TrainLoss: 0.39255; TestAcc: 0.88800;  
 TestLoss: 0.29792  
 Epoch: [ 12 / 20 ]; TrainAcc: 0.84148; TrainLoss: 0.38762; TestAcc: 0.87920;  
 TestLoss: 0.31755  
 Epoch: [ 13 / 20 ]; TrainAcc: 0.84342; TrainLoss: 0.38365; TestAcc: 0.88770;  
 TestLoss: 0.29869  
 Epoch: [ 14 / 20 ]; TrainAcc: 0.84450; TrainLoss: 0.37899; TestAcc: 0.88520;  
 TestLoss: 0.29808  
 Epoch: [ 15 / 20 ]; TrainAcc: 0.84840; TrainLoss: 0.37333; TestAcc: 0.88120;  
 TestLoss: 0.29977  
 Epoch: [ 16 / 20 ]; TrainAcc: 0.86108; TrainLoss: 0.33813; TestAcc: 0.89720;  
 TestLoss: 0.27954  
 Epoch: [ 17 / 20 ]; TrainAcc: 0.86385; TrainLoss: 0.33086; TestAcc: 0.89350;  
 TestLoss: 0.27931  
 Epoch: [ 18 / 20 ]; TrainAcc: 0.86533; TrainLoss: 0.32790; TestAcc: 0.89670;  
 TestLoss: 0.27869  
 Epoch: [ 19 / 20 ]; TrainAcc: 0.86557; TrainLoss: 0.32558; TestAcc: 0.89570;  
 TestLoss: 0.27975  
 Done, epsilon 0.05  
 epsilon 0.2  
 Epoch: [ 0 / 20 ]; TrainAcc: 0.51630; TrainLoss: 1.19701; TestAcc: 0.75040;  
 TestLoss: 0.63336  
 Epoch: [ 1 / 20 ]; TrainAcc: 0.65197; TrainLoss: 0.83879; TestAcc: 0.80350;  
 TestLoss: 0.55514  
 Epoch: [ 2 / 20 ]; TrainAcc: 0.69597; TrainLoss: 0.73323; TestAcc: 0.80630;  
 TestLoss: 0.52417  
 Epoch: [ 3 / 20 ]; TrainAcc: 0.71938; TrainLoss: 0.68434; TestAcc: 0.81840;  
 TestLoss: 0.51014  
 Epoch: [ 4 / 20 ]; TrainAcc: 0.73388; TrainLoss: 0.65096; TestAcc: 0.82030;  
 TestLoss: 0.47191  
 Epoch: [ 5 / 20 ]; TrainAcc: 0.74407; TrainLoss: 0.62970; TestAcc: 0.82060;  
 TestLoss: 0.47322  
 Epoch: [ 6 / 20 ]; TrainAcc: 0.75050; TrainLoss: 0.61232; TestAcc: 0.82620;  
 TestLoss: 0.46323  
 Epoch: [ 7 / 20 ]; TrainAcc: 0.75450; TrainLoss: 0.59980; TestAcc: 0.83120;  
 TestLoss: 0.45106  
 Epoch: [ 8 / 20 ]; TrainAcc: 0.76132; TrainLoss: 0.58800; TestAcc: 0.82830;  
 TestLoss: 0.44952  
 Epoch: [ 9 / 20 ]; TrainAcc: 0.76327; TrainLoss: 0.57910; TestAcc: 0.83140;

TestLoss: 0.44403  
 Epoch: [ 10 / 20 ]; TrainAcc: 0.76572; TrainLoss: 0.57348; TestAcc: 0.82830;  
 TestLoss: 0.44780  
 Epoch: [ 11 / 20 ]; TrainAcc: 0.77102; TrainLoss: 0.56074; TestAcc: 0.82750;  
 TestLoss: 0.44293  
 Epoch: [ 12 / 20 ]; TrainAcc: 0.77337; TrainLoss: 0.55642; TestAcc: 0.83260;  
 TestLoss: 0.43581  
 Epoch: [ 13 / 20 ]; TrainAcc: 0.77522; TrainLoss: 0.54891; TestAcc: 0.83300;  
 TestLoss: 0.43869  
 Epoch: [ 14 / 20 ]; TrainAcc: 0.77975; TrainLoss: 0.54332; TestAcc: 0.83750;  
 TestLoss: 0.42460  
 Epoch: [ 15 / 20 ]; TrainAcc: 0.78157; TrainLoss: 0.53704; TestAcc: 0.83940;  
 TestLoss: 0.43287  
 Epoch: [ 16 / 20 ]; TrainAcc: 0.79227; TrainLoss: 0.51197; TestAcc: 0.84470;  
 TestLoss: 0.41035  
 Epoch: [ 17 / 20 ]; TrainAcc: 0.79607; TrainLoss: 0.50512; TestAcc: 0.84740;  
 TestLoss: 0.40755  
 Epoch: [ 18 / 20 ]; TrainAcc: 0.79582; TrainLoss: 0.50423; TestAcc: 0.84590;  
 TestLoss: 0.40835  
 Epoch: [ 19 / 20 ]; TrainAcc: 0.79747; TrainLoss: 0.50031; TestAcc: 0.84600;  
 TestLoss: 0.40717  
 Done, epsilon 0.2  
 epsilon 0.4  
 Epoch: [ 0 / 20 ]; TrainAcc: 0.25575; TrainLoss: 1.89715; TestAcc: 0.57650;  
 TestLoss: 1.13337  
 Epoch: [ 1 / 20 ]; TrainAcc: 0.54922; TrainLoss: 1.03925; TestAcc: 0.65180;  
 TestLoss: 1.01517  
 Epoch: [ 2 / 20 ]; TrainAcc: 0.62892; TrainLoss: 0.89012; TestAcc: 0.69420;  
 TestLoss: 0.82792  
 Epoch: [ 3 / 20 ]; TrainAcc: 0.64958; TrainLoss: 0.84977; TestAcc: 0.69720;  
 TestLoss: 0.83903  
 Epoch: [ 4 / 20 ]; TrainAcc: 0.67225; TrainLoss: 0.80242; TestAcc: 0.69730;  
 TestLoss: 0.83425  
 Epoch: [ 5 / 20 ]; TrainAcc: 0.68480; TrainLoss: 0.77602; TestAcc: 0.70230;  
 TestLoss: 0.86794  
 Epoch: [ 6 / 20 ]; TrainAcc: 0.69162; TrainLoss: 0.76109; TestAcc: 0.71500;  
 TestLoss: 0.76351  
 Epoch: [ 7 / 20 ]; TrainAcc: 0.69753; TrainLoss: 0.74598; TestAcc: 0.70360;  
 TestLoss: 0.79814  
 Epoch: [ 8 / 20 ]; TrainAcc: 0.70518; TrainLoss: 0.73418; TestAcc: 0.72430;  
 TestLoss: 0.75679  
 Epoch: [ 9 / 20 ]; TrainAcc: 0.70942; TrainLoss: 0.71949; TestAcc: 0.71130;  
 TestLoss: 0.76780  
 Epoch: [ 10 / 20 ]; TrainAcc: 0.70678; TrainLoss: 0.72502; TestAcc: 0.73820;  
 TestLoss: 0.69110  
 Epoch: [ 11 / 20 ]; TrainAcc: 0.70462; TrainLoss: 0.72950; TestAcc: 0.74060;  
 TestLoss: 0.68806  
 Epoch: [ 12 / 20 ]; TrainAcc: 0.71888; TrainLoss: 0.70329; TestAcc: 0.73550;

```


TestLoss: 0.67693
Epoch: [ 13 / 20 ]; TrainAcc: 0.72423; TrainLoss: 0.68603; TestAcc: 0.74520;
TestLoss: 0.66706
Epoch: [ 14 / 20 ]; TrainAcc: 0.72882; TrainLoss: 0.67622; TestAcc: 0.74920;
TestLoss: 0.67956
Epoch: [ 15 / 20 ]; TrainAcc: 0.73053; TrainLoss: 0.67330; TestAcc: 0.74790;
TestLoss: 0.65314
Epoch: [ 16 / 20 ]; TrainAcc: 0.74838; TrainLoss: 0.62920; TestAcc: 0.75840;
TestLoss: 0.65737
Epoch: [ 17 / 20 ]; TrainAcc: 0.74638; TrainLoss: 0.62725; TestAcc: 0.75700;
TestLoss: 0.65066
Epoch: [ 18 / 20 ]; TrainAcc: 0.74917; TrainLoss: 0.62426; TestAcc: 0.75880;
TestLoss: 0.64633
Epoch: [ 19 / 20 ]; TrainAcc: 0.74810; TrainLoss: 0.62980; TestAcc: 0.76130;
TestLoss: 0.64954
Done, epsilon 0.4
Done!

```

```

[10]: _, _ = plot_lab3_bonus_epoch_vs_test_acc(n_epochs=num_epochs,
        ↪in_dict=lab3_bonus_AT_dict, eps_list=eps_lab3_bonus, save=True)

```



```

[16]: def plot_lab3_bonus_eps_vs_acc(epsilon_arr, accuracy_dict, attack_list,
        ↪save=False):
    fig, ax = plt.subplots(1, len(accuracy_dict.keys()), figsize=(16, 5/1.2))

    for iii, model_eps in enumerate(accuracy_dict):
        for jjj, attack_name in enumerate(attack_list):
            ax[iii].plot(epsilon_arr, accuracy_dict[model_eps][attack_name],
            ↪label='%s attack' % attack_name)
            ax[iii].set_xlabel('eps')
            ax[iii].set_ylabel('Accuracy')
            ax[iii].set_title('Model Trained with PGD Attack\nmodel_eps=%g' %
            ↪model_eps)
            ax[iii].legend()

```



```

fig.tight_layout()
if save:
    plt.savefig('Figures/lab3e_acc_vs_eps.pdf', dpi=500,
↳bbox_inches='tight')
    return fig, ax

```

```

[ ]: ## lab 3 testing model with adversarial data
f_ptr = open('lab3_bonus_eps_vals.txt', 'r')
lines = f_ptr.readlines()
f_ptr.close()

eps_lst_from_file = [float(line) for line in lines]
eps_lst_from_file.append(0.1)
eps_lst_from_file = sorted(eps_lst_from_file)
EPS_list_lab3_bonus = np.linspace(0, 0.44, 23)
lab3_bonus_checkpoint_base = 'netA_advtrain_pgd_eps_'

# outer label: name of attack that was used to train the model; inner label:
↳name of attack
attack_list_lab3_bonus = ['FGSM', 'rFGSM', 'PGD']
lab3_bonus_attack_acc_dict = {eps: {at: [] for at in attack_list_lab3_bonus}
↳for eps in eps_lst_from_file}

print(EPS_list_lab3_bonus)
for epsilon in EPS_list_lab3_bonus: # attack EPS
    print('attack epsilon', epsilon)
    for model_eps in eps_lst_from_file: # EPS used to train model
        if model_eps != 0.1:
            model_checkpoint = lab3_bonus_checkpoint_base + '%g.pt' % model_eps
↳ # name of checkpoint
        else:
            model_checkpoint = 'netA_advtrain_pgd0p1.pt' # name of checkpoint

    print('    which eps used to train model', model_eps)
    whitebox = models.NetA()
    whitebox.load_state_dict(torch.load(model_checkpoint)) # TODO: Load
↳your robust models
    whitebox = whitebox.to(device)
    whitebox.eval();

    test_acc, _ = test_model(whitebox, test_loader, device)
    print("    Initial Accuracy of Whitebox Model: ", test_acc)

    ## Test the model against an adversarial attack

    # TODO: Set attack parameters here

```

```

ATK_EPS = epsilon
ATK_ITERS = 10 # for testing, ATK_ITERS = 10
ATK_ALPHA = 1.85 * (ATK_EPS/ATK_ITERS)
for which_attack in attack_list_lab3_bonus:
    print('        which attack', which_attack)
    whitebox_correct = 0.
    running_total = 0.
    for batch_idx, (data, labels) in enumerate(test_loader):
        data = data.to(device)
        labels = labels.to(device)

        # TODO: Perform adversarial attack here
        adv_data = lab3_ab_return_adv_data(model=whitebox,
↪device=device, dat=data, lbl=labels, eps=ATK_EPS, alpha=ATK_ALPHA,
                                                iters=ATK_ITERS,
↪rand_start=True, which_method=which_attack)
        # Sanity checking if adversarial example is "legal"
        assert(torch.max(torch.abs(adv_data-data)) <= (ATK_EPS + 1e-5) )
        assert(adv_data.max() == 1.)
        assert(adv_data.min() == 0.)

        # Compute accuracy on perturbed data
        with torch.no_grad():
            whitebox_outputs = whitebox(adv_data)
            _,whitebox_preds = whitebox_outputs.max(1)
            whitebox_correct += whitebox_preds.eq(labels).sum().item()
            running_total += labels.size(0)

        # Print final
        whitebox_acc = whitebox_correct/running_total
        lab3_bonus_attack_acc_dict[model_eps][which_attack] .
↪append(whitebox_acc)
        print("        Attack Epsilon: {}; Whitebox Accuracy: {}".
↪format(ATK_EPS, whitebox_acc))


    print("    Done with this model!")
    print("Done with this epsilon!")

```

```

[17]: _, _ = plot_lab3_bonus_eps_vs_acc(epsilon_arr=EPS_list_lab3_bonus,
↪accuracy_dict=lab3_bonus_attack_acc_dict,
                                                attack_list=attack_list_lab3_bonus, save=True)

```



## 1.7 Lab-3f, saliency maps (non-AT and PGD-AT models)

```
[5]: def plot_lab3_bonus_saliency(save=False):
    f_ptr = open('lab3_bonus_eps_vals.txt', 'r')
    lines = f_ptr.readlines()
    f_ptr.close()
    eps_lst_from_file = [float(line) for line in lines]
    eps_lst_from_file.append(0.1)
    eps_lst_from_file.append(-1)
    eps_lst_from_file = sorted(eps_lst_from_file)
    model_checkpoints = {'netA_advtrain_pgd_eps_%g.pt' % eps: eps for eps in
    ↪eps_lst_from_file if eps != 0.1}
    model_checkpoints['netA_standard.pt'] = -1
    model_checkpoints['netA_advtrain_pgd0p1.pt'] = 0.1

    inv_model_checkpoints = {v: k for k, v in model_checkpoints.items()}

    num_examples = 6
    inds = 0
    data, labels = None, None
    for d, l in test_loader:
        data, labels = d.to(device), l.to(device)
        inds = random.sample(list(range(data.size(0))), num_examples) # which
    ↪data points in batch to plot
        break
    fig, ax = plt.subplots(len(eps_lst_from_file) + 1, num_examples,
    ↪figsize=(13, 14))
    for jj in range(num_examples):
        ax[0, jj].imshow(data[inds[jj], 0].cpu().numpy(), cmap='gray',
    ↪interpolation='nearest')
        ax[0, jj].axis("off")
        if jj == 0:
            ax[0, jj].set_title('data')
    for i, eps in enumerate(eps_lst_from_file):
        model_checkpoint = inv_model_checkpoints[eps]
```

```

whitebox = models.NetA()
whitebox.load_state_dict(torch.load(model_checkpoint)) # TODO: Load
↳ your robust models
whitebox = whitebox.to(device)
whitebox.eval()

grad_wrt_data = attacks.gradient_wrt_data(whitebox, device, data,
↳ lbl=labels)

for jj in range(num_examples):
    ax[i+1, jj].imshow(grad_wrt_data[inds[jj], 0].cpu().numpy(),
↳ cmap='gray', interpolation='nearest')
    ax[i+1, jj].axis("off")
    # plt.title("clean. pred={}".format(classes[clean_preds[inds[jj]]]))
    if jj == 0:
        ax[i+1, jj].set_title('Saliency map,eps=%g' % eps)
        if eps == -1:
            ax[i+1, jj].set_title('netA_standard')

fig.suptitle("Saliency maps for different models")
fig.tight_layout()
if save:
    plt.savefig('Figures/lab3f_saliency.pdf', dpi=500, bbox_inches='tight')
return

plot_lab3_bonus_saliency(save=True)

```